Frontier and Localhost: A Survey of How Production AI Learns Outside the Weights

Mikhail L. Arbuzov
Independent Researcher
mike.arbuzov54@gmail.com

Lee Mosbacker
Independent Researcher
lee.mosbacker@gmail.com

Sisong Bei
Independent Researcher
qurining@gmail.com

Ziwei Dong
Independent Researcher
ziwei.dong@alumni.emory.edu

Dmitri Kalaev
Independent Researcher
kalaevdr@gmail.com

Alexey A. Shvets
Palo Alto Networks
ashvets@paloaltonetworks.com

Abstract

Production systems built on large language models now adapt mostly outside the model weights, and this article asks what optimization discipline that adaptation needs. After a failure, teams edit instructions, skills, memories, tools, orchestration graphs, evaluation suites, and governance rules far faster than any weight update. Two earlier results of this series suggest why that layer is the right place: errors concentrate in a few key decisions that fail in a small, recurring catalogue of modes. We read the layer as one surface governed by one loop: a recurring failure is traced to a scaffold coordinate, one edit is proposed, a gate accepts or rejects it on measured loss, and a second decision sets how far the accepted edit may travel. Five conditions make this loop an optimizer rather than patchwork; the decisive one, a measurable loss, depends on the feedback signal the gate receives: an exact check (code, mathematics), an observable outcome (completed tasks), a judgment (writing, design), or a delayed reward (research). The signal's class fixes what the gate can be, which pairings of algorithmic and human gates are reachable, and which failure dominates. Across 256 production and research systems, scored on six substrates, with the scores released, the local loop is routine and the governed second loop is undocumented: no surveyed system combines automated local adaptation with promotion across organizations under versioned lineage, access control, and rollback. We map the failure surface this gap leaves open and state the open problem each signal class poses.

1 Introduction

The standard mental model of a system built on a large language model (LLM) is the model. Pre-training fixes the weights, fine-tuning nudges them, and deployment wraps a thin layer of prompts and tools around the result. Production in 2025 and 2026 has stopped behaving this way. Frontier weights change on a cycle of months; the systems built on them change daily, in instruction files checked into repositories, memories that persist across sessions, skills loaded on demand, registered tools, evaluation suites that block a merge, agent topologies, and pull-request gates on shared rule repositories. A growing share of the adaptation that makes a deployed system reliable now happens outside the weights, in a fast-changing local scaffold wrapped around a slow-changing frontier model.

Read one at a time, these artifacts look like unrelated engineering. A skill file is a procedure the model did not have; a repository rule is a convention for one project; a retrieval connector is memory access for one workflow; a tool registry provisions a capability the base model cannot reach; a pull-request review against a shared rule repository is a promotion gate. Read together, they are one adaptation surface, the non-model part of what Zaharia et al. call a compound AI system [1]. This survey asks what optimization discipline that surface needs, and finds that the field has built the layer and, as far as public evidence shows, not the optimizer that should govern it. Most deployments maintain the layer as patchwork: a failure happens, someone adds a rule, the pile grows, nobody knows which addition helped, and fixes leak into contexts that never validated them.

The disciplined alternative is one loop with two halves. Its first half turns a recurring failure into one tested edit to the scaffold and keeps the edit only if the deployment's own failure rate falls; the second decides how far the edit may travel, to a team, to an organization, or to every customer, and lets it travel only as far as evidence supports. The first half is routine: in most research systems and a growing number of production ones it closes under an automatic gate. None of the 256 systems surveyed has a second half governed by versioned lineage, access control, and rollback. What separates the two halves is the feedback the gate receives. A test or an observed outcome can gate an edit automatically; a judgment can only once the judge is calibrated against people; a reward that arrives months later cannot gate edits one at a time. The class of the signal decides which half of the loop a deployment can close.

This article makes five contributions.

1. The loop that turns patchwork into optimization, in plain language and as a gated update rule, with the five conditions under which it is an optimizer (Section 3).
2. Four classes of feedback signal, from an exact check to a delayed reward, each fixing what the gate can be and which failure dominates (Section 4).
3. A survey of 256 production and research systems, scored on six substrates against a published rubric, with the scores released (Sections 5 and 6).
4. A design space indexed by the gate on each half of the loop, and the convergence of five research families on the same loop (Sections 7 and 8).
5. A four-property audit criterion for governed cross-organization scaffold optimization, which no surveyed system meets, with the failure surface its absence leaves open and the open problem each signal class poses (Sections 9 to 11).

The survey runs no experiments. Its evidence is the published behavior of 256 systems, scored on a released table.

2 Why adaptation lives outside the weights

2.1 Breadth in the weights, depth in the scaffold

Frontier training optimizes for breadth. The weights are tuned to generalize across millions of unseen deployments, and the objective smooths over exactly what one deployment cares about: local naming conventions, contradictory preferences between customers, rare workflow-specific failures. Reinforcement learning from human feedback reduces output diversity and homogenizes preferred style across users [2, 3], and instruction-tuning data rewards patterns that generalize across the annotator pool, so the frontier model is trained to disfavor one customer's idiosyncrasies. Tiwari et al. [4] reach the same place from the optimization side: in-context adaptation is cheap but capacity-bounded, in-weights adaptation is expressive but induces forgetting and loss of plasticity

[5], so a system that learns slowly in the weights needs a cheap outer layer that learns fast outside them.

Production reliability is won in depth, inside one deployment, and it turns on what global training had to smooth away. Encoding every deployment in the weights is ruled out three times over: fine-tuning per customer multiplies the training, evaluation, and serving lifecycle by the number of deployments; weights ship quarterly at best, while the deployment wants daily updates; and overfitting the weights to one customer degrades the neighbors whose preferences contradict. In principle, per-customer adapters [6], distilled models, and learned routing could absorb some of the residual, the failures that remain after training. In practice they are rare. Adaptation lands in the scaffold, the one substrate that is cheap, inspectable, reversible, scoped to one customer, and updatable daily.

2.2 Why the scaffold is the right architecture: the two earlier results

The case for learning outside the weights does not rest on cost alone. The two papers that precede this one in the series supply the hypothesis that explains the pattern the corpus shows, and each states the scope on which it rests. *Beyond Exponential Decay* [7] argues that long outputs do not fail at the rate a uniform per-token error model predicts, because reliability hinges on a sparse set of key decisions. Its evidence is twofold: a direct measurement of long-range dependence on one natural-text corpus, where under a tenth of tokens depend on distant context, and the accuracy that majority voting recovers on reasoning benchmarks. The argument stops where the number of decisions grows with the task, as in long-horizon agentic work. For a deployment that operates inside a bounded slice of recurring tasks, tools, and evaluator expectations, *The Architecture of Errors* [8] argues that those decisions fail in a small recurring catalogue of modes whose size grows only slowly with the number of observed failures. A targeted intervention budget that grows polylogarithmically in sequence length, and becomes constant once the catalogue saturates, then covers the residual. The paper's own negative result marks the limit: no finite catalogue covers an unbounded domain.

If those results hold in general, a deployment's errors are few enough to fix one at a time, and a small catalogue of fixable modes is exactly what a cheap, inspectable, reversible artifact layer is for. This article asks where the fixes live, and the corpus answers: in persistent artifacts outside the weights, one recurring mode at a time, which is the pattern the hypothesis predicts. On that hypothesis the gated and scoped layer is not a workaround for an inadequate model. It is the architecture the error structure calls for: a frontier model that generalizes, a local scaffold that specializes, and reliability that comes from governing the specialization. The survey's own findings, in Sections 6 to 9, stand without it.

2.3 Relation to existing surveys

This survey classifies systems on four axes: whether a feedback loop closes at all, whether its gate is a program or a person, what the gate can observe, and whether promotion beyond one deployment is governed. Surveys of self-evolving agents, agent memory, context engineering, skills, and harness design catalogue mechanisms by what they edit and when; none uses all four axes, and none covers promotion across organizations (Table 1). The one closest in spirit, Ding et al.'s survey of persistent-state agents, finds that literature concentrated on accumulating and retrieving state rather than on governing, recovering, or relinquishing it [9].

A reader will expect the prompt-optimization lineage here: OPRO, EvoPrompt, Promptbreeder, DSPy, and TextGrad [18–22]. That lineage optimizes a prompt or program once against a develop-

Table 1: Related surveys. The last column says which of this survey's four axes each one carries, as a classification axis or only in passing.

Survey	Scope	Organizes by	Of the four axes here
Gao et al. [10]	self-evolving agents	what, when, how, where to evolve	loop closure and gate type, in passing
Fang et al. [11]	self-evolving agents	inputs, agent system, environment, optimizers	loop closure, in passing
Zhou et al. [12]	self-evolving coding agents (108 methods)	evolution target, timing, code signals	all four in passing; reversibility as a governance axis
Ding et al. [9]	persistent-state agents (435 works)	authority, scope, mutability, provenance, recoverability	governance (provenance, rollback) as an axis; gate type in passing
Lin et al. [13]	memory security (~90 works)	lifecycle phases × security objectives	governance (principal scoping, rollback) as an axis; gate type in passing
Guo et al. [14]	agent system and harness design	six runtime responsibilities	all four, in passing
Xia et al. [15]	skill-mediated agents (8 systems)	four responsibility layers, ten patterns	loop closure and governance, in passing
Ding et al. [16]	skill evaluation and evolution	four evolution paradigms × six benchmark classes	all four, in passing
Mei et al. [17]	context engineering (1,411 papers)	components × system implementations	none
This survey	256 scaffold-adaptive systems	loop closure, gate type, signal class, governance	all four, as classification axes

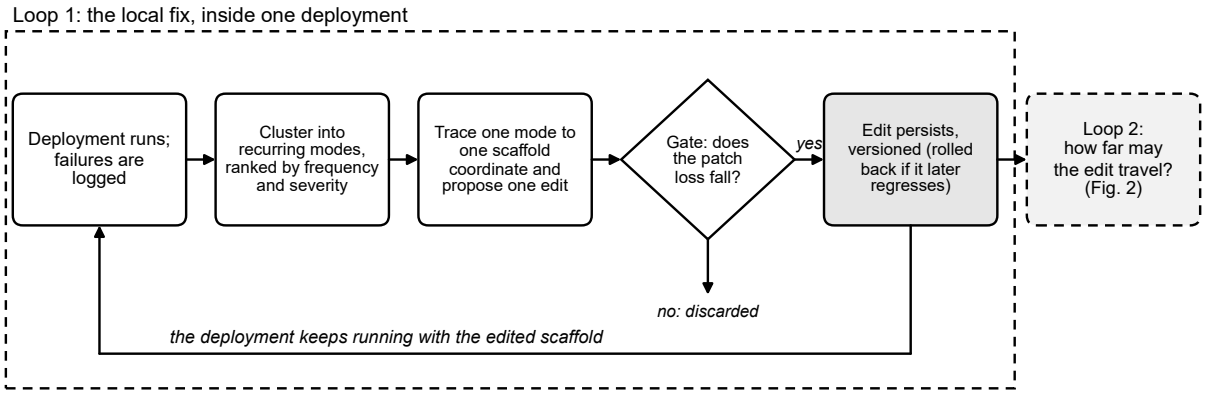


Figure 1: Loop 1, the local half of artifact-layer descent. A recurring failure is traced to one scaffold coordinate, one edit is proposed, and a gate keeps the edit only if the deployment's patch loss falls. An accepted edit then enters Loop 2 (Fig. 2).

ment set and is not a deployed loop. Agentic context engineering [23] does persist its artifact, an evolving playbook, but keeps it inside one context and never promotes it. The closest precedents optimize the scaffold itself: agent symbolic learning updates prompts, tools, and pipeline as symbolic parameters from textual feedback [24]; in STOP, a scaffolding program recursively improves its own improvement code [25]; and a coding agent edits its own harness to raise its benchmark score [26]. Each runs one deployment's loop under a benchmark score. What changes here is that the edited artifact persists in a deployment and must clear a second gate before it earns any scope beyond that deployment (Appendix H).

3 The loop: artifact-layer descent

3.1 The loop in plain language

We call the disciplined version of scaffold maintenance artifact-layer descent. It is one loop with two halves, run again and again over the life of a deployment: Loop 1 fixes a recurring failure inside one deployment, and Loop 2 decides how far the fix may travel. Fig. 1 draws the first half. The deployment runs and logs what goes wrong: failed outputs, user corrections, rejected answers, evaluation regressions, tool-call errors. The loop does not react to single failures. It clusters the log into recurring modes and takes up the ones that matter most, ranked by how often each recurs and how badly it fails. For such a mode it locates the scaffold coordinate responsible, which may be a stale memory entry, an under-specified instruction, a missing tool argument, a brittle retrieval rule, or a misrouted agent edge, and it proposes one candidate edit to that coordinate.

Then it measures. A deployment serves one **patch**: the slice of users, tasks, and conventions it handles. The failure rate on that patch's recurring tasks is the **patch loss**. A gate accepts the edit if the patch loss falls and rejects it otherwise. An accepted edit stays in the scaffold, versioned; one that a later check shows made things worse is rolled back. That is the whole of Loop 1, and it is what most research systems and a growing number of production systems now do.

Loop 2 is the step most systems skip. An engineer notices that a coding agent keeps misnaming a class in one repository. The fix is a one-line rule in that repository's instruction file. That same

rule could be promoted to a team-wide template, to the company's core instructions, or to a skill registry shared across companies. The edit did not change; its scope did. And the gate that should let it through gets stricter at each step, because the loss it is tested against is now an average over more deployments, and a fix that helps one repository can hurt another.

The scopes nest (Fig. 2). USER is one person's account or fork; PROJECT is one repository, workflow, or product workspace; STACK is a curated bundle of projects under one team; a **tenant** is one organization or customer; CORE is the scope every tenant shares. Each level's loss is the average failure rate over the deployments it contains. Promotion that reaches CORE crosses tenants, which is what cross-tenant promotion means here. Loop 2 promotes an edit upward only as far as the evidence reaches; we call that limit the edit's **evidence radius**. The check that belongs at each rung is **receiving-side evaluation**: a check run where the edit arrives, not where it was made. In optimizer terms, Loop 1 sets the step (how much one accepted edit changes the local scaffold) and Loop 2 sets the target (which average failure rate the scaffold is tuned to reduce). Governance is therefore not administrative overhead. It is the regularization that keeps a small fix from being aimed at the wrong target.

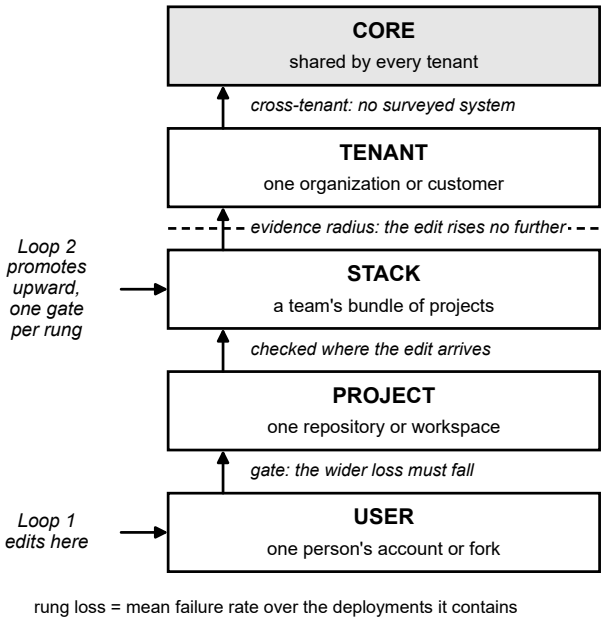


Figure 2: Loop 2, the promotion half. Scopes nest from USER to CORE, each rung's loss averages the deployments it contains, and an edit rises one rung at a time, only while a check at the receiving side shows the wider loss falling and only as far as its evidence radius. Promotion into CORE crosses tenants.

3.2 When the loop is an optimizer

The loop is an optimization process, rather than a description of good practice, when five conditions hold:

1. a measurable patch loss, such as a failure rate, a correction rate, or an evaluation score;
2. captured residuals, so that failures are logged rather than lost;

3. credit assignment, so that a failure can be traced, at least sometimes, to the scaffold coordinate that caused it;
4. a synthesized candidate edit to that coordinate; and
5. a gate that checks whether the edit lowers the loss before the edit is kept or promoted.

The conditions presuppose frozen frontier weights, edits that persist as inspectable artifacts rather than transient prompts, and promotion that widens an artifact's scope under a gate rather than averaging comparable updates as weight-space federated learning does.

Written out, the patch loss of deployment D is

$$L_D(\theta_s) = \mathbb{E}_{x\sim D}[\ell(\pi(\theta_M, \theta_s, x))]\,,$$

where θ_M are the frontier weights, fixed during deployment, θ_s is the scaffold, a structured object over the six substrate coordinates of Section 5, and ℓ is an operational loss. A failure yields no analytic gradient. It yields a candidate edit z_t, proposed by reflecting over the accumulated failure history. The gate accepts the edit when its noisy estimate of the change in loss clears a margin,

$$G_t(z_t) = 1 \iff \widehat{\Delta}_{z_t} L_D(\theta_s^t) + \lambda\,\Omega(z_t) \le -\tau_t,$$

where $\Omega \ge 0$ penalizes the complexity of the edit and the breadth of its promotion, and $\tau_t \ge 0$ is the required margin. Production systems implement this inequality without ever writing it down: an evaluation blocks a merge on regression, a reviewer rejects a vague rule, a governance system holds an artifact back until evidence accumulates. Under a bounded-bias assumption on the gate, the update satisfies a one-step descent inequality. When that bias is summable over the run, the loss converges and improving edits stop clearing the gate; the scaffold is then gate-stable, which is a statement about the gate, not about optimality (Appendix A). Where a condition fails, the picture is an analogy, and the survey marks which systems meet which conditions.

Each half of the loop has a gate, and we name three kinds by who decides. A gate is **Auto** when a program applies a check before the edit goes through: an evaluation score, a verifier, a reinforcement-learning threshold, or a language-model judge [27] with a quantified agreement rate. It is **Human** when a person decides. It is **None** when nothing decides, either because that half of the loop is absent or because the step runs unchecked, as when a system pushes every edit to its neighbors without a test. Who writes the candidate is recorded separately and is not a gate: the author may be an agent reflecting over failures, an engineer at a desk, or a contributor registering an artifact. We write a system's two gates as a pair, Loop 1 first: [Auto × Human] means an algorithmic local gate and a human promotion gate. These pairs index the design space of Section 7.

Condition 1 says the loss must be measurable. It does not say what measures it, and that turns out to be the question that sorts the corpus.

3.3 From patchwork to descent

Three stages of scaffold maintenance recur across the corpus, and the five conditions separate them. At the first stage, patchwork, edits are made by intuition after a failure, credit assignment is implicit, nothing is promoted on evidence, and the scaffold bloats. At the second stage, random artifact search, many variants are tried under an evaluation and the best is merged; the evaluation is the only signal, no failure is traced to a coordinate, and the dominant failure is overfitting to the in-distribution patch. At the third stage, artifact-layer descent, a residual is traced to a coordinate, one edit is made, promotion is bounded by evidence radius and reversible. Its dominant failure

becomes evaluation fragility, because the suite that gates the edits is itself a scaffold artifact that drifts and can be gamed (Section 10). A loop reaches the third stage once the five conditions hold. Production deployments that reach it on Loop 1 then promote in one of two ways: behind a human gate, as Spotify [28] and SkillForge do, or, in Letta's Context Repositories alone, by an ungated push that carries lineage but no evaluation.

4 Where the signal comes from: four classes of feedback

Four scaffolds all run the loop of Section 3. A coding agent changes one line in its repository rule file; the test suite runs and reports pass or fail within seconds. A scheduling assistant changes how it phrases a booking request; the appointment either appears on the calendar or it does not, and the system learns which an hour later. A marketing assistant changes the tone rule in its instructions; nobody can say whether the next draft is better, only whether a reviewer prefers it. A research assistant changes how it chooses which direction to pursue; whether the choice was good becomes visible months later, after many other edits have landed. All four produce a candidate edit. They differ in what the gate can know about it.

That difference is the whole of this section. The gate needs an estimate of one quantity: whether the edit lowered the patch loss. That estimate comes from a feedback signal. Signals differ in three ways: how soon they arrive, how exact they are, and whether there is a fact of the matter at all. Those three differences sort every surveyed system into four classes (Table 2), and the class fixes what the gate can be, which half of the loop can close automatically, and which failure dominates.

Table 2: The four feedback-signal classes. Reading across a row: what the gate receives, what it can therefore be, and what it cannot fix.

Class	The signal the gate receives	What the gate can be	What stays hard	Corpus examples
C1 exact	A check that returns pass or fail, or an exact score, at once and repeatably: tests, compilers, answer keys	An automatic check; one run per case, repeated only to sample the task mix	The check names the failed output, not the coordinate that caused it; a check that is also the target can be gamed	DGM, AlphaEvolve, SkillOpt
C2 observable outcome	A real result one episode later: a booked slot, a merged pull request, a closed ticket; noisy and confounded	An automatic threshold over batched episodes; a person at promotion	The outcome says *that* the episode failed, not *why*; the channel can be poisoned	Devin, Sierra, SAGE, SkillRL, PACE, Voyager
C3 judged	A rater, a rubric, or a model judge; no ground truth	A judge calibrated against people; a person at promotion	The rubric is itself a scaffold coordinate, so the gate judges edits to itself; the proxy gets optimized past the point where it tracks quality	SkillForge, CASCADE, skills-repository review
C4 delayed	An outcome weeks or months later, once per many edits	No per-edit gate; proposals are pruned and milestones stand in	Credit across many edits; drift; the loop cannot be told from random search	NanoResearch, CORAL, the AI Scientist

The classes are a partition, not three independent axes: a system is assigned by the property that

binds its gate first. Delay binds before anything else, so an outcome that arrives months later is C4 even when it is objective; the absence of ground truth binds next, so an immediate judgment is C3; and among signals that are prompt and objective, exactness separates C1 from C2.

C1, exact. A deterministic checker returns pass or fail, or an exact score, immediately and repeatably: unit tests, compilers, type checkers, exact-match answer keys, proof checkers, schema validators. DGM accepts a self-modification only after it passes in a sandbox on coding benchmarks [29]; AlphaEvolve scores candidate programs with a cascade of exact evaluators [30]; SkillOpt gates a skill edit on a held-out exact-match split of a math benchmark [31]. Because the check is exact for the artifact and the test it runs, and repeating it returns the same answer, one run settles the checked case. What the check does not settle is the patch loss: the suite is a sample of the patch's task mix, and the agent that runs under the edited scaffold is stochastic, so the gate's estimate of the loss change still carries sampling error. The gate can be a program and its margin can be small, but the margin absorbs sampling error, not error in the check. Two things stay hard. A failing test localizes the output, not the scaffold coordinate that caused it, so credit assignment remains open; and a checker that is also the training target can be gamed. The dominant failures are test overfitting and scaffold bloat, because cheap acceptances accumulate.

C2, observable outcome. The environment or the user emits an outcome after an episode: a booked appointment, a closed ticket, a merged pull request, a tool call that returns without error, a retrieved document that contains the answer. The signal is real but one episode late, noisy, and confounded, since a user can accept a wrong answer and reject a right one. Gates of this class include Devin's merged pull requests [32], SAGE's task success on AppWorld [33], SkillRL's success-rate threshold [34], and PACE's failure signals from a deployed small model [35]. The gate can still be a program (a reinforcement-learning threshold or a held-out task suite), but it needs a positive margin, and it absorbs the noise by batching episodes or by requiring agreement across samples (Appendix B). Credit assignment takes its most stubborn form here: the outcome says *that* the episode failed, not *why*. Most production deployments live in this class, and its signal arrives through a channel an attacker can control, which is why poisoned memory and tool-output injection dominate.

C3, judged. There is no fact of the matter. An essay, a marketing page, a design, a pitch deck, or a code review for style can be scored only by a judge: a human rater, a language-model judge, a rubric, or a comparison against references. SkillForge accepts a skill diff when its judge agrees with human engineers more than nine times in ten [36]; CASCADE consolidates memory under a human-and-agent gate [37]; Anthropic's skills repository promotes a skill when a reviewer approves the pull request. The gate can be a program only after the judge has been calibrated against people, because the judge's bias is now the estimator's bias. Residuals, meanwhile, are captured only when someone objects, so condition 2 depends on complaint channels rather than instrumentation. The dominant failure is Goodhart's law in the form measured for reward models: optimizing a proxy past a point lowers the true score [38]. Once the judge is the target, the scaffold learns the judge. Nor is a human judge stable: Shankar et al. find that evaluators' criteria drift as they grade more outputs [39].

C4, delayed and sparse. The outcome arrives after weeks or months, once per many edits, in a world that has moved on. Research programs, product strategy, and open-ended discovery are of this kind; the clearest instances are automated research pipelines such as the AI Scientist, which generates ideas, runs experiments, and writes the paper [40]. NanoResearch merges candidate skills by semantic compatibility rather than by a measured loss change [41], and CORAL promotes findings into shared memory on a heartbeat rather than on an evaluation [42]. AlphaEvolve's evolutionary

database can look like an instance, and is not one: it balances improving the best programs against diversity only after every candidate has been scored by its exact evaluators, so it is a search strategy inside the exact class, not a gate for a delayed signal [30]. With no per-edit estimate to gate on, the gate degenerates to proposal pruning plus proxy milestones, and the descent argument covers only the composition of pruning and whatever gate exists at the horizon (Appendix B). Temporal credit assignment across many edits is open. The dominant failure is drift: the scaffold ages faster than the signal that would correct it. An optimizer in this class also cannot be told from random search by its outcomes alone; Ong et al. [43] report the same indistinguishability for harness optimizers.

The judge is part of the scaffold. One feature of the judged class deserves its own paragraph, because it returns in the findings and in the open problems. The rubric that sets the loss is itself a scaffold coordinate. Suppose a pitch-review system keeps overrating decks that claim a large market without numbers. It traces the failure to its rubric and edits "evaluate market size" into "require a numeric estimate and a source". That edit changes the gate that will judge the next edit. We call this the *reflexive gate*, and nothing in the corpus calibrates one: three scaffold systems report judge calibration [36, 44, 45], and none calibrates a judge that is itself being edited. The weight side faces the same question. Reinforcement learning from verifiable rewards lies outside the scored corpus, but pushing it into writing and other open-ended tasks has forced researchers to build judged gates. One group holds the gate fixed while the policy learns, as a checklist, a pairwise judge, a reference-anchored judge, or a fixed rubric [46–49]; a second group lets the gate move with the policy, scoring by the policy's own probability of a reference or updating the rubric or judge during training [50–54]. Fixed versus co-evolved is exactly the reflexive-gate question, and of the co-evolving methods only RLPR, whose reward is a probability rather than a rubric, reports agreement with human labels (Appendix G).

What the class decides. Read down the table and three consequences follow, each a prediction about the corpus that Section 7 checks. First, which half of the loop can close automatically. An exact check or an observed outcome can gate a local edit by program. A judgment can do so only once the judge is calibrated, and a delayed reward cannot gate edits one at a time. Promotion is stricter. Only the exact class can gate it by program without a person on the receiving side, and only when the receiving side runs the same check, because only then does the sending side's evidence mean the same thing on arrival. In every other case the receiving-side evaluation of Section 3 is still owed. Second, which gate pairings a deployment can reach. [Auto × Auto] belongs to exact signals and simulated outcomes; [Auto × Human] belongs to observed outcomes and judgments, where the person on Loop 2 does the work the missing exact signal cannot; and the delayed class has no per-edit gate on either half. Third, which failure dominates, as the last column of the table records and Section 10 develops.

5 The survey: substrates, corpus, and scoring

To see which half of the loop the field has built, we scored 256 systems on where their scaffold adapts, whether a loop closes there, and under what gate.

5.1 Six substrates

A deployment's scaffold is composed of six substrates (Table 3). They are the coordinates of θ_s, and each can be edited independently of the others; the three kinds of memory on the memory substrate follow the cognitive-architecture taxonomy of Sumers et al. [55]. Other partitions are possible, for instance separating evaluation suites from governance. The argument of Section 3 is

invariant to the partition so long as the coordinates stay independently editable. This partition is the one that recurs in practice: every integrated system in the corpus has a recognizable instance of each substrate under some name.

Table 3: The six scaffold substrates, what each holds, and what a system must do on it to score 5.

Substrate	What persists	A score of 5 requires
S1 Instructions	Project, team, or organization rules attached to a workspace	Versioned rules promoted under an automatic gate
S2 Skills	Procedures, as code or text, loaded on demand	A skill bank that edits itself on failure, test-gated, versioned, and promoted
S3 Memory	Experience kept across sessions: episodic, semantic, procedural	Provenance, versions, and a correction path for every entry
S4 Tools	The vetted action surface and its schemas	Evaluation-gated tool versions; a curated bundle as the product
S5 Orchestration	The multi-role graph and its routing	Topology updated from an outer-loop signal
S6 Governance	Versioning, promotion, audit, rollback	Staged promotion with evaluation thresholds, an audit log, and rollback

Every substrate is scored on the same 0 to 5 scale: 0, ephemeral; 1, persistent and local; 2, persistent with tool attachment but no feedback; 3, multi-scope or feedback-driven but without automated promotion; 4, automated feedback updates the scaffold, so one half of the loop is closed; 5, the substrate's own top mark in Table 3. Systems that adapt several substrates at once receive an additional integrated row. The top mark means different things on two kinds of substrate. On those that hold adaptable content (S1, S2, S5) it requires both halves of the loop. On those that hold governance machinery (S3 provenance and correction, S4 evaluation-gated tool versions, S6 evaluation-gated promotion with audit and rollback) it requires the machinery a second loop would need, not the loop itself. That is why tool registries and the prompt-management platforms of Section 6 can score 5 with no inner loop at all and still sit outside the two-loop cells of the design space.

5.2 How the corpus was assembled

A system enters the corpus on either of these grounds: it wraps a foundation model with at least one persistent artifact updated from deployment signal, or it is described in public as self-improving, scaffolded, or agentic. All systems were found in arXiv, vendor documentation, and public repositories between January 2023 and the October 2026 cutoff. The corpus was assembled in two passes, the second adding 126 systems up to the cutoff, and each system was scored from its own method section or documentation, never from a secondary summary (Appendix C). Two kinds of system are excluded: those whose only update mechanism is weight-level fine-tuning or reinforcement learning, which are model-side and belong to the self-evolution literature this survey positions against, and those with no public evidence of deployment or release.

Two filters structure the corpus (Fig. 3). The patch-local discriminator admits a system to the core corpus when it scores at least 3 on at least one substrate, that is, when some artifact persists across scope or is updated from feedback; systems that carry the labels but fail it form a contrast corpus, kept so that the boundary can be audited rather than asserted. Of the 256 systems, 204 form the core corpus and 52 the contrast corpus.

A stricter composite two-loop audit then asks two questions of each core system: does Loop 1 modify

a persisted artifact from session signal, and does Loop 2 have an explicit mechanism that promotes the artifact beyond the deployment that produced it, with versioning or lineage? The audit was applied to the 204 core systems and to ten further systems that the text treats at mechanism level. Those ten were classified from their papers after the released table closed; they are not rows of it, and the corpus counts above exclude them (Appendix C). The audit admits 14 research systems and 3 industry exemplars. Of the research systems, 13 are rows of the released table and the fourteenth, SkillOpt, comes from the mechanism-level group; the industry exemplars are Sierra Agent OS 2.0, Cognition Devin, and Letta's Context Repositories. Only these are placed in the design space. Passing the audit is a weaker thing than having the governed second loop of the thesis: every passing system promotes across users, teams, or repositories inside one operator, none reaches another organization, and one of them, Letta, promotes with no gate at all (Section 9).

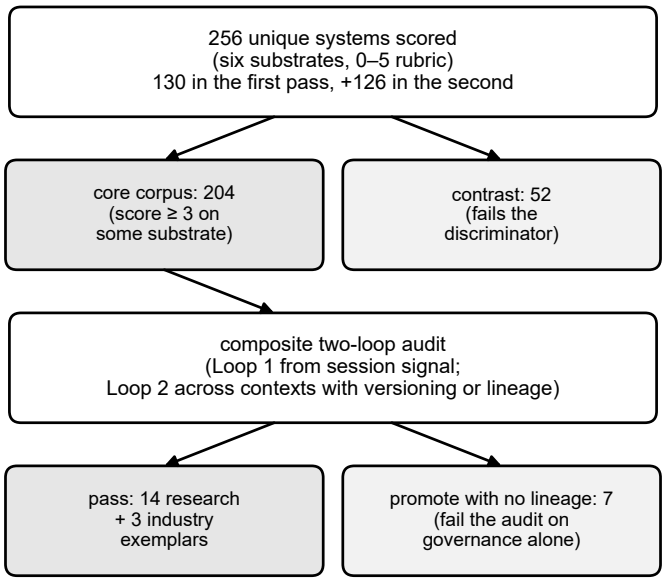


Figure 3: Corpus flow. Every system is scored on the substrates it adapts, with an integrated row for systems that adapt several, giving 269 substrate-level rows for 256 systems; the discriminator splits core from contrast; the composite two-loop audit selects the systems placed in the design space of Section 7.

5.3 Evidence and scoring discipline

Every classification rests on public evidence of uneven strength and on one rater. Each system carries one or more evidence tiers (T1 to T5 in the released table: peer-reviewed publication, arXiv preprint, production claim with metrics, open-source artifact, industry blog or interview). Peer-reviewed papers and open-source artifacts weigh most in cluster-level statements; a production claim counts when it names a benchmark or integration target; a blog post or interview counts only when a higher tier corroborates it; and production claims are never pooled with peer-reviewed evidence in a cluster-level statistic. Industry systems evaluated from vendor documentation rather than a paper are flagged wherever they appear. One author scored every system, used half-scores only when one criterion was clearly met and another partially, and rounded ambiguous cases down; the scores have not been independently re-rated, and that check is future work. The complete score table, with a one-line justification per row, is released so that every scored classification can be re-derived (Appendix I). The counts do move under natural relaxations of the composite audit (for

instance, counting rollback alone as Loop-2 governance, asking for a promotion mechanism with no governance at all, or admitting a Loop 1 that persists only optimizer state); the qualitative findings survive every one of them (Appendix C).

6 What the field has built

The corpus answers the question of Section 1 in one sentence: the local half of the loop is routine, and no surveyed system documents the governed second half. The evidence comes in three cuts: per substrate, where the top mark is reached and where it is not; per missing piece, which parts of the second half are absent and in which signal class; and per cluster, which kinds of system have which half.

6.1 The local loop is routine, and no substrate is finished

Nine in ten of the most recently added core systems close the local loop under an automatic gate, and none of them reaches the top mark on any substrate (Table 4).

Table 4: Substrate maturity across the 256-system corpus. Column n counts scored rows on each substrate, Mean is the average score, and Score-5 systems reach the substrate's top mark; the integrated row counts systems scored across several substrates as one composite row. Of the most recently added 126 systems, 31 are skill systems and 102 are core; 92 of the core systems close Loop 1 under an automatic gate, and none of the 126 scores 5 on any substrate.

Substrate	n	Mean	Score-5 systems	Survey finding
S1 Instructions	15	3.07	(none)	Rules persist across IDEs and agents; none closes both loops; ceiling 4
S2 Skills	50	3.34	SAGE, COSPLAY	Most active research substrate; Loop 1 routinely automated; production repositories human-merged
S3 Memory	50	3.06	Cuadros et al.	Lineage shipped in products; enforced correction still absent
S4 Tools	20	3.65	MCP, Harvey, Hippocratic AI	Most mature production substrate; registries govern admission, not evaluation
S5 Orchestration	37	3.08	(none)	Topologies evolve under evaluation; one structurally gated instance; no two-loop system
S6 Governance	45	3.02	Braintrust, Vellum, LangSmith Hub, AGENTS.md/AAIF	Mature prompts-as-code; registries add lineage and rollback; no inner loop
Integrated	52	3.81	NanoResearch, AutoAgent, SkillRL	High-water mark; research systems without enterprise governance

On the three substrates that hold adaptable content, only two systems reach the top mark, and both reach it inside one research environment. Instruction files persist across IDEs and agents (Claude Code, Cursor, AGENTS.md [56], Windsurf, Continue.dev) and appear in a fifth to a quarter of coding-agent projects [57, 58], yet in a controlled study they raise inference cost by about a fifth and buy nothing in task success [59].

Skills are the most active research substrate: SAGE and COSPLAY close both halves of the loop [33, 60], Loop 1 is routinely automated, and SkillChain's judge-approved edits pass a human

reflection gate before they enter its skill bank, a two-stage local gate rather than a promotion loop [44]. Eleven skill systems treat retirement as a gated operation [61–64], and the public skill repositories are merged by hand [65].

On orchestration, automated workflow search came first, in GPTSwarm's optimizable agent graphs [66] and AFlow's tree search over code-represented workflows [67]. Topologies evolve under evaluation [68–70]; TOCOMAS alone keeps a topology change only when it satisfies a structural invariant [71], and Evo2Team promotes routing skills across teams [72].

Memory carries lineage without enforced correction. Production memory records its own history, in Letta's git-backed Context Repositories [73], Mem0's per-memory history [74], and Zep's bitemporal facts [75]. Research memory splits into write-gating and provenance on one side [76–78] and verifier-gated failure loops on the other [79–81]. Among the memory systems that gate writes or correct from failures, only Cuadros et al. does both, which is why it alone reaches the top mark.

Tools and governance hold the machinery of a second loop without the first, which is how they reach the top mark with no inner loop (Section 5). MCP registries split into catalogues (namespaces and version records, but no scanning) and control planes (scanning, access control, and quarantine) [82–84]; none has evaluation-gated promotion. The governance registries bring lineage and rollback [85–87], and none of them has an inner loop.

No system matures all six substrates at once. Research systems learn fast and lack enterprise governance; production systems govern well and learn slowly; open standards solve artifact portability but not adaptation. This pattern is not a quirk of the corpus: it tracks which constraints each kind of system was built to satisfy.

6.2 Where the second loop should be, five pieces are missing

Walk the second half of the loop from the failure that starts it to the promotion that ends it, and five pieces are missing or rare across the core corpus. Each is missing in a particular feedback-signal class, which is the first sign that the gap is not one gap but several.

- **Failures are not captured** (C2). Deployed memory systems record preferences and successful facts, not residuals, so the loop has nothing to cluster. The one field report of stores written from dead-ends to prevent recurrence describes a single, self-reported deployment [88].
- **Which sessions inform the next edit is left to chance** (C2). Nearly every system takes failures in the order they arrive. Task-CoEvolve's variance-weighted task sampling and AutoSaddler's mini-batch failure signals are the exceptions [89, 90].
- **Memory is versioned, but correction is not enforced** (C2, C3). Git-backed and bitemporal memory ship in products, and five production backends still return facts that were revoked [91].
- **Promotion is rarely checked where the edit arrives** (C1, C2). ShareMem gives evaluation tasks to a different user, XRepoSkill verifies in at least two repositories, Evo2Team confirms at the target team, and VACE evaluates the candidate with the receiving model held fixed [92]; even Evo2Team's confirmation lets nearly a fifth of its transfers fail on new graphs. EvolveNet has no such check, and the federated optimization framework that proposes a receiving-client gate has not implemented it [93].
- **No promotion crosses an organization** (all classes). ShareMem, EvolveNet, XRepoSkill, and RASO promote across users, clients, or repositories inside one operator [94–97]. No system promotes across customers or organizations with evidence.

The last piece is the thesis of this article. The four before it are the reasons a system that tried to add it would have little to promote and no way to check what it promoted.

6.3 Four kinds of system, each with half the loop

Sort the core corpus by which half of the loop it automates and four clusters appear (Table 5). Each has one half of the loop; none has both halves governed.

Table 5: Four clusters of the core corpus.

Cluster	Representative systems	Loop 1 / Loop 2	What the cluster shows
Research full-auto	NanoResearch, SkillOpt, SkillRL, AlphaEvolve, DGM, ADAS, Voyager, EvolveNet, Evo2Team, XRepoSkill	Auto / Auto	Both loops can run under algorithmic gates and benchmark gains compound; enterprise governance is uniformly absent
Production hybrid	SkillForge, CASCADE, Sierra Agent OS 2.0, Cognition Devin, Spotify's recommendation loop [28]	Auto / Human	An algorithmic Loop-1 gate with a reviewer on Loop 2 is the production-viable cell; Spotify adds online A/B gating
Governance-first	Braintrust, Vellum, LangSmith Hub, EDDOps, Agentic Configuration Management, Governed Capability Evolution	None / Auto or Human	Evaluation-gated promotion, lineage, and rollback exist; candidate generation is manual and no inner loop fires from session signal
Open standards and registries	AGENTS.md, MCP, Anthropic Agent Skills, Cursor Rules, the MCP control planes	None / Human, or None / None	Cross-vendor convergence on artifact format and tool attachment; registries govern admission and access but do not evaluate

Research full-auto systems show that the architecture works when both gates are algorithmic [29, 30, 41, 98–101]. AlphaEvolve is the clearest case: it pairs an evaluator with an evolutionary archive, as FunSearch did before it, and reports production gains at Google. EvolveNet, Evo2Team, and XRepoSkill extend the loop from a single deployment to a federation of deployments. What every member lacks is governance: no access control and no audited promotion.

Production hybrid systems show that the production-viable pairing is an algorithmic gate on Loop 1 with a reviewer on Loop 2. SkillForge illustrates it most plainly: a judge that agrees with people more than nine times in ten passes each skill diff, and support engineers decide whether it ships [36]. Sierra routes agent-generated insights through human-authored expert answers [102], Devin's merged pull requests are its promotion step [32], and Spotify launches automatically optimized prompts and plans only behind an online A/B test. In every member, a person on Loop 2 stands in for the exact signal the cluster lacks.

Governance-first systems automate Loop 2 alone. The prompts-as-code platforms [103–105] and the governance registries [85–87] keep immutable versions, evaluation thresholds, and rollback, but they have no Loop 1: engineers author the candidates, and session feedback never triggers one. This is MLOps for scaffolds, distinct from MLOps for weights.

Open standards and registries converge across vendors on format and attachment. Claude Agent Skills, Cursor Rules, GitHub Copilot instructions, Windsurf, Continue.dev, Zed, and Replit all keep a git-tracked markdown artifact attached to a workspace with an MCP-style tool layer [58, 106], a convergence that Section 8 reads as evidence that the architecture is real. Adaptation is what the

cluster lacks: its registries govern admission and access but do not evaluate, and in the repositories that run autonomous loops, the runtime state of those loops sits outside version control [107].

7 The two-loop design space

Index each system by the gate on each half of the loop and the field collapses into two gated cells, one half-cell, and one ungated passer (Fig. 4).

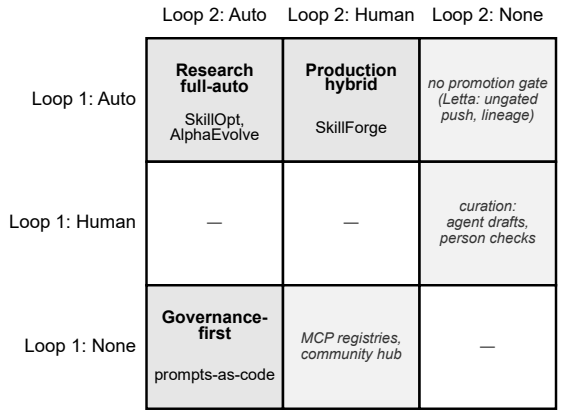


Figure 4: The two-loop design space. Rows give the Loop 1 gate and columns the Loop 2 gate, so a cell [Loop 1 × Loop 2] such as [Auto × Human] pairs an algorithmic local gate with a human promotion gate; a gate is None when the step is absent or runs unchecked. Two cells hold every audit passer whose promotion has a gate; a third holds the governance-first platforms, which have the gated second loop but no inner loop; Letta's Context Repositories pass the audit on lineage and promote with no gate at all. The human-gated row holds only curation workflows, and no surveyed system puts a person on the local gate and then promotes. No populated system adds governed cross-tenant promotion.

[Auto × Auto] holds the research full-auto systems, machine-evaluated on both loops. [Auto × Human], the production-viable cell, holds SkillForge, CASCADE, Sierra Agent OS 2.0, and Cognition Devin: Loop 1 fires at machine speed under an algorithmic criterion, and Loop 2 waits for a person. The governance-first platforms sit in [None × Auto]: they gate promotion on evaluation but generate no candidates from failures, so they pass only the Loop-2 half of the audit and enter the matrix as contrast entries (Appendix E).

One audit passer sits in neither gated cell. Letta's Context Repositories push shared memory across agents automatically, with git lineage and no evaluation of any kind, so the audit admits them on lineage and the matrix places them in [Auto × None]: a local loop with a gate, and a promotion with none. The cell makes the point the pair notation exists for. A push that carries every edit is not a gate, and the column records that nothing decides.

Seven systems promote across users, clients, teams, or repositories with no versioning or lineage on what they promote, and fail the audit on that alone. If the audit asked for a promotion mechanism and nothing more, six of them would enter, all in [Auto × Auto] except Spotify's production loop, which would land in [Auto × Human]; none would open a new cell. The seventh, Future AGI's

prompt lifecycle, rests on a vendor blog alone, which the tier rule does not count, and is not admitted under any rule (Appendix C).

The human-gated row holds only curation. In [Human × None] an agent drafts and a person decides what is kept: VerificAgent's agent-drafted memories pass a human fact-check and are frozen [108], and SkillChain's judge-approved skill edits re-enter a human reflection gate before they are deployed. Nothing in the row promotes. [Human × Auto] and [Human × Human] are empty: no surveyed system puts a person on the local gate of a failure-driven loop and then promotes the result, by program or by person. The systems that look like human-gated promotion on a first reading are registries, in which a person registers an artifact and a person, or an admission scan, approves its use, and no failure-driven loop runs at all. They sit in [None × Human]: the MCP registries and a community hub [82–84, 109], and, in practice, the five public skill repositories of one field study, where every substantive edit is human-authored and human-merged and one deprecation appears in 254 edits [65].

The occupancy bears out the three consequences of Section 4. Every [Auto × Auto] system whose gate estimates the loss change per edit runs on an exact signal or a simulated outcome, and those that work on delayed signals prune proposals instead. [Auto × Human] holds only systems whose signal is an observed outcome or a judgment; no exact-signal system sits there. Nor does any system in the matrix calibrate a judge that is itself under edit.

Two further patterns cut across the cells. Earlier systems accepted an edit on one evaluation and absorbed estimator noise through the per-edit margin alone; the newer ones batch evidence before acceptance. A paired per-item test commits an edit only when its wins are reliable [110]; EvolveNet accepts a merge only when it fixes at least as many held-out items as it breaks [95]; in VaG, three heterogeneous critics must all pass [111]; and AutoSaddler feeds mini-batch failure signals into validation-based selection [90]. Requiring agreement among independent proposals and averaging over them are different operations that reduce error at different rates (Appendix B). Multi-layer hierarchies remain rare: AlphaEvolve is the only system that combines a cascade evaluator with a diversity archive, and the other systems that keep archives use them for lineage, not for staged evaluation.

8 Five research families converge on the loop

The loop is not a prescription the field has yet to adopt. Five research families, starting from different substrates and under different names, have arrived independently at parts of it, which we read as evidence that the loop is a shared attractor rather than one group's local optimum. The vendors' convergence on one artifact format and tool attachment (Section 6) is independent evidence of the same kind. Convergence does not show that the design is optimal, only that the same constraints keep producing the same shape.

The clearest single instance is SkillOpt [31]. It treats one markdown skill document as the trainable external state of a frozen agent and edits it with explicit optimizer controls: rollout and reflection batches, add/delete/replace edits, a bounded textual learning rate with cosine decay, a held-out selection gate that accepts only a strictly better score, a rejected-edit buffer kept as negative feedback, and an epoch-wise slow update. Across six benchmarks, seven target models, and three harnesses it is best or tied on every cell its authors evaluate. It also reports transfer across models, harnesses, and benchmarks, which reads as an empirical measurement of evidence radius rather than as a governance mechanism. Its own follow-up, SkillOpt-Lite [112], removes the buffer, the

slow update, and the schedule, keeps the held-out gate and a consensus filter across rollouts, and matches or beats it. So the gate and a consensus form of credit assignment are load-bearing; the rest, it turns out, are stabilizers.

The five families instantiate complementary parts of the loop (Table 6), and each was pushed into its part by the constraint it started from. Reflective evolution started from trajectories and proposes edits by reflecting over them, with a two-timescale argument for keeping those edits outside the weights [4, 113]. Reflective context learning, the most direct independent formalization, started from the observation that credit assignment, overfitting, forgetting, and high-variance signals persist whether the learned object lies in parameter space or context space; its generalization penalty is the complexity term Ω of the gate [114]. Harness optimization started from orchestration rather than skills and names credit assignment as its bottleneck; it is also the largest family in the corpus: roughly two hundred papers cite its founding entries, and a dozen are scored here [70, 115–120]. The lifecycle family started from artifact pools that bloat and found hygiene to be a precondition for descent [121]: the complexity penalty becomes a constraint frontier once platform limits bind [122], each outcome is attributed to skill, exploration, or environment before an update is admitted [123], and eleven systems make retirement a gated operation [61–64]. The transfer family started from how far an edit travels and instruments Loop 2, from measuring evidence radius one task at a time to federating skill patches across clients or deployments [35, 42, 95, 124–126].

Table 6: Five research families converging on artifact-layer descent, with the part of the loop each instantiates.

Family	Representative systems	What it instantiates	What it leaves open
Reflective evolution	GEPA; Tiwari et al.	proposal by reflection over trajectory evidence; two-timescale motivation	persistent artifact; bounded step; held-out gate
Reflective context	RCL	the learned object in context space; generalization penalty as Ω	cross-deployment Loop 2
Harness optimization	Meta-Harness; AHE; MOSS; CANTANTE; HARNESSEVO; HarnessForge; Growing Harness	credit assignment over orchestration; source-level edits; slot attribution	governed Loop 2; access control; rollback
Lifecycle and multi-objective	Ratchet; MOCHA; SkillsVote; Library Drift; SkillSmith; ASSAY	hygiene as precondition; constraint frontier as the shape of Ω; gated retirement	gradient-analogue closure; retirement of a promoted artifact's descendants
Transfer and federation	ICT; Combee; CORAL; PACE; FederatedSkill; EvolveNet; Evo2Team; XRepoSkill	empirical evidence radius; cross-client aggregation; receiving-side confirmation	cross-organization gate, lineage, rollback

Read against the loop, the families support four conclusions.

1. The scaffold-versus-weights decomposition is independently confirmed. GEPA, Meta-Harness, MOSS, RCL, PACE, and the harness-optimization papers all begin from the observation that weights and the external artifact layer are distinct adaptation surfaces with different cost, reversibility, and scope, and a system that updates both keeps them as separate levers [127].
2. Credit assignment is the central engineering bottleneck. AHE, CANTANTE, Meta-Harness, HARNESSEVO, SkillsVote, and the harness-optimizer evaluation of Ong et al. [43] all name

the same step as what separates descent from trial-and-error: attributing a system-level failure to one scaffold coordinate. That step remains unsolved in the general multi-substrate case.

3. Lifecycle hygiene is a precondition for descent, not an alternative to it. A dirty artifact pool starves the gate of signal (Table 8), gated retirement is the field's answer, and the one negative result cuts deeper: once a contaminated skill has descendants, rolling it back recovers little, so gating must happen before commit [111].
4. No family documents the governed second loop. FederatedSkill and EvolveNet each narrow the gap on different properties, and neither meets the four that Section 9 sets out.

9 The missing architecture: governed cross-tenant scaffold optimization

What is missing is specific. Four properties define an architecture that no surveyed system instantiates:

- (a) **auto-gated Loop 1**: the commit criterion is an algorithmic threshold;
- (b) **cross-organization promotion**: Loop 2 routes artifacts between organizations (we use *cross-tenant* and *cross-organizational* interchangeably);
- (c) **versioned lineage with access control**: promoted artifacts carry versions, audit trails, and role-based access policies;
- (d) **rollback**: a bad promotion can be reverted without redeploy.

The closest approaches split these properties among them, and no row of Table 7 is full on promotion across organizations.

Table 7: The closest approaches to governed cross-organization scaffold optimization, scored against the four properties.

System	(a) auto-gated Loop 1	(b) cross-organization promotion	(c) versioned lineage with access control	(d) rollback
FederatedSkill [126]	full	partial: privacy-preserving aggregation of skill patches across heterogeneous clients that sit under one operator	partial: auditable per-modification decision log	none
EvolveNet [95]	full	partial: held-out merge gate at the aggregator, but clients are benchmark partitions, not organizations	none: no lineage, no access control	full: rollback to the previous shared harness after a failed round
SkillForge [36]	full	partial: cross-team within one enterprise	partial	not documented
Governance-first registries [85–87]	none: no inner loop	none	partial: lineage without access control	full, with evidence
MCP control planes [82]	none: no inner loop	none	partial: access control without evaluation-gated promotion	not reported

Row by row, the partition is tidy: the systems with a local loop carry the automatic gate and part of the promotion, the registries carry lineage and rollback without an inner loop, and the control planes carry access control alone. Rollback appears inside one federated round and inside single deployments, as edit-level reversal [62, 128] and as an admission test that rejects edits which cannot be undone [129]. No system rolls back a promotion that has crossed an organization boundary, because none makes such a promotion (Appendix E).

In machine-learning terms, what is missing is a cross-tenant aggregator with an automated evaluation gate at the scaffold layer. It would carry five parts: per-tenant artifact pools, each tagged with provenance; an aggregation step across tenants that respects tenant isolation; a held-out evaluation at the receiving tenant that blocks a promotion whenever it does not lower that tenant's patch loss; versioned lineage with access control; and a rollback envelope on every promoted artifact. The federated-learning analogy [130–132] is useful only in the limit where artifacts are mergeable update representations; federated prompt tuning, in which participants learn a shared prompt rather than shared weights [133], is the weight-side precedent closest to that limit. In surveyed practice Loop 2 is selection and merge: the nearest systems aggregate program edits with a language model, not by averaging.

Whether the architecture *should* be built depends on the feedback-signal class, not on the application domain (Section 4). Where the signal is exact, automatically gated cross-tenant promotion of a test-validated skill would be low-risk, and receiving-side evaluation is cheap to add. Where it is an observable outcome, the sending side's outcome does not transfer, so promotion should be automated only behind receiving-side evaluation, the check that would catch a scaffold over-fitted to one deployment before it propagates to another. Even that check is necessary rather than sufficient: some of Evo2Team's transfers still fail on new graphs after target-side confirmation [72]. Where it is judged, a human gate on Loop 2 is a design feature, and governance means lineage and rollback around a human decision. With a delayed signal, property (a) cannot be met; what remains is lineage and rollback without an automatic gate. Safety-critical domains, the usual case for human gates, are judged or delayed domains with high stakes. Naming the four properties turns a vague gap into an audit criterion that any future system can be measured against; naming the signal class says which of the four a given deployment should expect to meet.

10 The failure surface

Until the governed second loop exists, scaffold maintenance carries failure modes that weight-only systems do not have; they are the scaffold-layer entries in the technical-debt ledger that Sculley et al. opened for machine-learning systems a decade ago [134]. Each is a missing piece of the optimizer, and each dominates in a particular feedback-signal class (Table 8). The entries are documented attacks and degradations, not a new threat model.

Bloat is not only slow; it is a signal-quality problem. Ratchet reports what that looks like: LLM-authored skills gain nothing over a no-skill baseline, human-curated ones show a double-digit gain, and the authors attribute the gap to lifecycle management [121]. In the vocabulary of Section 3, that means residual capture and mode discovery cannot work on a messy artifact pool.

Evaluation fragility compounds because the evaluation suite is itself a patch-local artifact rather than a fixed instrument. It is optimized against: once a metric becomes the target, it is gamed. It is read non-independently, because the gate scores path-dependent agent trajectories whose readings do not compose like ordinary test-set statistics. It goes stale, because the suite adapts on the same

Table 8: The failure surface of patchwork scaffold maintenance, the measured evidence, the architectural response, and the feedback-signal class in which each failure dominates.

Failure	What has been measured	Architectural response	Dominant class
Scaffold bloat	Instruction-following accuracy falls and latency grows sharply as instruction count scales [135]; LLM-authored skills add nothing over no skill until lifecycle hygiene is applied [121]	Complexity penalty Ω in the gate; pruning and expiry as first-class gated operations	C1, C3
Poisoned memory	High attack success against standard retrieval-augmented generation [136]; backdoors planted in long-term memory or knowledge bases [137]; memory control-flow attacks across major agent frameworks [138]	Provenance on every entry; write gating; a rollback envelope [139, 140]	C2
Injection via tools and MCP	A tool's output or a connected server can smuggle instructions into context [141, 142]; one approved MCP configuration could later be modified without re-approval (CVE-2025-54136) [143]	Narrow tool permissions; a privilege hierarchy over instructions [144]; tool bindings pass the same promotion gate as any other coordinate	C2
Evaluation fragility	Optimizing a proxy reward past a point lowers the true score [38]; access asymmetries alone produce large relative gains that reflect exposure to the test [145]; harness optimizers cannot be told from trial-and-error by downstream score alone [43]	Held-out evaluations scoped within evidence radius, weighted toward recurring modes, re-read as the target drifts, and checked against an update's attribution	C3 (primary), C1
Coordination failures	Fourteen failure modes in multi-agent systems, dominated by role confusion and context loss [146]	Orchestration topology treated as a scaffold coordinate under the same gate and rollback discipline	C2
Patch overfitting	A scaffold tuned to one deployment over-fits it by construction (Appendix D)	Evidence-radius discipline at promotion; receiving-side held-out evaluation	C2
Plasticity loss	Context rot (long-context degradation [147]), instruction conflict, memory pollution, tool-version drift	Expiry and rollback as explicit loop operations rather than retrospective hygiene	C4

cadence as the scaffold it gates. And because failures concentrate in a finite catalogue of recurring modes, a suite that samples tasks at random under-weights exactly the families that drive the failure rate. The response in the table follows from those four facts: held-out and scoped, mode-weighted, re-read, and checked against attribution rather than only against score.

The failure surface is real, quantified in the cited measurements, and largely unmitigated in production. Every row names a piece of the optimizer that current systems do not yet have.

11 Open problems and conclusion

Five open problems follow from the survey, one per feedback-signal class and one for the architecture.

1. **Exact signals: credit assignment from a failing check to a scaffold coordinate.** A test tells the loop that an output failed, not which instruction, skill, or tool binding caused it. Slot attribution [119], intervention-based attribution [69], and contrastive multi-agent attribution [118] each solve one substrate's case; a general method is the bottleneck the harness-optimization family names.
2. **Observable outcomes: receiving-side evaluation before promotion.** The outcome that justified an edit on the sending side does not transfer, and even where a promoted edit is checked on arrival, some transfers still fail (Section 6). The gate a federated scaffold optimizer needs, a held-out test at the receiving tenant, has been proposed but not built [93].
3. **Judged signals: calibrating a gate that is itself being edited.** When the rubric is a scaffold coordinate, each accepted edit to it moves the instrument that judges the next edit. No method that co-evolves a rubric or judge reports agreement with human judgment, and no scaffold system calibrates a judge under edit (Section 4); the nearest measurement is on the human side, where evaluators' criteria drift as they grade [39].
4. **Delayed signals: telling descent from search.** Where the outcome arrives after many edits, an optimizer's gains cannot be distinguished from trial-and-error by outcome alone [43]; evaluating harness optimizers against their attribution claims rather than their scores is the open method.
5. **The governed cross-organization loop.** The architecture of Section 9, properties (a) through (d), is documented by no surveyed system; the receiving-side gate of problem 2 is the piece that would make cross-organization promotion safe in the two classes where it should be automatic.

The field has built an external adaptation layer for language models but not, as far as public evidence shows, the optimizer that should govern it. Both earlier papers of this series supply the hypothesis for why the layer belongs where it is: errors concentrate in a few key decisions; inside a bounded deployment, those decisions fail in a small and recurring catalogue of modes; and such a catalogue can be fixed one mode at a time in a persistent, inspectable, reversible artifact outside the weights. The corpus is consistent with that hypothesis. If it holds, the gated and scoped artifact layer is not a workaround but the architecture the error structure calls for.

The survey shows how far practice has come toward that optimizer, and the progress is lopsided. The local loop closes automatically in most research systems and in a growing number of production ones; evidence is batched before acceptance; retirement and expiry are first-class operations; production memory carries lineage. What has not moved is the second loop. No surveyed system promotes a validated edit across an organization boundary under an automatic gate with lineage, access control, and rollback. Read by signal class, this is not one gap but four, each with its own reachable answer.

The frontier model generalizes. The local scaffold specializes. Reliability comes from governing that specialization. Because the instruments that judge each edit are themselves part of the moving scaffold, the same evidence-radius and lifecycle discipline the optimizer needs must extend to the gates that judge it.

Acknowledgments

The authors used an AI assistant (Claude, Anthropic) to help search the literature, check citations, and edit prose; all claims, scores, and classifications were reviewed and approved by the authors.

References

[1] M. Zaharia, O. Khattab, L. Chen, J. Q. Davis, H. Miller, C. Potts, J. Zou, M. Carbin, J. Frankle, N. Rao, and A. Ghodsi, "The shift from models to compound AI systems," Berkeley Artificial Intelligence Research (BAIR) Blog, https://bair.berkeley.edu/blog/2024/02/18/compound-ai-systems/, 2024.

[2] R. Kirk, I. Mediratta, C. Nalmpantis *et al.*, "Understanding the effects of RLHF on LLM generalisation and diversity," in *ICLR 2024*, 2024.

[3] V. Padmakumar and H. He, "Does writing with language models reduce content diversity?" *arXiv preprint*, 2023.

[4] R. Tiwari, K. Sareen, L. A. Agrawal, J. E. Gonzalez, M. Zaharia, K. Keutzer, I. S. Dhillon, R. Agarwal, and D. Khatri, "Learning, fast and slow: Towards LLMs that adapt continually," *arXiv preprint*, 2026.

[5] S. Dohare *et al.*, "Loss of plasticity in deep continual learning," *Nature*, vol. 632, no. 8026, pp. 768–774, 2024.

[6] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, "LoRA: Low-rank adaptation of large language models," *arXiv preprint arXiv:2106.09685*, 2021.

[7] M. L. Arbuzov, S. Bei, Z. Dong, D. Kalaev, and A. A. Shvets, "Beyond exponential decay: Rethinking error accumulation in large language models," *arXiv preprint*, 2025.

[8] M. L. Arbuzov, L. Mosbacker, S. Bei, Z. Dong, D. Kalaev, and A. Shvets, "The architecture of errors: From universal impossibility to patch-local LLM reliability," *arXiv preprint arXiv:2605.30628*, 2026.

[9] T. Ding, A. Nannapaneni, B. Liu, and L. Zhang, "Always-on agents: A survey of persistent memory, state, and governance in LLM agents," *arXiv preprint arXiv:2606.30306*, 2026.

[10] H.-a. Gao, J. Geng, W. Hua, M. Hu, X. Juan *et al.*, "A survey of self-evolving agents: What, when, how, and where to evolve on the path to artificial super intelligence," *Transactions on Machine Learning Research*, 2026, arXiv:2507.21046v4, 16 Jan 2026.

[11] J. Fang, Y. Peng, X. Zhang, Y. Wang, X. Yi, G. Zhang, Y. Xu, B. Wu, S. Liu, Z. Li, Z. Ren, N. Aletras, X. Wang, H. Zhou, and Z. Meng, "A comprehensive survey of self-evolving AI agents: A new paradigm bridging foundation models and lifelong agentic systems," *arXiv preprint arXiv:2508.07407*, Aug. 2025.

[12] H. Zhou, H. Hu, T. Luo, Y. Shang, C. Fang, Z. Chen, L. Xiao, and Q. Zhang, "Self-evolving coding agents," *arXiv preprint arXiv:2608.03392*, Aug. 2026.

[13] Z. Lin, X. Hao, R. Fu, S. Cui, K. Chen, C. Li, Z. Li, and F. Xiong, "A survey on long-term memory security in LLM agents: Attacks, defenses, and governance across the memory lifecycle," in *Proceedings of EMNLP 2026*, 2026, arXiv:2604.16548v3, 22 Sep 2026.

[14] J. Guo, Z. Hao, C. Wang, C. Fan, T. Luo, H. Li, Y. Gao, H. Mei, J. Peng, R. Xu, M. Dong, H. Wu, M. Zheng, K. Han, S. Wang, C. Xu, and Y. Wang, "From question answering to task completion: A survey on agent system and harness design," *arXiv preprint arXiv:2606.20683*, Jun. 2026.

[15] B. Xia, L. Zhu, Z. Xing, Q. Lu, D. Sejdinovic, and X. Xu, "Harnessing agent skills: Architectural patterns and a reference architecture for skill-mediated LLM agents," *arXiv preprint arXiv:2606.20631*, May 2026.

[16] K. Ding, Y. Zhou, C. Jin, F. Tong, M. Zhou, and D. N. Metaxas, "Agent skill evaluation and evolution: Frameworks and benchmarks," *arXiv preprint arXiv:2606.11435*, Jun. 2026.

[17] L. Mei, J. Yao, Y. Ge, Y. Wang, B. Bi, Y. Cai, J. Liu, M. Li, Z.-Z. Li, D. Zhang, C. Zhou, J. Mao, T. Xia, J. Guo, and S. Liu, "A survey of context engineering for large language models," *arXiv preprint arXiv:2507.13334*, Jul. 2025.

[18] C. Yang, X. Wang, Y. Lu, H. Liu, Q. V. Le, D. Zhou, and X. Chen, "Large language models as optimizers," in *International Conference on Learning Representations (ICLR)*, 2024, arXiv:2309.03409.

[19] Q. Guo, R. Wang, J. Guo, B. Li, K. Song, X. Tan, G. Liu, J. Bian, and Y. Yang, "EvoPrompt: Connecting LLMs with evolutionary algorithms yields powerful prompt optimizers," in *International Conference on Learning Representations (ICLR)*, 2024, arXiv:2309.08532.

[20] C. Fernando, D. Banarse, H. Michalewski, S. Osindero, and T. Rocktäschel, "Promptbreeder: Self-referential self-improvement via prompt evolution," *arXiv preprint arXiv:2309.16797*, 2023.

[21] O. Khattab, A. Singhvi, P. Maheshwari, Z. Zhang, K. Santhanam, S. Vardhamanan, S. Haq, A. Sharma, T. T. Joshi, H. Moazam, H. Miller, M. Zaharia, and C. Potts, "DSPy: Compiling declarative language model calls into self-improving pipelines," *arXiv preprint arXiv:2310.03714*, 2023.

[22] M. Yuksekgonul, F. Bianchi, J. Boen, S. Liu, Z. Huang, C. Guestrin, and J. Zou, "TextGrad: Automatic "differentiation" via text," *arXiv preprint*, 2024.

[23] Q. Zhang, C. Hu, S. Upasani *et al.*, "Agentic context engineering: Evolving contexts for self-improving language models," arXiv preprint arXiv:2510.04618, 2025.

[24] W. Zhou, Y. Ou, S. Ding, L. Li, J. Wu, T. Wang, J. Chen, S. Wang, X. Xu, N. Zhang, H. Chen, and Y. E. Jiang, "Symbolic learning enables self-evolving agents," *arXiv preprint arXiv:2406.18532*, 2024.

[25] E. Zelikman, E. Lorch, L. Mackey, and A. T. Kalai, "Self-taught optimizer (STOP): Recursively self-improving code generation," in *First Conference on Language Modeling (COLM 2024)*, 2024, arXiv:2310.02304.

[26] M. Robeyns, M. Szummer, and L. Aitchison, "A self-improving coding agent," *arXiv preprint arXiv:2504.15228*, 2025.

[27] L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, and I. Stoica, "Judging LLM-as-a-judge with MT-Bench and Chatbot Arena," in *Advances in Neural Information Processing Systems (NeurIPS 2023), Datasets and Benchmarks Track*, 2023, arXiv:2306.05685.

[28] E. Palumbo, A. Tamborrino, V. Ode, B. Lacker, A. Casas Escoda, J. Hopple, M. Better, J. Leoni, H. Galvão, H. Bouchard, M. Lalmas, J. L. Redondo García, A. Abebe, A. Clifton, A. Blomberg, H. Lindström, D. Doro, and C. Doig Cardet, "Bootstrapping conversational recommendation agents at spotify: Synthetic data generation and self-improvement loops," *arXiv preprint arXiv:2609.30297*, 2026.

[29] J. Zhang *et al.*, "Darwin Gödel machine: Open-ended evolution of self-improving agents," *arXiv preprint*, 2025.

[30] A. Novikov *et al.*, "AlphaEvolve: A coding agent for scientific and algorithmic discovery," arXiv preprint arXiv:2506.13131, 2025.

[31] Y. Yang, Z. Gong, W. Huang, Q. Yang, Z. Zhou, Z. Huang, Y. Li, X. Gao, Q. Dai, B. Liu, K. Qiu, Y. Yang, D. Chen, X. Yang, and C. Luo, "SkillOpt: Executive strategy for self-evolving agent skills," arXiv preprint arXiv:2605.23904, 2026. [Online]. Available: https://microsoft.github.io/SkillOpt

[32] Cognition, "Devin's 2025 performance review: Learnings from 18 months of agents at work," https://cognition.com/blog/devin-annual-performance-review-2025, 2025.

[33] J. Wang *et al.*, "Reinforcement learning for self-improving agent with skill library," *arXiv preprint*, 2025.

[34] P. Xia *et al.*, "SkillRL: Evolving agents via recursive skill-augmented reinforcement learning," *arXiv preprint*, 2026.

[35] C. Ling, P. Chen, A. Guan, J. Qu, S. A. Akbar, M. Gopinathan, and E. Cornejo, "PACE: Two-timescale self-evolution for small language model agents," *arXiv preprint*, 2026.

[36] X. Liu *et al.*, "SkillForge: Forging domain-specific, self-evolving agent skills in cloud technical support," *arXiv preprint*, 2026, accepted at ACM SIGIR 2026 Industry Track.

[37] X. Huang *et al.*, "CASCADE: Cumulative agentic skill creation through autonomous development and evolution," *arXiv preprint*, 2025.

[38] L. Gao, J. Schulman, and J. Hilton, "Scaling laws for reward model overoptimization," *arXiv preprint arXiv:2210.10760*, 2022.

[39] S. Shankar, J. Zamfirescu-Pereira, B. Hartmann, A. G. Parameswaran, and I. Arawjo, "Who validates the validators? aligning LLM-assisted evaluation of LLM outputs with human preferences," *arXiv preprint arXiv:2404.12272*, 2024.

[40] C. Lu, C. Lu, R. T. Lange, J. Foerster, J. Clune, and D. Ha, "The AI scientist: Towards fully automated open-ended scientific discovery," *arXiv preprint arXiv:2408.06292*, 2024.

[41] J. Xu *et al.*, "NanoResearch: Co-evolving skills, memory, and policy for personalized research automation," *arXiv preprint*, 2026.

[42] A. Qu, H. Zheng, Z. Zhou, Y. Yan, Y. Tang, S. Y. Ong, F. Hong, K. Zhou, C. Jiang, M. Kong, J. Zhu, X. Jiang, S. Li, C. Wu, B. K. H. Low, J. Zhao, and P. P. Liang, "CORAL: Towards autonomous multi-agent evolution for open-ended discovery," *arXiv preprint*, 2026.

[43] K. T.-i. Ong, M. Kang, D. Choi, J. Cho, S. Kim, S. Lim, G. Jang, M. Oh, B. Jeong, S. Kim, T. Kwon, and J. Yeo, "Towards direct evaluation of harness optimizers via priority ranking," *arXiv preprint*, 2026.

[44] Y. Hu, M. Xu, H. Guo, Y. Song, X. Zhu, and B. Zheng, "Skillchain: Closing the loop on skill evolution for image-based e-commerce ai assistants," *arXiv preprint arXiv:2606.12984*, 2026.

[45] J. Zhu, K. Mao, Y. Guo, D. He, S. Xu, S. Gu, and Y. Yue, "Skillcoach: Self-evolving rubrics for evaluating and enhancing agentic skill-use," *arXiv preprint arXiv:2607.01874*, 2026.

[46] A. Gunjal, A. Wang, E. Lau, V. Nath, Y. He, B. Liu, and S. Hendryx, "Rubrics as rewards: Reinforcement learning beyond verifiable domains," *arXiv preprint arXiv:2507.17746*, 2025.

[47] R. Jia, Y. Yang, Y. Gai, K. Luo, S. Huang, J. Lin, X. Jiang, and G. Jiang, "Writing-zero: Bridge the gap between non-verifiable tasks and verifiable rewards," *arXiv preprint arXiv:2506.00103*, 2025.

[48] K. Shi, Y. Liu, P. Wang, A. R. Fabbri, S. Joty, and A. Cohan, "References improve LLM alignment in non-verifiable domains," *arXiv preprint arXiv:2602.16802*, Feb. 2026.

[49] X. Jin, J. Zhang, T. Cao, and B. Yang, "Beyond score prediction: LLM-based essay scoring and feedback generation via reinforcement learning with rubric rewards," *arXiv preprint arXiv:2607.19219*, Jul. 2026.

[50] T. Yu, B. Ji, S. Wang, S. Yao, Z. Wang, G. Cui, L. Yuan, N. Ding, Y. Yao, Z. Liu, M. Sun, and T.-S. Chua, "RLPR: Extrapolating RLVR to general domains without verifiers," *arXiv preprint arXiv:2506.18254*, Jun. 2025.

[51] R. Xu, T. Liu, Z. Dong, T. Yu, I. Hong, C. Yang, L. Zhang, T. Zhao, and H. Wang, "Alternating reinforcement learning for rubric-based reward modeling in non-verifiable LLM post-training," *arXiv preprint arXiv:2602.01511*, 2026.

[52] H. Ding, B. Huang, Y. Fang, W. Liao, Z. Li, J. Zhang, Z. Wu, J. Zhao, and Y. Wang, "EvoRubrics: Dynamic rubrics as rewards via adversarial co-evolution for LLM reinforcement learning," *arXiv preprint arXiv:2606.23038*, Jun. 2026.

[53] H. Zhu, C. Cai, Y. Song, X. Chen, S. Han, and Y. Guo, "Self-evolving deep research via joint generation and evaluation," *arXiv preprint arXiv:2606.04507*, Jun. 2026.

[54] G. Li, B. Dalvi Mishra, Z. Wang, J. Yan, Y. Chen, C.-L. Li, L. T. Le, R. Han, G. Lee, H. Tong, C.-Y. Lee, and T. Pfister, "RubricEM: Meta-RL with rubric-guided policy decomposition beyond verifiable rewards," *arXiv preprint arXiv:2605.10899*, May 2026.

[55] T. R. Sumers, S. Yao, K. Narasimhan, and T. L. Griffiths, "Cognitive architectures for language agents (CoALA)," *Transactions on Machine Learning Research (TMLR)*, 2023.

[56] AGENTS.md, "AGENTS.md: A cross-vendor open standard for agent instructions," 2025. [Online]. Available: https://agents.md

[57] R. Robbes, T. Matricon, T. Degueule, A. Hora, and S. Zacchiroli, "Agentic much? adoption of coding agents on github," *arXiv preprint arXiv:2601.18341*, 2026.

[58] M. Galster, S. Mohsenimofidi, J. L. Lulla, M. A. Abubakar, C. Treude, and S. Baltes, "Harness engineering for agentic ai coding tools: An exploratory study," *arXiv preprint arXiv:2602.14690*, 2026.

[59] T. Gloaguen, N. Mündler-Sasahara, M. N. Müller, V. Raychev, and M. Vechev, "Evaluating agents.md: Are repository-level context files helpful for coding agents?" *arXiv preprint arXiv:2602.11988*, 2026.

[60] X. Wu *et al.*, "Co-evolving LLM decision and skill bank agents for long-horizon tasks," *arXiv preprint*, 2026.

[61] X. Zhang, Y. Cui, G. Wang, Z. Li, W. Qiu, B. Zhu, and P. He, "Library drift: Diagnosing and fixing a silent failure mode in self-evolving llm skill libraries," *arXiv preprint arXiv:2605.19576*, 2026.

[62] Y. Wei, Z. Huang, S. Lu, J. Qian, Q. Wang, C. Wu, and L. He, "Skillsmith: Co-evolving skills and tools for self-improving agent systems," *arXiv preprint arXiv:2606.01314*, 2026.

[63] H. Pu, X. Song, and L. Zhao, "Skillops: Managing llm agent skill libraries as self-maintaining software ecosystems," *arXiv preprint arXiv:2605.13716*, 2026.

[64] Y. Wang, Y. Zhou, Y. Liang, C. Zhang, F. Liu, J. Zhou, and H. Yao, "Not all skills help: Measuring and repairing agent knowledge," *arXiv preprint arXiv:2606.15390*, 2026.

[65] C. Shen and E. Hruschka, "Who maintains agent skills? A longitudinal study of human-governed, AI-assisted skill maintenance," *arXiv preprint arXiv:2609.05677*, 2026.

[66] M. Zhuge, W. Wang, L. Kirsch, F. Faccio, D. Khizbullin, and J. Schmidhuber, "Language agents as optimizable graphs," in *Forty-first International Conference on Machine Learning (ICML 2024)*, 2024, arXiv:2402.16823.

[67] J. Zhang, J. Xiang, Z. Yu, F. Teng, X. Chen, J. Chen, M. Zhuge, X. Cheng, S. Hong, J. Wang, B. Zheng, B. Liu, Y. Luo, and C. Wu, "AFlow: Automating agentic workflow generation," *arXiv preprint arXiv:2410.10762*, 2024.

[68] Y. Cheng, C. Liu, H. Yu, H. Tan, T. Wu, W. Jin, J. Ran, K. Wang, X. Li, H. Bai, G. Chesi, and N. Wong, "Optimas: Automatically optimize multi-agent system," *arXiv preprint arXiv:2608.21918*, 2026.

[69] J. Chu, J. Seo, J. Cho, J. Na, Y. Xiong, Y. Kim, and H. J. Kim, "Agentgrad: Intervention-guided prompt optimization for multi agent systems," *arXiv preprint arXiv:2609.08572*, 2026.

[70] M. Chen, C. Lv, G. Zhang, H. Chang, and S. Zhou, "Harnessforge: Joint harness and policy evolution for adaptive agent systems," *arXiv preprint arXiv:2606.01779*, 2026.

[71] S. Zhao, R. Kong, Z. Zhang, L. Shen, X. He, X. Zhang, and Q. Zhang, "Topological coherence for self-evolving multi-agent systems," *arXiv preprint arXiv:2609.37953*, 2026.

[72] R. Wang and J. Cui, "Evo2team: When do evolved skills transfer? from selection to deployment," *arXiv preprint arXiv:2609.34135*, 2026.

[73] Letta, "MemFS; memory & dreaming; introducing context repositories: Git-based memory for coding agents," https://docs.letta.com/letta-code/memfs, https://docs.letta.com/letta-agent/memory, https://www.letta.com/blog/context-repositories/, 2026.

[74] Mem0, "Mem0 platform: Memory history, feedback, and changelog," https://docs.mem0.ai/api-reference/memory/history-memory, https://docs.mem0.ai/api-reference/memory/feedback, https://docs.mem0.ai/changelog, 2026.

[75] Zep, "Zep documentation: Facts; february 2026 deprecation wave; changelog," https://help.getzep.com/facts, https://help.getzep.com/february-2026-deprecation-wave, https://help.getzep.com/changelog/2026/9/24, 2026.

[76] Z. Wang, "TOKI: A bitemporal operator algebra for contradiction resolution in LLM-agent persistent memory," *arXiv preprint arXiv:2606.06240*, Jun. 2026.

[77] J. Xu, Y. Xiao, W. Shao, H. Liu, and X. Li, "Memory provenance laundering in LLM agents: A non-amplification firewall for persistent memory," *arXiv preprint arXiv:2607.29167*, Jul. 2026.

[78] Y. Margalit, N. Cohen-Inger, E. Avram, R. Taig, and O. Margalit, "Governed shared memory for multi-agent LLM systems," *arXiv preprint arXiv:2606.24535*, Jun. 2026.

[79] Z. Yu, Y. Wu, Z. Yin, K. Chen, Z. Zhao, M. Wang, S. Yan, and L. Yang, "Recursive experiential–working memory evolution for long-horizon agent harnesses," *arXiv preprint arXiv:2608.24876*, Aug. 2026.

[80] H. Wang, G. Dong, H. Liang, Z. Zhang, J. Luo, C. Liu, C. Xue, and H. Tang, "MemGuard: Persisting verifier signals for LLM-agent memory governance," *arXiv preprint arXiv:2608.21867*, Aug. 2026.

[81] X. Sun, Q. Wang, H. Li, and L. Chen, "Verifiable memory: Learning unified memory management with local and global verifiers for large language model agents," *arXiv preprint arXiv:2608.03137*, Aug. 2026.

[82] agentic-community, "Mcp gateway & registry," https://github.com/agentic-community/mcp-gateway-registry, 2026.

[83] Microsoft, "Inventory and discover mcp servers in your api center," https://learn.microsoft.com/en-us/azure/api-center/register-discover-mcp-server, 2026.

[84] TrueFoundry, "Mcp gateway authentication and security," https://www.truefoundry.com/docs/ai-gateway/mcp-gateway-auth-security, 2026.

[85] R. Kang and V. Wang, "Registry-governed agent lifecycle: Completing EDDOps with evaluation-driven registration, promotion, and retirement on AWS AgentCore," *arXiv preprint arXiv:2607.00345*, Jul. 2026.

[86] A. Quessada-Vial, "Agentic configuration management (ACM): A reference configuration model for governed agentic systems," *arXiv preprint arXiv:2608.11166*, Aug. 2026.

[87] X. Qin, S. Luan, J. See, Z. Boukhers, C. Yang, and Z. Li, "Governed capability evolution: Lifecycle-time compatibility checking and rollback for AI-component-based systems, with a proof-of-concept evaluation on embodied agents," *arXiv preprint arXiv:2604.08059*, Apr. 2026.

[88] M. Helwig, "Memory as infrastructure: Reliability engineering for persistent agent memory in months-long LLM-assisted development," *arXiv preprint arXiv:2609.05510*, Aug. 2026.

[89] A. Miyai, K. Aizawa, and T. Yamasaki, "Task-coevolve: Efficient harness optimization via adaptive validation task selection," *arXiv preprint arXiv:2608.20169*, 2026.

[90] S. Park, W. Kim, R. Tan, J. Zhang, W.-S. Han, P. Gao, C. Park, Y. Yao, R. Fu, E. Nallipogu, Q. Lin, S. Rajmohan, and D. Zhang, "Autosaddler: Automatic harness optimization with durable updates from agent execution traces," *arXiv preprint arXiv:2608.23041*, 2026.

[91] Y. T. Shen, K. Toyoda, and A. Leung, "Revoked but still authoritative: An empirical study of revocation enforcement in agent-memory systems," *arXiv preprint arXiv:2609.08258*, Sep. 2026.

[92] J. Qi, Y. He, J. Liu, Q. Huang, S. Hu, Z. Dang, G. Chen, R. Yang, W. Jiang, Y. Liu, T. Lyu, and F. Li, "Vace: Validation-gated alternating co-evolution of agent models and harnesses," *arXiv preprint arXiv:2609.37105*, 2026.

[93] Q. Yang, Z. Kou, X. Zhang, D.-D. Wu, H. Gu, J. Guo, Y. Liu, D. Jiang, and Q. Xu, "Federated agent optimization," *arXiv preprint arXiv:2610.01195*, 2026.

[94] J. Hu, H. Zhao, Q. Jia, D. Chen, T. Zhao, S. Duan, and G. Liu, "Learning from others, acting for you: Cross-user memory sharing for llm agents," *arXiv preprint arXiv:2609.32511*, 2026.

[95] J. Nie, Y. Zhang, Q. Cai, Y.-m. Cheung, X. Tian, and B. Han, "Evolvenet: Collaborative harness evolution for agent self-improvement," *arXiv preprint arXiv:2608.04968*, 2026.

[96] Y. Guo, H. Zhou, J. Zhang, Y. Zhang, Y. Liu, Q. Chen, Q. Lin, H. Dong, J. M. Zhang, and Z. Chen, "Xreposkill: Learning transferable skills for software engineering agents," *arXiv preprint arXiv:2609.36807*, 2026.

[97] J. Chu, J. S. Lee, J. Park, D. Ko, J. Na, S. Lee, T. Lee, M. Yoon, M. Joo, Y. Xiong, and H. J. Kim, "Retrieval-augmented skill optimization via cross-harness adaptation," *arXiv preprint arXiv:2609.38024*, 2026.

[98] X. Wang *et al.*, "AutoAgent: Evolving cognition and elastic memory orchestration for adaptive agents," *arXiv preprint*, 2026.

[99] B. Romera-Paredes, M. Barekatain, A. Novikov, M. Balog, M. P. Kumar, E. Dupont, F. J. R. Ruiz, J. S. Ellenberg, P. Wang, O. Fawzi, P. Kohli, and A. Fawzi, "Mathematical discoveries from program search with large language models," *Nature*, vol. 625, pp. 468–475, 2024.

[100] S. Hu, C. Lu, and J. Clune, "Automated design of agentic systems (ADAS)," *arXiv preprint*, 2024.

[101] G. Wang *et al.*, "Voyager: An open-ended embodied agent with large language models," *arXiv preprint*, 2023.

[102] Sierra, "Agent OS 2.0: From answers to memory and action," 2025. [Online]. Available: https://sierra.ai/blog/agent-os-2-0

[103] Braintrust, "Use prompts in code," Braintrust documentation, https://www.braintrust.dev/docs/guides/prompts, 2026, accessed 2026-10-03.

[104] Vellum, "Release tags," Vellum documentation, https://docs.vellum.ai/product/deployments/release-tags, 2026, accessed 2026-10-03.

[105] LangChain, "Prompt engineering concepts," LangSmith documentation, https://docs.langchain.com/langsmith/prompt-engineering-concepts, 2026, accessed 2026-10-03.

[106] P. Barbaste, T. Darrigol, G. Vu, and T. Wiltberger, "Harness engineering: Anatomy, architecture, and evolution of coding agents – a source-code study of eleven systems," *arXiv preprint arXiv:2609.00006*, 2026.

[107] J. L. Lulla, V. Nersesyan, S. Mohsenimofidi, C. Treude, and S. Baltes, "Loop engineering: Building blocks, adoption, and impact," *arXiv preprint arXiv:2608.21884*, 2026.

[108] T. Q. Nguyen, S. Desai, R. H. Anwar, F. Shaik, V. Suryanarayanan, and V. Chowdhary, "VerificAgent: Domain-specific memory verification for scalable oversight of aligned computer-use agents," *arXiv preprint arXiv:2506.02539*, Jun. 2025.

[109] Z. Jiang and B. Dong, "Buildrix: An open platform for sharing and benchmarking agentic ai skills in building engineering," *arXiv preprint arXiv:2606.25139*, 2026.

[110] Y. Wang, L. Xia, R. Liu, Z. Zhang, H. Wu, Y. Yang, J. He, Y. Guo, and K. Lei, "Sage: A statistical acceptance gate for self-evolving agents," *arXiv preprint arXiv:2609.36043*, 2026.

[111] L. Shang, M. Xu, Y. Sun, T. Xia, L. Hu, L. Xu, and N. Zheng, "When self-evolution backfires: Pre-commit gating against skill contamination in llm agents," *arXiv preprint arXiv:2608.05810*, 2026.

[112] Y. Shen, B. Li, and X. Zhang, "SkillOpt-Lite: Better and faster agent self-evolution via one line of vibe," *arXiv preprint arXiv:2607.03451*, 2026.

[113] L. A. Agrawal, S. Tan, D. Soylu, N. Ziems, R. Khare, K. Opsahl-Ong, A. Singhvi, H. Shandilya, M. J. Ryan, M. Jiang, C. Potts, K. Sen, A. G. Dimakis, I. Stoica, D. Klein, M. Zaharia, and O. Khattab, "GEPA: Reflective prompt evolution can outperform reinforcement learning," *arXiv preprint*, 2025.

[114] N. Vassilyev, W. Berrios, R. Zhang, B. Han, D. Kiela, and S. Mehri, "Reflective context learning: Studying the optimization primitives of context space," *arXiv preprint*, 2026.

[115] Y. Lee, R. Nair, Q. Zhang, K. Lee, O. Khattab, and C. Finn, "Meta-Harness: End-to-end optimization of model harnesses," *arXiv preprint*, 2026.

[116] J. Lin, S. Liu, C. Pan, L. Lin, S. Dou, Z. Xi, X. Huang, H. Yan, Z. Han, T. Gui, and Y.-G. Jiang, "Agentic harness engineering: Observability-driven automatic evolution of coding-agent harnesses," *arXiv preprint*, 2026.

[117] Q. Cai, Y. Zhang, X. Jia, H. Zheng, W. Xue, J. Song, X. Tian, and Y. Guo, "MOSS: Self-evolution through source-level rewriting in autonomous agent systems," *arXiv preprint*, 2026.

[118] T. Zehle, "CANTANTE: Optimizing agentic systems via contrastive credit attribution," *arXiv preprint*, 2026.

[119] M. Nguyen, W. C. Tan, N. A. Hassan, A. Raman, L. H. Lim, and A. F. Razak, "Where does harness-optimization value live? Localized gains and the budget-splitting trap in self-evolving LLM agents," *arXiv preprint arXiv:2609.02889*, 2026.

[120] L. Li, J. Li, J. Zhao, K. Ye, Y. Li, C.-z. Xu, and X. Gao, "Grow the harness, not the context: From strategy-free scaffolds to reusable specialist agents," *arXiv preprint arXiv:2609.26760*, 2026.

[121] X. Zhang, Y. Cui, G. Wang, Z. Li, W. Qiu, B. Zhu, and P. He, "Ratchet: A minimal hygiene recipe for self-evolving LLM agents," *arXiv preprint*, 2026.

[122] M. M. Tanjim, J. Subramanian, X. Chen, B. Kveton, S. Mukherjee, A. Zhang, S. Kim, S. Sarkhel, and S. Choudhury, "MOCHA: Multi-objective chebyshev annealing for agent skill optimization," *arXiv preprint*, 2026.

[123] H. Liu, H. Yang, T. Jiang, B. Tang, F. Xiong, Y. Luo, and Z. Li, "SkillsVote: Lifecycle governance of agent skills from collection, recommendation to evolution," *arXiv preprint arXiv:2605.18401*, 2026.

[124] Y. Shalev, Z. Ding, and M. Jamnik, "Training language agents to learn from experience," *arXiv preprint*, 2026.

[125] H. Li, R. He, Q. Zhang, C. Ji, Q. Mang, X. Chen, L. A. Agrawal, W.-L. Liao, E. Yang, A. Cheung, J. Zou, K. Olukotun, I. Stoica, and J. E. Gonzalez, "Combee: Scaling prompt learning for self-improving language model agents," *arXiv preprint*, 2026.

[126] J. Yang, G. Yao, Y. Zhang, R. R. Kompella, G. Liu, and S. Chang, "FederatedSkill: Federated learning for agentic skill evolution," *arXiv preprint arXiv:2606.03143*, 2026.

[127] P. Hebbar, Y. Manawat, S. Verboomen, A. Ivanova, S. Palanimalai, K. Bhatia, and V. Baskaran, "Sia: Self improving ai with harness & weight updates," *arXiv preprint arXiv:2605.27276*, 2026.

[128] H. Gao, H. Chen, C. Wang, S. Guo, L. Pang, Z. Liu, H. Shen, and X. Cheng, "Skill-laudit: Ground-truth-free skill evolution via paired trajectory auditing," *arXiv preprint arXiv:2606.14239*, 2026.

[129] T. Sah, D. Sah, H. Jain, and T. Sah, "Evoundo: Recoverability-constrained self-evolution for llm agent harnesses," *arXiv preprint arXiv:2608.28363*, 2026.

[130] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. Agüera y Arcas, "Communication-efficient learning of deep networks from decentralized data," in *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, ser. Proceedings of Machine Learning Research, vol. 54, 2017, pp. 1273–1282.

[131] H. B. McMahan, D. Ramage, K. Talwar, and L. Zhang, "Learning differentially private recurrent language models," in *International Conference on Learning Representations (ICLR)*, 2018, arXiv:1710.06963.

[132] P. Kairouz, H. B. McMahan, B. Avent, A. Bellet, M. Bennis *et al.*, "Advances and open problems in federated learning," *arXiv preprint arXiv:1912.04977*, 2019, published in Foundations and Trends in Machine Learning, vol. 14, no. 1–2, 2021.

[133] T. Guo, S. Guo, J. Wang, and W. Xu, "PromptFL: Let federated participants cooperatively learn prompts instead of models – federated learning in age of foundation model," *arXiv preprint arXiv:2208.11625*, 2022.

[134] D. Sculley, G. Holt, D. Golovin, E. Davydov, T. Phillips, D. Ebner, V. Chaudhary, M. Young, J.-F. Crespo, and D. Dennison, "Hidden technical debt in machine learning systems," in *Advances in Neural Information Processing Systems 28 (NIPS 2015)*, 2015.

[135] D. Jaroslawicz *et al.*, "How many instructions can LLMs follow at once?" *arXiv preprint*, 2025.

[136] W. Zou *et al.*, "PoisonedRAG: Knowledge corruption attacks to retrieval-augmented generation of large language models," *arXiv preprint*, 2024, accepted at USENIX Security 2025.

[137] Z. Chen, Z. Xiang, C. Xiao, D. Song, and B. Li, "AgentPoison: Red-teaming LLM agents via poisoning memory or knowledge bases," *arXiv preprint arXiv:2407.12784*, 2024.

[138] Z. Xu *et al.*, "From storage to steering: Memory control flow attacks on LLM agents," *arXiv preprint*, 2026.

[139] Y. Su, W. Xu, C. Zuo, and E. Bertino, "ChronoMem: Version control and semantic rollback for large language model agent memory," *arXiv preprint arXiv:2607.27773*, 2026.

[140] C. Yu, Y. Wang, J. Zhang, Y. Duan, M. Zheng, Z. Wu, K. Shi, and T. Cai, "From faulty memories to corrected actions: Dependency-guided rollback repair for memory-augmented agents," *arXiv preprint arXiv:2608.10502*, 2026.

[141] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, "Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection," *arXiv preprint arXiv:2302.12173*, 2023.

[142] E. Debenedetti, J. Zhang, M. Balunović, L. Beurer-Kellner, M. Fischer, and F. Tramèr, "AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents," *arXiv preprint arXiv:2406.13352*, 2024.

[143] Check Point Research, "MCPoison: Cursor IDE's MCP vulnerability (CVE-2025-54136)," https://research.checkpoint.com/2025/cursor-vulnerability-mcpoison/, 2025.

[144] E. Wallace, K. Xiao, R. Leike, L. Weng, J. Heidecke, and A. Beutel, "The instruction hierarchy: Training LLMs to prioritize privileged instructions," *arXiv preprint arXiv:2404.13208*, 2024.

[145] S. Singh *et al.*, "The leaderboard illusion," *arXiv preprint*, 2025.

[146] M. Cemri *et al.*, "Why do multi-agent LLM systems fail?" *arXiv preprint*, 2025.

[147] N. F. Liu *et al.*, "Lost in the middle: How language models use long contexts," *Transactions of the Association for Computational Linguistics (TACL)*, 2023.

[148] H. Robbins and D. Siegmund, "A convergence theorem for non negative almost supermartingales and some applications," in *Optimizing Methods in Statistics*. Academic Press, 1971, pp. 233–257.

[149] A. Madaan *et al.*, "Self-refine: Iterative refinement with self-feedback," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

[150] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, "ReAct: Synergizing reasoning and acting in language models," in *International Conference on Learning Representations (ICLR)*, 2023.

[151] D. Tran and D. Kiela, "Single-agent LLMs outperform multi-agent systems on multi-hop reasoning under equal thinking token budgets," *arXiv preprint*, 2026.

[152] N. Shinn *et al.*, "Reflexion: Language agents with verbal reinforcement learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

[153] Z. Ma, S. Yang, Y. Ji, X. Wang, Y. Wang, Y. Hu, T. Huang, and X. Chu, "SkillClaw: Let skills evolve collectively with agentic evolver," *arXiv preprint arXiv:2604.08377*, 2026.

[154] J. Ni *et al.*, "Trace2Skill: Distill trajectory-local lessons into transferable agent skills," *arXiv preprint*, 2026.

[155] S. Alzubi *et al.*, "EvoSkill: Automated skill discovery for multi-agent systems," *arXiv preprint*, 2026.

[156] W. Song, J. Yue, and Z. Pang, "ABSTRAL: Automatic design of multi-agent systems through iterative refinement and topology optimization," *arXiv preprint*, 2026.

[157] Y. He *et al.*, "EvoTest: Evolutionary test-time learning for self-improving agentic systems," *arXiv preprint*, 2025.

[158] A. Plesner, C. Northcutt, F. Guzmán, and A. Athalye, "A cheap verifier is good enough: LLM post-training is robust to erroneous rewards," *arXiv preprint arXiv:2609.33467*, Sep. 2026.

[159] M. Bhattarai, I. Boureima, N. R. Ranasinghe, S. Pakin, and D. O'Malley, "Rubric-grounded RL: Structured judge rewards for generalizable reasoning," *arXiv preprint arXiv:2605.08061*, May 2026.

[160] R. Jia, Y. Yang, W. Wang, Y. Wu, Y. Gai, S. Tao, M. Zhou, J. Lin, X. Jiang, and G. Jiang, "Open rubric system: Scaling reinforcement learning with pairwise adaptive rubric," *arXiv preprint arXiv:2602.14069*, Feb. 2026.

[161] C. Lv, J. Zhou, W. Zhao, J. Xu, S. Dou, Z. Huang, M. Tian, X. Wang, Z. Guo, Y. Liu, P. Zhou, T. Gui, L. Tian, X. Zhou, X. Zheng, X. Huang, and J. Zhou, "Learning query-specific rubrics from human preferences for DeepResearch report generation," *arXiv preprint arXiv:2602.03619*, Feb. 2026.

[162] H. Chen, Z. Han, Y. Yan, Q. Zhu, M. Sun, and W. Che, "From holistic evaluation to structured criteria: Rubrics across the evolving LLM landscape," *arXiv preprint arXiv:2606.08625*, 2026.

[163] Z. Ren, Y. Chen, D. Guo, G. Rong, T. Li, R. B. Xiong, Q. Lan, W. Wang, L. Nanbo, Y. Yang, M. Zhuge, and J. Schmidhuber, "Self-improvements in modern agentic systems: A survey," *arXiv preprint arXiv:2607.13104*, 2026.

[164] P. Du, "Memory for autonomous LLM agents: Mechanisms, evaluation, and emerging frontiers," *arXiv preprint arXiv:2603.07670*, Mar. 2026.

[165] W.-C. Huang, W. Zhang, Y. Liang, Y. Bei, Y. Chen, T. Feng *et al.*, "A survey of agent memory in the second half: Towards self-evolving and long-horizon agents," *arXiv preprint arXiv:2602.06052*, 2026.

[166] J. Luo, Y. Tian, C. Cao, Z. Luo, H. Lin, K. Li, C. Kong, R. Yang, and J. Ma, "From storage to experience: A survey on the evolution of LLM agent memory mechanisms," *arXiv preprint arXiv:2605.06716*, May 2026.

[167] Q. Hua, L. Ye, D. Fu, Y. Xiao, X. Cai, Y. Wu, J. Lin, J. Wang, and P. Liu, "Context engineering 2.0: The context of context engineering," *arXiv preprint arXiv:2510.26493*, Oct. 2025.

[168] Y. Zhou, W. Shu, Y. Su, W. Du, Y. Fang, and X. Lin, "A comprehensive survey on agent skills: Taxonomy, techniques, and applications," *arXiv preprint arXiv:2605.07358*, 2026.

[169] Y. Zhou, A. I. Muresanu, Z. Han, K. Paster, S. Pitis, H. Chan, and J. Ba, "Large language models are human-level prompt engineers," *arXiv preprint arXiv:2211.01910*, 2022.

[170] R. Pryzant, D. Iter, J. Li, Y. T. Lee, C. Zhu, and M. Zeng, "Automatic prompt optimization with "gradient descent" and beam search," in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP 2023)*, 2023, arXiv:2305.03495.

[171] K. Opsahl-Ong, M. J. Ryan, J. Purtell, D. Broman, C. Potts, M. Zaharia, and O. Khattab, "Optimizing instructions and demonstrations for multi-stage language model programs," in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP 2024)*, 2024, arXiv:2406.11695.

[172] C.-A. Cheng, A. Nie, and A. Swaminathan, "Trace is the next AutoDiff: Generative optimization with rich feedback, execution traces, and LLMs," *arXiv preprint arXiv:2406.16218*, 2024.

[173] J. S. Park *et al.*, "Generative agents: Interactive simulacra of human behavior," in *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)*, 2023.

[174] A. Zhao, D. Huang, Q. Xu, M. Lin, Y.-J. Liu, and G. Huang, "ExpeL: LLM agents are experiential learners," in *Proceedings of the 38th Annual AAAI Conference on Artificial Intelligence (AAAI-24)*, 2024, arXiv:2308.10144.

[175] C. Packer, S. Wooders, K. Lin, V. Fang, S. G. Patil, I. Stoica, and J. E. Gonzalez, "MemGPT: Towards LLMs as operating systems," *arXiv preprint arXiv:2310.08560*, 2023.

[176] Z. Z. Wang, J. Mao, D. Fried, and G. Neubig, "Agent workflow memory," *arXiv preprint arXiv:2409.07429*, 2024.

[177] W. Xu, Z. Liang, K. Mei, H. Gao, J. Tan, and Y. Zhang, "A-MEM: Agentic memory for LLM agents," in *Advances in Neural Information Processing Systems (NeurIPS 2025)*, 2025, arXiv:2502.12110.

[178] M. Suzgun, M. Yuksekgonul, F. Bianchi, D. Jurafsky, and J. Zou, "Dynamic cheatsheet: Test-time learning with adaptive memory," *arXiv preprint arXiv:2504.07952*, 2025.

A Formal statement of artifact-layer descent

This appendix states the setup behind Section 3, a one-step descent inequality under a noisy gate, an informal convergence proposition, the optimization-analogy correspondence, and four worked cases. No new theorem is claimed; the proposition is the ordinary target of local stochastic optimization under noisy estimators, restated for a discrete, versioned parameter.

A.1 Setup

Let the frontier model have fixed parameters θ_M and let the scaffold be a structured object $\theta_s \in \Theta_s = \Theta_1 \times \cdots \times \Theta_6$ over the six substrate coordinates of Section 5; Θ_s is discrete, symbolic, and versioned. For a deployment D, $L_D(\theta_s) = \mathbb{E}_{x\sim D}[\ell(\pi(\theta_M, \theta_s, x))]$. A candidate edit $z_t \sim Q_t(z \mid H_t, \theta_s^t)$ is drawn from a proposal process conditioned on the accumulated history H_t of failures, corrections, rejected outputs, evaluation regressions, and prior accepted edits. Q_t is not uniform: for a mode m that recurs with frequency f_m and carries per-occurrence loss s_m, its expected contribution to the failure rate is $f_m s_m$, so high-frequency or high-severity modes are proposed first. The finite directional difference is $\Delta_z L_D(\theta_s) = L_D(\theta_s \oplus z) - L_D(\theta_s)$; the gate sees only a noisy estimate $\widehat{\Delta}_z L_D$. The update is

$$\theta_s^{t+1} = \theta_s^t \oplus z_t \text{ if } G_t(z_t) = 1, \quad \theta_s^{t+1} = \theta_s^t \text{ otherwise.}$$

Table 9: Optimization concepts and their scaffold analogues.

Optimization concept	Scaffold analogue
Parameter vector	Instructions, skills, memories, tools, routing graphs, evaluation gates, governance rules
Loss	Failure, user correction, evaluation regression, retrieval miss, tool error, schema violation
Gradient estimate	Credit assignment identifying the coordinate responsible for a failure
Step	Accepted scaffold edit; step size is edit magnitude under $\oplus$
Aggregated update	Several recurring failures combined before an update, by averaging or by agreement across k proposals
Update radius	Promotion scope: which loss is descended (USER, PROJECT, STACK, TENANT, CORE)
Validation loss	Held-out patch evaluation before merge
Regularization	Complexity penalty $\Omega(z)$, bloat control, cross-patch regression checks
Overfitting	Scaffold improves on this deployment and degrades elsewhere
Rollback	Reverting a harmful accepted update

A.2 Descent inequality

Assume (A1) positive descent bias on accepted edits, $\mathbb{E}[\widehat{\Delta}_{z_t} L_D \mid G_t = 1, \theta_s^t] \le -\gamma_t - \tau_t$ for some $\gamma_t > 0$, and (A2) bounded estimator bias on accepted edits, $|\mathbb{E}[\Delta_{z_t} L_D - \widehat{\Delta}_{z_t} L_D \mid G_t = 1, \theta_s^t]| \le \varepsilon_t$. With $p_t = \Pr(G_t = 1 \mid \theta_s^t)$,

$$\mathbb{E}[L_D(\theta_s^{t+1}) \mid \theta_s^t] \le L_D(\theta_s^t) - p_t(\gamma_t + \tau_t) + p_t\,\varepsilon_t.$$

Per-step expected improvement is the gate's margin minus its estimator bias, weighted by acceptance probability. In the feedback-signal classes of Section 4, ε_t is near zero for the check itself in C1, with the sampling error of a finite suite over the patch remaining, bounded by batching for an observable outcome (C2), equal to the judge's bias and bounded only by calibration for a judged signal (C3), and unavailable per edit for a delayed signal (C4).

A.3 Convergence (informal)

Add (A3) proposal coverage with acceptance: at every non-gate-stable scaffold, Q_t places probability at least $q > 0$ on an admissible edit with $\mathbb{E}[\Delta_z L_D] \le -\delta$, and the gate accepts such an edit with probability at least $q' > 0$, so that $p_t(\gamma_t + \tau_t) \ge qq'\delta$ there; (A4) effective compactness of the reachable orbit; and (A5) summable bias, $\sum_t p_t \varepsilon_t < \infty$; and let $0 \le L_D \le B$. The one-step inequality is then of the Robbins–Siegmund form, a non-negative process that decreases in expectation up to a summable perturbation $p_t\varepsilon_t$, so $L_D(\theta_s^t)$ converges almost surely and $\sum_t p_t(\gamma_t + \tau_t) < \infty$ almost surely [148]. The finite sum carries the content: a bounded loss cannot pay for infinitely many accepted edits with a positive expected margin, so under (A3) the process visits non-gate-stable scaffolds only finitely often, and with (A4) the limit set of $\{\theta_s^t\}$ is gate-stable: no admissible edit clears the gate with negative expected directional loss in the limit. This is convergence of the loss and stability of the gate decision, not convergence to an optimum; a globally suboptimal gate-stable scaffold is consistent with the result, because Q_t may never propose the improving edit or $\widehat{\Delta}$ may be biased to reject it. In a non-stationary deployment ($D = D_t$) the same mechanism tracks a moving optimum rather than converging once, which is why recency weighting, pruning, expiry, and rollback are the scaffold-level analogues of continual-learning machinery rather than optional hygiene.

A.4 Correspondence

Table 9 pairs each optimization concept with its scaffold analogue. Credit assignment here is attribution, not backpropagation. Scaffold systems do not differentiate through instruction files. They trace an output error to an upstream coordinate, manually or with a judge or verifier. The descent argument needs only that this attribution succeeds with non-zero probability per failure, which is the credit-assignment condition of Section 3.

A.5 Worked cases

Retrieval. Failures recur because retrieved context omits the decisive clause. Attribution traces the error to the retrieval coordinate: chunking, query rewriting, metadata filter, embedding model, reranker, or context-packing rule. The candidate edit tightens chunk boundaries or adds a metadata constraint, and the gate checks that recall and answer accuracy improve on the patch evaluation without degrading adjacent cases. Signal class C1 or C2.

Pitch review. The system repeatedly overrates decks that claim a large market without numbers, and the failure cluster points at an under-specified rubric coordinate. The edit changes "evaluate market size" into "require a numeric estimate and a source; if absent, cap the score"; the gate runs the revised rubric against held-out decks with human scores. What the edit changed was the gate that will judge the next edit: the reflexive gate of Section 4, class C3.

Tool use. Calls fail because optional parameters are under-specified. The responsible coordinate may be the tool schema, the call policy, or the clarification rule. The edit adds required fields or a pre-call uncertainty check; the gate checks that tool-call failures fall without raising refusals or latency. Class C2.

Memory. A stale customer preference is repeatedly retrieved and causes bad decisions. The responsible coordinate is a memory entry together with its provenance, priority, and expiry policy. The edit corrects, expires, or version-tags the entry; the gate checks that future behavior improves without erasing useful history. Class C2, with the correction itself observable only when a later outcome reveals it.

B Gate families and feedback-signal classes

Section 3 collapses every production gate into one per-edit test G_t on an estimate of the loss change. Surveyed systems instantiate that test in several distinct ways, and some of them are better read as proposal pruning that precedes the per-edit gate than as estimators of $\widehat{\Delta}L_D$. Table 10 lists the families, the slot each occupies, representative systems, and the feedback-signal class (Section 4) each family serves.

The first five families populate G_t directly and, under the assumptions of Appendix A, inherit its descent guarantee. Semantic merge and archive selection are not estimators of $\widehat{\Delta}L_D$; they constrain the proposal distribution, and the descent argument applies to the composition of pruning and gate. Multi-sample agreement is listed because it is the correct analogue of "wait for k pieces of evidence before accepting an edit", a different operation from mini-batch averaging and one that should not be conflated with it. Read down the signal-class column. The only families available to a judged (C3) deployment are the two whose estimator is a judge, and a delayed (C4) deployment has no family that estimates the loss change per edit at all: what it has is semantic merge before the gate and whatever proxy milestone stands at the horizon. That is the structural content of Section 4.

Table 10: Gate families, their slot in the formalism, and the feedback-signal class each serves.

Gate family	What it estimates	Slot in Section 3	Signal class	Representative systems
Surrogate verifier	Pass or fail from tests, a symbolic checker, or an exact match	Per-edit gate G_t	C1	DGM (sandboxed code), SkillOpt (held-out exact match), Voyager (game-state check: C2)
Evaluation cascade	$\widehat{\Delta}L_D$ on a held-out task suite, cheap-to-expensive evaluators	Per-edit gate G_t	C1 when the suite is exact; C2 when it is a task suite	AlphaEvolve, Braintrust, Vellum, LangSmith Hub
RL threshold	Reward improvement against a deployment threshold	Per-edit gate G_t	C2	SkillRL, SAGE, EvolveR, PACE
LLM judge with calibration	Pairwise or rubric score, calibrated against human agreement	Per-edit gate G_t	C3	SkillForge (Loop 1), CASCADE (consolidation gate)
Reviewer judgment	Human credit assignment by pull-request review or expert sign-off	Per-edit gate G_t, human-instantiated	C3 (and C2 when the reviewer has outcome data)	SkillForge (Loop 2), Sierra Agent OS 2.0, Anthropic skills repository
Semantic merge	Whether a candidate duplicates or conflicts with existing artifacts	Proposal pruning before G_t	C4 (also a pre-filter in C2, C3)	NanoResearch orchestrator, FederatedSkill server agent
Archive selection	Whether a candidate improves on, or diversifies, an archive of already-scored programs	Proposal pruning before G_t	Any; inherits the class of the evaluator that scored the archive (C1 in AlphaEvolve)	AlphaEvolve (MAP-Elites-style database), ADAS
Multi-sample agreement	Whether k independent proposals or evaluations agree; false-accept rate decays as p^k rather than variance as $1/k$	Aggregation filter on the estimate	C2, C3 noise absorption	VaG (three heterogeneous critics must all pass); SkillOpt-Lite (consensus filter across rollouts)

C Data appendix

Corpus and window. The corpus holds 256 unique systems from January 2023 to the October 2026 cutoff, released as two files. `p3_master_scores.csv` holds the 130 first-pass systems in 140 rows; `p3_refresh_2026-10.csv` holds the 126 second-pass systems in 133 rows, four of which re-audit a first-pass system and replace its row rather than add a system. The corpus therefore has 140 + 133 - 4 = 269 substrate-level rows. A system may be scored on several substrates, so rows outnumber systems, and Table 4 counts rows.

Two passes. Every system in either pass is scored against one rubric, from its own method section or documentation. First-pass systems were found in arXiv, vendor documentation, and public repositories through May 2026. Second-pass systems, found from June to the October 2026 cutoff, come from six query families: self-evolving skills and skill lifecycle; harness and context optimization; agent memory governance; prompt-as-code and evaluation-gated promotion; federated or multi-tenant skill sharing; and multi-agent topology evolution. Some were reached instead through the citations of the first-pass systems closest to the missing architecture of Section 9. The second pass was a systematic search of the field, not only a search around the known systems, so the per-substrate statistics pool both passes.

Inclusion and exclusion. A system enters the corpus if it either wraps a foundation model with at least one persistent artifact updated from deployment signal or is described in public as "self-improving," "scaffolded," or "agentic." The contrast set enters through the second criterion and is kept on purpose, so that the discriminator has something to mark against (Self-Refine, ReAct, TextGrad, OPRO, engineer-authored convention files; Appendix E.1.1). Two kinds of system are excluded: those whose only update mechanism is weight-level fine-tuning or reinforcement learning, which are model-side and belong to the self-evolution literature this survey positions against, and those with no public evidence of deployment or release, such as white papers without code, demos, or a production claim.

Two filters. The patch-local discriminator admits a system to the core corpus when it scores at least 3 on at least one substrate. Of the 256 systems, 204 pass it, 99 of them first-pass. The other 52, 31 of them first-pass, form the contrast corpus: systems described in public as agentic or self-improving that fail the discriminator. Each second-pass contrast row gives its exclusion reason in the released CSV. The stricter composite two-loop audit then asks whether Loop 1 modifies a persisted artifact from session signal and whether Loop 2 has an explicit cross-context promotion mechanism with versioning or lineage. Representative first-pass members of both groups, contrast systems and core systems that fail the composite audit, are listed in Appendix E.1.

Systems outside the released table. Fifteen named systems have no row in the released table. Ten are treated at mechanism level in the main text and this supplement: SkillOpt, GEPA, Meta-Harness, CORAL, PACE, MOSS, RCL, FederatedSkill, SkillClaw, and ChronoMem. Each was classified from its paper after the first pass closed and is placed by that reading, in the matrix or among the research families. The other five are contrast or boundary entries, classified from their papers and marked with a dagger in Appendix E.1: ReAct, MetaGen, MetaReflection, CLIN, and ACE. The corpus counts (256, 204, 52) exclude all fifteen. SkillOpt is nonetheless one of the fourteen research systems that pass the composite audit, because the audit was applied to it from its paper.

Evidence tiers. Each system carries one or more of five tiers: (T1) peer-reviewed publication; (T2) arXiv preprint; (T3) production claim with metrics; (T4) open-source artifact; (T5) industry

blog or interview. T1 and T4 weigh most in cluster analysis. T3 counts where the claim is concrete, naming a benchmark or an integration target; T5 counts only when T1–T4 evidence corroborates it elsewhere. Tier flags reach cluster claims through the per-row evidence column of the released CSV, and T3 production claims are never pooled with T1 or T2 evidence in a cluster-level statistic.

Tier distribution. In the first-pass core corpus, about a quarter of the systems carry T1 evidence as their primary tier, roughly half carry T2, and the rest T3 or T4. Preprint-only systems in either pass carry T2, recorded per row in the released CSV. Some T2 work is recent enough that its ablations may yet be revised; the cluster-level claims survive dropping any single T2 system.

Scoring discipline. Each system is scored substrate by substrate on the survey's 0–5 rubric and also receives an integrated patch-local-adaptation score in {0, 1, 2, 2.5, 3, 3.5, 4, 4.5, 5}. Half-scores are used sparingly, where one criterion is clearly met and another only partly, and ambiguous cases are rounded down. Every score comes from a single rater. The full score table is therefore published, each row with a one-line justification in the released CSVs, so that any ambiguous classification can be checked independently (Appendix I).

Composite two-loop criterion (sensitivity analysis). The strict composite audit passes fourteen research systems, all first-pass except SkillOpt, and three industry exemplars. Two natural relaxations change that count.

The governance relaxation counts any of versioning, lineage, or rollback as Loop-2 governance. It re-admits MAE, MetaGen, MetaReflection, and CLIN from the first pass, raising the research count to approximately eighteen, and of the second-pass promoters below it admits EvolveNet alone, whose rollback to the previous shared harness counts.

Seven second-pass systems promote across users, clients, teams, or repositories with no versioning or lineage on the promoted artifact: EvolveNet, Evo2Team, XRepoSkill, ShareMem, RASO, Spotify's recommendation loop, and Future AGI's prompt lifecycle. They fail the audit on governance alone. A weaker relaxation that asks only for an explicit promotion mechanism, with no governance property at all, would admit six of them, all landing in [Auto × Auto] except Spotify's loop, which would land in [Auto × Human]; none would open a new cell. The seventh, Future AGI's prompt lifecycle, rests on T5 evidence alone, which the tier rule above does not count without corroboration, so it is not admitted under any rule.

The optimizer-state relaxation admits as Loop 1 a session signal that updates optimizer state without persisting an artifact distinct from that state. It re-admits DSPy and TextGrad from the first pass, raising the research count to approximately sixteen. It admits no second-pass system. The second-pass contrast rows are benchmarks, evaluation and attack studies, registries, audits, and conceptual frameworks, plus one orchestrator whose adaptation lives in the weights, and none of them persists optimizer state as its only artifact.

The qualitative findings hold under every relaxation: [Auto × Auto] is well populated; [Human × Auto] and [Human × Human] are empty; [Human × None] holds only curation workflows and [None × Human] only registries, all of which fail the composite audit; the governance-first cluster sits in [None × Auto]; Letta alone passes the audit with an ungated promotion; and no system adds governed cross-tenant promotion.

D The mechanism stack and continual-learning notes

D.1 The mechanism stack

The update rule of Section 3 expands into six operations that every patch-local system performs, whether by explicit engineering or by accident.

1. **Residual capture.** Log failures, corrections, rejected outputs, user edits, evaluation regressions, and tool-call failures. This is the signal source of Loop 1; without it, z_t has no provenance.
2. **Mode discovery and credit assignment.** Cluster repeated failures into the deployment's own error modes (the catalogue of *The Architecture of Errors* [8], instantiated locally for one patch) and attribute each mode to the scaffold coordinate responsible for it (the correspondence note of Appendix A).
3. **Delta synthesis.** Turn a recurring mode into a candidate scaffold edit: an instruction rewrite, a new skill, a memory correction, a tool binding, an orchestration change, or an evaluation case. This produces z_t.
4. **Delta validation.** Estimate whether the edit lowers L_D without collateral damage on adjacent inputs. The gate G_t is Auto when the check is an algorithmic criterion and Human when it is reviewer judgment.
5. **Promotion control.** Decide whether the accepted edit stays in the local fork or climbs the scope hierarchy: PROJECT, STACK, TENANT, CORE. This is Loop 2.
6. **Drift and bloat control.** Expire stale rules, prune conflicting memories, roll back regressed skills, and detect overfitting. Without this operation, accumulated edits produce the scaffold-level analogue of plasticity loss.

Each operation is the locus of a separate engineering discipline, and surveyed systems implement them partially and unevenly. Research full-auto systems excel at operations 3 and 4, governance-first systems at operation 5, and production hybrids at most of operations 1 to 5 but rarely at operation 6. Implementation patterns for each operation lie outside this survey.

D.2 Loop 1 and Loop 2 as descent steps of different radius

Loop 1 proposes and validates local descent steps on a single L_D. Loop 2 widens the update radius: it selects, merges, and promotes accepted local edits so that they apply to a larger set of patch losses averaged together. Acceptance at USER scope descends on a single L_D. PROJECT promotion descends on a loss averaged over the project, and STACK promotion widens the average further. TENANT promotion descends on a tenant-wide average; CORE promotion, the cross-tenant level, descends on an average over tenants. The two-loop design space is therefore more than a taxonomy of engineering patterns: it records *which loss each surveyed system is descending on*, and which estimator gates the step.

D.3 Overfitting in scaffold adaptation

In weight training, overfitting is usually a defect. In scaffold adaptation, local overfitting is partly the point. Frontier weights are trained to preserve generality across patches, so they must smooth over local conventions and rare workflow-specific failures. A deployment's patch needs the opposite: selective specialization to its own residuals. The scaffold is therefore a layer meant to overfit, but one whose overfitting is scoped, inspectable, versioned, gated, and reversible. A scaffold edit that improves one repository, hospital workflow, or legal diligence process may degrade performance

elsewhere. That is not a contradiction; it is patch adaptation. The mistake is promoting that edit beyond its evidence radius.

In non-stationary patches ($D = D_t$), the same descent mechanism tracks a moving optimum rather than converging once. Recency weighting, replay, pruning, expiry, and rollback are therefore not optional hygiene; they are the scaffold-level analogues of continual-learning machinery.

E Contrast set, matrix membership, and closest approaches (May 2026 corpus)

The tables in this section describe the first-pass corpus, which closed in May 2026; the second-pass systems that occupy a cell of the Section 7 matrix are named in the cell-by-cell discussion that follows the matrix.

E.1 Contrast set

Two tables mark the two boundaries of the survey. The first lists representative members of the first-pass contrast corpus: 31 systems that fail the patch-local discriminator, scoring below 3 on every substrate. The second lists core systems, each scoring at least 3 on some substrate, that still fail the stricter composite two-loop criterion, so that the asterisked entries of the matrix in Appendix E.2.1 can be checked.

Two marks recur. An asterisk, used in the matrix, marks a system that fails at least one criterion of the composite audit and is included for contrast. A dagger has two uses: in the two contrast tables it marks a system classified from its paper because the score table has no row for it (ReAct, MetaGen, MetaReflection, CLIN, ACE); in the matrix it marks an industry exemplar evaluated with vendor documentation as the primary evidence (Sierra Agent OS 2.0, Cognition Devin).

E.1.1 Contrast corpus: fails the patch-local discriminator

System	Year	Discriminator failure	Brief reason
Self-Refine [149]	2023	No persistent artifact	In-context iteration only
ReAct [150]†	2023	No persistent artifact	Thought-act-observe loop is intra-session
TextGrad [22]	2024	No artifact distinct from optimizer state	Textual-gradient optimization; SkillOpt [31] (Section 8; Appendix F) adds the persistent `best_skill.md`, held-out validation gate, rejected-edit memory, and textual learning-rate schedule that TextGrad lacks.
OPRO, Toolformer	2023	No persistent deployment artifact	Optimization or tool use inside a single training or inference run
Aider conventions, Lovable knowledge	2024	No feedback loop	Engineer-authored convention files

System	Year	Discriminator failure	Brief reason
AGENTS.md (as a standard)	2025	No feedback loop	Cross-vendor artifact format as a workspace file; adoption measured, adaptation absent (Section 6). The specification's governance under the AAIF is a separate S6 entry, "AGENTS.md / AAIF"
LangChain Skills, HuggingFace Agent Skills, CrewAI Skills	2025–2026	No feedback loop	Curated skill packs without failure-triggered update
Tran & Kiela [151]	2026	Not a system	Single-agent baselines match or beat multi-agent under matched thinking-token budgets; constrains S5 claims (see Section 6)

E.1.2 Core corpus, fails the composite two-loop criterion

System	Year	Composite-criterion failure	Brief reason
Reflexion [152]	2023	No Loop 2	Per-agent episodic memory without promotion
Generative Agents	2023	No Loop 2	Per-agent reflection + importance memory
ExpeL	2023	Loop boundary blurred	Experience bank built from training tasks
Mem0 (vanilla)	2024	No Loop 2	Per-user memory; no cross-context promotion
Zep	2024	No Loop 2	Per-agent/session knowledge graph
ChatGPT Memory	2024	No Loop 2	Per-user persistence
Claude Memory	2025	No Loop 2	Per-conversation-thread persistence
MemGPT	2023	No Loop 2	Virtual-context management per agent
DSPy	2023	No artifact distinct from optimizer state	Prompt-template optimization; SkillOpt [31] (Section 8; Appendix F) realizes the artifact-bearing form.
MetaGen†	2026	No Loop 2; no cross-session persistence	Role pool dynamic within single inference call
MAE (Multi-Agent Evolve)	2025	No Loop 2	Co-evolution within Proposer/Solver/Judge triplet
MetaReflection†	2024	No Loop 2	Per-agent semantic memory
CLIN†	2023	No cross-instance promotion	Causal abstractions per-agent-per-environment

System	Year	Composite-criterion failure	Brief reason
LangGraph, AutoGen, CrewAI	2023–2024	Framework only	Orchestration frameworks; not deployed systems
W&B Weave, Agenta, Langfuse	2024	No Loop 1	Observability / eval-as-code
Mobile-Agent-E, Cradle, OS-Copilot	2024–2025	No Loop 2	Self-evolving but per-session/per-deployment
Agent S, Agent S2	2024–2025	No Loop 2	Computer-use agentic framework
ACE [23]†	2025	No Loop 2 in surveyed form	Single-context evolving playbook
Cursor Rules (vanilla)	2024	No Loop 2	Per-project artifacts
GitHub Copilot custom instructions, Windsurf, Continue.dev, Zed AI rules	2024	No Loop 1	Engineer-authored at repo base branch
Replit `replit.md`	2025	Loop 2 absent	Self-writing scaffold but per-workspace

The contrast systems fail the discriminator in three main ways: they iterate in context only, they have no feedback loop, or they are not systems at all. Core systems fail the composite criterion mostly for want of a Loop 2.

E.2 Matrix membership and closest approaches

E.2.1 Two-loop matrix: cell-by-cell discussion

Both axes record who decides, as Section 3 defines: a gate is Auto when a program applies a check, Human when a person decides, and None when the step is absent or runs unchecked. An engineer-authored or registrant-supplied artifact with no failure trigger therefore has Loop 1 None in both passes, and who supplies each candidate is listed separately in the occupant table at the end of this section.

The first-pass matrix, with the marks of Appendix E.1:

	Loop-2 Auto-gated	Loop-2 Human-gated	Loop-2 None
Loop-1 Auto-gated	NanoResearch, SkillOpt, SkillRL, AlphaEvolve, ADAS, DGM, Voyager, SAGE, EvolveR, CoEvoSkills, AutoAgent, AutoManual	SkillForge, CASCADE, Sierra Agent OS 2.0†, Cognition Devin†	MAE*, MetaGen*, MetaReflection*, CLIN*, DSPy*, TextGrad*
Loop-1 Human-gated	**[empty by engineering logic]**	**[empty in the first pass]**	*(none)*
Loop-1 None	ExpeL* (weak), Braintrust*, LangSmith Hub*, Vellum*	Anthropic skills repository*	Self-Refine*, ReAct*, Mem0*

The matrix places the first-pass systems audited under the composite criterion, with SkillOpt, and the contrast entries marked with an asterisk. Every matrix entry that passes the audit sits in [Auto × Auto] or [Auto × Human]: the twelve research systems of the first cell, with SkillForge and CASCADE, are the fourteen research systems of Section 5, and Sierra Agent OS 2.0 and Cognition Devin are two of its three industry exemplars. The third exemplar, Letta's Context Repositories, is a second-pass system.

Two placements need a word. DSPy and TextGrad sit in [Auto × None] because their Loop 1 is automated but persists only optimizer state, which only the optimizer-state relaxation of Appendix C would count as a persisted artifact. DGM runs in two modes depending on deployment configuration, archive-based evolution in [Auto × Auto] and sandboxed code-modification review in [Auto × Human], and is listed once, in the dominant cell.

COSPLAY and Cuadros et al., which are first-pass rows, and PACE, CORAL, GEPA, Meta-Harness, MOSS, and RCL, which were classified from their papers after the first pass closed, carry no composite audit, so they are neither placed in the matrix nor counted in Section 5. Second-pass systems record their gate pairs in the released score table, and the paragraphs below name those that occupy a cell.

[Auto × Auto] is the most populated cell: twelve systems, each with a machine-evaluated criterion on both loops. Speed on both loops is what its members share; cross-organization governance is what every one of them lacks.

[Auto × None] holds the systems whose local loop is gated and whose promotion has no gate, whether because nothing is promoted or because everything is. One audit passer sits here: Letta's Context Repositories push shared memory across agents automatically, with git lineage and no evaluation of any kind, so the composite audit admits them on lineage and the matrix records a promotion that nothing decides (Section 7). Mem0 Platform, TOKI, and Learning on the Job sit here with no promotion at all.

[Auto × Human] is the production-viable cell. Loop 1 fires at machine speed under an algorithmic criterion, and Loop 2 waits for a person before an edit propagates; the pairing reconciles the speed of automation with a reviewer's accountability. SkillForge is the clearest instance: an LLM judge

gates each skill diff on Loop 1, and reviewers decide on Loop 2 whether it ships. CASCADE, Sierra Agent OS 2.0, and Cognition Devin are variants of the same pairing.

Seven second-pass systems promote across users, clients, teams, or repositories with no versioning or lineage on the promoted artifact, so they sit outside the strict audit: EvolveNet, Evo2Team, XRepoSkill, ShareMem, RASO, Spotify's recommendation loop, and Future AGI's prompt lifecycle. Were the audit to ask for a promotion mechanism and no governance property at all, six of them would enter, all landing in [Auto × Auto] except Spotify's loop in [Auto × Human], without opening a new cell; the seventh, Future AGI's prompt lifecycle, rests on T5 evidence alone and is not admitted under any rule (Appendix C).

[**None** × **Auto**] holds the governance-first prompts-as-code systems: Braintrust, LangSmith Hub, and Vellum. Their Loop 2 is real, since evaluation-gated CI blocks a promotion on regression. They have no Loop 1 in the patch-local sense, because engineers author the candidate artifacts at their desks and session feedback never triggers one, and the asterisks record exactly that: governance without inner-loop adaptation. The cell is informative because it shows what survives when only the outer loop is automated. ExpeL, starred and also marked weak, sits here for a different reason: its experience bank is built from training tasks, which blurs the boundary of its loop, so it fails the audit (Appendix E.1.2).

[**Human** × **Auto**] is empty in both passes: no surveyed system puts a person on the local gate of a failure-driven loop and then promotes the result by program. A plausible occupant would ship person-approved skills through an automatic held-out evaluation; no published system does.

[**Human** × **Human**] is empty too. The registries that a first reading puts here (Azure API Center, the agentic-community MCP Gateway and Registry, TrueFoundry, Buildrix) have no failure-driven loop: a person registers an artifact and a person, or an admission scan, approves its use, which is a promotion gate on a registrant-supplied artifact. They sit in [**None** × **Human**] with the Anthropic skills repository and Agentic Configuration Management, whose lifecycle states a governance role assigns, and the five public skill repositories of the field study (Section 7) show the same shape in practice. Enoch's checkpoint-and-validate protocol moves one agent's state between hosts and is neither loop. The released table records these re-codings (Appendix G.7).

[**Human** × **None**] is the curation cell, also filled only from the second pass and single-loop by construction. In VerificAgent an agent drafts, a person fact-checks, and the result is frozen; in SkillChain a judge filters the candidate edits and a person decides, at a reflection gate, which refined skills are deployed; in the Davis et al. case study one engineer both writes and accepts the governance mechanisms that failures call for (Section 7). Curation stays patch-local, rather than collapsing into an engineer editing a configuration file, because the candidate answers a failure and a person decides whether it stays.

The occupants, with the variable the matrix does not show:

Cell	Systems	Candidate source	Loop 1 gate	Promotion mechanism	Loop 2 gate
[Auto × Auto]	NanoResearch, SkillOpt, SkillRL, AlphaEvolve, ADAS, DGM, Voyager, SAGE, EvolveR, CoEvoSkills, AutoAgent, AutoManual	agent proposals from trajectories or failures	evaluator, verifier, or threshold	merge into a shared bank or archive	the same evaluator, on the shared bank
[Auto × Human]	SkillForge, CASCADE, Sierra Agent OS 2.0, Cognition Devin	agent proposals	calibrated judge or outcome threshold	release to the team or product	reviewer, support engineer, or merge approval
[Auto × None]	Letta Context Repositories; Mem0 Platform, TOKI, Learning on the Job	agent (Letta: sleep-time subagents)	automatic write gate	Letta: unconditional git push to shared memory; the others: none	none
[None × Auto]	Braintrust, LangSmith Hub, Vellum, Governed Capability Evolution; EDDOps, with a sign-off after the threshold	engineers	none (no failure-driven loop)	evaluation-gated release	evaluation threshold
[None × Human]	Azure API Center, MCP Gateway and Registry, TrueFoundry, Buildrix, Anthropic skills repository, Agentic Configuration Management	registrants, contributors, engineers	none	admission to a registry, merge into a shared repository, or a lifecycle state	approval by a person (the MCP Gateway adds an automated admission scan)
[Human × None]	VerificAgent, SkillChain, Davis et al.	agent drafts (VerificAgent; SkillChain after a judge); the engineer (Davis et al.)	a person	none	none

E.2.2 Closest approaches to governed cross-tenant scaffold optimization

Section 9 defines the four properties: (a) an auto-gated Loop 1, (b) cross-organization promotion, (c) versioned lineage with access control, and (d) rollback. The closest approaches split them among

themselves:

- **AlphaEvolve** has (a) and a partial (b) confined to one organization; it has no mechanism for promotion across organizations.
- **NanoResearch** has (a) at the research level, and its paper describes no enterprise governance layer.
- **SkillOpt** has (a) fully and a partial (c), through its exported `best_skill.md` artifact and protected slow-update field. Its reported positive transfer across models, harnesses, and benchmarks is an empirical measurement of evidence radius. It implements neither (b) nor explicit rollback (d). The walkthrough in Appendix F maps each of its design decisions onto the loop.
- **CORAL** [42] has a partial (a) through reflection-written shared memory and an automatic Loop 2 through heartbeat-based promotion; it has no (b), no formal access control for (c), and no (d). It is the only system classified before the second pass whose Loop 2 runs asynchronously across multiple agents inside one organization, through shared persistent memory. Evo2Team and ShareMem, second-pass systems, promote across teams and across users without versioning.
- **PACE** [35] has (a) at small-model production timescales and a fast/slow two-timescale pair of loops that consolidates across one deployment's own episodes, not across deployments; it has no (b) and no formal (c) or (d). It is the closest demonstration that the two-loop architecture runs under small-frozen-model constraints.
- **Meta-Harness** [115] has (a) over the harness (substrate S5), with a filesystem of candidates that records credit assignment and serves as the accumulated history H_t of Appendix A. It has no cross-organization promotion, no access control, and no rollback.
- **SkillForge** has (a), a partial (b) that is cross-team within one enterprise, and a partial (c) through VFS commits; (d) is undocumented in the public report.
- **Sierra Agent OS 2.0** has a partial (a). Its cross-organization promotion is claimed in vendor documentation only and not counted as evidence under the tier rule of Section 5, and its GitHub-style Workspaces give some lineage but not access control with rollback in the strict sense.
- **FederatedSkill** [126] has (a) and, with EvolveNet (Table 7), the strongest partial (b) among the surveyed systems: heterogeneous clients, each a backbone-plus-harness profile under one federation operator, exchange semantic skill diffs, and a server-side evolution agent applies, adapts, or rejects each patch per client. An auditable decision log, recording source patch, reward, and justification for each modified path, gives a partial (c); there is no access control and no (d).
- **SkillClaw** [153] has a partial (a), through a nightly validation step, and a partial (b) across the users of one shared, synchronized skill repository; it reports no (c) or (d).
- **ChronoMem** [139] has (d) and the versioning half of (c) for the memory substrate alone: whole-memory snapshots, version histories, and natural-language rollback inside Google's Agent Development Kit. It has no Loop 1, no (b), and no access control.
- **Vellum, Braintrust, and LangSmith Hub** have (c) and (d) explicitly but neither (a) nor (b) in the strict sense; they are governance-first by construction.

No surveyed system combines (b), (c), and (d) under any gate type. The cross-organization governance triad is unfilled on its own, whether or not Loop 1 is auto-gated.

E.2.3 Whether the missing architecture should be built

Whether the missing architecture should be built depends on the feedback-signal class, not on the application domain: safety-critical domains such as medical decision support, legal advice, and financial advice, where a reviewer's gate on promotion is a design feature rather than a defect, are judged or delayed domains with high stakes. Section 9 gives the argument class by class.

F The convergent research families, system by system

This appendix walks the convergent research families of Section 8 system by system, at the level of mechanism, so that each conclusion of that section can be checked against what the systems do. The walkthroughs were written before the second pass and updated where the audit of Appendix G revised a claim; the members the second pass adds to each family appear in Table 6 and in Appendix I.

F.1 The convergent landscape

SkillOpt [31] is the clearest instance of artifact-layer descent we found, and it is one of at least eighteen independent research threads from 2025–2026 that converge on the same disciplined, gradient-like update for the scaffold layer. They converge from different starting points: GEPA from prompt optimization and Tiwari et al. from continual learning; RCL from cognitive-agent learning; Meta-Harness, AHE, MOSS, CANTANTE, and Ong et al. from harness engineering; MOCHA and Ratchet from skill-deployment constraints; ICT, Combee, CORAL, and PACE from cross-task, multi-agent, and production deployment. Each instantiates a different subset of the five conditions of Section 3, and each exposes a different gap.

By substrate (Section 5), GEPA and RCL work on the prompt and context layer (S1–S2); the harness family on orchestration (S5), with MOSS reaching into source code (S2+S5); MOCHA and Ratchet on skills under deployment constraints (S2); the transfer family on integrated memory-plus-orchestration stacks (S3+S5); and SkillOpt on an integrated stack led by S2.

No member of any family combines the four properties of Section 9: (a) auto-gated Loop 1, (b) cross-organizational Loop 2, (c) versioned lineage with role-based access control (RBAC), and (d) rollback. The loop of Section 3 is realized in fragments across the families, and every member leaves the Section 9 gap open.

F.2 SkillOpt: full correspondence

The table maps each component of SkillOpt to its quantity in Section 3, with the role that quantity plays and a pointer into the published procedure (Algorithm 1 of the SkillOpt paper).

Section 3 quantity	SkillOpt instantiation	Section 3 role explained	Algorithm pointer
θ_s (scaffold)	`best_skill.md` — a 379–1,995-token markdown document	the trainable external state of the frozen target model; persistent and exported as the deployment artifact	s_{best} (Algorithm 1, lines 1, 22, 37)
L_D (patch loss)	benchmark-native hard score / exact-match on a held-out test split	scalar quality signal that the optimizer is implicitly minimizing	EVALUATE(M, h, s, D) (lines 2, 17, 37)
Q_t (proposal)	optimizer-model reflection over failure and success minibatches; failures and successes analyzed separately, then merged with priority on failure fixes	proposal distribution conditioned on accumulated rollout evidence and current scaffold	optimizer calls O for failure analysis, success analysis, and merge (lines 9–11)
z_t (edit)	structured add/delete/replace patch on the skill (or rewrite-from-suggestions in rewrite mode)	the candidate update applied via $\oplus$ to the current scaffold	rank/clip to η_t edits, apply to obtain $\tilde{s}$ (lines 12–13)
η_t (textual learning rate)	bounded per-step edit budget; constant / linear / cosine / autonomous schedules; default cosine with floor $\eta_t = 2$	edit-magnitude controller; the step size of Section 3 in textual form	rank/clip step (line 12); schedule applied per epoch
G_t (gate)	held-out selection-split accept gate, strictly greater than current selection score	the per-edit gate G_t — accepts only when $\widehat{\Delta}_{z_t} L_D$ improves the selection score	comparison at line 20; cache lookup at lines 14–18
rejected-edit buffer	epoch-local store of failed proposals and the score drop they caused	converts rejected steps into negative feedback for later proposals within the same epoch; zero deployment cost	$\mathcal{B}$ updated at line 26; reset at line 5
slow/meta update	epoch-end longitudinal guidance written to a markup-fenced protected slow-update field, still passed through D_{sel}	second-order consolidation across adjacent epochs; step-level edits cannot overwrite the protected slow-update field	lines 28–31; gate applied to injected guidance
optimizer-side meta skill	m_{meta} — a prompt-side summary of which edits helped, which failed, and which failures persisted	training-time guidance for Q_t; never shipped with the deployed artifact	lines 33–34 (training-side only)

The mapping is one-to-one at the level of Algorithm 1: every Section 3 quantity has a named, code-level counterpart, and every SkillOpt control surface maps onto a Section 3 quantity. SkillOpt is therefore the clearest instance we found in which the loop of Section 3 is *mechanism* rather than *metaphor*: all five conditions are instantiated.

F.3 What SkillOpt does that prior textual-optimization systems do not

SkillOpt is compared with six baselines under one protocol (same target model, held-out test split, and scorer), and each baseline lacks a specific piece of the loop.

- **TextGrad** [22] backpropagates textual feedback through a compound system's variables but leaves no artifact distinct from optimizer state, which is why Appendix E.1 places it outside the core corpus; it has no step bound, held-out gate, or rejected-edit buffer. SkillOpt leads it on direct chat by +5.9 to +34.0 points across benchmarks at GPT-5.5 scale, and at all seven target-model scales.
- **GEPA** [113] evolves prompts by reflection under Pareto-style selection. Its target is a prompt rather than a persistent skill, its selection is not a strictly-greater held-out gate, and nothing bounds an accepted edit, so one edit can rewrite the prompt wholesale. SkillOpt matches or beats it in every reported cell of its main results table [31].
- **Trace2Skill** [154] distils trajectory-local lessons into reusable skills offline. No held-out loop accepts or rejects each candidate, so delta validation (operation 4 of Appendix D.1) and the rejected-edit buffer are missing. SkillOpt leads it by +5.0 to +19.0 points at GPT-5.5 direct chat.
- **EvoSkill** [155], the strongest baseline inside the Codex and Claude Code harnesses, evolves a skill folder from failure analysis. Its credit assignment is genuine, but per-step folder edits are unbounded and rejected edits are not kept. SkillOpt leads it by +14.0 inside Codex and +3.2 inside Claude Code on the GPT-5.5 five-benchmark average.
- **A human-written skill** (145–516 tokens per benchmark) persists but never updates: the no-Loop-1 ceiling.
- **A one-shot LLM skill**, generated once by GPT-5.5 from the benchmark description, never updates and proposes without rollout evidence: the untrained-scaffold baseline.

ABSTRAL [156] and EvoTest [157], cited by SkillOpt but not measured, work at the multi-agent-design and test-time-configuration layer and fall outside the comparison.

The live difference across the six is four controls: a persistent edit target distinct from optimizer state (θ_s); a held-out gate on every acceptance (G_t, with estimator $\widehat{\Delta}$ and threshold τ_t); rejected edits kept as negative feedback for Q_t; and a bounded, scheduled edit budget (η_t). SkillOpt is the first system we found that has all four, run as one offline training loop whose deployed artifact, a 379–1,995-token markdown file, needs no optimizer calls at inference.

Its follow-up, SkillOpt-Lite [112], asks which controls are necessary. Recasting skill optimization as zeroth-order search, it keeps three principles: trajectory exploration through the file system, consensus attribute mining (attributes common across heterogeneous rollouts, single-sample quirks discarded), and a held-out gate strictly disjoint from the training failures. It drops minibatch reflection pooling, the rejected-edit buffer, the slow update, and the learning-rate schedule. It still converges faster and scores higher than full SkillOpt (LiveMath +8.8 points at GPT-5.5, +25.4 at GPT-5.4-nano). The gate and a consensus form of credit assignment are load-bearing; negative-feedback retention and the bounded step are stabilizers whose value depends on the benchmark. None of this touches the governance properties of Section 9.

F.4 GEPA, Tiwari, and the reflective-evolution family

The immediate predecessor to SkillOpt's optimizer framing is **GEPA** [113]. It shows that natural-language reflection over execution trajectories can provide a richer adaptation signal than sparse scalar rewards for prompt evolution in compound AI systems. GEPA samples trajectories containing reasoning steps, tool calls, and tool outputs; reflects on them to diagnose failures and propose prompt updates; and maintains a Pareto frontier of candidate prompts rather than greedily mutating only the current best. Across six tasks, GEPA outperforms GRPO by 6% on average and by up to 20% while using up to 35× fewer rollouts, and also outperforms MIPROv2 in the reported aggregate comparison. In Section 3 vocabulary, GEPA strongly instantiates the proposal side, Q_t, through trajectory-conditioned reflection, and uses Pareto-based candidate selection as a proposal and pruning mechanism. What it does not instantiate is the full SkillOpt-style artifact lifecycle: a bounded edit budget on a reusable skill document, a strictly-greater held-out gate on each accepted skill update, a rejected-edit buffer used as negative feedback, and an exported deployment artifact with governance hooks.

Tiwari et al. [4] supply the broader framing from the optimization side, arguing that restricting learning to *either* in-context adaptation (prompt optimization) *or* in-weights adaptation (RL or fine-tuning) is unnecessary and costly. In-context adaptation is cheap and rapid but capacity-bounded; in-weights adaptation is expressive but induces catastrophic forgetting and loss of plasticity [5]. The scaffold layer of Section 2 is the substrate that absorbs what in-context adaptation can carry cheaply, leaving the weights free to generalize. Tiwari et al. reach the two-timescale motivation from optimization and Section 2 reaches it from production reliability; the two framings are independent and reinforce each other.

F.5 Reflective Context Learning: a parallel formalization

The most direct independent convergence on the Section 3 formalization comes from **Reflective Context Learning** [RCL; [114]]. Their opening observation is that "the fundamental problems of learning, including credit assignment, overfitting, forgetting, local optima, and high-variance learning signals, persist whether the learned object lies in parameter space or context space," and that current methods are "fragmented and ad hoc" precisely because they address these challenges one at a time rather than as a unified optimization problem over the context. The parallel to Section 3's argument, that scaffold maintenance needs the discipline of weight-space optimization applied to an external artifact layer, is exact and was arrived at independently.

RCL maps onto Section 3 as follows. The *learned context* is θ_s. The *reflection signal* is a form of credit assignment identifying which context coordinate caused a failure (Appendix D.1, operation 2). The *generalization penalty* that RCL introduces to prevent narrow sample-specific rules is structurally the complexity penalty $\Omega(z)$ of the gate condition. Where RCL and Section 3 differ is in Loop-2 scope: RCL is framed around single-agent cross-task generalization, while Section 3's Loop 2 is cross-deployment, cross-tenant, and cross-organizational promotion. The convergence is nonetheless significant: two research threads starting from different motivations (production reliability in Section 3, cognitive agent learning in RCL) arrive at the same formal observation that the classical learning problems do not disappear when the learned object moves outside the weights.

F.6 The harness-optimization cluster

A second family optimizes the *harness*, the code and configuration that wrap a frozen model, rather than skill documents. It names credit assignment as the central unsolved problem of scaffold

evolution, and its solutions instantiate two operations of the mechanism stack in Appendix D.1: mode discovery with credit assignment (operation 2) and delta validation (operation 4).

Meta-Harness [115] starts from the observation that system performance "depends not only on model weights, but also on their harness", the scaffold-versus-weights decomposition stated for this substrate. Existing text optimizers, it finds, "compress feedback too aggressively" for credit assignment, so it gives an outer-loop proposer filesystem access to the source code, scores, and execution traces of every prior candidate. The edit history becomes the credit-assignment signal. In Section 3 terms the filesystem is the accumulated history H_t that conditions the proposal Q_t, and the outer loop closes Loop 1 over orchestration (S5) rather than skills (S2).

Agentic Harness Engineering [AHE; [116]] comes at the problem from observability. It names three obstacles: a heterogeneous action space across editable components, voluminous trajectories that bury the actionable signal, and edits whose effect is hard to attribute. They are, in order, coordinate identification (operation 2), residual capture (operation 1), and the credit-assignment half of operation 2. AHE's answer, three observability pillars that give every editable harness component a file-level representation, builds the coordinate decomposition that Section 3 assumes. Descent needs that instrumentation as much as it needs the loop.

MOSS [117] observes that earlier self-evolving systems "confine evolution to text-mutable artifacts" and leave the harness untouched. Routing, hook ordering, state invariants, and dispatch live in code, so a whole class of structural failure is, in MOSS's phrase, "physically unreachable from the text layer". MOSS therefore extends Loop 1 to rewriting source code, evolving skills (S2) and orchestration (S5) together. The consequence for Section 9: a governed optimizer that covers only text artifacts leaves the code substrate ungoverned.

CANTANTE [118] takes up credit assignment in multi-agent scaffolds, where "scores are available only at the system level, whereas the parameters governing agent behavior are local". It contrasts rollouts of several jointly acting configurations to split a system-level reward into per-agent update signals. That gives the credit-assignment step a concrete mechanism when the scaffold is spread across agents (S5): the attribution, succeeding with non-zero probability per failure, that the correspondence note of Appendix A requires.

Harness-optimizer evaluation [43] finds that judging scaffold optimizers "solely by observing target agents' performance gains" misses intermediate erroneous edits that hinder performance, and that it is "unclear whether harness optimization is driven by optimizers' informed update actions or simply trial-and-error". That is the line between Stage 2 and Stage 3 of Section 3, measured. An optimizer that cannot tell informed updates from lucky ones operates in Stage 2, whatever its benchmark gains.

HARNESSEVO [119] turns the diagnosis into a measurement. It splits the harness into four separately evolvable slots (role, task strategy, tool and format rules, reflection and control) and attributes value by leave-one-in and leave-one-out evolution. On ALFWorld with a 7B target model nearly all of the gain sits in the reflection and control slot (leave-one-in +0.119; the other slots individually null). Spread evenly, a fixed 64-rollout budget leaves every slot below the optimizer's effective search floor, so each freezes at its seed; concentrated on the high-value slot, it recovers the gain (0.642 to 0.761). On WebShop the effect is absent. Credit assignment over coordinates has to precede budget allocation, and which coordinate carries the value depends on the task.

F.7 Multi-objective and lifecycle constraints

Descent on a single loss misses two kinds of constraint: the platform limits a deployed skill must satisfy, and the hygiene of the artifact pool from which the gate draws its signal.

MOCHA [122] observes that agent skills are "multi-field artifacts subject to hard platform constraints": descriptions are truncated for routing, instruction bodies are compacted by progressive disclosure, and co-resident skills compete for the context window. A skill must maximize task performance and satisfy those limits at once, and optimizing performance alone yields skills that pass benchmarks but degrade on real deployment infrastructure. MOCHA navigates the constraint frontier with *Chebyshev annealing*. It operationalizes the complexity penalty $\Omega(z)$ of the gate, and it shows that the scalar Ω of Section 3 simplifies what is, in production, a Pareto frontier.

Ratchet [121] places the bottleneck elsewhere. LLM-authored skills deliver +0.0 pp over no-skill baselines while human-curated ones deliver +16.2 pp, which points away from the optimizer, the mechanism that proposes and accepts edits, and toward *lifecycle management*: whether skills are written, retrieved, curated, and retired correctly. Four hygiene mechanisms (outcome-driven retirement, a bounded active cap, meta-skill authoring, and pattern canonicalization) lift the +0.0 pp floor substantially before any optimizer runs.

The finding clarifies a precondition of descent rather than refuting it. The gate needs a meaningful signal: if retrieval injects the wrong skill into every rollout, the proposal process conditions on corrupted evidence and descent degrades to random search. Ratchet's mechanisms keep residual capture and mode discovery (operations 1 and 2) working. Retirement keeps stale artifacts out of the residual signal, the active cap keeps scaffold bloat from making retrieval unreliable, and canonicalization collapses near-duplicates so that proposals condition on one clean artifact set. A system with strong hygiene and no optimizer sits at Stage 1 with a clean scaffold; one with a strong optimizer and poor hygiene sits at Stage 2 with a noisy loss. Stage 3 needs both.

SkillsVote [123] carries the lifecycle argument to an open ecosystem. It profiles a million-scale open-source skill corpus for environment requirements, quality, and verifiability, synthesizes tasks for the verifiable subset, and exposes skills through agentic library search before execution. After execution it splits each trajectory into skill-linked subtasks and attributes their outcomes to skill guidance, agent exploration, environment, or result signals; only successful, reusable discoveries pass into evidence-gated updates. Attribution thus happens at the admission gate rather than inside the optimizer, credit assignment used as a filter on residual capture. The redundancy and environment sensitivity SkillsVote measures are the bloat and drift of Section 10, arising here from many authors rather than one loop.

F.8 Cross-task transfer and evidence radius

Several papers measure what Section 3 calls the *evidence radius* of a scaffold artifact, how far an edit validated in one context keeps helping when redeployed elsewhere, without using the term.

In-context Training [ICT; [124]] tests cross-task self-improvement: whether experience distilled from one set of tasks transfers to unseen ones. Its finding that "current reflection-based methods can only self-correct within a single task instance" measures the radius at its lower bound, one task. ICT tries to widen it by distilling reflections into reusable lessons, which instruments Loop 2 empirically: does an edit that passed the gate on one task's L_D still help across a distribution of unseen tasks?

Combee [125] reaches the same measurement through scale. ACE [23] showed that a system prompt which accumulates and organizes strategies through agentic interaction improves on a static prompt for a single agent; Combee learns from the traces of many agents running in parallel. Each proposal then draws on a wider deployment distribution, and the promoted artifact is tested against a larger average of patch losses. This is not cross-organizational promotion. Combee measures how much of the evidence-radius benefit parallelism inside one organization captures.

CORAL [42] runs long-lived agents that explore, reflect, and collaborate through shared persistent memory, asynchronous execution, and heartbeat-based interventions. A finding one agent writes to the shared store becomes available to all, so the store is the cross-agent promoted artifact and the heartbeat a form of auto-gated Loop-2 promotion; there is no versioned lineage, RBAC, or rollback. CORAL is the only fully asynchronous multi-agent Loop-2 promotion mechanism among the systems classified before the second pass; in the full corpus, Evo2Team and ShareMem add cross-team and cross-user promotion without versioning (Appendix G.3).

PACE [35] is a production-oriented two-timescale design for small language-model agents. The fast timescale adapts prompts, parsers, validators, and control logic from failure signals within a deployment, as Loop 1 does; the slow timescale consolidates across the deployment's own episodes, a Loop 2 whose radius is one deployment. The two-loop structure holds on small, frozen production models, not only under frontier-model assumptions, and persists even when the optimizer and the target are the same frozen small model.

SkillClaw [153] and **FederatedSkill** [126] are the two systems of these walkthroughs whose Loop 2 crosses users rather than tasks or agents; in the full corpus, EvolveNet, ShareMem, XRepoSkill, and RASO also promote across users, clients, or repositories, each inside one operator (Section 6). SkillClaw aggregates many users' trajectories, turns recurring behavioral patterns into skill-set updates, and synchronizes the updated skills across users; its evaluation runs a day-night cycle in which simulated users generate trajectories by day and the system reflects, proposes, validates, and deploys by night. FederatedSkill exchanges semantic skill diffs, structured patches over a local library shipped with the task reward that produced them, so that clients share what they learned without sharing raw trajectories. A server-side evolution agent aggregates the diffs into a per-client capability matrix (covered, absorbing, broken, gap) and applies, adapts, or rejects each peer patch per client, logging source patch, reward, and justification for every modified path. Across 20 task families it reports up to +44.4% success and -37.5% compute against isolated self-evolution. FederatedSkill is the first surveyed system in which Loop 2 aggregates artifact updates rather than promoting a single artifact, and its decision log is the first lineage record at an aggregator. Its acceptance, though, uses the source client's reward rather than a held-out evaluation at the receiving client.

F.9 Synthesis: what the convergence establishes

Section 8 draws four conclusions from the families, and the walkthroughs support them as follows.

1. **The scaffold-versus-weights decomposition is independently confirmed.** GEPA, Meta-Harness, MOSS, RCL, and PACE all begin from the observation that model weights and the external artifact layer are distinct adaptation surfaces with different cost, reversibility, and scope properties. No surveyed paper disputes this decomposition; the variation is in which substrate each paper targets (skill document, harness code, system prompt, shared memory).
2. **Credit assignment is the central engineering bottleneck.** AHE, CANTANTE, Meta-Harness, the harness-optimizer evaluation of Ong et al. [43], HARNESSEVO, and SkillsVote

all identify credit assignment, attributing a system-level failure to a specific scaffold coordinate, as the problem that separates disciplined Loop-1 evolution from trial-and-error. It is operation 2 of Appendix D.1, the attribution that the correspondence note of Appendix A requires, and it remains unsolved in the general multi-substrate case.

3. **Lifecycle hygiene is a precondition for descent, not an alternative to it.** Ratchet establishes that the gradient-analog mechanism of Section 3 requires a maintained scaffold. A dirty artifact pool (stale, bloated, miscurated) keeps the gate G_t from receiving a meaningful signal. Hygiene and optimization are sequential dependencies, not competing paradigms.

4. **No surveyed system combines Loop-1 automation with cross-organizational Loop-2 governance.** SkillOpt has the strongest Loop-1 closure and the most systematic evidence-radius measurement, but it does not implement cross-organizational promotion, versioned lineage with RBAC, or rollback. CORAL implements asynchronous cross-agent promotion within one system but not across organizations. PACE implements two-timescale consolidation within one deployment. Meta-Harness implements Loop-1 closure over the harness substrate but no Loop-2 promotion mechanism. FederatedSkill federates skill patches across clients under one operator, with a decision log but no RBAC or rollback, and EvolveNet adds a held-out merge gate and round rollback without lineage or RBAC. The four-property gap of Section 9 is narrowed on (b) and (d) by different systems and closed on none, and the independent convergence of several research threads on the same single-organization ceiling confirms it.

F.10 SkillOpt against Section 9's four properties

Mapped onto Section 9's four properties of governed cross-tenant scaffold optimization:

- (a) **Auto-gated Loop 1.** *Satisfied.* The held-out selection-split gate at strictly-greater margin instantiates G_t.
- (b) **Multi-tenant cross-org promotion.** *Absent.* SkillOpt is a single-organization research system; there is no governed cross-tenant aggregator.
- (c) **Versioned lineage with RBAC.** *Partial.* The exported `best_skill.md` is a versioned artifact, and the protected slow-update field separates fast and slow consolidation. There is no documented RBAC layer or audit-trail policy, and `edit_apply_report.json` is an operational log rather than an access-controlled lineage system.
- (d) **Rollback.** *Absent.* The training loop selects the best validation-gated skill and exports it; no rollback envelope on deployed artifacts is reported.

The cross-model, cross-harness, and cross-benchmark transfer experiments reported by SkillOpt are *empirical evidence-radius measurements*: a skill optimized in one (model, benchmark, harness) cell continues to improve performance when redeployed in nearby cells without further optimization. In Section 3 vocabulary, this is a measured Loop-2 *radius* for the artifact, not a *governance mechanism* for promoting it. Both readings are useful: SkillOpt is the strongest Loop-1 proof and supplies the first systematic transfer evidence at the scaffold layer, while the Section 9 cross-tenant governance triad ((b), (c), (d)) remains unfilled.

F.11 Cost and bloat

SkillOpt reports 0.6M to 46.4M training tokens per absolute test-point gain (their Table 6), with the procedural benchmarks (SpreadsheetBench, OfficeQA, LiveMath) at the cheap end and search/visual

benchmarks (SearchQA, DocVQA) an order of magnitude more expensive per point. The deployed artifact remains 379–1,995 tokens regardless. This is, to our knowledge, the first published cost-per-point figure for scaffold-layer descent, and it complements the Section 10 scaffold-bloat discussion (and the Ratchet hygiene-precondition finding walked in Appendix F.7): the textual learning rate η_t with cosine decay keeps the deployed skill compact even when the optimizer explores many more edits than are accepted (only 1–4 edits survive into the final skill across the six benchmarks). The mechanism that prevents the deployed artifact from bloating is exactly the complexity penalty $\Omega(z)$ of Section 3, instantiated as a hard cap on accepted edits per step rather than a soft regularizer.

G Audit of the first-pass draft's absolute claims against the full corpus

The first-pass draft of this survey, written on the 130 systems of the first pass, made absolute claims of the form *no surveyed system*, *only X*, *first*, *absent*, and *empty cell*; each row below sets one of them against the full corpus of 256 systems and records what that corpus shows. A claim **holds** when the full corpus bears it out as worded, is **sharpened** when it is true and the corpus supports a more precise statement, is **revised** when the corpus contradicts its wording, and is **dropped** when the article no longer makes it; for sharpened and revised claims the Evidence column also states the claim as the full corpus supports it. Evidence names the systems that decide each verdict, and parenthetical numbers such as 2608.04968 are arXiv identifiers.

G.1 The four properties of Section 9

Claim	Verdict	Evidence
No public system combines (a) auto-gated Loop 1, (b) cross-organization promotion, (c) versioned lineage with role-based access control (RBAC), (d) rollback	**holds**	Across the 126 second-pass systems: no system demonstrates promotion across organizations, and EvolveNet and RASO have a partial (b) inside one operator (EvolveNet's clients are benchmark partitions, "organizations" only in motivation; RASO draws on a public corpus as its source); MemClaw is single-tenant in its evaluation, and Buildrix is an open community hub with no gate; (c) appears as lineage without RBAC everywhere (EDDOps, Agentic Configuration Management (ACM), MemClaw, Letta, Mem0 history) and as RBAC without versioned promotion in the MCP control planes (agentic-community gateway, TrueFoundry); (d) is demonstrated by Governed Capability Evolution and EDDOps and claimed by EvolveNet (round rollback) and EvoUndo (recoverability as admission), none of which has (a)+(b).

Claim	Verdict	Evidence
"No surveyed system meets the FedAvg limit; FederatedSkill comes closest"	**sharpened**	EvolveNet (2608.04968) comes nearer: it broadcasts a shared harness, optimizes it locally at each client, merges on the server only when the merge fixes at least as many held-out items as it breaks, and rolls back to the previous shared harness after a failed round. It is still selection and merge over program edits, not averaging, so neither it nor FederatedSkill meets the FedAvg limit.
Fourth conclusion on the research families: "FederatedSkill comes closest ... narrowed on one property, closed on none"	**revised**	The two closest approaches are FederatedSkill (decision log; no gate at the receiving client) and EvolveNet (held-out merge gate and round rollback; no lineage or RBAC). The gap is narrowed on (b) and (d) by different systems and closed on none (Table 7).
"no member of the Section 8 research families provides (d), and only FederatedSkill provides even a partial (b)"	**revised**	(d) appears in EvolveNet (round rollback), SkillSmith and SkillAudit (edit-level rollback inside one agent), Growing Harness and WDA (rollback within an optimization run), and EvoUndo (recoverability as an admission test); EvolveNet and RASO provide a partial (b). Rollback appears inside single deployments and inside one federated round; no system rolls back a promotion that has crossed an organization boundary, because none makes such a promotion.

G.2 Matrix cells

Claim	Verdict	Evidence
"Three cells are populated; the rest are empty for structural reasons"	**revised**	Single-loop and governance-only systems occupy these cells, under the definition of Section 3 that a gate records who decides and not who supplies the candidate: [None × Human] (MCP registries: Azure API Center, agentic-community MCP Gateway & Registry, TrueFoundry; Buildrix; the Anthropic skills repository; ACM), [Human × None] (VerificAgent, SkillChain's second-stage reflection gate, Davis et al. governance conversion), [Auto × None] (Letta's ungated push; Mem0 Platform, TOKI, Learning on the Job), [None × Auto] (Governed Capability Evolution; EDDOps, with a sign-off after the threshold), [None × None] (Official MCP Registry; Enoch's migration protocol; Continual Harness's unchecked self-edit; the revocation guard). Two cells hold every audit passer whose promotion has a gate, a third holds the governance-first platforms, Letta passes with no promotion gate, and the human-gated row holds only curation (Section 7).
"[Human × Auto] is empty by engineering logic . . . no published example exists"	**holds**	No system in the full corpus puts a person on the local gate of a failure-driven loop and promotes by program. MCP Gateway & Registry adds an automated security scan to a human admission, but it has no failure-driven local loop, so it sits in [None × Human].
"[Human × Human] is empty in the surveyed corpus"	**holds, after re-coding**	The registries and the community hub that a first reading places here have no failure-driven loop and sit in [None × Human] (first row). The occupant the first-pass draft predicted, solo or small-team workflows, is documented by "Who Maintains Agent Skills?" (2609.05677): across five public skill repositories every substantive edit is human-authored and human-merged, 62% are AI co-authored, and there is 1 deprecation in 254 edits; that shape is [None × Human] too, because no loop fires from session signal (Section 7).
"[Human × None] is structurally improbable"	**revised**	VerificAgent (agent-drafted memories, human fact-check, frozen bank), SkillChain's second-stage human reflection gate, and the Davis et al. governance conversion occupy the cell. It is the human-curation cell, single-loop by construction.

Claim	Verdict	Evidence
"Nearly every auto-gated system fires at n = 1" (accepts an edit on a single evaluation)	**revised**	Batched or statistical acceptance is common, and has been since mid-2026: SAGE-gate (paired per-item test), Future AGI (paired bootstrap CI), AutoSaddler (mini-batch failure signals), MGM, ModularRSI (evidence across tasks), VaG (3-of-3 critics), EvolveNet (fixed $\geq$ broken on held-out), VACE, HackProbe, PROCTOR. Evidence is batched before acceptance, with SAGE-gate and VaG as instances (Section 7).
"no surveyed system has formalized agreement as a k-of-n acceptance rule"	**revised**	VaG requires all three heterogeneous critics to pass (3-of-3), and SkillSpec is consensus-gated, so agreement-based acceptance exists. The distinction between agreement and averaging stands: under agreement the false-accept rate decays as p^k, under averaging the variance decays as $1/k$ (Appendix B).
"Only AlphaEvolve demonstrates explicit multi-layer depth"	**sharpened**	HarnessBank, JIT-Agent, and MGM keep archives for lineage, and VaG adds tiered subset selection; none pairs a cascade evaluator with a diversity archive. AlphaEvolve is the only system that combines the two.

G.3 "Only" and "first"

Claim	Verdict	Evidence
CORAL "the only asynchronous multi-agent Loop-2 mechanism in the corpus"	**revised**	Evo2Team (2609.34135) promotes routing and communication skills across teams of 4–32 agents with target-side confirmation, and ShareMem pools memory across users. CORAL is the only such mechanism among the systems classified before the second pass; in the full corpus Evo2Team and ShareMem add cross-team and cross-user promotion without versioning (Appendix F.8).
AlphaEvolve Strassen claim	**holds** (attributed)	The claim is AlphaEvolve's reported improvement on Strassen's algorithm for 4×4 matrix multiplication. The result is taken from the AlphaEvolve report and was not independently measured; with that attribution in the same sentence, the claim holds.
"CASCADE reports the largest within-benchmark ablation gain in the corpus"	**dropped**	Not established for the 126 second-pass systems, and ablation gains measured on different benchmarks are not comparable. The article does not make the claim.

Claim	Verdict	Evidence
ChronoMem and rollback repair "the first surveyed instances of the rollback envelope"	**revised**	Memory rollback also appears in VerMem (soft-delete and restore operations), Mi-Memory (as a design), and Letta (git history), and checkpoint-rollback studies (2608.29381, 2608.15939) show rollback unsafe. ChronoMem and rollback repair are among the first instances, and the checkpoint-rollback studies are the counter-evidence.
Cuadros "only surveyed memory system with full provenance + versioning + correction"	**holds**	The memory-governance systems of the second pass (TOKI, PPMF, TMA-NM, MemClaw, He & Yu) each carry provenance or lineage without correction loops or versioning, and the auto-gated memory loops (Recuris, MemGuard, VerMem, Compiled Memory) carry no lineage.
SkillOpt "cleanest published instance" / "first to instantiate all four simultaneously"	**sharpened**	SkillOpt-Lite (2607.03451) removes two of the four controls and matches or beats SkillOpt; VALVE, SkillSmith, and SkillAudit instantiate held-out gates with rollback. SkillOpt is the first system we found that instantiates all four, and its own follow-up shows that two of them are stabilizers rather than necessities (Appendix F.3).
FederatedSkill "first … federated aggregation … first lineage record at the aggregator"	**holds**	EvolveNet (August 2026) postdates it and keeps no lineage record.
SkillOpt "first published cost-per-point figure"	**holds**	The full corpus contains no earlier figure.

G.4 Absences

Claim	Verdict	Evidence
"production deployments sit at Stage 1 or 2; none reaches Stage 3"	**sharpened**	Spotify (2609.30297) runs automated prompt and planning optimization gated by human launch and A/B metrics; SkillForge passes each skill diff through a judge that agrees with engineers more than nine times in ten, and support engineers decide whether it ships. Both are Stage 3 on Loop 1 with a human Loop 2: production deployments reach Stage 3 on the local loop only behind a human promotion gate (Section 3).
No S1 system scores 5; "no instruction system closes both loops"	**holds**	On S1, n=15 and the maximum score is 4. Adoption studies point the same way: context files raise cost by more than 20% without raising success (Gloaguen), and context files dominate configuration while skills are rare (Galster).

Claim	Verdict	Evidence
"Anthropic's skills repository is one-pool with PR review but no automated eval gate"	**holds**	The study of public skill repositories (2609.05677) observes human merges and no evaluation gate.
"correctable, versioned memory is rare; production memory lacks failure-triggered correction and versioned lineage"	**revised**	Letta Context Repositories (git-backed memory, every edit committed, shared memory auto-pushed), Mem0's per-memory history and feedback endpoint, and Zep's bitemporal invalidation make lineage common in production memory. Failure-triggered correction is not common, and 2609.08258 shows revoked facts still outranking their replacements in five backends. Production memory is versioned, but correction is not enforced (Section 6).
No S5 system scores 5; "topology evolution is rare"	**sharpened**	With MAE and Memento-Skills corrected to 4 (Appendix G.7), S5 has no score of 5 in either pass. Gated topology evolution exists: TOCOMAS keeps only candidates that satisfy structural invariants; Evo2Team promotes across teams; OptiMAS, AgentGrad, and HarnessForge update S5 under evaluation. Topology evolution is rare and, where it exists, gated by evaluation alone except in TOCOMAS.
"no system matures all six at once"	**holds**	None of the 126 second-pass systems scores 5 on any substrate.
"Failure-triggered memory is rare in production"	**holds**	One field report exists (SIx Harness, N=1, self-reported); Letta's dreaming writes from conversation history, not from failures.
"Cross-customer skill promotion is nearly absent"	**holds**	ShareMem, EvolveNet, XRepoSkill, and RASO promote across users, clients, or repositories inside one operator; none promotes across customers or organizations with evidence.
"Memory versioning is mostly unsolved outside Cuadros et al. and ChronoMem"	**revised**	The evidence is the same as for the claim that production memory lacks versioned lineage (fourth row of this table): Letta, Mem0 history, Zep bitemporal facts, TOKI audit rows, and MemClaw supersession chains all version memory. Memory versioning is common, not confined to Cuadros et al. and ChronoMem.
"Scaffold-level curriculum is unspecified in nearly every surveyed system"	**sharpened**	Task-CoEvolve (variance-weighted task sampling) and AutoSaddler (mini-batch failure signals) specify it; they remain the exception.

Claim	Verdict	Evidence
"Held-out evals at Loop-2 promotion are absent"	**revised**	Receiving-side evaluation exists: ShareMem (evaluation tasks belong to a different user), XRepoSkill (independent verification in $\geq$2 repositories), Evo2Team (target-side confirmation; 4 of 22 transfers still fail on new graphs), VACE (candidate evaluated with the receiving model fixed). EvolveNet gates the merge on a server-side held-out slice but runs no check at the receiving client, and Federated Agent Optimization (2610.01195) proposes a receiving-client gate without implementing it. Receiving-side evaluation at promotion is rare, not absent (Section 6).
Research full-auto cluster: "enterprise governance is uniformly absent"	**holds**	No second-pass [Auto $\times$ Auto] system reports RBAC or audited promotion.
Governance-first cluster: "no patch-local inner loop"	**holds**	EDDOps, ACM, Governed Capability Evolution, and the MCP control planes all lack an inner loop.
"AGENTS.md in over 60,000 repositories"	**revised**	Adoption is measured directly rather than by one repository count: a 2,853-repository study (Galster), a 128,018-project adoption study (Robbes: 22–29% of coding-agent projects), and 36,710 repositories mined for loop engineering (Lulla, 2608.21884), alongside a controlled study in which context files raise cost without raising success (Gloaguen; Section 6).
"no surveyed system treats pruning or expiration as a first-class gated operation (outside Ratchet, SkillOpt, ChronoMem)"	**revised**	Retirement is a gated operation in at least eleven skill systems (Library Drift: retire when n $\geq$ 100 and contribution $\leq$ -0.10 with a cap of 50; SkillSmith; SkillOps; ASSAY; HiSME; SkeMex; MUSE-Autoskill; SkillMAS; SafeEvolve; Double Ratchet; SkillGen), and expiry is a product control in Mem0, Zep, and LangGraph Platform. Pruning and expiry are first-class operations; what remains rare is pruning of a promoted artifact's descendants, which VaG shows is structurally irreversible.
The failure surface is "largely unmitigated in production"	**holds**	Revocation is unenforced in five production backends (2609.08258), checkpoint rollback is unsafe in three frameworks (2608.29381), and 16 deployed frameworks keep no evidentiary record (Hearsay, 2609.32495).
Multi-sample agreement is "under-explored; SkillOpt-Lite nearest"	**revised**	The evidence is the same as for the two acceptance-rule claims of Appendix G.2, single-evaluation acceptance and k-of-n agreement: SAGE-gate, VaG, and SkillSpec. Agreement-based acceptance is instantiated, not under-explored (Appendix B).

Claim	Verdict	Evidence
"No surveyed paper disputes the scaffold-vs-weights decomposition"	**holds**	SIA (2605.27276) updates harness and weights together but keeps them as distinct levers.
Credit assignment "remains unsolved in the general multi-substrate case"	**holds**	HARNESSEVO (slot attribution), AgentGrad (intervention-based), and CANTANTE (contrastive) each solve one substrate's case.

G.5 The three predictions of Section 4

Section 4 predicts from the signal classes which cells of the design space each class can reach, and Section 7 confirms the predictions; the rows give the systems behind each verdict, in the order Section 4 states them.

Claim	Statement	Verdict	Evidence
First	An [Auto × Auto] system whose gate estimates the loss change per edit should run on a C1 signal or a simulated C2 environment; one that works on a C4 signal should prune proposals rather than estimate	**holds**	Before the second pass: DGM, Voyager, AlphaEvolve, SkillOpt, SkillRL, SAGE, EvolveR, CoEvoSkills, AutoAgent, AutoManual (C1 or simulated C2); NanoResearch (semantic merge, C4) and CORAL (heartbeat, C4) prune. Second pass: EvolveNet, Evo2Team, XRepoSkill (C1 code or benchmark signals); ShareMem (C2, Mind2Web); RASO (target execution feedback, C1/C2); Letta (no gate).
Second	[Auto × Human] should hold only C2 and C3 systems, because the human on Loop 2 substitutes for the exact signal that Loop 1 lacks	**holds**	SkillForge (C3 judge, human Loop 2), CASCADE (C3), Sierra (C2/C3), Devin (C2 merge outcome), Spotify (C2 A/B behind human launch). No C1 system occupies the cell.
Third	No scaffold system should yet calibrate a judged gate under the reflexive condition	**sharpened**	Three scaffold systems report calibration: SkillForge (>90% agreement), SkillChain (κ 0.72–0.91), and SkillCoach (trajectory-filter consistency 82→96% against human gold); none calibrates a judge that is itself being edited. On the weight side, no method that co-evolves a rubric or judge reports agreement with people (Appendix G.6).

G.6 The rubric-reward comparison

Reinforcement learning from verifiable rewards trains weights on exact signals, and extending it to writing and other open-ended tasks has produced the judged gates that a C3 scaffold loop can borrow. These are weight-side methods and the corpus does not score them; Section 4 reads one distinction from them, whether the rubric or judge is held fixed while the policy learns or moves with it. The comparison covers fourteen papers, thirteen methods and one survey. Of the methods,

eight hold the rubric or judge fixed during policy optimization and five co-evolve it. Five report a human-agreement figure. Four of them hold the gate fixed, and the fifth, RLPR, co-evolves a probability reward rather than a rubric, checked against human labels. No method that co-evolves a rubric or judge reports one.

Gate during policy optimization	Papers	Members	Agreement with human judgment
Held fixed	8 methods	Rubrics as Rewards, an instance-specific checklist scored by a fixed judge [46]; Writing-Zero, a pairwise judge trained on writing principles and frozen for policy training [47]; reference-guided evaluation [48]; an essay scorer trained against a fixed rubric of binary items [49]; a study of cheap training-time verifiers [158]; Rubric-Grounded RL, the Open Rubric System, and query-specific rubrics learned from human preferences [159–161]	Reference-guided evaluation reaches 79.1% agreement with human annotations; the essay scorer reaches a quadratic weighted kappa of 0.803 against a single human rater's 0.805; the cheap-verifier study calibrates against language-model golden judges, not people
Co-evolved with the policy	5 methods	RLPR, scoring by the policy's own probability of a fixed reference [50]; Rubric-ARM, alternating updates of a rubric generator and a judge at the reward-model stage [51]; EvoRubrics, SCORE, and RubricEM, updating rubrics adversarially or from a rubric buffer within training [52–54]	RLPR only, for its probability reward against human labels; no method that co-evolves a rubric or judge reports a figure
Survey	1 survey	A survey of rubric evaluation, which names the dynamic regime as the trend [162]	—

The only co-evolving method with a human-agreement figure, RLPR, co-evolves a probability reward, not a rubric. No method that co-evolves a rubric or judge reports agreement with people. That is the reflexive-gate problem of Section 4, seen from the weight side.

G.7 Recount of the full corpus

The full corpus holds 256 unique systems, 130 from the first pass and 126 from the second; 204 form the core corpus and 52 the contrast corpus. None of the 126 second-pass systems scores 5 on any substrate. The composite two-loop audit passes 14 research systems and 3 industry exemplars: Sierra Agent OS 2.0, Cognition Devin, and Letta's Context Repositories, whose shared-memory push carries git lineage but no gate. The 7 second-pass systems that promote across contexts with no versioning or lineage on the promoted artifact are EvolveNet, Evo2Team, XRepoSkill, ShareMem, RASO, Spotify, and Future AGI. Of these, EvolveNet alone carries rollback and so enters under

the governance relaxation of Appendix C; the first six enter only if the audit drops its governance requirement altogether; Future AGI's evidence is T5 only, so it enters not at all.

Two first-pass rows are corrected before the recount. The row "Governed Collaborative Memory / SSGM" conflated Cuadros et al. (2605.04264) with SSGM (Lam et al., 2603.11768); it is split, and SSGM, a conceptual framework, scores 2. MAE and Memento-Skills were scored 5 on S5 in the released CSV while the first-pass draft's table said no orchestration system reaches 5; both close one loop only and score 4 under the uniform rubric, which moves the first-pass S5 mean from 2.68 to 2.58.

Nine second-pass rows had their gate labels re-coded on 2026-10-08 under the definition of Section 3, which records who decides and not who supplies the candidate: Letta's Loop-2 label moved from Auto to None (an ungated git push of shared memory; the row still passes the composite audit on its lineage); the Loop-1 label of Azure API Center, the agentic-community MCP Gateway & Registry, the Official MCP Registry, TrueFoundry, Buildrix, the revocation guard, and "Who Maintains Agent Skills?" moved from Human to None (artifacts registered or authored by people, with no failure-driven loop), and Enoch's checkpoint-and-validate protocol moved to None on both loops (a migration of one agent's state, not a promotion). The MCP Gateway's Loop-2 label now records its automated admission scan beside the human approval, and the field study's Loop-2 label records that maintainers merge. No score changed; the cells of Section 7 and Appendix E.2.1 follow the re-coded labels.

The recount runs over all substrate-level rows and gives the counts of Table 4; its last line is the first pass alone (systems / core / contrast) after the two corrections.

```
rows: 269  unique systems: 256  core: 204  contrast: 52
substrate            n   mean  score-5 systems
INTEGRATED           52   3.81  AutoAgent, NanoResearch, SkillRL
S1_Instructions      15   3.07  (none)
S2_Skills            50   3.34  COSPLAY, SAGE
S3_Memory            50   3.06  Governed Collaborative Memory (Cuadros et al.)
S4_Tools             20   3.65  Harvey, Hippocratic AI, MCP
S5_Orchestration     37   3.08  (none)
S6_Governance        45   3.02  AGENTS.md / AAIF, Braintrust, LangSmith Hub, Vellum
First pass alone (after corrections): 130 / 99 / 31; S5 mean 2.58.
```

H Extended related work

The main article makes four design commitments that another survey could make differently: it classifies systems by whether a feedback loop closes over a persisted artifact, rather than by which mechanism is edited; it labels the gate on each of two loops as algorithmic, human, or absent; it classifies the signal that feeds the gate by verifiability; and it audits governance against four named properties. Each commitment meets a separate line of prior work. This section traces what each line established, the setting it assumed, and the condition that changes here, so that a reader can see the differences without the article having to argue for them. Every description below is taken from the cited work's own abstract, method section, or documentation; the evidence record is released with the article.

H.1 Loop closure as the unit of classification

The surveys of self-evolving agents organize the field by what is evolved and when. Gao et al. [10] ask what, when, how, and where to evolve, with a timing axis that separates adaptation inside

a test-time episode from adaptation between episodes, and they note that a human in the loop can validate a newly created tool. Fang et al. [11] give a unified frame of system inputs, agent system, environment, and optimizers, and cover single-agent, multi-agent, and domain-specific settings. Zhou et al. [12] taxonomize 108 self-evolving coding methods by evolution target, timing, and code-specific signal, and they come closest to a gate: their admission rule accepts an update when it improves a metric by a margin and passes a check, and they observe that a passing test "checks only selected behavior". Ren et al. [163] describe an agent as a foundation model times an operational scaffold and treat self-improvement as a self-induced update operator organized by target and driving signal. Surveys of agent memory [164–166], context engineering [17, 167], and agent skills [15, 16, 168] do the same for their substrate: each catalogues mechanisms and, in the skills surveys, lifecycle stages.

What this line established is the inventory. The substrates this article scores are the ones these surveys name, and the mechanisms it finds in the corpus are the ones they catalogue. The setting they assume is a mechanism-centered one: a system is described by what it edits. The condition that changes here is the question asked of each system. Whether a loop closes at all, whether it closes under a gate, and whether a second loop promotes the result are properties that the mechanism inventory does not determine, and they are what the two-loop audit of Section 5 measures. Zhou et al.'s admission rule is the nearest precedent for the gate; it is stated for one domain, code, where the check is a test, and this article asks what the rule becomes when the check is a judge or a delayed outcome.

H.2 The gate as algorithmic, human, or absent

Three lines bear on gate type. The first is prompt and program optimization, which a reader will expect to see named. APE treats the instruction as a program and selects among LLM-proposed candidates [169]; ProTeGi forms natural-language "gradients" from minibatches that criticize the current prompt and searches edits with beam search [170]; OPRO has the language model propose new instructions from a prompt that lists previous candidates with their scores and reports gains of up to 8% on GSM8K and up to 50% on Big-Bench Hard over human-designed prompts [18]; EvoPrompt applies evolutionary operators to a prompt population and improves it against a development set, with gains of up to 25% on BBH across 31 datasets [19]; Promptbreeder evolves task prompts and the mutation prompts that modify them, evaluated on arithmetic and commonsense reasoning and on hate-speech classification [20]. DSPy compiles pipelines of language-model calls into text-transformation graphs and optimizes each pipeline to a metric by creating and collecting demonstrations [21], MIPRO proposes instructions and demonstrations jointly for multi-stage programs under a stochastic mini-batch surrogate [171], TextGrad backpropagates textual feedback through the variables of a compound system [22], and Trace generalizes the pattern to any computation whose execution trace can be shown to an optimizer [172]. This line established that a development-set score is a usable gate for an algorithmic proposer and that the proposer can be the model itself. The setting it assumes is a single optimization run against a fixed task distribution. The condition that changes here is persistence and scope: these optimizers produce a prompt or program as the optimizer's output, not a versioned deployment artifact edited from session signal and promoted under a second gate. That is why OPRO and TextGrad score below the core-corpus threshold and DSPy, which persists only optimizer state, closes no second loop; Appendix E.1 lists all three. SkillOpt is the point where the lineage crosses into the corpus, by adding a persistent skill document, a held-out acceptance gate, and a rejected-edit buffer to the same proposer (Section 8).

The second line optimizes the scaffold as a whole rather than one prompt, and it is the nearest precedent for the loop of Section 3. Agent symbolic learning treats prompts, tools, and the pipeline that connects them as symbolic parameters and updates them with textual losses and gradients [24]; STOP has a scaffolding program recursively rewrite the code that improves it [25]; a self-improving coding agent edits its own harness and raises its benchmark score as a result [26]; and workflow search optimizes the orchestration graph itself, by node and edge optimization over agent graphs [66] or by tree search over code-represented workflows [67]. These established that the whole scaffold is an optimizable object and that textual feedback can drive its updates. Each runs within one deployment under a benchmark score, with no promotion to other deployments and no governance of the artifact it produces; the two-loop audit of Section 5 is the measurement of exactly those two absences.

The third line is the classic single-loop agents and their memories: Reflexion's verbal reinforcement into episodic memory [152], the reflection-and-importance memory of generative agents [173], Voyager's growing skill library [101], the experience bank ExpeL extracts from training tasks [174], MemGPT's virtual context management [175], Agent Workflow Memory's induced and reused workflows [176], the Zettelkasten-style organization of A-MEM [177], the persistent test-time memory of Dynamic Cheatsheet [178], the evolving context playbook of agentic context engineering [23], and the in-context iteration of Self-Refine and ReAct [149, 150]. They established that a frozen model can improve across episodes through an artifact it writes, and the cognitive-architecture framing of Sumers et al. [55] gave that artifact the episodic, semantic, and procedural division this article's memory substrate uses. Reflexion, generative agents, Voyager, ExpeL, MemGPT, Agent Workflow Memory, A-MEM, Dynamic Cheatsheet, and the context playbook persist the artifact and close Loop 1; Self-Refine and ReAct do not persist it. None promotes the artifact beyond the agent that wrote it, which is the absence the second loop of this article is defined to measure.

H.3 Verifiability of the feedback signal

No surveyed work classifies systems by the verifiability of the signal that feeds the gate, but three lines come near it. Guo et al. [14] note that software engineering provides mechanical feedback, in the form of builds, unit tests, linters, and execution logs, that most other domains lack; Zhou et al. [12] note that even a passing test is incomplete feedback. The judged end of the spectrum has its own measurements: language-model judges agree with human preference at rates that make them usable evaluators of open-ended answers [27], a proxy reward optimized past a point lowers the true score [38], and people who grade model outputs revise their own criteria as they grade, so that the evaluator drifts along with the evaluated [39]. The second line is the weight-side effort to extend reinforcement learning from verifiable rewards into domains without verifiers. Rubrics as Rewards generates an instance-specific checklist before training and scores responses with a fixed judge [46]; Writing-Zero trains a pairwise generative judge on writing principles and holds it fixed during policy training [47]; reference-guided evaluation reaches 79.1% agreement with human annotations across the datasets its authors test [48]; an essay scorer trained against a fixed rubric reaches a quadratic weighted kappa of 0.803 against a single human rater's 0.805 [49]. A second group moves the gate with the policy: RLPR scores by the policy's own probability of a fixed reference [50], Rubric-ARM alternates updates of a rubric generator and a judge [51], and EvoRubrics, SCORE, and RubricEM update rubrics adversarially or from a rubric buffer within training [52–54]. The groups part on human calibration: among the methods that move the gate with the policy, only RLPR, whose reward is a probability rather than a rubric, reports agreement with human labels, and no method that co-evolves a rubric or judge does (Appendix G.6 gives the tally). One study

calibrates its verifier against language-model golden judges rather than people [158], and a survey of rubric evaluation names the dynamic regime as the trend [162].

This line established the gate constructions for judged signals: checklists, references, pairwise judges, and rubric buffers. The setting it assumes is weight-space optimization with the policy as the learner. The condition that changes here is the learner: when the scaffold is what learns, the rubric is itself a scaffold coordinate, so the fixed-versus-co-evolved distinction becomes the question of whether a gate under edit can be calibrated at all. The article borrows the constructions and the distinction; it does not score the methods.

H.4 Governance: lineage, access control, rollback, and promotion across organizations

Three lines bear on the four-property audit. The first is governance surveys. Ding et al. [9] code 435 works on six dimensions per state item, including provenance and recoverability, and find the literature "concentrates more heavily on accumulating and retrieving state than on governing, recovering, or relinquishing it". Lin et al. [13] organize memory security by lifecycle phase and name five primitives, among them principal-scoped retrieval and rollbackability. Neither carries a promotion dimension, and neither scores whether an inner loop fires from session signal.

The second line is the systems that come closest to the four properties. FederatedSkill exchanges semantic skill diffs between clients that are agent instances under one operator, and a server-side evolution agent applies, adapts, or rejects each patch per client while writing an auditable decision log of source patch, reward, and justification; it reports up to 44.4% higher success than isolated self-evolution across 20 task families [126]. EvolveNet broadcasts a shared harness, optimizes it locally at each client, and merges only when a merge fixes at least as many held-out items as it breaks, rolling back to the previous shared harness after a failed round; its authors describe clients as spanning users and organizations, evaluate on benchmark partitions, and note that a program edit can itself encode client information [95]. The registries supply the remaining properties without an inner loop: EDDOps promotes through automatic evaluation thresholds and human sign-off with supersession lineage and emergency rollback, and states that it does not address role-based access control [85]; Agentic Configuration Management keeps immutable revisions and baselines that enable rollback [86]; Governed Capability Evolution gates activation and rolls back on monitoring, recovering in 72.4% of drift cases in its proof of concept [87]; and the MCP control planes add role-based access and quarantine without an evaluation gate [82–84]. Letta's Context Repositories commit every memory edit to git and push shared memory across agents without a gate [73].

The third line is federated learning. FedAvg averages client updates in parameter space [130], DP-FedAvg adds user-level differential privacy [131], the field's open problems are catalogued in Kairouz et al. [132], and PromptFL moves the federated object from weights to a shared prompt that participants learn cooperatively [133]. The setting is numeric updates that are comparable by construction. The condition that changes here is the artifact: scaffold promotion is selection and merge over symbolic objects, and the nearest systems aggregate program edits with a language model rather than by averaging, which is why the article uses the federated analogy only as a limit.

H.5 What is and is not new

The ingredients have precedent, and the article says so. The six substrates are the inventory of the self-evolving-agent, memory, context, skill, and harness surveys, and the memory taxonomy is CoALA's. The gate as a development-set score is the prompt-optimization lineage; the scaffold as

an optimizable object is agent symbolic learning, STOP, and workflow search; the persisted artifact that closes Loop 1 is Reflexion, Voyager, and the agent-memory line; the admission rule with a margin and a check is Zhou et al.'s; the gate constructions for judged signals are the rubric-reward line; provenance, recoverability, and rollback as governance dimensions are Ding et al. and Lin et al.; federated aggregation of artifact patches is FederatedSkill and EvolveNet. What is new is the combination and the measurement: a corpus of 256 systems scored on one rubric and released; a classification of each by which of two loops it closes and under which gate type; a classification of the signal that feeds the gate into four verifiability classes, with the demonstration that the class fixes the reachable cells of the design space and the dominant failure; and an audit of every system against four governance properties, with the finding that none meets all four and that the two nearest, FederatedSkill and EvolveNet, each narrow different properties without meeting the set. These are the claims of the abstract and the introduction. Three limits bound them: scoring is single-rater over public evidence of uneven strength (Section 5), the corpus closes at the October 2026 cutoff, and the evidence is the published behavior of the 256 systems rather than new experiments (Section 1).

I Full score table

Every scored (substrate, system) row of the corpus: 130 systems scored by May 2026 and 126 added by the October 2026 sweep. Scores follow the 0–5 rubric of the main text (Table 3); Loop 1 / Loop 2 gate types and signal classes are recorded for the October rows (and for May rows re-audited in Section 7). Substrate codes: INT integrated, S1 instructions, S2 skills, S3 memory, S4 tools, S5 orchestration, S6 governance. A year of "—" marks a classical system that predates the survey window. The CSV files with per-row justifications are released as ancillary files with this preprint.

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
INT	ADAS / Meta Agent Search	4	2024	arXiv:2408.08435				May 26
INT	Agent S / S2	4.5	2024	arXiv:2410.08164				May 26
INT	Agent Workflow Memory (AWM)	3.5	2024	arXiv:2409.07429				May 26
INT	AlphaEvolve	4	2025	arXiv:2506.13131				May 26
INT	AutoAgent	5	2026	arXiv:2603.09716				May 26
INT	AutoManual	3.5	2024	NeurIPS 2024				May 26
INT	AutoSaddler	4	2026	arXiv:2608.23041	Auto	None	C2	Oct 26
INT	CaptainAgent	2.5	2024	arXiv:2405.19425				May 26
INT	CASCADE	4.5	2025	arXiv:2512.23880				May 26
INT	ChatDev	3	2023	arXiv:2307.07924				May 26
INT	Co-Harness	4	2026	arXiv:2607.22688	Auto	None	C2	Oct 26
INT	CoEvoSkills	3.5	2026	arXiv:2604.01687				May 26
INT	Cognition Devin	3	2024	cognition.ai				May 26
INT	Continual Harness	4	2026	arXiv:2605.09998	None	None	C4	Oct 26
INT	Cradle	4	2024	arXiv:2403.03186				May 26
INT	Darwin Gödel Machine	4	2025	arXiv:2505.22954				May 26
INT	Double Ratchet	4	2026	arXiv:2607.12790	Auto	None	C3 anchored by C1	Oct 26
INT	EvolveNet	4	2026	arXiv:2608.04968	Auto	Auto	C1,C2	Oct 26
INT	EvolveR	4	2025	arXiv:2510.16079				May 26
INT	ExpeL	3.5	2023	arXiv:2308.10144				May 26
INT	Generative Agents	4	2023	arXiv:2304.03442				May 26
INT	GenericAgent	4	2026	arXiv:2604.17091	Auto	None	C2	Oct 26
INT	GITM	3	2023	arXiv:2305.17144				May 26
INT	HarnessBank	4	2026	arXiv:2607.13683	Auto	None	C2	Oct 26
INT	HarnessX	4	2026	arXiv:2606.14249	Auto	None	C2	Oct 26
INT	Hippocratic AI Polaris	3.5	2024	hippocraticai.com				May 26
INT	JARVIS-1	3.5	2023	arXiv:2311.05997				May 26
INT	JIT-Agent	4	2026	arXiv:2608.25593	Auto	None	C2	Oct 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
INT	Life-Harness	4	2026	arXiv:2605.22166	Auto	None	C2	Oct 26
INT	Mendel Gödel Machine (MGM)	4	2026	arXiv:2608.07645	Auto	None	C1	Oct 26
INT	MetaGPT	3.5	2023	arXiv:2308.00352				May 26
INT	Mobile-Agent-E	4.5	2025	arXiv:2501.11733				May 26
INT	ModularRSI	4	2026	arXiv:2609.14857	Auto	None	C2	Oct 26
INT	NanoResearch	5	2026	arXiv:2605.10813				May 26
INT	OpenCore	4	2025	github.com/sibmike/ opencore				May 26
INT	OpenHands / CodeAct	3	2024	arXiv:2407.16741				May 26
INT	OS-Copilot / FRIDAY	4	2024	arXiv:2402.07456				May 26
INT	Plan4MC	2.5	2023	GitHub/arXiv				May 26
INT	SAGE	4	2025	arXiv:2512.17102				May 26
INT	Sakana AI Scientist	3	2024	arXiv:2408.06292				May 26
INT	ScholarEvolve	4	2026	arXiv:2609.40169	Auto	None	C2	Oct 26
INT	SIA	4	2026	arXiv:2605.27276	Auto	None	C2	Oct 26
INT	Sierra Agent OS	4	2024	sierra.ai				May 26
INT	Sierra Agent OS 2.0	3	2025	sierra.ai/blog/ agent-os-2-0				May 26
INT	SkillForge	4	2026	arXiv:2604.08618				May 26
INT	SkillMAS	4	2026	arXiv:2605.09341	Auto	None	C2	Oct 26
INT	SkillOS	3	2026	SemanticScholar: 2653e4c7				May 26
INT	SkillRL	5	2026	arXiv:2602.08234				May 26
INT	TTHE	4	2026	arXiv:2607.08124	Auto	None	C3	Oct 26
INT	VACE	4	2026	arXiv:2609.37105	Auto	None	C2	Oct 26
INT	Voyager	4	2023	arXiv:2305.16291				May 26
INT	Workspace Optimization (DreamTeam)	3	2026	arXiv:2605.09650	Auto	None	C1	Oct 26
S1	AGENTS.md	2	2026	arXiv:2601.20404				May 26
S1	Aider Conventions	2	2024	aider.chat				May 26
S1	Claude Code CLAUDE.md	4	2025	docs.anthropic.com				May 26
S1	Cody / Sourcegraph	2	2024	sourcegraph.com				May 26
S1	Continue.dev Rules	3	2024	docs.continue.dev				May 26
S1	Cursor Rules	3	2024	docs.cursor.com/rules				May 26
S1	GitHub Copilot Custom Instructions	3	2024	docs.github.com				May 26
S1	GRACE	4	2026	arXiv:2607.09175	Auto	None	C2	Oct 26
S1	Lovable Knowledge	2	2024	docs.lovable.dev				May 26
S1	Meta Context Engineering (MCE)	4	2026	arXiv:2601.21557	Auto	None	C2	Oct 26
S1	Replit replit.md	4	2025	docs.replit.com				May 26
S1	Spotify conversational-recommendation loop	4	2026	arXiv:2609.30297	Auto	Human	C3	Oct 26
S1	VS Code Copilot Instructions	3	2024	code.visualstudio.com				May 26
S1	Windsurf Cascade Rules	3	2024	docs.windsurf.com				May 26
S1	Zed AI Rules	3	2024	zed.dev/docs				May 26
S2	AgentVerse	1	2023	arXiv:2308.10848				May 26
S2	Anthropic Claude Agent Skills	3	2025	anthropic.com/ engineering				May 26
S2	Anthropic/skills GitHub repo	2	2025	github.com/ anthropics/skills				May 26
S2	ASSAY	4	2026	arXiv:2606.15390	Auto	None	C2	Oct 26
S2	AutoSkill	3	2026	arXiv:2603.01145	Auto	None	C3	Oct 26
S2	Buildrix	3	2026	arXiv:2606.25139	None	Human	C3	Oct 26
S2	ChatGPT Plugins (legacy)	1	2023	openai.com				May 26
S2	Claude Code Skills	3	2025	docs.anthropic.com				May 26
S2	COSPLAY	5	2026	arXiv:2604.20987				May 26
S2	CrewAI Skills	2	2025	docs.crewai.com				May 26
S2	HiSME	4	2026	arXiv:2605.28390	Auto	None	C1,C3	Oct 26
S2	HuggingFace Agent Skills	2	2025	huggingface.co				May 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S2	LangChain Skills	2	2025	langchain.com				May 26
S2	Library Drift curator	4	2026	arXiv:2605.19576	Auto	None	C1	Oct 26
S2	Meta Context Engineering (MCE)	4	2026	arXiv:2601.21557	Auto	None	C2	Oct 26
S2	MetaGPT (SOP-as-skill)	3	2023	arXiv:2308.00352				May 26
S2	MetaSkill-Evolve	4	2026	arXiv:2607.05297	Auto	None	C2	Oct 26
S2	Microsoft Semantic Kernel Plugins	2	2024	learn.microsoft.com				May 26
S2	MUSE-Autoskill	4	2026	arXiv:2605.27366	Auto	None	C1,C2	Oct 26
S2	NanoResearch Skill Bank	4	2026	arXiv:2605.10813				May 26
S2	OpenAI Custom GPTs / Actions	2	2023	platform.openai.com				May 26
S2	Personalized developer skills (Huang-Du-Lan)	3	2026	arXiv:2608.10319	Auto	None	C3	Oct 26
S2	RASO	4	2026	arXiv:2609.38024	Auto	Auto	C2	Oct 26
S2	SAGE	5	2025	arXiv:2512.17102				May 26
S2	SCAFFOLD	4	2026	arXiv:2609.05511	Auto	None	C2	Oct 26
S2	SE-GoS	4	2026	arXiv:2609.08228	Auto	None	C2	Oct 26
S2	SkeMex	4	2026	arXiv:2606.09365	Auto	None	C2	Oct 26
S2	Skill Issues (Bauplan)	4	2026	arXiv:2606.01185	Auto	None	C1,C3	Oct 26
S2	SKILL.md smells (Hong-Imani-Ahmed)	1	2026	arXiv:2607.01456	None	None	-	Oct 26
S2	SkillAudit	4	2026	arXiv:2606.14239	Auto	None	C3,C1-partial	Oct 26
S2	SkillBoost	4	2026	arXiv:2607.26643	Auto	None	C1	Oct 26
S2	SkillChain	4	2026	arXiv:2606.12984	Auto->Huma	None	C3,C2	Oct 26
S2	SkillComposer	4	2026	arXiv:2606.06079	Auto	None	C1,C2	Oct 26
S2	SkillEvoReg	4	2026	arXiv:2609.30861	Auto	None	C1,C2	Oct 26
S2	SkillFoundry	3	2026	arXiv:2604.03964	Auto	None	C1	Oct 26
S2	SkillGen	4	2026	arXiv:2605.10999	Auto	None	C2	Oct 26
S2	SkillGLoW	4	2026	arXiv:2609.02217	Auto	None	C2	Oct 26
S2	SkillLift	4	2026	arXiv:2609.15396	Auto	None	C3 inner, C2 outer	Oct 26
S2	SkillOps	4	2026	arXiv:2605.13716	Auto	None	C2	Oct 26
S2	SkillSmith	4	2026	arXiv:2606.01314	Auto	None	C1,C2	Oct 26
S2	SkillSpec (consensus-gated)	4	2026	arXiv:2610.00704	Auto	None	C1,C2	Oct 26
S2	SuperAGI Toolkits	1	2024	github.com/TransformerOptimus				May 26
S2	TRUSS	3	2026	arXiv:2608.17588	Auto	None	C1,C2	Oct 26
S2	VALVE	4	2026	arXiv:2609.32990	Auto	None	C1	Oct 26
S2	VCE-Skill	4	2026	arXiv:2608.16544	Auto	None	C1,C2	Oct 26
S2	Voyager	4	2023	arXiv:2305.16291				May 26
S2	Web2BigTable	3	2026	arXiv:2604.27221	Auto	None	C1,C2	Oct 26
S2	Who Maintains Agent Skills? (empirical)	2	2026	arXiv:2609.05677	None	Human	C3	Oct 26
S2	WikiSkill	4	2026	arXiv:2608.27454	Auto	None	C2	Oct 26
S2	XRepoSkill	4	2026	arXiv:2609.36807	Auto	Auto-parti	C1	Oct 26
S3	A-MEM	3	2025	NeurIPS 2025				May 26
S3	Agent-native memory system evaluation (Zhou et a	2	2026	arXiv:2606.24775	n/a	n/a	n/a	Oct 26
S3	AIM privacy-aware interoperable memory (Johnson	3	2026	arXiv:2609.12320	Auto	None	C3	Oct 26
S3	Beyond the Harness (enterprise Text-to-SQL cards	3	2026	arXiv:2608.22830	Auto	None	C2	Oct 26
S3	ChatGPT Memory	3	2024	openai.com				May 26
S3	Claude Memory	3	2025	anthropic.com				May 26
S3	Cognee	3	2024	github.com/topoteretes/cognee				May 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S3	Compiled Memory (Rhodes & Kang)	4	2026	arXiv:2603.15666	Auto-gated	None	C1	Oct 26
S3	Control-plane placement / ForgetEval (Yang)	2	2026	arXiv:2606.15903	n/a	n/a	n/a	Oct 26
S3	FRESH failure-aware heterogeneous graph memory (	4	2026	arXiv:2609.28003	Auto	None	C2	Oct 26
S3	GateMem multi-principal memory-governance benchm	2	2026	arXiv:2606.18829	n/a	n/a	n/a	Oct 26
S3	Gemini Code Assist Memory	4	2025	cloud.google.com				May 26
S3	Generative Agents	3	2023	arXiv:2304.03442				May 26
S3	Governed Collaborative Memory (Cuadros et al.)	5	2026	arXiv:2605.04264				May 26
S3	HippoRAG	2	2024	arXiv				May 26
S3	Honest Lying mitigation / programmatic failure s	4	2026	arXiv:2605.29463	Auto	None	C1/C2	Oct 26
S3	KV-cache rollback consistency (Zhang & Yang)	2	2026	arXiv:2608.15939	n/a	n/a	n/a	Oct 26
S3	LangChain Primitives (legacy)	1	2023	langchain.com				May 26
S3	LangGraph Long-Term Memory	3	2024	langchain.com				May 26
S3	LangGraph Store / LangMem (2026 refresh)	3	2026	docs.langchain.com/ oss/python	None	None	n/a	Oct 26
S3	Learning on the Job	4	2026	arXiv:2607.22157	Auto+ Human	None	C2,C3	Oct 26
S3	Letta MemFS + Context Repositories + Dreaming (2	4	2026	docs.letta.com/ letta-code	Auto	None	C3	Oct 26
S3	Mem0	3	2025	arXiv				May 26
S3	Mem0 Platform (2026 refresh: history, feedback,	3	2026	docs.mem0.ai/ api-reference	Auto	None	C3	Oct 26
S3	MemAudit post-hoc poisoned-memory audit (Tan et	3	2026	arXiv:2605.23723	Auto	None	C2	Oct 26
S3	MemClaw / Governed Shared Memory (Margalit et al	3	2026	arXiv:2606.24535	Auto	None	C3	Oct 26
S3	MemGPT / Letta	3	2023	arXiv:2310.08560				May 26
S3	MemGuard persisted verifier signals (Wang et al.	4	2026	arXiv:2608.21867	Auto-gated	None	C3	Oct 26
S3	MemoryBank	3	2023	Zhong et al.				May 26
S3	Mi-Memory lifecycle framework (Liu et al.; Xiaom	3	2026	arXiv:2607.18975	Auto	None	C3	Oct 26
S3	MIRIX	3	2025	arXiv				May 26
S3	MPBench memory-poisoning study (Dash et al.)	2	2026	arXiv:2606.04329	n/a	n/a	n/a	Oct 26
S3	NanoResearch Memory Module	3	2026	arXiv:2605.10813				May 26
S3	PPMF provenance-preserving memory firewall (Xu e	3	2026	arXiv:2607.29167	Auto	None	C1	Oct 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S3	REALM retrieval-driven memory reconsolidation (S	4	2026	arXiv:2609.16053	Auto	None	C2/C3	Oct 26
S3	Recuris recursive experiential-working memory (Y	4	2026	arXiv:2608.24876	Auto-gated	None	C2	Oct 26
S3	Reflexion	4	2023	arXiv:2303.11366				May 26
S3	Revocation guard / Revoked but Still Authoritati	3	2026	arXiv:2609.08258	None	None	C1	Oct 26
S3	Safe to Resume? checkpoint-rollback continuity (	2	2026	arXiv:2608.29381	n/a	n/a	n/a	Oct 26
S3	ShareMem	4	2026	arXiv:2609.32511	Auto	Auto-parti	C3	Oct 26
S3	SIx Harness / Memory as Infrastructure (Helwig)	3	2026	arXiv:2609.05510	Auto	None	C2	Oct 26
S3	SOAR Cognitive Architecture	4	—	Newell/Laird				May 26
S3	SSGM stability-and-safety governed memory (Lam e	2	2026	arXiv:2603.11768	None	None	n/a	Oct 26
S3	Stale constraints in inherited memory (Nakayashi	2	2026	arXiv:2608.25553	n/a	n/a	n/a	Oct 26
S3	TMA-NM origin-bound memory authority (Louck)	3	2026	arXiv:2606.24322	Auto	None	C1	Oct 26
S3	TOKI bitemporal operator algebra (Wang)	3	2026	arXiv:2606.06240	Auto or Hu	None	C3	Oct 26
S3	Trust-scored memory sanitization (Devarangadi Su	2	2026	arXiv:2601.05504	Auto gate	None	n/a	Oct 26
S3	Typed provenance + assertion guardrails / Stored	3	2026	arXiv:2609.02127	Auto	None	C1	Oct 26
S3	Utility Under Attack (Karunanidhi)	2	2026	arXiv:2608.21230	n/a	n/a	n/a	Oct 26
S3	VerificAgent human-verified frozen memory (Nguye	3	2025	arXiv:2506.02539	Human	None	C3	Oct 26
S3	VerMem verifiable memory management (Sun et al.)	4	2026	arXiv:2608.03137	Auto-gated	None	C2/C3	Oct 26
S3	Voyager Skill Library	4	2023	arXiv:2305.16291				May 26
S3	Zep / Graphiti	3	2025	arXiv				May 26
S3	Zep / Graphiti bitemporal facts (2026 refresh)	3	2026	help.getzep.com	Auto	None	C3	Oct 26
S4	Anthropic Tool Use API	4	2024	docs.anthropic.com				May 26
S4	Browser Tools (Playwright/Computer Use)	3	2024	anthropic.com				May 26
S4	Claude Code Tool API	4	2024	docs.anthropic.com				May 26
S4	Code Interpreter	3	2023	openai.com				May 26
S4	Cursor MCP Integration	4	2024	docs.cursor.com				May 26
S4	Devin / SWE-agent Tool Harness	4	2024	cognition.ai				May 26
S4	Gemini Function Calling	3	2024	ai.google.dev				May 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S4	GPT Actions	3	2023	platform.openai.com				May 26
S4	Harvey (Legal Vertical)	5	2024	harvey.ai				May 26
S4	Hippocratic AI	5	2024	hippocraticai.com				May 26
S4	LangChain Tools Abstraction	3	2023	langchain.com				May 26
S4	LangGraph Tool Nodes	4	2024	langchain.com				May 26
S4	LlamaIndex Tool Ecosystem	4	2023	llamaindex.ai				May 26
S4	Model Context Protocol (MCP)	5	2024	modelcontextprotocol.io				May 26
S4	OpenAI Assistants API	3	2023	platform.openai.com				May 26
S4	OpenAI Function Calling	3	2023	platform.openai.com				May 26
S4	Sierra	4	2024	sierra.ai				May 26
S4	SkillSmith	4	2026	arXiv:2606.01314	Auto	None	C1	Oct 26
S4	Toolformer	2	2023	arXiv:2302.04761				May 26
S4	ToolLLM / ToolBench	3	2023	arXiv:2307.16789				May 26
S5	AgentConductor	0	2026	arXiv:2602.17100	None	None	C2	Oct 26
S5	AgentGrad	4	2026	arXiv:2609.08572	Auto	None	C2	Oct 26
S5	AgentVerse	2	2023	arXiv:2308.10848				May 26
S5	AutoGen / MagenticOne	3	2024	arXiv:2308.08155				May 26
S5	AWS Bedrock Multi-Agent	2	2024	aws.amazon.com				May 26
S5	Bilevel Autoresearch	4	2026	arXiv:2603.23420	Auto	None	C1	Oct 26
S5	CAMEL	1	2023	arXiv				May 26
S5	ChatDev	2	2023	arXiv:2307.07924				May 26
S5	ContextEvo	4	2026	arXiv:2609.34649	Auto	None	C2	Oct 26
S5	Control System for Frozen Agents (Paul)	4	2026	arXiv:2607.25415	Auto	None	C2	Oct 26
S5	CrewAI	3	2025	docs.crewai.com				May 26
S5	Evo2Team	4	2026	arXiv:2609.34135	Auto	Auto	C2	Oct 26
S5	Growing Harness	4	2026	arXiv:2609.26760	Auto	None	C2	Oct 26
S5	HARBOR	4	2026	arXiv:2604.20938	Auto	None	C2	Oct 26
S5	HarnessForge	4	2026	arXiv:2606.01779	Auto	None	C2	Oct 26
S5	Inspect AI	2	2024	ukgovernmentbeis.github.io				May 26
S5	Institutional MAS architectures (Fei et al.)	1	2026	arXiv:2604.27691	None	None	C2	Oct 26
S5	LangGraph	3	2024	langchain.com				May 26
S5	Letta	3	2025	letta.com				May 26
S5	LlamaIndex Workflows	2	2024	llamaindex.ai				May 26
S5	Memento-Skills	4	2026	arXiv				May 26
S5	MetaGPT	2	2023	arXiv:2308.00352				May 26
S5	Microsoft Semantic Kernel	2	2024	learn.microsoft.com				May 26
S5	Multi-Agent Evolve (MAE)	4	2025	arXiv				May 26
S5	OpenAI AgentKit	4	2025	openai.com				May 26
S5	OpenAI Swarm	1	2024	openai.com				May 26
S5	OpenHands / CodeAct	3	2024	arXiv:2407.16741				May 26
S5	OptiMAS	4	2026	arXiv:2608.21918	Auto	None	C2	Oct 26
S5	PluginRSI	4	2026	arXiv:2609.32423	Auto	None	C2	Oct 26
S5	Pydantic AI	2	2024	pydantic.dev				May 26
S5	Recursive Harness Self-Improvement (RHI)	4	2026	arXiv:2607.15524	Auto	None	C3	Oct 26
S5	Skill-MAS (meta-skill)	4	2026	arXiv:2606.18837	Auto	None	C2	Oct 26
S5	TACO	4	2026	arXiv:2604.19572	Auto	None	C2	Oct 26
S5	Task-CoEvolve	4	2026	arXiv:2608.20169	Auto	None	C2	Oct 26
S5	TOCOMAS	4	2026	arXiv:2609.37953	Auto	None	C2	Oct 26
S5	Vertex AI Agent Builder	4	2024	cloud.google.com				May 26
S5	Workflow-Designing Agents (WDA)	4	2026	arXiv:2609.32458	Auto	None	C2	Oct 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S6	Agenta	3	2024	agenta.ai				May 26
S6	Agentic Configuration Management ACM (Quessada-V	3	2026	arXiv:2608.11166	None	Human	n/a	Oct 26
S6	AGENTS.md / AAIF	5	2025	agents.md				May 26
S6	Anthropic skills repo	4	2025	github.com/ anthropics/skills				May 26
S6	Azure API Center MCP registry	3	2026	learn.microsoft.com	None	Human	-	Oct 26
S6	Braintrust	5	2024	braintrust.dev				May 26
S6	Cleric.ai	2	2025	cleric.ai				May 26
S6	Competence Not Accuracy	2	2026	arXiv:2608.18719	None	None	-	Oct 26
S6	Constitutional AI	2	2022	Anthropic				May 26
S6	Cursor Rules	3	2024	cursor.com				May 26
S6	Decagon	2	2026	decagon.ai				May 26
S6	DSPy compile	3	2024	dspy.ai				May 26
S6	EvolveNet	4	2026	arXiv:2608.04968	Auto	Auto	C2	Oct 26
S6	EvoUndo	4	2026	arXiv:2608.28363	Auto	None	C1	Oct 26
S6	From Prompts to Contracts	2	2026	arXiv:2607.08028	None	None	C1	Oct 26
S6	Future AGI prompt lifecycle	4	2026	futureagi.com/blog	Auto	Auto	C3	Oct 26
S6	GitHub Copilot enterprise	4	2026	github.com				May 26
S6	Governance conversion (Davis et al. case study)	3	2026	arXiv:2607.01087	Human	None	C2	Oct 26
S6	Governed Capability Evolution (Qin et al.)	3	2026	arXiv:2604.08059	None	Auto	C2	Oct 26
S6	HackProbe	3	2026	arXiv:2609.04665	Auto	None	C1	Oct 26
S6	Inspect AI	3	2024	ukgovernmentbeis. github.io				May 26
S6	LangFuse	3	2024	langfuse.com				May 26
S6	LangSmith Hub	5	2024	smith.langchain.com				May 26
S6	MCP Gateway & Registry (agentic-community)	3	2026	github.com/ agentic-community	None	Human	C1	Oct 26
S6	NanoResearch	2	2026	arXiv:2605.10813				May 26
S6	Official MCP Registry	2	2026	modelcontextprotocol. io	None	None	-	Oct 26
S6	OpenAI AgentKit Evals	3	2025	openai.com				May 26
S6	OPRO	2	2023	arXiv				May 26
S6	Organizational Control Layer (Shi et al.)	2	2026	arXiv:2606.04306	None	None	n/a	Oct 26
S6	Registry-Governed Agent Lifecycle / EDDOps on AW	3	2026	arXiv:2607.00345	None	Auto+ Human	C3	Oct 26
S6	Reward-Free Pairwise Validator	4	2026	arXiv:2607.14408	Auto	None	C3	Oct 26
S6	Runtime Skill Audit	2	2026	arXiv:2606.11671	None	None	-	Oct 26
S6	Runtime-Independent Persistent Agents (Enoch)	3	2026	arXiv:2609.00546	None	None	C1	Oct 26
S6	SafeEvolve	4	2026	arXiv:2608.12851	Auto	None	C3,C2	Oct 26
S6	SAGE statistical acceptance gate ($\neq$ SAGE 2512.17	4	2026	arXiv:2609.36043	Auto	None	C2	Oct 26
S6	Self-Refine	1	2023	Madaan et al.				May 26
S6	Sierra Ghostwriter	2	2026	sierra.ai				May 26
S6	SkillCoach	3	2026	arXiv:2607.01874	Auto	None	C3,C1	Oct 26
S6	SkillProbe	2	2026	arXiv:2603.21019	None	None	-	Oct 26
S6	SkillSpec (intent-masked correctness)	2	2026	arXiv:2609.06052	None	None	-	Oct 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S6	TextGrad	2	2024	arXiv				May 26
S6	TrueFoundry MCP Gateway	3	2026	truefoundry.com/docs	None	Human	-	Oct 26
S6	Vellum	5	2024	vellum.ai				May 26
S6	Verifier-as-Gatekeeper (VaG)	4	2026	arXiv:2608.05810	Auto	None	C2	Oct 26
S6	W&B Weave (Prompts)	3	2024	weave-docs.wandb.ai				May 26