Frontier and Localhost: A Survey of How Production AI Learns Outside the Weights

Mikhail L. Arbuzov
Independent Researcher
mike.arbuzov54@gmail.com

Lee Mosbacker
Independent Researcher
lee.mosbacker@gmail.com

Sisong Bei
Independent Researcher
qurining@gmail.com

Ziwei Dong
Independent Researcher
ziwei.dong@alumni.emory.edu

Dmitri Kalaev
Independent Researcher
kalaevdr@gmail.com

Alexey A. Shvets
Palo Alto Networks
ashvets@paloaltonetworks.com

Abstract

Deployed AI systems change through instructions, skills, memories, tools, and orchestration long before their model weights are updated. These changes accumulate experience, but their persistence alone does not establish that the system is improving. This survey examines how 256 production and research systems turn experience into lasting changes outside the weights. We distinguish two decisions: whether an edit improves the deployment that produced it, and whether there is enough evidence to share it more widely. Together, these decisions define a local improvement loop and a promotion loop. Automatic acceptance tests are present in nearly every system with failure-driven local adaptation, while failure capture and attribution remain less developed. Promotion is more fragmented: the surveyed systems provide evaluation, shared artifact pools, lineage, access control, and rollback in different combinations, but none documents their integration into governed promotion across organizations. The feedback available to a deployment helps explain this pattern. Exact checks, observed outcomes, judgments, and delayed rewards support different acceptance rules and expose different limitations. In every case, an accepted edit is supported only for the measure and task distribution on which it was evaluated. We use this observation to develop five conditions for artifact-layer optimization, a classification of feedback and gate types, and four properties for auditing cross-organization promotion. The resulting framework identifies what existing systems can already automate and where new evidence or human judgment is still required.

1 Introduction

A coding agent repeatedly uses the wrong naming convention in a repository. An engineer adds a line to its instructions, and the next session follows the convention. The model has not changed, but the deployed system has. Similar changes happen when an assistant corrects a stored preference, acquires a reusable skill, or learns which tool to call. Much of the experience that makes an AI system useful in one setting now accumulates in these persistent artifacts outside the model's weights.

We call this layer the **scaffold**. It includes instructions, skills, memory, tools, orchestration, and the

evaluation and governance rules that control their use. In the language of compound AI systems [1], it is the part surrounding the model that a deployment can inspect and revise without retraining it. The distinction matters because the two parts change on different schedules. A team can revise an instruction or tool binding as soon as a problem is understood, while a new model requires a separate training, evaluation, and release process.

The ease of making a change also makes it easy to accumulate changes whose value is unknown. A rule added after one incident may prevent a recurring error, conflict with another rule, or merely make the prompt longer. A corrected memory may remain in the history while the system continues to retrieve the outdated entry. Even a demonstrably useful change can become harmful when copied into a deployment with different tasks or conventions.

These situations raise two questions. What evidence justifies keeping a local change? What additional evidence justifies sharing it? This survey uses those questions to examine 256 public production and research systems. The first defines a **local loop**: identify a recurring failure, propose a change to the responsible artifact, and test whether the change helps. The second defines a **promotion loop**: decide whether the change should be adopted by another project, team, or organization, and preserve enough history to reverse that decision.

The systems surveyed have developed these loops unevenly. Automatic tests for accepting a local edit are present in nearly every system that adapts from failures. The steps that should supply those tests with well-founded candidates are weaker: failures are often not retained in a usable form, and a failed output rarely identifies which instruction, memory, skill, or tool binding caused it. Meanwhile, production platforms have built substantial versioning and release machinery without any process that proposes changes from deployment experience.

Sharing adds a further limitation. A change that improves one deployment has been evaluated against that deployment's tasks and measure of success. Another deployment may need different behavior. Shared formats and registries make artifacts portable, but they do not establish that an artifact will help its recipient. Some surveyed systems check transferred changes within one operator. None documents an automatic local improvement loop connected to promotion across organizations with versioned lineage, access control, and rollback.

The feedback available to each system helps explain these differences. A compiler or test suite can check a candidate immediately. A booked appointment provides a later, noisier outcome. A writing task may depend on a judge, whose preferences can drift. A research decision may take months to evaluate. These signals support different kinds of automation, and none establishes an improvement beyond what it actually measures.

The survey develops a framework for reading this evidence. We first define the two loops and the conditions under which artifact maintenance constitutes optimization. After describing the corpus and scoring method, we examine local adaptation, the feedback that constrains it, and the evidence needed for promotion. The later sections consider failures that arise when these checks are missing, specify the unassembled architecture for governed sharing, and connect the findings to the research lines that produced its components. The evidence consists of published methods and public documentation; the survey introduces no new experiments. Its scores and classifications are released for inspection.

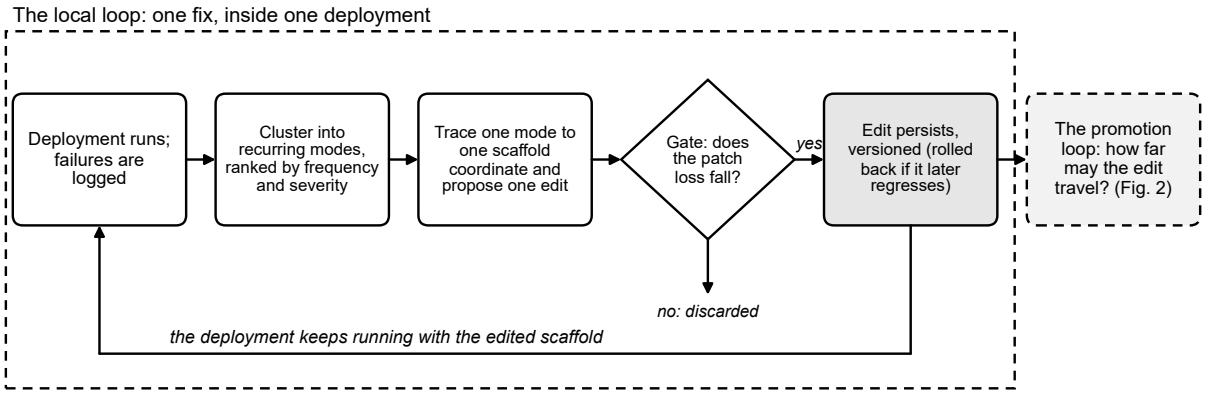


Figure 1: The local improvement loop. Recurring failures guide a proposed scaffold edit. An acceptance gate tests the change before it persists, and later regressions can trigger rollback. Sharing the accepted edit requires a separate promotion decision (Fig. 3).

2 From a local fix to a measured improvement

Adding a naming rule after an error is a reasonable response. To learn whether it works, however, the team must run the agent again on relevant tasks and compare the results. If the change helps, it should persist. If it causes regressions, the team needs a way to reject or reverse it. Repeating this process turns a growing collection of fixes into a procedure for improving the deployment.

2.1 The local loop

The procedure starts with a record of failures: rejected answers, user corrections, failed tool calls, evaluation regressions, or unsuccessful tasks. Repeated failures are grouped into modes and prioritized by frequency and severity. For each selected mode, the system tries to identify the artifact responsible. A stale memory entry calls for a different intervention from a missing tool argument or an ambiguous instruction.

We call this step **credit assignment**. It links an observed failure to an editable part of the scaffold, or **coordinate**. The system then proposes a candidate edit and evaluates it against the current scaffold. The acceptance test is the **gate**. An accepted edit is retained as a versioned artifact; subsequent evidence of regression can trigger rollback. Fig. 1 shows the complete local loop.

The evaluation must be tied to the deployment's actual workload. A deployment serves a particular distribution of users, tasks, and conventions, which we call its **patch**. Its **patch loss** measures failure on that distribution. Depending on the application, this may be a task-failure rate, correction rate, retrieval error, or another operational measure. A gate can establish an improvement only to the extent that its evaluation estimates this loss.

2.2 The second decision: whether to share the fix

Suppose the naming rule improves the original repository. The team may now consider adding it to a shared template. That decision changes the population on which the rule will act. Other repositories may follow a different convention, and the local evaluation says little about them.

The promotion loop therefore asks for evidence at the receiving scope. We call this **receiving-side**

evaluation: a check conducted on the tasks of the deployment that would adopt the edit. The **evidence radius** is the scope for which the available evaluation supports the change. A local success does not automatically enlarge that radius. Section 6 develops the scopes and examines systems that already transfer artifacts between them.

2.3 Five conditions for artifact-layer optimization

We use **artifact-layer descent** for the form of maintenance that connects recurring failures to measured improvements through five conditions:

1. **A measurable patch loss:** the deployment has a defined measure of success or failure.
2. **A failure log:** relevant failures are retained as evidence for later changes.
3. **Credit assignment:** failures can, at least sometimes, be traced to the scaffold coordinates responsible.
4. **A candidate edit:** the process proposes a change to those coordinates.
5. **An acceptance gate:** an edit is retained or promoted only after a check of whether it lowers the relevant loss.

The model weights remain fixed during this process, and the edits persist beyond the current interaction. These conditions distinguish three recurring forms of maintenance. In *patchwork*, an engineer adds a rule after an incident without measuring the result. In *random artifact search*, a system tests many variants and keeps a winner without tracing failures to their causes. Artifact-layer descent uses that trace to direct the proposed change and tests it before acceptance. Where the conditions are only partly met, the optimization interpretation is correspondingly limited.

For deployment D, write the patch loss as

$$L_D(\theta_s) = \mathbb{E}_{x\sim D}[\ell(\pi(\theta_M, \theta_s, x))]\,,$$

where θ_M denotes the fixed model weights, θ_s the editable scaffold, and ℓ the operational loss. A proposal process uses the failure history to produce an edit z_t. There is no analytic gradient through an instruction file or memory entry. Instead, the gate estimates the change in loss and accepts when that estimate clears a required margin:

$$G_t(z_t) = 1 \iff \widehat{\Delta}_{z_t} L_D(\theta_s^t) + \lambda\Omega(z_t) \le -\tau_t.$$

Here $\Omega \ge 0$ penalizes edit complexity and promotion breadth, and $\tau_t \ge 0$ is the acceptance margin. An evaluation that blocks a regressing merge is a practical instance of this rule. A reviewer or governance process can also instantiate the gate by requiring stronger evidence for a broader or more complex change.

Appendix A gives the update rule and a one-step descent inequality under assumptions on the gate's estimation bias. With additional coverage, compactness, and summable-bias assumptions, improving edits eventually stop clearing the gate. This is **gate stability**, not a claim of global optimality. A limited proposal process or biased evaluator can leave useful changes undiscovered or rejected.

2.4 Who proposes and who decides

The author of an edit and the authority that accepts it are separate. An agent may draft a revision that a person approves; an engineer may write a revision that an automatic evaluation rejects. We classify the gate by who makes the acceptance decision:

- **Auto:** a program applies a check, including a verifier, evaluation threshold, or language-model judge [2]. Whether that judge has been calibrated against people is a separate question, which Section 5 takes up.
- **Human:** a person decides whether to accept the change.
- **None:** no acceptance decision is made, either because the loop is absent or because every proposed edit is retained.

A pair such as [Auto × Human] denotes an automatic local gate followed by a human promotion gate. These labels describe the decision mechanism. They do not, by themselves, establish that the decision is well supported.

The loss definition imposes two limits on every such loop. First, an accepted edit is evidence of improvement only on the measure and tasks evaluated. Promotion must preserve that meaning or establish the improvement again. Second, the loop optimizes a defined objective. It does not discover a better objective simply by accumulating edits. Revising the evaluator changes the objective itself, a complication we return to in Section 7.2.

3 Corpus and method

We read each system's public method or documentation for four things: the persistent artifact it changes, the feedback that prompts the change, the decision that accepts it, and the scope in which the result is used. This allows systems with different names and application domains to be compared by their behavior.

3.1 Scope and selection

The corpus contains 256 unique systems from January 2023 through the October 2026 cutoff. The first search pass identified 130 systems by May 2026; the second added 126. Sources include research papers, public repositories, product documentation, and supported production reports; the search families, evidence tiers, and inclusion rules are recorded with the released data (Appendix C).

A system enters the survey if it adapts a foundation model through a persistent artifact updated from deployment signal, or if it is publicly described as self-improving, scaffolded, or agentic. Including the latter group provides a contrast set: the survey can test whether those descriptions correspond to persistent adaptation. Systems whose only update mechanism is weight-level training are excluded, as are systems without public evidence of release or deployment.

We apply two filters. The first admits a system to the **core corpus** if it scores at least 3 on at least one scaffold substrate under the rubric below. Of the 256 systems, 204 form the core corpus and 52 the contrast corpus. A system may have one integrated row or several substrate-specific rows, so the 256 systems occupy 269 scored rows.

The second filter is a **two-loop audit**. It asks whether session feedback changes a persistent artifact and whether that artifact is subsequently promoted beyond the deployment that produced it, with versioning or lineage. Applied to the core systems and the additional paper-based classifications (Appendix C), it admits 11 research systems and 3 industry exemplars. Fig. 2 summarizes the selection.

The research systems passing the audit are NanoResearch, SkillOpt, AlphaEvolve, ADAS, DGM, Voyager, SAGE, EvolveR, CoEvoSkills, AutoManual, and CASCADE. The industry exemplars are Sierra Agent OS 2.0, Cognition Devin, and Letta's Context Repositories. SkillOpt is classified

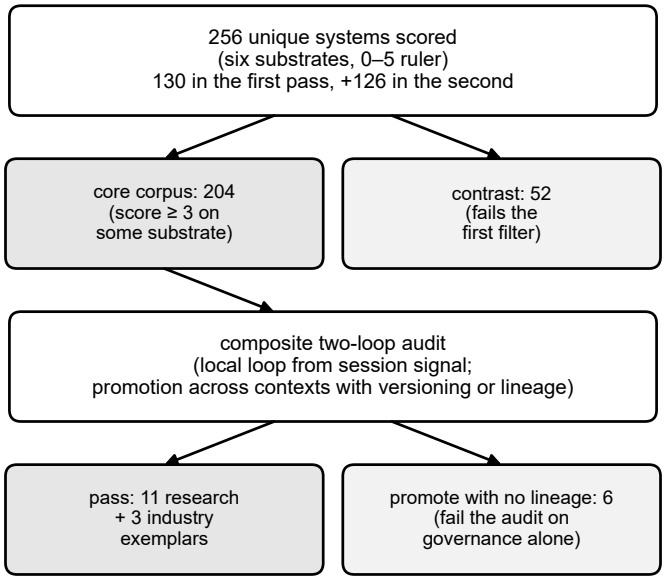


Figure 2: Corpus selection. The 256 systems occupy 269 scored rows. The first filter separates core from contrast systems; the two-loop audit then checks persistent local adaptation and promotion with lineage. The audit includes the additional paper-based classifications described in the text.

from its paper and is not a row in the 256-system score table. The remaining paper-based systems are identified in the data notes (Appendix C) and the audit's boundary cases in the contrast set (Appendix E).

Passing this audit does not establish governed sharing across organizations. All passing systems promote within one operator, into a shared bank or archive that later runs draw from, or across the agents and sessions of one deployment. Letta passes on the lineage of its shared-memory updates even though those updates have no acceptance gate. Six further systems promote across contexts without lineage on the promoted artifact; their treatment under relaxed audit rules is reported separately.

3.2 Substrates and scoring

The scaffold has six substrates: instructions, skills, memory, tools, orchestration, and governance. These are the possible targets of a proposed edit. Table 1 describes them; the episodic, semantic, and procedural memory categories follow Sumers et al. [3].

The common scoring scale runs from 0 to 5. At 0 an artifact is ephemeral. At 1 it persists locally; at 2 it persists with tool attachment but without feedback. At 3 it either spans scopes without a gate or is updated from feedback by a person. At 4 a program updates it from feedback. The score-5 requirement depends on the substrate, as specified in the table.

This last distinction is important when interpreting the scores. Instructions and skills require both loops to reach 5. Orchestration requires a topology updated from deployment feedback by a loop outside the agent's own execution. Memory, tools, and governance are scored for their corresponding governance capabilities, so a registry can score 5 without generating local edits. Integrated rows measure breadth and depth of adaptation, not closure of either loop. The scale therefore supports substrate-specific comparisons; a top score is not a universal certificate of complete optimization.

Table 1: The six scaffold substrates, what each holds, and what a system must do on it to score 5.

Substrate	What persists	A score of 5 requires
S1 Instructions	Project, team, or organization rules attached to a workspace	Versioned rules promoted under an automatic gate
S2 Skills	Procedures, as code or text, loaded on demand	A skill bank that edits itself on failure, test-gated, versioned, and promoted
S3 Memory	Experience kept across sessions: episodic, semantic, procedural	Provenance, versions, and a correction path for every entry
S4 Tools	The vetted action surface and its schemas	Evaluation-gated tool versions; a curated bundle as the product
S5 Orchestration	The multi-role graph and its routing	Topology updated from deployment feedback by a loop outside the agent's own run
S6 Governance	Versioning, promotion, audit, rollback	Staged promotion with evaluation thresholds, an audit log, and rollback

Table 2 gives the corpus distribution. Research systems hold the top marks on skills, memory, and the integrated row. Production platforms and standards hold them on tools and governance. No system reaches the top score across all six substrates, and none reaches it on instructions or orchestration.

For the 126 most recently added systems, gate types, feedback classes, and governance properties are recorded as explicit columns of the released table, so counts based on those columns apply to those systems. For example, 89 of their 105 core systems close the local loop under a program. This is not a gate count over all 256 systems.

3.3 Evidence and limitations

The evidence ranges from peer-reviewed studies and open-source artifacts to vendor claims. Every row of the released table carries its source and a one-line justification, and the rows for the 126 most recently added systems also carry an evidence tier. Peer-reviewed papers and open-source artifacts receive the greatest weight. A production claim must name a benchmark or integration target, and a blog or interview is counted only when stronger evidence corroborates it. Production claims are not pooled with peer-reviewed evidence in corpus-level statistics; vendor-based exemplars are identified as such.

One author scored all systems. Ambiguous cases were rounded down, and half-scores were used only when one criterion was clearly met and another partly met. Independent re-rating remains future work. The classifications should be read as an auditable interpretation of public evidence, not a complete inventory of private production behavior.

We also test the audit under alternative boundaries: counting rollback as sufficient promotion governance, allowing promotion without governance, and admitting loops that persist only optimizer state. Membership changes under these relaxations, but the principal structural findings remain: [Auto × Auto] stays well populated, the human-local cells hold only curation and registries, Letta alone passes with an ungated promotion, and no relaxation adds governed promotion across organizations (Section 8). The sensitivity analysis is released with the data notes (Appendix C) and the full score table with the released table (Appendix I).

Table 2: Substrate maturity across the 256-system corpus. Column n counts scored rows on each substrate, Mean is the average score, and Score-5 systems reach the substrate's top mark; the integrated row counts systems scored across several substrates as one row. Of the 126 most recently added systems, 105 are core and 32 are skill systems; 89 of those core systems close the local loop under a program, and none of the 126 scores 5 on any substrate.

Substrate	n	Mean	Score-5 systems	What the corpus holds on it
S1 Instructions	15	3.07	(none)	Rules persist across IDEs and agents and nothing measures them; no system closes both loops; ceiling 4
S2 Skills	50	3.34	SAGE, COSPLAY	The local loop closes under a program in the research systems; the production repositories merge by hand
S3 Memory	50	3.06	Cuadros et al.	The ledger ships in products; the correction a gate would enforce is absent
S4 Tools	20	3.65	MCP, Harvey, Hippocratic AI	The most mature production substrate; registries govern admission, not evaluation
S5 Orchestration	37	3.08	(none)	Topologies change under an evaluation gate; one structurally gated instance; no two-loop system with lineage
S6 Governance	45	3.02	Braintrust, Vellum, LangSmith Hub, AGENTS.md under the Agentic AI Foundation (AAIF)	Prompts-as-code mature; registries add lineage and rollback; no local loop generates a candidate
Integrated	52	3.81	NanoResearch, AutoAgent, SkillRL	Breadth and depth of adaptation, not the loop; reached only by research systems without enterprise governance

4 What local adaptation actually closes

A system can detect a failure, explain it, and write a plausible correction without ever establishing that the correction helps. SkillForge illustrates this distinction. Its judge agrees with people more than nine times in ten and identifies replies that should be rewritten. Those replies drive updates to the skill, but the updated skill is committed without an acceptance test [4]. The calibration concerns which replies are inadequate. It does not measure the effect of the resulting skill changes.

SkillRL exposes the same separation through a different mechanism. A success-rate threshold triggers proposals for new skills, and the skill bank accepts each proposal unconditionally [5]. Both systems respond to experience, yet neither has the local acceptance gate defined in Section 2. This is why the survey records who accepts an edit separately from who generates it or which signal triggers it.

Among the 105 core systems most recently added to the corpus, automatic local acceptance is nevertheless the dominant pattern: 89 use a program to retain or reject changes. Looking inside those loops reveals a less even picture. Evaluation is comparatively mature, while the failure record and the link from a failure to the responsible artifact are often incomplete.

4.1 Finding the right thing to change

A useful failure history must retain what went wrong in a form that supports later comparison. Production memory more often records preferences, successes, or conversation history. Only one field report in the corpus describes a store of dead ends intended to prevent their recurrence, and that evidence comes from a single self-reported deployment [6]. The persistent store may therefore be extensive while the particular record needed for failure-driven adaptation is missing.

Even when failures are retained, nearly every system uses them in arrival order. Few explicitly choose which recurring modes deserve attention. Task-CoEvolve's variance-weighted task sampling and AutoSaddler's mini-batch failure signals are exceptions [7, 8]. They weight the experience used for the next update rather than letting the most recent session determine it. The general loop described earlier would prioritize both recurrence and severity; this remains uncommon in the surveyed implementations.

The next step is harder. A failed output does not reveal whether its cause was an instruction, a retrieved memory, a skill, or the route through a group of agents. Harness-optimization papers repeatedly identify this attribution problem as the bottleneck [9–13]. An evaluation of harness optimizers reaches the same conclusion: by downstream score alone, a harness optimizer cannot be told from trial and error [14].

Existing solutions usually operate within one substrate. Meta-Harness preserves candidate harnesses and the record of which edits helped [11]. SkillsVote attributes outcomes to the skill, exploration, or environment before admitting an update [13]. SkillOpt-Lite keeps only the attributes that recur across rollouts, and its ablation finds this consensus mechanism load-bearing [15]. These methods give the proposal process more useful evidence. None solves the general attribution problem across all six scaffold substrates.

4.2 Testing the candidate

Research systems commonly generate candidates by reflecting on trajectories. In production prompt platforms, an engineer writes the candidate and session feedback never triggers one, while public

skill repositories depend on contributors. A field study of five such repositories finds that every substantive edit was committed under a named human account and merged by a person, most carrying an AI co-author trailer [16]. Where proposal generation depends on a person noticing a problem, the evaluation machinery can run automatically while the improvement process itself remains intermittent.

On skills, SAGE and COSPLAY close both local and promotion loops inside a research environment [17, 18]. SkillChain adds a human reflection decision after the judge, on which refined skills are deployed [19]. SkillRL and SkillForge, described above, leave the proposed changes unchecked. These differences would be obscured by classifying all three simply as systems that learn from feedback.

Orchestration has followed a similar progression. Early systems searched for agent topologies under an evaluation score [20, 21]. More recent systems evaluate proposed topology changes before keeping them; TOCOMAS also requires a structural invariant [22–24]. OptiMAS is an exception: its optimizer reads trajectories and validation scores to write the next configuration, but the score informs the edit rather than supplying a keep-if-better decision [25]. No orchestration system closes both loops with lineage on the promoted artifact.

The most widely deployed artifact remains a useful counterpoint. Instruction files are present across Claude Code, Cursor, AGENTS.md [26], Windsurf, and Continue.dev, and adoption studies find them in roughly a fifth to a quarter of coding-agent projects [27, 28]. Their prevalence does not settle their effectiveness. One controlled study reports higher inference cost without improved task success [29]. Another toggles the root AGENTS.md in paired runs on 124 pull requests and reports about 29% lower median wall-clock time and 17% fewer median output tokens, with total tokens approximately unchanged; it does not measure correctness [30]. The studies measure different outcomes and disagree on cost. Neither establishes a general task-success benefit from accumulating instruction files.

Human local gates appear only in curation. VerificAgent uses fact-checked, frozen memories [31], and in SkillChain a person decides, at a reflection gate, which refined skills are deployed. Prompt platforms and registries without a failure-driven proposal process have no local gate in this classification. Their checks act on release or admission to a shared pool, which belongs to promotion.

4.3 Gathering enough evidence to accept

An acceptance test can itself be too noisy to support a decision from one run. Several newer systems address this directly. SAGE-gate, a statistical acceptance method unrelated to SAGE, uses paired per-item comparisons before committing an edit [32]. VaG requires three heterogeneous critics to pass [33]. AutoSaddler uses mini-batch failure signals in validation-based selection [8].

These mechanisms should not be conflated. Averaging k independent noisy scores reduces the variance of their mean by a factor of $1/k$. Requiring agreement among k independent checks, each with false-acceptance probability p, gives the idealized probability p^k. The independence assumptions matter, and actual critics need not satisfy them. Appendix B distinguishes the gate families and their roles. The shared development is that acceptance increasingly depends on a body of evidence rather than one favorable observation.

4.4 Maintaining the artifacts that survive

The evaluation of a new edit also depends on the artifacts already present. Stale or conflicting skills can cause failures that obscure the effect of the candidate under test. Ratchet reports no gain from model-written skills over a no-skill baseline until lifecycle hygiene is applied, while human-curated skills provide a double-digit gain [34]. Eleven surveyed skill systems now treat removal as an operation that must pass a check, just as admission does [35–38] (Appendix G).

Memory products have made substantial progress on a different part of maintenance: recording an entry's history. Letta's Context Repositories use git, Mem0 retains per-memory history, and Zep represents facts bitemporally [39–41]. These records can explain what the system remembers and how that state arose.

History does not guarantee correction. A study of five production backends finds that all five still return facts that were revoked [42]. The record contains the change, but retrieval does not reliably enforce it. Research systems likewise divide between provenance and write controls [43–46] and verifier-gated failure correction [47, 48]. Cuadros et al.'s governed collaborative memory combines provenance, versioning, and correction, earning the top memory score [49].

Reversal has its own limits. SkillSmith and SkillAudit support edit-level rollback [36, 50], and EvoUndo rejects edits that cannot be undone [51]. VaG finds that reversing a contaminated skill recovers little once other skills have been derived from it [33]. The consequences of an edit may therefore outlast the edit itself. Section 7 returns to this problem when considering promotion and artifact descendants.

4.5 A complete local example

SkillOpt provides a concrete instance of all five conditions [52]. Its persistent state is one markdown skill document used by a frozen agent. Rollout and reflection batches supply the failure evidence. A consensus filter helps attribute recurring errors, and bounded add, delete, and replace operations propose changes. A held-out exact-match split supplies the loss, and a candidate is accepted only when its score strictly improves.

Across six benchmarks, seven target models, and three harnesses, SkillOpt is best or tied in every cell reported by its authors. Its follow-up, SkillOpt-Lite, removes a rejected-edit buffer and slow update schedule while retaining the held-out gate and consensus filter; it matches or exceeds the original [15]. This comparison supports the importance of the gate and this form of credit assignment in that setting, while identifying the removed controls as stabilizers rather than necessities.

SkillOpt also evaluates transfer across models, harnesses, and benchmarks. Those experiments measure how far the learned artifact remains useful. They are distinct from the acceptance decision that admits an edit to the shared skill bank. This distinction becomes central once the same local procedure is used in settings where correctness is harder to measure.

5 What the feedback allows a gate to know

A coding agent changes a repository rule and receives a test result within seconds. A scheduling assistant changes its booking procedure and learns later whether an appointment was made. A marketing assistant changes its tone instructions and receives a reviewer's preference. A research assistant changes its strategy and may wait months to learn whether the direction was productive.

All four systems can propose a persistent edit. What differs is the evidence available for accepting it. The signal's timing, exactness, and dependence on judgment determine how well a gate can estimate the change in loss. We divide the surveyed systems into four feedback classes, summarized in Table 3.

Table 3: The four feedback-signal classes. Reading across a row: what the gate receives, what it can therefore be, and what it cannot fix.

Class	The signal the gate receives	What the gate can be	What stays hard	Corpus examples
C1 exact	A check that returns pass or fail, or an exact score, at once and repeatably: tests, compilers, answer keys	An automatic check; one run per case, repeated only to sample the task mix	The lift is bounded and specific, an improvement in what the check measures on the cases it ran; the check names the failed output, not the coordinate that caused it; a check that is also the target can be gamed	DGM, AlphaEvolve, SkillOpt
C2 observable outcome	A real result one episode later: a booked slot, a merged pull request, a closed ticket; noisy and confounded	An automatic threshold over batched episodes; a person at promotion	The outcome says *that* the episode failed, not *why*; the channel can be poisoned	Devin, Sierra, SAGE, SkillRL, PACE, Voyager
C3 judged	A rater, a rubric, or a model judge; no ground truth	A judge calibrated against people; a person at promotion	The rubric is itself a scaffold coordinate, so the gate judges edits to itself; the proxy gets optimized past the point where it tracks quality	SkillForge, CASCADE, skills-repository review
C4 delayed	An outcome weeks or months later, once per many edits	No per-edit gate; proposals are pruned and milestones stand in	Credit across many edits; drift; the loop cannot be told from random search	NanoResearch, CORAL, the AI Scientist

The classes are assigned by the constraint that binds first. A long delay makes a signal C4 even when the eventual outcome is objective. Among timely signals, the absence of an independent ground truth makes the signal C3. Among timely objective signals, repeatable exact checks distinguish C1 from noisy observed outcomes in C2. A system may use several signals, but its operative acceptance decision must be read against the signal that supports it.

5.1 Exact checks: C1

A compiler, proof checker, schema validator, or exact-match answer key can decide a specified property immediately and repeatably. This makes an automatic gate straightforward to implement. AlphaEvolve scores candidates through a cascade of exact evaluators [53], and SkillOpt uses held-out exact-match evaluation [52]. DGM makes a narrower acceptance decision: it retains self-modifications that pass a sandboxed compile-and-edit functionality check on ten benchmark tasks. The broader benchmark score weights parent selection from the archive rather than determining retention [54].

Exactness applies to the check, not to the deployment as a whole. Repeating a deterministic check on the same output yields the same answer, but an agent produces stochastic outputs, and a finite suite samples only part of the task distribution. Estimating the change in patch loss still requires accounting for that variation. Passing the suite also establishes only the properties the suite tests.

Two problems remain even with a reliable checker. A failed test may not identify the instruction or skill that caused the failure, and repeated optimization against the same suite can overfit it. Cheap acceptance can also allow the artifact pool to grow without sufficient pruning. Exact feedback makes evaluation easier; it does not supply attribution or lifecycle management.

5.2 Observed outcomes: C2

A booked appointment, closed ticket, merged pull request, or successful tool call provides an observable result after an episode. The result is meaningful but may be noisy or confounded. A user can accept an incorrect answer or reject a correct one, and a merged pull request reflects more than the agent's technical performance.

Devin reads merged pull requests as feedback [55]. SAGE measures task success in AppWorld [17], and PACE uses failure signals from a deployed small model [56]. SkillRL uses a success-rate threshold only to trigger proposals, without testing whether the proposed skills should remain [5]. The same kind of outcome can therefore support different operations in different systems.

Simulated outcomes, such as those in Voyager's game world or SAGE's AppWorld, remove the user confound and can be sampled repeatedly. They are treated as a distinguishable case within C2. For observed outcomes more generally, an automatic gate needs a positive margin and enough episodes or independent evidence to absorb noise. Even then, the result usually reveals that the task failed more readily than why it failed.

Promotion is harder because another deployment may differ in users, tasks, and acceptance behavior. A sender's successful outcome is not an evaluation of the recipient. This class also depends on channels that can carry poisoned memories or injected tool instructions, connecting the reliability of feedback to the integrity of its source.

5.3 Judgments: C3

For style, persuasion, design, or other preference-dependent tasks, a gate relies on a person, rubric, reference comparison, or model judge. There is no independent exact answer to the preference being measured. CASCADE uses a human-and-agent gate for memory consolidation [57], and a reviewer approves promotion into Anthropic's skills repository. SkillForge uses its calibrated judge earlier in the process, to identify replies that motivate a rewrite [4].

A program can make these decisions, but its score must be calibrated against the human judgment it is intended to represent before the score is treated as evidence of that preference. The gate labels deliberately do not record this distinction: an uncalibrated model judge is still an Auto gate, though the meaning of its acceptances is less established. SkillForge, SkillChain, and SkillCoach report calibration measurements [4, 19, 58]; no other surveyed judge does.

Optimization introduces another risk. A scaffold can become better at satisfying the judge while becoming worse at the underlying task. Reward-model research measures this form of proxy overoptimization [59]. Human evaluation is not fixed either: Shankar et al. find that evaluators'

criteria change as they grade more outputs [60]. If the system also edits the rubric, it changes the instrument used to assess subsequent edits. This **reflexive gate** is developed in Section 7.2.

5.4 Delayed and sparse outcomes: C4

Research programs and open-ended discovery may produce meaningful outcomes only after weeks or months and many intervening edits. The AI Scientist illustrates the broad setting by generating ideas, conducting experiments, and writing a paper [61]. The eventual outcome cannot usually be assigned to each preceding scaffold change.

Systems in this class consequently use intermediate decisions. NanoResearch merges skills by semantic compatibility [62]. CORAL promotes findings into shared memory on a heartbeat [63]. These operations can prune or consolidate proposals, but they do not estimate each edit's contribution to the eventual loss. AlphaEvolve's archive diversity does not make it a delayed-signal example: its candidates have already been scored by exact evaluators, so archive selection is a search strategy within C1 [53].

For C4, the descent interpretation applies only to whatever proxy gate is available or to an evaluation at the longer horizon. Temporal credit assignment remains unresolved, and the environment can change before the result arrives. Outcome gains alone are also insufficient to distinguish directed improvement from a successful search, a limitation emphasized in the evaluation of harness optimizers [14].

5.5 Consequences for the two loops

The classification yields three expectations that can be checked against the corpus. First, local automatic acceptance is most directly supported by exact checks and observed outcomes. Judged acceptance depends on calibration, while delayed feedback cannot evaluate edits individually. Second, automatic promotion that estimates a loss change should occur only where a comparable check can be applied at the recipient, in exact and re-runnable simulated settings, and human promotion should hold the observed outcomes and judgments. Delayed feedback supports no per-edit gate on either half. Third, the main failure risks should differ: overfitting to a suite and bloat under exact feedback, corrupted outcome channels under C2, proxy optimization under C3, and drift under C4.

These are expectations about the evidence a gate can establish, not restrictions on whether software can execute a decision. A program can merge two judged artifacts or prune a delayed proposal without estimating a loss change. Section 6 makes this qualification explicit in the gate grid. In every class, sharing remains an additional empirical question: even the same exact checker must be applied to the recipient's relevant cases before broader performance is established.

6 When a local improvement is shared

A rule can be useful in one repository and inappropriate in the next. Copying the file transfers the instruction, but not the evidence that justified it. The recipient must either share the relevant evaluation conditions or test the change on its own workload. This is the additional decision that separates local adaptation from promotion.

6.1 Scope and receiving-side evaluation

We distinguish five nested scopes. **USER** covers one person's account or fork. **PROJECT** covers a repository, workflow, or product workspace. **STACK** is a team's curated bundle of projects. A **TENANT** is an organization or customer, and **CORE** is the scope shared by all tenants. A wider scope aggregates the losses of the deployments it contains. Promotion changes which population the edit is intended to improve, as shown in Fig. 3.

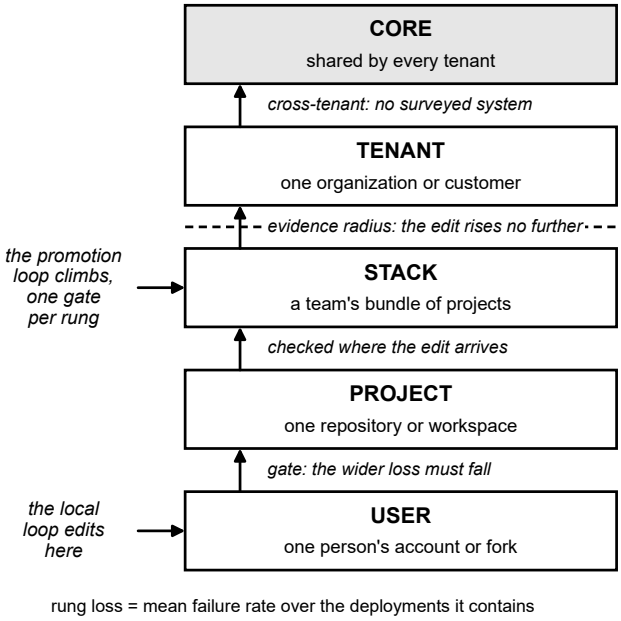


Figure 3: Promotion scopes, from a user account to the shared core. Wider adoption changes the population on which an edit acts. Receiving-side evaluation checks whether the improvement is supported for the recipient.

The nesting makes an important limit visible. A lower average loss at a broad scope can conceal harm to an individual recipient. The proposed architecture in Section 8 therefore includes a held-out check at the receiving tenant. Versioning and access control govern whether an edit may be delivered and reversed; receiving-side evaluation tests whether it should be used there.

Several systems already evaluate transfers within one operator. ShareMem consolidates memories from one user for another and evaluates them in the receiving context [64]. XRepoSkill verifies a rule in at least two repositories [65]. Evo2Team promotes routing skills between teams and confirms them at the target [66]. These are concrete instances of evidence being gathered beyond the source deployment.

Federated systems place the check elsewhere. EvolveNet merges a shared harness only when the merge fixes at least as many held-out items as it breaks, but that check occurs at the aggregator [67]. FederatedSkill's server agent applies, adapts, or rejects each client's edit and records the decision [68]. Their clients are partitions or agent instances under one operator, not independently governed customer organizations. They do not instantiate a receiving-organization gate.

The location of the evaluation matters, and its coverage still matters after it is correctly located. ShareMem, XRepoSkill, and Evo2Team check the receiving context; EvolveNet checks the shared

merge; FederatedSkill records a server decision. Even in Evo2Team, nearly a fifth of confirmed transfers fail on graphs not covered by confirmation. A receiving-side check provides necessary evidence for adoption, but finite evaluation does not guarantee performance on all future tasks.

6.2 Automatic, human, and unchecked promotion

Research systems often promote by admitting a candidate to a bank or archive that later runs reuse. SkillOpt, DGM, ADAS, Voyager, and SAGE apply programmatic decisions within this shared environment [17, 52, 54, 69, 70]. AlphaEvolve and its predecessor FunSearch score candidates before archive admission [53, 71]. AlphaEvolve has produced gains used within Google, but it is a research loop within one operator, rather than a product for governed sharing between organizations. That pool is as far as any of them travels, and none has access control, audited promotion, or a gate at a receiving organization.

Not every automatic admission has the same evidentiary meaning. NanoResearch's semantic merge, AutoManual's judged rewriting and consolidation, and ShareMem's judged consolidation are programmatic decisions, but none of them estimates a candidate's reduction in patch loss. The distinction anticipated by the feedback classes is therefore essential: an automatic promotion decision can exist where a per-edit improvement estimate does not.

Production systems more often retain a person at the promotion decision. Sierra routes generated insights through human-authored expert answers [72]. In Devin, the same pull-request merge serves two roles: it supplies an observed outcome to local adaptation and records the reviewer's approval for the persistent repository [55]. Spotify's optimized prompts and plans enter production after online A/B testing and a launch decision [73]. CASCADE also combines human and agent approval for memory consolidation [57]. In each, the person at the promotion gate stands in for the exact signal production lacks, which makes a program on the local gate with a person on the promotion gate the production-viable pairing rather than a compromise.

Letta's Context Repositories take a different path. Shared-memory changes are pushed with git lineage but without an acceptance check. The record explains what was shared, while leaving whether it should have been shared unresolved. In many other systems, promotion is absent altogether. AutoAgent, for example, revises the tool and peer descriptions of one agent without propagating those revisions to another deployment [74]. Both cases receive None at the promotion gate, but for different reasons: one shares without checking; the other does not share.

6.3 What formats, registries, and release platforms provide

Much of the infrastructure needed for sharing already exists independently of these learning loops. Claude Agent Skills, Cursor Rules, GitHub Copilot instructions, Windsurf, Continue.dev, Zed, and Replit use git-tracked markdown artifacts attached to workspaces, with a tool layer organized around MCP-style interfaces [28, 75]. This convergence makes artifacts easier to carry between runtimes. A study of repository-level loops nevertheless finds that the runtime state of the adaptation process remains outside version control [76]. The portable file does not necessarily carry the history that produced it.

Public repositories provide a shared scope with human admission. In a study of five skill repositories, every substantive change was committed under a named human account and merged by a person; only one of 254 edits deprecated a skill [16]. MCP registries provide namespaces and version records, while control planes add scanning, access control, or quarantine [77–79]. Governance registries add

lineage and rollback [80–82]. EDDOps gates promotion on evaluation thresholds before a human sign-off, but none of these registries generates candidate artifacts from failures.

Braintrust, Vellum, and LangSmith Hub likewise provide immutable versions, evaluation thresholds, and rollback [83–85]. Engineers author the candidates, and evaluations can block a regressing release. In the surveyed descriptions, session feedback does not trigger candidate generation or the release process. These platforms supply much of the machinery for controlled change without closing the failure-driven local loop.

6.4 The two decisions viewed together

Fig. 4 places the local and promotion gates on a common grid. Table 4 summarizes the occupants and their signals. The grid distinguishes a system that generates and tests improvements from one that distributes human-authored artifacts under admission control, while allowing their complementary capabilities to be compared.

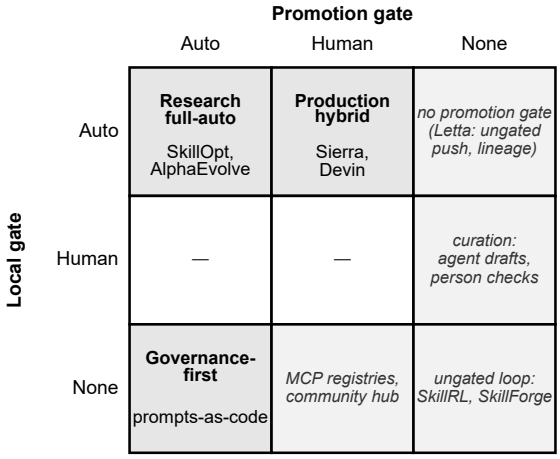


Figure 4: Local and promotion gates. Each axis records who accepts a change: a program, a person, or no gate. None includes both absent promotion and unchecked sharing; the text distinguishes these cases.

Among [Auto × Auto] systems whose promotion gate estimates a loss change, the operative signals are exact checks or simulated outcomes. The judged and delayed cases use pruning, rewriting, or merging rather than that estimate. [Auto × Human] contains observed outcomes and judgments, with no exact-signal system in the surveyed cell. These patterns match the feedback-based expectations once the meaning of automatic admission is made explicit.

The [Human × Auto] and [Human × Human] cells are empty in the surveyed set: no system combines a human local gate on a failure-driven loop with subsequent promotion. Human local decisions occur in curation. The bottom row includes comparison systems without a failure-driven local loop, such as prompt platforms, registries, and community hubs including Buildrix [86]. SkillRL and SkillForge occupy [None × None] because their failure-driven skill updates, new skills in SkillRL and rewrites in SkillForge, are kept unchecked.

The audit's boundary cases do not change this structure. Of six systems that promote without artifact lineage, five have sufficient evidence under the tier rule and enter already occupied cells

Table 4: The gate grid read cell by cell: who keeps a local edit and who lets it climb, which systems sit in the cell, and the feedback signal they gate on. [Human × Auto] and [Human × Human] hold no surveyed system.

Cell	Who keeps the local edit, who lets it climb	Who is there	Signal found there
[Auto × Auto]	A program keeps it; a program lets it climb	Research systems that close both halves inside one environment or one operator	Exact and simulated outcomes; delayed and judged only under pruning or merging
[Auto × Human]	A program keeps it; a person lets it climb	Production systems with a reviewer on promotion, and CASCADE	Observed outcomes and judgments
[Auto × None]	A program keeps it; nothing lets it climb	Letta's ungated push; systems that promote nothing	As on the local gate
[Human × None]	A person keeps it; nothing lets it climb	Curation	Judged
[None × Auto]	No failure triggers a candidate; a threshold gates release	Prompts-as-code platforms and governance registries	No failure-driven loop
[None × Human]	No failure triggers a candidate; a person admits	Open standards, MCP registries, community hubs	No failure-driven loop
[None × None]	Nothing checks or promotes the edit	SkillRL and SkillForge, unchecked skill loops	Outcomes and judgments

when the audit is relaxed; Spotify is one of them. The sixth, Future AGI's prompt lifecycle, has only uncorroborated T5 evidence. No relaxation yields governed promotion across organizations. What the grid reveals is a division of capabilities: no system has both the learning of the top row and the governance of the bottom row.

7 How adaptation can undermine its own evidence

An improvement loop changes the context in which its next decision is made. New rules can conflict with old ones, a corrected memory can still be retrieved in its earlier form, and a frequently reused evaluator can become a target. These failures extend the technical-debt concerns identified for machine-learning systems by Sculley et al. [87]. The literature documents six of these failures in a measured case; a seventh, overfitting to one deployment, follows from the scope of the loss rather than from a measurement. Table 5 connects each to the control that addresses it.

7.1 Accumulation and contamination

As instructions accumulate, following them becomes harder and latency increases [88]. Ratchet's finding that model-written skills need lifecycle hygiene before producing gains makes the consequence visible at the system level [34]. Conflicting or obsolete artifacts also complicate attribution: an evaluation may detect failure without revealing whether the cause is the new candidate or an older artifact. Gated retirement, pruning, and expiry preserve the conditions under which a new edit can be meaningfully assessed.

Table 5: The failure surface of patchwork scaffold maintenance: what goes wrong, the piece of the loop that answers it, and the feedback-signal class in which it dominates.

Failure	What goes wrong	The loop piece that answers it	Dominant class
Scaffold bloat	Rules and skills accumulate until the model follows them worse and the gate has no clean signal	Complexity penalty Ω in the gate; pruning and expiry as first-class gated operations	C1, C3
Poisoned memory	An attacker writes an entry that later retrieval trusts	Provenance on every entry; write gating; a rollback envelope	C2
Injection via tools and MCP	A tool's output or a connected server carries instructions into context	Narrow tool permissions; a privilege hierarchy over instructions; tool bindings pass the same promotion gate as any other coordinate	C2
Evaluation fragility	The suite that gates edits is itself edited, gamed, and read non-independently	Held-out evaluations scoped within evidence radius, weighted toward recurring modes, re-read as the target drifts, checked against attribution	C3 (primary), C1
Coordination failures	Multi-agent roles blur and context is lost between agents	Orchestration topology treated as a scaffold coordinate under the same gate and rollback discipline	C2
Patch overfitting	A scaffold tuned to one deployment degrades another	Evidence-radius discipline at promotion; receiving-side held-out evaluation	C2
Plasticity loss	Old rules, polluted memory, and drifting tool versions crowd out adaptation	Expiry and rollback as explicit loop operations rather than retrospective hygiene	C4

Persistent memory creates another route to lasting error. PoisonedRAG demonstrates attacks on retrieved knowledge [89], AgentPoison plants backdoors in long-term memory or knowledge bases [90], and memory control-flow attacks operate across major agent frameworks [91]. Tool outputs and connected servers can also insert instructions into the model's context [92, 93]. In CVE-2025-54136, a previously approved MCP configuration could be changed without renewed approval [94].

These failures motivate provenance, write gating, and reversible memory changes [95, 96], together with restricted tool permissions and an instruction privilege hierarchy [97]. The tool binding itself must be subject to controlled revision. Checking only the initial registration leaves later changes outside the decision that authorized it.

Rollback also needs to account for derivation. If a contaminated skill has been used to create other skills, restoring its earlier version does not remove those descendants. VaG's limited recovery after such rollback [33] supports checking contamination before commitment and tracking consequences after it. A version history is a prerequisite for this work, but it does not perform the recovery by itself.

7.2 When the evaluator changes

Consider a pitch-review assistant that consistently overrates decks claiming a large market without supporting numbers. The team revises its rubric from "evaluate market size" to "require a numeric estimate and a source." The edit may improve reviews, but it also changes the standard used to evaluate subsequent behavior. If the same rubric judges future rubric revisions, the improvement process is modifying its own measuring instrument.

This is the reflexive-gate problem introduced with C3 feedback. The survey finds three scaffold systems reporting calibration of a judge against people: SkillForge, SkillChain, and SkillCoach [4, 19, 58]. None reports calibration for a judge that is itself being edited. The missing result is therefore more specific than an absence of calibrated judges. It concerns maintaining calibration while the criterion changes.

Weight-level reinforcement learning has encountered a related problem in extending verifiable rewards to open-ended tasks. Some methods hold the judge fixed, using checklists, pairwise preferences, references, or rubrics [98–101]. Others update the reward, rubric, or judge alongside the policy [102–106]. Among the co-evolving methods reviewed, only RLPR reports agreement with human labels, and its reward is a probability rather than a rubric. The comparison is tabulated in the audit notes (Appendix G); these weight-side methods are not included in the scaffold scores.

Even a fixed evaluator can become unreliable through repeated use. Optimizing a proxy beyond a point can reduce the underlying quality [59]. Unequal exposure to evaluation data can create apparent performance gains [107]. Evaluations of agent trajectories are also path-dependent, and deployment tasks drift while the suite remains fixed. Human criteria themselves can shift during grading [60]. Consequently, a gate needs its own held-out checks, calibration where judgment is involved, and revision history.

The failure-mode hypothesis discussed in Section 9.1 adds a conditional implication: if deployment failures concentrate in recurring modes, uniformly sampled tasks may underrepresent the modes that most need correction. Evaluation should then include those modes while preserving a held-out assessment. The requirement is broader than improving a benchmark score: score alone cannot tell a harness optimizer from trial and error [14], so the evaluation must also test the claimed link between failures and edits.

7.3 Scope, coordination, and drift

Some failures arise from the structure of the deployment. A study of multi-agent systems identifies fourteen failure modes, prominently including role confusion and context lost between agents [108]. Treating routing and role assignments as editable scaffold coordinates makes them subject to evaluation and rollback alongside instructions and skills.

Specialization creates a related boundary. A scaffold adapted to one deployment can degrade another because the tasks and preferred behavior differ. Receiving-side evaluation limits this risk at promotion; the evidence radius records where the improvement has been tested. This does not eliminate uncertainty outside the evaluated cases.

Finally, the scaffold can become less adaptable as old rules, polluted memories, and drifting tool versions accumulate. Long-context degradation [109] provides one related mechanism. Expiry, pruning, and rollback must remain available throughout the system's lifetime, especially when delayed feedback arrives too slowly to keep the artifacts current.

The feedback classes help organize these risks without making them exclusive to one class. Exact checks invite suite overfitting and cheap accumulation; observable outcomes depend on channels vulnerable to poisoning and injection; judged feedback is exposed to proxy optimization and moving criteria; delayed feedback makes drift difficult to correct promptly. Production systems provide individual controls, but the surface stays largely unmitigated because no surveyed system runs them together.

8 What governed sharing would require

The evidence so far identifies a practical integration problem. Local optimizers can propose and test changes. Release platforms can version artifacts and reverse deployments. Registries can govern admission and access. The surveyed systems do not connect these capabilities into a process that learns in one organization and shares the resulting improvement with another under equivalent controls.

We audit that missing combination through four properties:

(a) **An auto-gated local loop:** an algorithmic acceptance criterion tests a proposed artifact change.
(b) **Cross-organization promotion:** artifacts can be promoted between independently governed organizations or customers.
(c) **Versioned lineage with access control:** a promoted artifact carries its history and audit trail, with policies controlling who may access it.
(d) **Rollback:** a harmful promotion can be reversed without redeploying the whole system.

The terms *cross-tenant* and *cross-organization* refer to the same boundary here. Sharing among agents, users, or benchmark partitions within one operator does not establish property (b).

Table 6 compares the closest systems. FederatedSkill combines local adaptation with semantic skill exchange and an auditable server decision log [68]. EvolveNet evaluates shared merges and can return to the previous shared harness [67]. Both remain within one operator. Governance registries supply lineage and rollback without a local learning loop, and MCP control planes supply access control alone. Rollback exists inside one federated round and inside single deployments; no system

rolls back a promotion that has crossed an organization boundary, because none makes such a promotion.

Table 6: The closest approaches to governed cross-organization scaffold optimization, scored against the four properties.

System	(a) auto-gated local loop	(b) cross-organization promotion	(c) versioned lineage with access control	(d) rollback
FederatedSkill [68]	full	partial: privacy-preserving aggregation of skill edits across heterogeneous clients that sit under one operator	partial: auditable per-modification decision log	none
EvolveNet [67]	full	partial: held-out merge gate at the aggregator, but clients are benchmark partitions, not organizations	none: no lineage, no access control	full: rollback to the previous shared harness after a failed round
SkillForge [4]	none: a calibrated judge selects the failures, nothing tests the rewrite	none: skill versions stay in the scenario deployment	partial	not documented
Governance-first registries [80–82]	none: no inner loop	none	partial: lineage without access control	full, with evidence
MCP control planes [77]	none: no inner loop	none	partial: access control without lineage	not reported

Receiving-side evaluation is an additional requirement on the promotion decision. It is kept separate from the four-property audit, which records loop, scope, and governance capabilities. A system could satisfy those capabilities while still making an unsupported decision about whether an artifact benefits its recipient.

An architecture that addresses both would contain five connected parts: per-tenant artifact pools with provenance; an aggregation or selection process that respects tenant isolation; held-out evaluation at the receiving tenant; versioned lineage and access control; and a defined reversal path for each promoted artifact. The recipient's check would reject a promotion that does not lower its own patch loss. VaG's result (Section 4) suggests that the rollback envelope would also need to track an artifact's descendants.

Federated learning offers a useful comparison, but the representation being shared is different. FedAvg combines comparable numeric updates, and its privacy and deployment problems have been extensively studied [110–112]. Federated prompt tuning moves the learned object to a shared prompt [113]. The nearest scaffold systems instead select, revise, or merge symbolic edits, often through a language model. Averaging is an appropriate analogy only where the artifact updates have a mergeable representation.

The extent of automation should depend on the feedback. Exact checks provide the most direct setting for testing a transferred skill, although the recipient still needs relevant test cases. Observed

outcomes require evaluation on the recipient's workload because the sender's outcome is not portable evidence of success. Where quality is judged, a human promotion gate is a design feature, supported by lineage and rollback. Delayed outcomes cannot support the per-edit automatic local gate in property (a); those systems still benefit from provenance and reversibility around proposals and milestones.

Thus the four properties define a target for the settings that can support an automatic local gate. They should not be interpreted as a requirement to automate every decision in every domain. The common requirement is that the evidence attached to an edit remain meaningful where it is used, and that the system preserve the ability to inspect and reverse its adoption.

9 How the research lines fit together

The uneven architecture has a history. Different research communities began with different problems: improving prompts, learning in context, selecting agent workflows, maintaining skill libraries, or transferring experience. Each developed the part of the loop most relevant to its initial problem. Table 7 places five such lines side by side.

Table 7: Five research lines converging on artifact-layer descent: the problem each started from, the part of the loop it built, and what it leaves open.

Research line	Started from	The part of the loop it built	What it leaves open
Reflective evolution	Prompt optimization and continual learning	Condition 4, the candidate: proposal by reflection over trajectory evidence; the two-timescale case for editing outside the weights	A persistent artifact; a bounded step; a held-out gate
Reflective context	Learning in context as a setting of its own	Condition 5, the penalty: the learned object in context space; a generalization penalty, the Ω of the gate	A cross-deployment promotion loop
Harness optimization	Orchestration and harness engineering	Condition 3, the trace: attribution of a failure to one slot of the orchestration graph; source-level edits	Attribution across substrates; a governed promotion loop; access control; rollback
Lifecycle and multi-objective	Skill pools that bloat under platform limits	The gate on removal: gated retirement; hygiene as a precondition; a hard frontier as the shape of Ω	Closing the loop as descent; retiring a promoted artifact's descendants
Transfer and federation	How far an edit travels	The promotion loop, the check on arrival: evidence radius measured one task at a time; cross-client aggregation; receiving-side confirmation	A cross-organization gate, lineage, and rollback

Prompt optimization and reflective evolution developed ways to propose edits from trajectory evidence, together with an argument for adapting artifacts on a faster schedule than weights [114, 115]. Reflective Context Learning formalizes learning in context and the need to control generalization, corresponding to the complexity penalty in the acceptance rule [116]. Harness

optimization treats source-level orchestration changes as the learned object and repeatedly identifies attribution as its difficult step [9–12, 23, 117, 118].

The lifecycle literature starts from a different constraint: skill pools can become too large or internally inconsistent to remain useful. Its contributions include evaluated retirement, hygiene, and optimization under platform limits [13, 34–38, 119]. Transfer and federation ask how far a learned artifact remains useful, progressing from cross-task measurements to receiving-team confirmation and shared client updates [56, 63, 66–68, 120, 121].

This history helps explain why evaluation and proposal mechanisms are more developed than failure capture. Each research line inherits its loss from its benchmark and begins with available trajectories. Building a record of recurring failures was none of their starting problems. Attribution receives substantial attention, but usually within the substrate a particular line studies. The local loop has consequently developed from evaluation backward toward the evidence needed to choose a useful edit.

The convergence supports a shared decomposition of the system into weights and editable artifacts. GEPA, Meta-Harness, MOSS, RCL, PACE, and the harness-optimization papers all use it. A system that updates both, SIA, still treats them as separate levers [122]. The convergence does not establish that one architecture is optimal, and none of these lines documents the complete cross-organization process. It shows that different starting problems repeatedly lead to comparable mechanisms for proposing, testing, retaining, and reusing changes.

9.1 Why specialization remains outside the weights

A general model and a particular deployment face different requirements. The model must work across many settings. A deployment must follow its own conventions, honor its users' preferences, and avoid failures that may be rare globally but frequent locally. The scaffold provides a place to encode those requirements with a scope narrower than the model's entire user population.

Work on reinforcement learning from human feedback reports reduced output diversity and more homogeneous preferences [123, 124]. This helps motivate the distinction between broad model behavior and customer-specific requirements. Tiwari et al. approach it from adaptation cost: in-context changes are cheap but capacity-bounded, while in-weight adaptation is more expressive and can introduce forgetting and loss of plasticity [115, 125].

Per-customer adapters [126], distilled models, and learned routing can absorb some local requirements. They also introduce separate training, evaluation, and serving lifecycles. In practice they are rare. Adaptation lands in persistent artifacts, the one substrate that is cheap, inspectable, reversible, scoped to one customer, and updatable daily.

Two earlier papers in this series offer a hypothesis for why targeted artifact changes may be effective. *Beyond Exponential Decay* argues that long-output reliability depends on a sparse set of key decisions rather than a uniform error probability at every token [127]. Its evidence includes long-range dependence measured on one natural-text corpus, where fewer than a tenth of tokens depend on distant context, and the accuracy recovered through majority voting on reasoning benchmarks. Its argument is limited when the number of consequential decisions grows with the task, as in long-horizon agentic work.

The Architecture of Errors considers a bounded deployment with recurring tasks, tools, and evaluator expectations [128]. It argues that failures there fall into a catalogue of modes whose size grows slowly

with observed experience. Under its assumptions, the intervention budget grows polylogarithmically with sequence length and becomes constant once the catalogue saturates. Its negative result is equally relevant: a finite catalogue cannot cover an unbounded domain.

If these hypotheses generalize, a deployment may be improved through a manageable set of targeted interventions. Persistent rules, memories, skills, and tool changes provide an inspectable representation for those interventions. On that hypothesis the gated and scoped artifact layer is not a workaround for an inadequate model but the architecture the error structure calls for: a frontier model that generalizes, a local scaffold that specializes, and reliability that comes from governing the specialization. The survey does not test the hypotheses; its findings about acceptance gates, feedback, and promotion follow from the corpus independently of them.

9.2 Relation to neighboring surveys

Existing surveys organize the field around agent evolution, persistent state, context engineering, skills, memory security, and harness design. This survey follows a different unit of analysis: the decision that accepts a persistent change, the feedback supporting that decision, and the scope in which the change is subsequently used. Table 8 records the distinction for each neighboring survey.

Table 8: Neighboring surveys and where this one departs from each. This survey classifies by loop closure, gate type, signal class, and governance, and is the only one to cover promotion across organizations.

Survey	Scope	How this survey differs
Gao et al. [129]	self-evolving agents	organizes by what, when, how, and where an agent evolves; the gate that accepts an edit, and the edit's scope, are mentioned but not classified
Fang et al. [130]	self-evolving agents	organizes by inputs, agent system, environment, and optimizer; loop closure appears in passing, gates and promotion do not
Zhou et al. [131]	self-evolving coding agents (108 methods)	organizes by evolution target, timing, and code signal; closest on reversibility, with no gate or promotion axis
Ding et al. [132]	persistent-state agents (435 works)	organizes by authority, scope, mutability, provenance, and recoverability of state; governance is an axis, the gate on an edit is not
Lin et al. [133]	memory security (~90 works)	organizes by lifecycle phase and security objective; principal scoping and rollback are axes, feedback and promotion are not
Guo et al. [134]	agent system and harness design	organizes by six runtime responsibilities; touches all four questions without classifying by any
Xia et al. [135]	skill-mediated agents (8 systems)	organizes by four responsibility layers and ten patterns; loop closure and governance appear as patterns, not as axes
Ding et al. [136]	skill evaluation and evolution	organizes by four evolution paradigms and six benchmark classes; evaluation is the subject, the gate it feeds is not
Mei et al. [137]	context engineering (1,411 papers)	organizes by components and system implementations; adaptation loops lie outside its scope

The closest governance precedents are Ding et al.'s survey of persistent-state agents and Lin et al.'s memory-security survey [132, 133]. The former finds greater attention to accumulating and retrieving state than to governing, recovering, or relinquishing it. The latter treats scoping and

rollback as security primitives. Neither asks who gates an edit, and neither covers promotion across organizations.

Prompt optimization supplies another important lineage: OPRO, EvoPrompt, Promptbreeder, DSPy, and TextGrad optimize a prompt or program once against development evidence and are not deployed loops [138–142]. ACE persists an evolving playbook but never promotes it [143]. Agent symbolic learning treats prompts, tools, and pipelines as symbolic parameters [144]; STOP improves its own scaffolding program [145]; and SICA edits its harness to improve benchmark performance [146]. These works establish many of the underlying mechanisms.

The contribution here is their joint reading at deployment scope: whether an artifact persists, whether a local acceptance decision is made, which signal supports it, and what authorizes use elsewhere. The 256-system corpus, two-loop classification, feedback classes, and governance audit make that comparison explicit. An extended discussion separates these contributions from their precedents (Appendix H).

10 Open problems and conclusion

The remaining problems become more specific when expressed in terms of the decisions a deployment must make.

Attributing a failure to an editable cause. An exact check can establish that an output failed without identifying the instruction, memory, skill, or tool binding responsible. No surveyed method solves this across all substrates. Slot attribution [12], intervention-based attribution [22], and contrastive multi-agent attribution [10] provide narrower starting points. Observed outcomes make the problem harder because their causes are more confounded.

Establishing improvement at the recipient. A transferred artifact must be evaluated against the receiving deployment's tasks and measure. Existing receiving-side checks show that this is feasible in selected settings, while Evo2Team's failures on unseen graphs show its coverage limits. Federated Agent Optimization proposes a receiving-client gate but does not implement it [147]. Connecting such checks to provenance, access control, and reversal remains an architectural problem.

Maintaining calibration while the criterion changes. For judged tasks, the open question is how to keep a judge aligned with people while revising the rubric it applies. No surveyed scaffold system reports this measurement. The neighboring reward literature also lacks a reported human-agreement result for a co-evolving rubric or judge, with RLPR's probability-based reward a distinct case. A useful solution must evaluate the moving criterion as well as the outputs it scores.

Evaluating improvement with delayed outcomes. When many edits occur before a meaningful result arrives, per-edit acceptance is unavailable. Proxy milestones and proposal pruning may still be useful, but downstream gains alone do not establish successful attribution. Harness optimizers should therefore be assessed against their claims about why particular edits help, as well as their final scores [14].

Connecting local learning to governed sharing. No surveyed system combines automatic local acceptance, promotion across organizations, versioned lineage with access control, and rollback. The components exist in different systems. Their integration must preserve tenant isolation, evaluate the recipient, and account for artifacts derived from a change that is later withdrawn. Where feedback cannot support automatic acceptance, the design must retain the corresponding human decision or delayed evaluation.

The survey's central observation is that persistent changes outside the weights have become an important form of adaptation, while the evidence supporting those changes remains fragmented. Many systems can test a local candidate. Fewer can reliably explain which recurring failure it addresses, and the machinery for transferring the resulting improvement to another organization has not been assembled in the public systems examined here.

The frontier model generalizes. The local scaffold specializes. Reliability comes from governing that specialization: knowing what each change improved, where that evidence applies, and how the change can be reversed. That discipline must extend to the evaluations themselves when they become part of the adapting system.

Acknowledgments

The authors used an AI assistant (Claude, Anthropic) to help search the literature, check citations, and edit prose; all claims, scores, and classifications were reviewed and approved by the authors.

References

[1] M. Zaharia, O. Khattab, L. Chen, J. Q. Davis, H. Miller, C. Potts, J. Zou, M. Carbin, J. Frankle, N. Rao, and A. Ghodsi, "The shift from models to compound AI systems," Berkeley Artificial Intelligence Research (BAIR) Blog, https://bair.berkeley.edu/blog/2024/02/18/compound-ai-systems/, 2024.

[2] L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, and I. Stoica, "Judging LLM-as-a-judge with MT-Bench and Chatbot Arena," in *Advances in Neural Information Processing Systems (NeurIPS 2023), Datasets and Benchmarks Track*, 2023, arXiv:2306.05685.

[3] T. R. Sumers, S. Yao, K. Narasimhan, and T. L. Griffiths, "Cognitive architectures for language agents (CoALA)," *Transactions on Machine Learning Research (TMLR)*, 2023.

[4] X. Liu *et al.*, "SkillForge: Forging domain-specific, self-evolving agent skills in cloud technical support," *arXiv preprint*, 2026, accepted at ACM SIGIR 2026 Industry Track.

[5] P. Xia *et al.*, "SkillRL: Evolving agents via recursive skill-augmented reinforcement learning," *arXiv preprint*, 2026.

[6] M. Helwig, "Memory as infrastructure: Reliability engineering for persistent agent memory in months-long LLM-assisted development," *arXiv preprint arXiv:2609.05510*, Aug. 2026.

[7] A. Miyai, K. Aizawa, and T. Yamasaki, "Task-coevolve: Efficient harness optimization via adaptive validation task selection," *arXiv preprint arXiv:2608.20169*, 2026.

[8] S. Park, W. Kim, R. Tan, J. Zhang, W.-S. Han, P. Gao, C. Park, Y. Yao, R. Fu, E. Nallipogu, Q. Lin, S. Rajmohan, and D. Zhang, "Autosaddler: Automatic harness optimization with durable updates from agent execution traces," *arXiv preprint arXiv:2608.23041*, 2026.

[9] J. Lin, S. Liu, C. Pan, L. Lin, S. Dou, Z. Xi, X. Huang, H. Yan, Z. Han, T. Gui, and Y.-G. Jiang, "Agentic harness engineering: Observability-driven automatic evolution of coding-agent harnesses," *arXiv preprint*, 2026.

[10] T. Zehle, "CANTANTE: Optimizing agentic systems via contrastive credit attribution," *arXiv preprint*, 2026.

[11] Y. Lee, R. Nair, Q. Zhang, K. Lee, O. Khattab, and C. Finn, "Meta-Harness: End-to-end optimization of model harnesses," *arXiv preprint*, 2026.

[12] M. Nguyen, W. C. Tan, N. A. Hassan, A. Raman, L. H. Lim, and A. F. Razak, "Where does harness-optimization value live? Localized gains and the budget-splitting trap in self-evolving LLM agents," *arXiv preprint arXiv:2609.02889*, 2026.

[13] H. Liu, H. Yang, T. Jiang, B. Tang, F. Xiong, Y. Luo, and Z. Li, "SkillsVote: Lifecycle governance of agent skills from collection, recommendation to evolution," *arXiv preprint arXiv:2605.18401*, 2026.

[14] K. T.-i. Ong, M. Kang, D. Choi, J. Cho, S. Kim, S. Lim, G. Jang, M. Oh, B. Jeong, S. Kim, T. Kwon, and J. Yeo, "Towards direct evaluation of harness optimizers via priority ranking," *arXiv preprint*, 2026.

[15] Y. Shen, B. Li, and X. Zhang, "SkillOpt-Lite: Better and faster agent self-evolution via one line of vibe," *arXiv preprint arXiv:2607.03451*, 2026.

[16] C. Shen and E. Hruschka, "Who maintains agent skills? A longitudinal study of human-governed, AI-assisted skill maintenance," *arXiv preprint arXiv:2609.05677*, 2026.

[17] J. Wang *et al.*, "Reinforcement learning for self-improving agent with skill library," *arXiv preprint*, 2025.

[18] X. Wu *et al.*, "Co-evolving LLM decision and skill bank agents for long-horizon tasks," *arXiv preprint*, 2026.

[19] Y. Hu, M. Xu, H. Guo, Y. Song, X. Zhu, and B. Zheng, "Skillchain: Closing the loop on skill evolution for image-based e-commerce ai assistants," *arXiv preprint arXiv:2606.12984*, 2026.

[20] M. Zhuge, W. Wang, L. Kirsch, F. Faccio, D. Khizbullin, and J. Schmidhuber, "Language agents as optimizable graphs," in *Forty-first International Conference on Machine Learning (ICML 2024)*, 2024, arXiv:2402.16823.

[21] J. Zhang, J. Xiang, Z. Yu, F. Teng, X. Chen, J. Chen, M. Zhuge, X. Cheng, S. Hong, J. Wang, B. Zheng, B. Liu, Y. Luo, and C. Wu, "AFlow: Automating agentic workflow generation," *arXiv preprint arXiv:2410.10762*, 2024.

[22] J. Chu, J. Seo, J. Cho, J. Na, Y. Xiong, Y. Kim, and H. J. Kim, "Agentgrad: Intervention-guided prompt optimization for multi agent systems," *arXiv preprint arXiv:2609.08572*, 2026.

[23] M. Chen, C. Lv, G. Zhang, H. Chang, and S. Zhou, "Harnessforge: Joint harness and policy evolution for adaptive agent systems," *arXiv preprint arXiv:2606.01779*, 2026.

[24] S. Zhao, R. Kong, Z. Zhang, L. Shen, X. He, X. Zhang, and Q. Zhang, "Topological coherence for self-evolving multi-agent systems," *arXiv preprint arXiv:2609.37953*, 2026.

[25] Y. Cheng, C. Liu, H. Yu, H. Tan, T. Wu, W. Jin, J. Ran, K. Wang, X. Li, H. Bai, G. Chesi, and N. Wong, "Optimas: Automatically optimize multi-agent system," *arXiv preprint arXiv:2608.21918*, 2026.

[26] AGENTS.md, "AGENTS.md: A cross-vendor open standard for agent instructions," 2025. [Online]. Available: https://agents.md

[27] R. Robbes, T. Matricon, T. Degueule, A. Hora, and S. Zacchiroli, "Agentic much? adoption of coding agents on github," *arXiv preprint arXiv:2601.18341*, 2026.

[28] M. Galster, S. Mohsenimofidi, J. L. Lulla, M. A. Abubakar, C. Treude, and S. Baltes, "Harness engineering for agentic ai coding tools: An exploratory study," *arXiv preprint arXiv:2602.14690*, 2026.

[29] T. Gloaguen, N. Mündler-Sasahara, M. N. Müller, V. Raychev, and M. Vechev, "Evaluating agents.md: Are repository-level context files helpful for coding agents?" *arXiv preprint arXiv:2602.11988*, 2026.

[30] J. L. Lulla *et al.*, "On the impact of AGENTS.md files on the efficiency of AI coding agents," *arXiv preprint*, 2026.

[31] T. Q. Nguyen, S. Desai, R. H. Anwar, F. Shaik, V. Suryanarayanan, and V. Chowdhary, "VerificAgent: Domain-specific memory verification for scalable oversight of aligned computer-use agents," *arXiv preprint arXiv:2506.02539*, Jun. 2025.

[32] Y. Wang, L. Xia, R. Liu, Z. Zhang, H. Wu, Y. Yang, J. He, Y. Guo, and K. Lei, "Sage: A statistical acceptance gate for self-evolving agents," *arXiv preprint arXiv:2609.36043*, 2026.

[33] L. Shang, M. Xu, Y. Sun, T. Xia, L. Hu, L. Xu, and N. Zheng, "When self-evolution backfires: Pre-commit gating against skill contamination in llm agents," *arXiv preprint arXiv:2608.05810*, 2026.

[34] X. Zhang, Y. Cui, G. Wang, Z. Li, W. Qiu, B. Zhu, and P. He, "Ratchet: A minimal hygiene recipe for self-evolving LLM agents," *arXiv preprint*, 2026.

[35] ——, "Library drift: Diagnosing and fixing a silent failure mode in self-evolving llm skill libraries," *arXiv preprint arXiv:2605.19576*, 2026.

[36] Y. Wei, Z. Huang, S. Lu, J. Qian, Q. Wang, C. Wu, and L. He, "Skillsmith: Co-evolving skills and tools for self-improving agent systems," *arXiv preprint arXiv:2606.01314*, 2026.

[37] H. Pu, X. Song, and L. Zhao, "Skillops: Managing llm agent skill libraries as self-maintaining software ecosystems," *arXiv preprint arXiv:2605.13716*, 2026.

[38] Y. Wang, Y. Zhou, Y. Liang, C. Zhang, F. Liu, J. Zhou, and H. Yao, "Not all skills help: Measuring and repairing agent knowledge," *arXiv preprint arXiv:2606.15390*, 2026.

[39] Letta, "MemFS; memory & dreaming; introducing context repositories: Git-based memory for coding agents," https://docs.letta.com/letta-code/memfs, https://docs.letta.com/letta-agent/memory, https://www.letta.com/blog/context-repositories/, 2026.

[40] Mem0, "Mem0 platform: Memory history, feedback, and changelog," https://docs.mem0.ai/api-reference/memory/history-memory, https://docs.mem0.ai/api-reference/memory/feedback, https://docs.mem0.ai/changelog, 2026.

[41] Zep, "Zep documentation: Facts; february 2026 deprecation wave; changelog," https://help.getzep.com/facts, https://help.getzep.com/february-2026-deprecation-wave, https://help.getzep.com/changelog/2026/9/24, 2026.

[42] Y. T. Shen, K. Toyoda, and A. Leung, "Revoked but still authoritative: An empirical study of revocation enforcement in agent-memory systems," *arXiv preprint arXiv:2609.08258*, Sep. 2026.

[43] Z. Wang, "TOKI: A bitemporal operator algebra for contradiction resolution in LLM-agent persistent memory," *arXiv preprint arXiv:2606.06240*, Jun. 2026.

[44] J. Xu, Y. Xiao, W. Shao, H. Liu, and X. Li, "Memory provenance laundering in LLM agents: A non-amplification firewall for persistent memory," *arXiv preprint arXiv:2607.29167*, Jul. 2026.

[45] Y. Margalit, N. Cohen-Inger, E. Avram, R. Taig, and O. Margalit, "Governed shared memory for multi-agent LLM systems," *arXiv preprint arXiv:2606.24535*, Jun. 2026.

[46] X. Sun, Q. Wang, H. Li, and L. Chen, "Verifiable memory: Learning unified memory management with local and global verifiers for large language model agents," *arXiv preprint arXiv:2608.03137*, Aug. 2026.

[47] Z. Yu, Y. Wu, Z. Yin, K. Chen, Z. Zhao, M. Wang, S. Yan, and L. Yang, "Recursive experiential–working memory evolution for long-horizon agent harnesses," *arXiv preprint arXiv:2608.24876*, Aug. 2026.

[48] H. Wang, G. Dong, H. Liang, Z. Zhang, J. Luo, C. Liu, C. Xue, and H. Tang, "MemGuard: Persisting verifier signals for LLM-agent memory governance," *arXiv preprint arXiv:2608.21867*, Aug. 2026.

[49] D. F. Cuadros, A.-A. Maiga, H. Meskhidze, and A. Curtis-Trudel, "Governed collaborative memory as artificial selection in LLM-based multi-agent systems," *arXiv preprint*, 2026.

[50] H. Gao, H. Chen, C. Wang, S. Guo, L. Pang, Z. Liu, H. Shen, and X. Cheng, "Skillaudit: Ground-truth-free skill evolution via paired trajectory auditing," *arXiv preprint arXiv:2606.14239*, 2026.

[51] T. Sah, D. Sah, H. Jain, and T. Sah, "Evoundo: Recoverability-constrained self-evolution for llm agent harnesses," *arXiv preprint arXiv:2608.28363*, 2026.

[52] Y. Yang, Z. Gong, W. Huang, Q. Yang, Z. Zhou, Z. Huang, Y. Li, X. Gao, Q. Dai, B. Liu, K. Qiu, Y. Yang, D. Chen, X. Yang, and C. Luo, "SkillOpt: Executive strategy for self-evolving agent skills," arXiv preprint arXiv:2605.23904, 2026. [Online]. Available: https://microsoft.github.io/SkillOpt

[53] A. Novikov *et al.*, "AlphaEvolve: A coding agent for scientific and algorithmic discovery," arXiv preprint arXiv:2506.13131, 2025.

[54] J. Zhang *et al.*, "Darwin Gödel machine: Open-ended evolution of self-improving agents," *arXiv preprint*, 2025.

[55] Cognition, "Devin's 2025 performance review: Learnings from 18 months of agents at work," https://cognition.com/blog/devin-annual-performance-review-2025, 2025.

[56] C. Ling, P. Chen, A. Guan, J. Qu, S. A. Akbar, M. Gopinathan, and E. Cornejo, "PACE: Two-timescale self-evolution for small language model agents," *arXiv preprint*, 2026.

[57] X. Huang *et al.*, "CASCADE: Cumulative agentic skill creation through autonomous development and evolution," *arXiv preprint*, 2025.

[58] J. Zhu, K. Mao, Y. Guo, D. He, S. Xu, S. Gu, and Y. Yue, "Skillcoach: Self-evolving rubrics for evaluating and enhancing agentic skill-use," *arXiv preprint arXiv:2607.01874*, 2026.

[59] L. Gao, J. Schulman, and J. Hilton, "Scaling laws for reward model overoptimization," *arXiv preprint arXiv:2210.10760*, 2022.

[60] S. Shankar, J. Zamfirescu-Pereira, B. Hartmann, A. G. Parameswaran, and I. Arawjo, "Who validates the validators? aligning LLM-assisted evaluation of LLM outputs with human preferences," *arXiv preprint arXiv:2404.12272*, 2024.

[61] C. Lu, C. Lu, R. T. Lange, J. Foerster, J. Clune, and D. Ha, "The AI scientist: Towards fully automated open-ended scientific discovery," *arXiv preprint arXiv:2408.06292*, 2024.

[62] J. Xu *et al.*, "NanoResearch: Co-evolving skills, memory, and policy for personalized research automation," *arXiv preprint*, 2026.

[63] A. Qu, H. Zheng, Z. Zhou, Y. Yan, Y. Tang, S. Y. Ong, F. Hong, K. Zhou, C. Jiang, M. Kong, J. Zhu, X. Jiang, S. Li, C. Wu, B. K. H. Low, J. Zhao, and P. P. Liang, "CORAL: Towards autonomous multi-agent evolution for open-ended discovery," *arXiv preprint*, 2026.

[64] J. Hu, H. Zhao, Q. Jia, D. Chen, T. Zhao, S. Duan, and G. Liu, "Learning from others, acting for you: Cross-user memory sharing for llm agents," *arXiv preprint arXiv:2609.32511*, 2026.

[65] Y. Guo, H. Zhou, J. Zhang, Y. Zhang, Y. Liu, Q. Chen, Q. Lin, H. Dong, J. M. Zhang, and Z. Chen, "Xreposkill: Learning transferable skills for software engineering agents," *arXiv preprint arXiv:2609.36807*, 2026.

[66] R. Wang and J. Cui, "Evo2team: When do evolved skills transfer? from selection to deployment," *arXiv preprint arXiv:2609.34135*, 2026.

[67] J. Nie, Y. Zhang, Q. Cai, Y.-m. Cheung, X. Tian, and B. Han, "Evolvenet: Collaborative harness evolution for agent self-improvement," *arXiv preprint arXiv:2608.04968*, 2026.

[68] J. Yang, G. Yao, Y. Zhang, R. R. Kompella, G. Liu, and S. Chang, "FederatedSkill: Federated learning for agentic skill evolution," *arXiv preprint arXiv:2606.03143*, 2026.

[69] S. Hu, C. Lu, and J. Clune, "Automated design of agentic systems (ADAS)," *arXiv preprint*, 2024.

[70] G. Wang *et al.*, "Voyager: An open-ended embodied agent with large language models," *arXiv preprint*, 2023.

[71] B. Romera-Paredes, M. Barekatain, A. Novikov, M. Balog, M. P. Kumar, E. Dupont, F. J. R. Ruiz, J. S. Ellenberg, P. Wang, O. Fawzi, P. Kohli, and A. Fawzi, "Mathematical discoveries from program search with large language models," *Nature*, vol. 625, pp. 468–475, 2024.

[72] Sierra, "Agent OS 2.0: From answers to memory and action," 2025. [Online]. Available: https://sierra.ai/blog/agent-os-2-0

[73] E. Palumbo, A. Tamborrino, V. Ode, B. Lacker, A. Casas Escoda, J. Hopple, M. Better, J. Leoni, H. Galvão, H. Bouchard, M. Lalmas, J. L. Redondo García, A. Abebe, A. Clifton, A. Blomberg, H. Lindström, D. Doro, and C. Doig Cardet, "Bootstrapping conversational recommendation agents at spotify: Synthetic data generation and self-improvement loops," *arXiv preprint arXiv:2609.30297*, 2026.

[74] X. Wang *et al.*, "AutoAgent: Evolving cognition and elastic memory orchestration for adaptive agents," *arXiv preprint*, 2026.

[75] P. Barbaste, T. Darrigol, G. Vu, and T. Wiltberger, "Harness engineering: Anatomy, architecture, and evolution of coding agents – a source-code study of eleven systems," *arXiv preprint arXiv:2609.00006*, 2026.

[76] J. L. Lulla, V. Nersesyan, S. Mohsenimofidi, C. Treude, and S. Baltes, "Loop engineering: Building blocks, adoption, and impact," *arXiv preprint arXiv:2608.21884*, 2026.

[77] agentic-community, "Mcp gateway & registry," https://github.com/agentic-community/mcp-gateway-registry, 2026.

[78] Microsoft, "Inventory and discover mcp servers in your api center," https://learn.microsoft.com/en-us/azure/api-center/register-discover-mcp-server, 2026.

[79] TrueFoundry, "Mcp gateway authentication and security," https://www.truefoundry.com/docs/ai-gateway/mcp-gateway-auth-security, 2026.

[80] R. Kang and V. Wang, "Registry-governed agent lifecycle: Completing EDDOps with evaluation-driven registration, promotion, and retirement on AWS AgentCore," *arXiv preprint arXiv:2607.00345*, Jul. 2026.

[81] A. Quessada-Vial, "Agentic configuration management (ACM): A reference configuration model for governed agentic systems," *arXiv preprint arXiv:2608.11166*, Aug. 2026.

[82] X. Qin, S. Luan, J. See, Z. Boukhers, C. Yang, and Z. Li, "Governed capability evolution: Lifecycle-time compatibility checking and rollback for AI-component-based systems, with a proof-of-concept evaluation on embodied agents," *arXiv preprint arXiv:2604.08059*, Apr. 2026.

[83] Braintrust, "Use prompts in code," Braintrust documentation, https://www.braintrust.dev/docs/guides/prompts, 2026, accessed 2026-10-03.

[84] Vellum, "Release tags," Vellum documentation, https://docs.vellum.ai/product/deployments/release-tags, 2026, accessed 2026-10-03.

[85] LangChain, "Prompt engineering concepts," LangSmith documentation, https://docs.langchain.com/langsmith/prompt-engineering-concepts, 2026, accessed 2026-10-03.

[86] Z. Jiang and B. Dong, "Buildrix: An open platform for sharing and benchmarking agentic ai skills in building engineering," *arXiv preprint arXiv:2606.25139*, 2026.

[87] D. Sculley, G. Holt, D. Golovin, E. Davydov, T. Phillips, D. Ebner, V. Chaudhary, M. Young, J.-F. Crespo, and D. Dennison, "Hidden technical debt in machine learning systems," in *Advances in Neural Information Processing Systems 28 (NIPS 2015)*, 2015.

[88] D. Jaroslawicz *et al.*, "How many instructions can LLMs follow at once?" *arXiv preprint*, 2025.

[89] W. Zou *et al.*, "PoisonedRAG: Knowledge corruption attacks to retrieval-augmented generation of large language models," *arXiv preprint*, 2024, accepted at USENIX Security 2025.

[90] Z. Chen, Z. Xiang, C. Xiao, D. Song, and B. Li, "AgentPoison: Red-teaming LLM agents via poisoning memory or knowledge bases," *arXiv preprint arXiv:2407.12784*, 2024.

[91] Z. Xu *et al.*, "From storage to steering: Memory control flow attacks on LLM agents," *arXiv preprint*, 2026.

[92] K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, and M. Fritz, "Not what you've signed up for: Compromising real-world LLM-integrated applications with indirect prompt injection," *arXiv preprint arXiv:2302.12173*, 2023.

[93] E. Debenedetti, J. Zhang, M. Balunović, L. Beurer-Kellner, M. Fischer, and F. Tramèr, "AgentDojo: A dynamic environment to evaluate prompt injection attacks and defenses for LLM agents," *arXiv preprint arXiv:2406.13352*, 2024.

[94] Check Point Research, "MCPoison: Cursor IDE's MCP vulnerability (CVE-2025-54136)," https://research.checkpoint.com/2025/cursor-vulnerability-mcpoison/, 2025.

[95] Y. Su, W. Xu, C. Zuo, and E. Bertino, "ChronoMem: Version control and semantic rollback for large language model agent memory," *arXiv preprint arXiv:2607.27773*, 2026.

[96] C. Yu, Y. Wang, J. Zhang, Y. Duan, M. Zheng, Z. Wu, K. Shi, and T. Cai, "From faulty memories to corrected actions: Dependency-guided rollback repair for memory-augmented agents," *arXiv preprint arXiv:2608.10502*, 2026.

[97] E. Wallace, K. Xiao, R. Leike, L. Weng, J. Heidecke, and A. Beutel, "The instruction hierarchy: Training LLMs to prioritize privileged instructions," *arXiv preprint arXiv:2404.13208*, 2024.

[98] A. Gunjal, A. Wang, E. Lau, V. Nath, Y. He, B. Liu, and S. Hendryx, "Rubrics as rewards: Reinforcement learning beyond verifiable domains," *arXiv preprint arXiv:2507.17746*, 2025.

[99] R. Jia, Y. Yang, Y. Gai, K. Luo, S. Huang, J. Lin, X. Jiang, and G. Jiang, "Writing-zero: Bridge the gap between non-verifiable tasks and verifiable rewards," *arXiv preprint arXiv:2506.00103*, 2025.

[100] K. Shi, Y. Liu, P. Wang, A. R. Fabbri, S. Joty, and A. Cohan, "References improve LLM alignment in non-verifiable domains," *arXiv preprint arXiv:2602.16802*, Feb. 2026.

[101] X. Jin, J. Zhang, T. Cao, and B. Yang, "Beyond score prediction: LLM-based essay scoring and feedback generation via reinforcement learning with rubric rewards," *arXiv preprint arXiv:2607.19219*, Jul. 2026.

[102] T. Yu, B. Ji, S. Wang, S. Yao, Z. Wang, G. Cui, L. Yuan, N. Ding, Y. Yao, Z. Liu, M. Sun, and T.-S. Chua, "RLPR: Extrapolating RLVR to general domains without verifiers," *arXiv preprint arXiv:2506.18254*, Jun. 2025.

[103] R. Xu, T. Liu, Z. Dong, T. Yu, I. Hong, C. Yang, L. Zhang, T. Zhao, and H. Wang, "Alternating reinforcement learning for rubric-based reward modeling in non-verifiable LLM post-training," *arXiv preprint arXiv:2602.01511*, 2026.

[104] H. Ding, B. Huang, Y. Fang, W. Liao, Z. Li, J. Zhang, Z. Wu, J. Zhao, and Y. Wang, "EvoRubrics: Dynamic rubrics as rewards via adversarial co-evolution for LLM reinforcement learning," *arXiv preprint arXiv:2606.23038*, Jun. 2026.

[105] H. Zhu, C. Cai, Y. Song, X. Chen, S. Han, and Y. Guo, "Self-evolving deep research via joint generation and evaluation," *arXiv preprint arXiv:2606.04507*, Jun. 2026.

[106] G. Li, B. Dalvi Mishra, Z. Wang, J. Yan, Y. Chen, C.-L. Li, L. T. Le, R. Han, G. Lee, H. Tong,

C.-Y. Lee, and T. Pfister, "RubricEM: Meta-RL with rubric-guided policy decomposition beyond verifiable rewards," *arXiv preprint arXiv:2605.10899*, May 2026.

[107] S. Singh *et al.*, "The leaderboard illusion," *arXiv preprint*, 2025.

[108] M. Cemri *et al.*, "Why do multi-agent LLM systems fail?" *arXiv preprint*, 2025.

[109] N. F. Liu *et al.*, "Lost in the middle: How language models use long contexts," *Transactions of the Association for Computational Linguistics (TACL)*, 2023.

[110] H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. Agüera y Arcas, "Communication-efficient learning of deep networks from decentralized data," in *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, ser. Proceedings of Machine Learning Research, vol. 54, 2017, pp. 1273–1282.

[111] H. B. McMahan, D. Ramage, K. Talwar, and L. Zhang, "Learning differentially private recurrent language models," in *International Conference on Learning Representations (ICLR)*, 2018, arXiv:1710.06963.

[112] P. Kairouz, H. B. McMahan, B. Avent, A. Bellet, M. Bennis *et al.*, "Advances and open problems in federated learning," *arXiv preprint arXiv:1912.04977*, 2019, published in Foundations and Trends in Machine Learning, vol. 14, no. 1–2, 2021.

[113] T. Guo, S. Guo, J. Wang, and W. Xu, "PromptFL: Let federated participants cooperatively learn prompts instead of models – federated learning in age of foundation model," *arXiv preprint arXiv:2208.11625*, 2022.

[114] L. A. Agrawal, S. Tan, D. Soylu, N. Ziems, R. Khare, K. Opsahl-Ong, A. Singhvi, H. Shandilya, M. J. Ryan, M. Jiang, C. Potts, K. Sen, A. G. Dimakis, I. Stoica, D. Klein, M. Zaharia, and O. Khattab, "GEPA: Reflective prompt evolution can outperform reinforcement learning," *arXiv preprint*, 2025.

[115] R. Tiwari, K. Sareen, L. A. Agrawal, J. E. Gonzalez, M. Zaharia, K. Keutzer, I. S. Dhillon, R. Agarwal, and D. Khatri, "Learning, fast and slow: Towards LLMs that adapt continually," *arXiv preprint*, 2026.

[116] N. Vassilyev, W. Berrios, R. Zhang, B. Han, D. Kiela, and S. Mehri, "Reflective context learning: Studying the optimization primitives of context space," *arXiv preprint*, 2026.

[117] Q. Cai, Y. Zhang, X. Jia, H. Zheng, W. Xue, J. Song, X. Tian, and Y. Guo, "MOSS: Self-evolution through source-level rewriting in autonomous agent systems," *arXiv preprint*, 2026.

[118] L. Li, J. Li, J. Zhao, K. Ye, Y. Li, C.-z. Xu, and X. Gao, "Grow the harness, not the context: From strategy-free scaffolds to reusable specialist agents," *arXiv preprint arXiv:2609.26760*, 2026.

[119] M. M. Tanjim, J. Subramanian, X. Chen, B. Kveton, S. Mukherjee, A. Zhang, S. Kim, S. Sarkhel, and S. Choudhury, "MOCHA: Multi-objective chebyshev annealing for agent skill optimization," *arXiv preprint*, 2026.

[120] Y. Shalev, Z. Ding, and M. Jamnik, "Training language agents to learn from experience," *arXiv preprint*, 2026.

[121] H. Li, R. He, Q. Zhang, C. Ji, Q. Mang, X. Chen, L. A. Agrawal, W.-L. Liao, E. Yang, A. Cheung, J. Zou, K. Olukotun, I. Stoica, and J. E. Gonzalez, "Combee: Scaling prompt learning for self-improving language model agents," *arXiv preprint*, 2026.

[122] P. Hebbar, Y. Manawat, S. Verboomen, A. Ivanova, S. Palanimalai, K. Bhatia, and V. Baskaran, "Sia: Self improving ai with harness & weight updates," *arXiv preprint arXiv:2605.27276*, 2026.

[123] R. Kirk, I. Mediratta, C. Nalmpantis *et al.*, "Understanding the effects of RLHF on LLM generalisation and diversity," in *ICLR 2024*, 2024.

[124] V. Padmakumar and H. He, "Does writing with language models reduce content diversity?" *arXiv preprint*, 2023.

[125] S. Dohare *et al.*, "Loss of plasticity in deep continual learning," *Nature*, vol. 632, no. 8026, pp. 768–774, 2024.

[126] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, "LoRA: Low-rank adaptation of large language models," *arXiv preprint arXiv:2106.09685*, 2021.

[127] M. L. Arbuzov, S. Bei, Z. Dong, D. Kalaev, and A. A. Shvets, "Beyond exponential decay: Rethinking error accumulation in large language models," *arXiv preprint*, 2025.

[128] M. L. Arbuzov, L. Mosbacker, S. Bei, Z. Dong, D. Kalaev, and A. Shvets, "The architecture of errors: From universal impossibility to patch-local LLM reliability," *arXiv preprint arXiv:2605.30628*, 2026.

[129] H.-a. Gao, J. Geng, W. Hua, M. Hu, X. Juan *et al.*, "A survey of self-evolving agents: What, when, how, and where to evolve on the path to artificial super intelligence," *Transactions on Machine Learning Research*, 2026, arXiv:2507.21046v4, 16 Jan 2026.

[130] J. Fang, Y. Peng, X. Zhang, Y. Wang, X. Yi, G. Zhang, Y. Xu, B. Wu, S. Liu, Z. Li, Z. Ren, N. Aletras, X. Wang, H. Zhou, and Z. Meng, "A comprehensive survey of self-evolving AI agents: A new paradigm bridging foundation models and lifelong agentic systems," *arXiv preprint arXiv:2508.07407*, Aug. 2025.

[131] H. Zhou, H. Hu, T. Luo, Y. Shang, C. Fang, Z. Chen, L. Xiao, and Q. Zhang, "Self-evolving coding agents," *arXiv preprint arXiv:2608.03392*, Aug. 2026.

[132] T. Ding, A. Nannapaneni, B. Liu, and L. Zhang, "Always-on agents: A survey of persistent memory, state, and governance in LLM agents," *arXiv preprint arXiv:2606.30306*, 2026.

[133] Z. Lin, X. Hao, R. Fu, S. Cui, K. Chen, C. Li, Z. Li, and F. Xiong, "A survey on long-term memory security in LLM agents: Attacks, defenses, and governance across the memory lifecycle," in *Proceedings of EMNLP 2026*, 2026, arXiv:2604.16548v3, 22 Sep 2026.

[134] J. Guo, Z. Hao, C. Wang, C. Fan, T. Luo, H. Li, Y. Gao, H. Mei, J. Peng, R. Xu, M. Dong, H. Wu, M. Zheng, K. Han, S. Wang, C. Xu, and Y. Wang, "From question answering to task completion: A survey on agent system and harness design," *arXiv preprint arXiv:2606.20683*, Jun. 2026.

[135] B. Xia, L. Zhu, Z. Xing, Q. Lu, D. Sejdinovic, and X. Xu, "Harnessing agent skills: Architectural patterns and a reference architecture for skill-mediated LLM agents," *arXiv preprint arXiv:2606.20631*, May 2026.

[136] K. Ding, Y. Zhou, C. Jin, F. Tong, M. Zhou, and D. N. Metaxas, "Agent skill evaluation and evolution: Frameworks and benchmarks," *arXiv preprint arXiv:2606.11435*, Jun. 2026.

[137] L. Mei, J. Yao, Y. Ge, Y. Wang, B. Bi, Y. Cai, J. Liu, M. Li, Z.-Z. Li, D. Zhang, C. Zhou, J. Mao, T. Xia, J. Guo, and S. Liu, "A survey of context engineering for large language models," *arXiv preprint arXiv:2507.13334*, Jul. 2025.

[138] C. Yang, X. Wang, Y. Lu, H. Liu, Q. V. Le, D. Zhou, and X. Chen, "Large language models as optimizers," in *International Conference on Learning Representations (ICLR)*, 2024, arXiv:2309.03409.

[139] Q. Guo, R. Wang, J. Guo, B. Li, K. Song, X. Tan, G. Liu, J. Bian, and Y. Yang, "EvoPrompt: Connecting LLMs with evolutionary algorithms yields powerful prompt optimizers," in *International Conference on Learning Representations (ICLR)*, 2024, arXiv:2309.08532.

[140] C. Fernando, D. Banarse, H. Michalewski, S. Osindero, and T. Rocktäschel, "Promptbreeder: Self-referential self-improvement via prompt evolution," *arXiv preprint arXiv:2309.16797*, 2023.

[141] O. Khattab, A. Singhvi, P. Maheshwari, Z. Zhang, K. Santhanam, S. Vardhamanan, S. Haq, A. Sharma, T. T. Joshi, H. Moazam, H. Miller, M. Zaharia, and C. Potts, "DSPy: Compiling declarative language model calls into self-improving pipelines," *arXiv preprint arXiv:2310.03714*, 2023.

[142] M. Yuksekgonul, F. Bianchi, J. Boen, S. Liu, Z. Huang, C. Guestrin, and J. Zou, "TextGrad: Automatic "differentiation" via text," *arXiv preprint*, 2024.

[143] Q. Zhang, C. Hu, S. Upasani *et al.*, "Agentic context engineering: Evolving contexts for self-improving language models," arXiv preprint arXiv:2510.04618, 2025.

[144] W. Zhou, Y. Ou, S. Ding, L. Li, J. Wu, T. Wang, J. Chen, S. Wang, X. Xu, N. Zhang, H. Chen, and Y. E. Jiang, "Symbolic learning enables self-evolving agents," *arXiv preprint arXiv:2406.18532*, 2024.

[145] E. Zelikman, E. Lorch, L. Mackey, and A. T. Kalai, "Self-taught optimizer (STOP): Recursively self-improving code generation," in *First Conference on Language Modeling (COLM 2024)*, 2024, arXiv:2310.02304.

[146] M. Robeyns, M. Szummer, and L. Aitchison, "A self-improving coding agent," *arXiv preprint arXiv:2504.15228*, 2025.

[147] Q. Yang, Z. Kou, X. Zhang, D.-D. Wu, H. Gu, J. Guo, Y. Liu, D. Jiang, and Q. Xu, "Federated agent optimization," *arXiv preprint arXiv:2610.01195*, 2026.

[148] H. Robbins and D. Siegmund, "A convergence theorem for non negative almost supermartingales and some applications," in *Optimizing Methods in Statistics*. Academic Press, 1971, pp. 233–257.

[149] A. Madaan *et al.*, "Self-refine: Iterative refinement with self-feedback," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

[150] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, "ReAct: Synergizing reasoning and acting in language models," in *International Conference on Learning Representations (ICLR)*, 2023.

[151] D. Tran and D. Kiela, "Single-agent LLMs outperform multi-agent systems on multi-hop reasoning under equal thinking token budgets," *arXiv preprint*, 2026.

[152] N. Shinn *et al.*, "Reflexion: Language agents with verbal reinforcement learning," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.

[153] J. Chu, J. S. Lee, J. Park, D. Ko, J. Na, S. Lee, T. Lee, M. Yoon, M. Joo, Y. Xiong, and H. J. Kim, "Retrieval-augmented skill optimization via cross-harness adaptation," *arXiv preprint arXiv:2609.38024*, 2026.

[154] J. Qi, Y. He, J. Liu, Q. Huang, S. Hu, Z. Dang, G. Chen, R. Yang, W. Jiang, Y. Liu, T. Lyu, and F. Li, "Vace: Validation-gated alternating co-evolution of agent models and harnesses," *arXiv preprint arXiv:2609.37105*, 2026.

[155] Z. Ma, S. Yang, Y. Ji, X. Wang, Y. Wang, Y. Hu, T. Huang, and X. Chu, "SkillClaw: Let skills evolve collectively with agentic evolver," *arXiv preprint arXiv:2604.08377*, 2026.

[156] J. Ni *et al.*, "Trace2Skill: Distill trajectory-local lessons into transferable agent skills," *arXiv preprint*, 2026.

[157] S. Alzubi *et al.*, "EvoSkill: Automated skill discovery for multi-agent systems," *arXiv preprint*, 2026.

[158] W. Song, J. Yue, and Z. Pang, "ABSTRAL: Automatic design of multi-agent systems through iterative refinement and topology optimization," *arXiv preprint*, 2026.

[159] Y. He *et al.*, "EvoTest: Evolutionary test-time learning for self-improving agentic systems," *arXiv preprint*, 2025.

[160] A. Plesner, C. Northcutt, F. Guzmán, and A. Athalye, "A cheap verifier is good enough: LLM post-training is robust to erroneous rewards," *arXiv preprint arXiv:2609.33467*, Sep. 2026.

[161] M. Bhattarai, I. Boureima, N. R. Ranasinghe, S. Pakin, and D. O'Malley, "Rubric-grounded RL: Structured judge rewards for generalizable reasoning," *arXiv preprint arXiv:2605.08061*, May 2026.

[162] R. Jia, Y. Yang, W. Wang, Y. Wu, Y. Gai, S. Tao, M. Zhou, J. Lin, X. Jiang, and G. Jiang, "Open rubric system: Scaling reinforcement learning with pairwise adaptive rubric," *arXiv preprint arXiv:2602.14069*, Feb. 2026.

[163] C. Lv, J. Zhou, W. Zhao, J. Xu, S. Dou, Z. Huang, M. Tian, X. Wang, Z. Guo, Y. Liu, P. Zhou, T. Gui, L. Tian, X. Zhou, X. Zheng, X. Huang, and J. Zhou, "Learning query-specific rubrics from human preferences for DeepResearch report generation," *arXiv preprint arXiv:2602.03619*, Feb. 2026.

[164] H. Chen, Z. Han, Y. Yan, Q. Zhu, M. Sun, and W. Che, "From holistic evaluation to structured criteria: Rubrics across the evolving LLM landscape," *arXiv preprint arXiv:2606.08625*, 2026.

[165] Z. Ren, Y. Chen, D. Guo, G. Rong, T. Li, R. B. Xiong, Q. Lan, W. Wang, L. Nanbo, Y. Yang, M. Zhuge, and J. Schmidhuber, "Self-improvements in modern agentic systems: A survey," *arXiv preprint arXiv:2607.13104*, 2026.

[166] P. Du, "Memory for autonomous LLM agents: Mechanisms, evaluation, and emerging frontiers," *arXiv preprint arXiv:2603.07670*, Mar. 2026.

[167] W.-C. Huang, W. Zhang, Y. Liang, Y. Bei, Y. Chen, T. Feng *et al.*, "A survey of agent memory in the second half: Towards self-evolving and long-horizon agents," *arXiv preprint arXiv:2602.06052*, 2026.

[168] J. Luo, Y. Tian, C. Cao, Z. Luo, H. Lin, K. Li, C. Kong, R. Yang, and J. Ma, "From storage to experience: A survey on the evolution of LLM agent memory mechanisms," *arXiv preprint arXiv:2605.06716*, May 2026.

[169] Q. Hua, L. Ye, D. Fu, Y. Xiao, X. Cai, Y. Wu, J. Lin, J. Wang, and P. Liu, "Context engineering 2.0: The context of context engineering," *arXiv preprint arXiv:2510.26493*, Oct. 2025.

[170] Y. Zhou, W. Shu, Y. Su, W. Du, Y. Fang, and X. Lin, "A comprehensive survey on agent skills: Taxonomy, techniques, and applications," *arXiv preprint arXiv:2605.07358*, 2026.

[171] Y. Zhou, A. I. Muresanu, Z. Han, K. Paster, S. Pitis, H. Chan, and J. Ba, "Large language models are human-level prompt engineers," *arXiv preprint arXiv:2211.01910*, 2022.

[172] R. Pryzant, D. Iter, J. Li, Y. T. Lee, C. Zhu, and M. Zeng, "Automatic prompt optimization with "gradient descent" and beam search," in *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP 2023)*, 2023, arXiv:2305.03495.

[173] K. Opsahl-Ong, M. J. Ryan, J. Purtell, D. Broman, C. Potts, M. Zaharia, and O. Khattab, "Optimizing instructions and demonstrations for multi-stage language model programs," in *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP 2024)*, 2024, arXiv:2406.11695.

[174] C.-A. Cheng, A. Nie, and A. Swaminathan, "Trace is the next AutoDiff: Generative optimization with rich feedback, execution traces, and LLMs," *arXiv preprint arXiv:2406.16218*, 2024.

[175] J. S. Park *et al.*, "Generative agents: Interactive simulacra of human behavior," in *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST)*, 2023.

[176] A. Zhao, D. Huang, Q. Xu, M. Lin, Y.-J. Liu, and G. Huang, "ExpeL: LLM agents are experiential learners," in *Proceedings of the 38th Annual AAAI Conference on Artificial Intelligence (AAAI-24)*, 2024, arXiv:2308.10144.

[177] C. Packer, S. Wooders, K. Lin, V. Fang, S. G. Patil, I. Stoica, and J. E. Gonzalez, "MemGPT: Towards LLMs as operating systems," *arXiv preprint arXiv:2310.08560*, 2023.

[178] Z. Z. Wang, J. Mao, D. Fried, and G. Neubig, "Agent workflow memory," *arXiv preprint arXiv:2409.07429*, 2024.

[179] W. Xu, Z. Liang, K. Mei, H. Gao, J. Tan, and Y. Zhang, "A-MEM: Agentic memory for LLM agents," in *Advances in Neural Information Processing Systems (NeurIPS 2025)*, 2025, arXiv:2502.12110.

[180] M. Suzgun, M. Yuksekgonul, F. Bianchi, D. Jurafsky, and J. Zou, "Dynamic cheatsheet: Test-time learning with adaptive memory," *arXiv preprint arXiv:2504.07952*, 2025.

A Formal statement of artifact-layer descent

This appendix states the setup behind Section 2, a one-step descent inequality under a noisy gate, an informal convergence proposition, the optimization-analogy correspondence, and four worked cases. No new theorem is claimed; the proposition is the ordinary target of local stochastic optimization under noisy estimators, restated for a discrete, versioned parameter.

Table 9: Optimization concepts and their scaffold analogues.

Optimization concept	Scaffold analogue
Parameter vector	Instructions, skills, memories, tools, routing graphs, and governance rules, evaluation gates among them
Loss	Failure, user correction, evaluation regression, retrieval miss, tool error, schema violation
Gradient estimate	Credit assignment identifying the coordinate responsible for a failure
Step	Accepted scaffold edit; step size is edit magnitude under $\oplus$
Aggregated update	Several recurring failures combined before an update, by averaging or by agreement across k proposals
Update radius	Promotion scope: which loss is descended (USER, PROJECT, STACK, TENANT, CORE)
Validation loss	Held-out patch evaluation before merge
Regularization	Complexity penalty $\Omega(z)$, bloat control, cross-patch regression checks
Overfitting	Scaffold improves on this deployment and degrades elsewhere
Rollback	Reverting a harmful accepted update

A.1 Setup

Let the frontier model have fixed parameters θ_M and let the scaffold be a structured object $\theta_s \in \Theta_s = \Theta_1 \times \cdots \times \Theta_6$ over the six substrate coordinates of Section 3; Θ_s is discrete, symbolic, and versioned. For a deployment D, $L_D(\theta_s) = \mathbb{E}_{x\sim D}[\ell(\pi(\theta_M, \theta_s, x))]$. A candidate edit $z_t \sim Q_t(z \mid H_t, \theta_s^t)$ is drawn from a proposal process conditioned on the accumulated history H_t of failures, corrections, rejected outputs, evaluation regressions, and prior accepted edits. Q_t is not uniform: for a mode m that recurs with frequency f_m and carries per-occurrence loss s_m, its expected contribution to the failure rate is $f_m s_m$, so high-frequency or high-severity modes are proposed first. The finite directional difference is $\Delta_z L_D(\theta_s) = L_D(\theta_s \oplus z) - L_D(\theta_s)$; the gate sees only a noisy estimate $\widehat{\Delta}_z L_D$. The update is

$$\theta_s^{t+1} = \theta_s^t \oplus z_t \text{ if } G_t(z_t) = 1, \quad \theta_s^{t+1} = \theta_s^t \text{ otherwise.}$$

A.2 Descent inequality

Assume (A1) positive descent bias on accepted edits, $\mathbb{E}[\widehat{\Delta}_{z_t} L_D \mid G_t = 1, \theta_s^t] \le -\gamma_t - \tau_t$ for some $\gamma_t > 0$, and (A2) bounded estimator bias on accepted edits, $|\mathbb{E}[\Delta_{z_t} L_D - \widehat{\Delta}_{z_t} L_D \mid G_t = 1, \theta_s^t]| \le \varepsilon_t$. With $p_t = \Pr(G_t = 1 \mid \theta_s^t)$,

$$\mathbb{E}[L_D(\theta_s^{t+1}) \mid \theta_s^t] \le L_D(\theta_s^t) - p_t(\gamma_t + \tau_t) + p_t\,\varepsilon_t.$$

Per-step expected improvement is the gate's margin minus its estimator bias, weighted by acceptance probability. In the feedback-signal classes of Section 5, ε_t is near zero for the check itself in C1, with

the sampling error of a finite suite over the patch remaining, bounded by batching for an observable outcome (C2), equal to the judge's bias and bounded only by calibration for a judged signal (C3), and unavailable per edit for a delayed signal (C4).

A.3 Convergence (informal)

Add (A3) proposal coverage with acceptance: at every non-gate-stable scaffold, Q_t places probability at least $q > 0$ on an admissible edit with $\mathbb{E}[\Delta_z L_D] \leq -\delta$, and the gate accepts such an edit with probability at least $q' > 0$, so that $p_t(\gamma_t + \tau_t) \geq qq'\delta$ there; (A4) effective compactness of the reachable orbit; and (A5) summable bias, $\sum_t p_t \varepsilon_t < \infty$; and let $0 \leq L_D \leq B$. The one-step inequality is then of the Robbins–Siegmund form, a non-negative process that decreases in expectation up to a summable perturbation $p_t\varepsilon_t$, so $L_D(\theta_s^t)$ converges almost surely and $\sum_t p_t(\gamma_t + \tau_t) < \infty$ almost surely [148]. The finite sum carries the content: a bounded loss cannot pay for infinitely many accepted edits with a positive expected margin, so under (A3) the process visits non-gate-stable scaffolds only finitely often, and with (A4) the limit set of $\{\theta_s^t\}$ is gate-stable: no admissible edit clears the gate with negative expected directional loss in the limit. This is convergence of the loss and stability of the gate decision, not convergence to an optimum; a globally suboptimal gate-stable scaffold is consistent with the result, because Q_t may never propose the improving edit or $\widehat{\Delta}$ may be biased to reject it. In a non-stationary deployment ($D = D_t$) the same mechanism tracks a moving optimum rather than converging once, which is why recency weighting, pruning, expiry, and rollback are the scaffold-level analogues of continual-learning machinery rather than optional hygiene.

A.4 Correspondence

Table 9 pairs each optimization concept with its scaffold analogue. Credit assignment here is attribution, not backpropagation. Scaffold systems do not differentiate through instruction files. They trace an output error to an upstream coordinate, manually or with a judge or verifier. The descent argument needs only that this attribution succeeds with non-zero probability per failure, which is the credit-assignment condition of Section 2.

A.5 Worked cases

Retrieval. Failures recur because retrieved context omits the decisive clause. Attribution traces the error to the retrieval coordinate: chunking, query rewriting, metadata filter, embedding model, reranker, or context-packing rule. The candidate edit tightens chunk boundaries or adds a metadata constraint, and the gate checks that recall and answer accuracy improve on the patch evaluation without degrading adjacent cases. Signal class C1 or C2.

Pitch review. The system repeatedly overrates decks that claim a large market without numbers, and the failure cluster points at an under-specified rubric coordinate. The edit changes "evaluate market size" into "require a numeric estimate and a source; if absent, cap the score"; the gate runs the revised rubric against held-out decks with human scores. What the edit changed was the gate that will judge the next edit: the reflexive gate of Section 5, class C3.

Tool use. Calls fail because optional parameters are under-specified. The responsible coordinate may be the tool schema, the call policy, or the clarification rule. The edit adds required fields or a pre-call uncertainty check; the gate checks that tool-call failures fall without raising refusals or latency. Class C2.

Memory. A stale customer preference is repeatedly retrieved and causes bad decisions. The

responsible coordinate is a memory entry together with its provenance, priority, and expiry policy. The edit corrects, expires, or version-tags the entry; the gate checks that future behavior improves without erasing useful history. Class C2, with the correction itself observable only when a later outcome reveals it.

B Gate families and feedback-signal classes

Section 2 collapses every production gate into one per-edit test G_t on an estimate of the loss change. Surveyed systems instantiate that test in several distinct ways, and some of them are better read as proposal pruning that precedes the per-edit gate than as estimators of $\widehat{\Delta}L_D$. Table 10 lists the families, the slot each occupies, representative systems, and the feedback-signal class (Section 5) each family serves.

The first five families populate G_t directly and, under the assumptions of Appendix A, inherit its descent guarantee. Semantic merge and archive selection are not estimators of $\widehat{\Delta}L_D$; they constrain the proposal distribution, and the descent argument applies to the composition of pruning and gate. Multi-sample agreement is listed because it is the correct analogue of "wait for k pieces of evidence before accepting an edit", a different operation from mini-batch averaging and one that should not be conflated with it. Read down the signal-class column. The only families available to a judged (C3) deployment are the two whose estimator is a judge. A delayed (C4) deployment has no family that estimates the loss change per edit at all: what it has is semantic merge before the gate and whatever proxy milestone stands at the horizon. That is the structural content of Section 5.

C Data appendix

Corpus and window. The corpus holds 256 unique systems from January 2023 to the October 2026 cutoff, released as two files. `p3_master_scores.csv` holds the 130 first-pass systems in 140 rows; `p3_refresh_2026-10.csv` holds the 126 second-pass systems in 133 rows, four of which re-audit a first-pass system and replace its row rather than add a system. The corpus therefore has 140 + 133 - 4 = 269 substrate-level rows. A system may be scored on several substrates, so rows outnumber systems, and Table 2 of the main article counts rows.

Two passes. Every system in either pass is scored against one rubric, from its own method section or documentation. First-pass systems were found in arXiv, vendor documentation, and public repositories through May 2026. Second-pass systems, found from June to the October 2026 cutoff, come from six query families: self-evolving skills and skill lifecycle; harness and context optimization; agent memory governance; prompt-as-code and evaluation-gated promotion; federated or multi-tenant skill sharing; and multi-agent topology evolution. Some were reached instead through the citations of the first-pass systems closest to the missing architecture of Section 8. The second pass was a systematic search of the field, not only a search around the known systems, so the per-substrate statistics pool both passes.

Inclusion and exclusion. A system enters the corpus if it either wraps a foundation model with at least one persistent artifact updated from deployment signal or is described in public as "self-improving," "scaffolded," or "agentic." The contrast set enters through the second criterion and is kept on purpose, so that the discriminator has something to mark against (Self-Refine, ReAct, TextGrad, OPRO, engineer-authored convention files; Appendix E.1.1). Two kinds of system are excluded: those whose only update mechanism is weight-level fine-tuning or reinforcement learning,

Table 10: Gate families, their slot in the formalism, and the feedback-signal class each serves.

Gate family	What it estimates	Slot in Section 2	Signal class	Representative systems
Surrogate verifier	Pass or fail from tests, a symbolic checker, or an exact match	Per-edit gate G_t	C1	DGM (sandboxed compile-and-edit check; score weights selection, not retention), SkillOpt (held-out exact match), Voyager (game-state check: C2)
Evaluation cascade	$\widehat{\Delta}L_D$ on a held-out task suite, cheap-to-expensive evaluators	Per-edit gate G_t	C1 when the suite is exact; C2 when it is a task suite	AlphaEvolve, Braintrust, Vellum, LangSmith Hub
RL threshold	Reward improvement against a deployment threshold	Per-edit gate G_t	C2	SAGE, PACE
LLM judge with calibration	Pairwise or rubric score, calibrated against human agreement	Per-edit gate G_t	C3	SkillChain (judge stage, before its human reflection gate)
Reviewer judgment	Human credit assignment by pull-request review, expert sign-off, or a user's acceptance	Per-edit gate G_t, human-instantiated	C3 (and C2 when the reviewer has outcome data)	Sierra Agent OS 2.0, Anthropic skills repository, CASCADE (a person and an agent approve consolidation into shared memory)
Semantic merge	Whether a candidate duplicates or conflicts with existing artifacts	Proposal pruning before G_t	C4 (also a pre-filter in C2, C3)	NanoResearch orchestrator, FederatedSkill server agent, EvolveR (admission: cosine pre-filter plus LLM equivalence judgment on each distilled principle; retention: periodic pruning on per-principle success rate)
Archive selection	Whether a candidate improves on, or diversifies, an archive of already-scored programs	Proposal pruning before G_t	Any; inherits the class of the evaluator that scored the archive (C1 in AlphaEvolve)	AlphaEvolve (MAP-Elites-style database), ADAS
Multi-sample agreement	Whether k independent proposals or evaluations agree; false-accept rate decays as p^k rather than variance as $1/k$	Aggregation filter on the estimate	C2, C3 noise absorption	VaG (three heterogeneous critics must all pass); SkillOpt-Lite (consensus filter across rollouts)

which are model-side and belong to the self-evolution literature this survey positions against, and those with no public evidence of deployment or release, such as white papers without code, demos, or a production claim.

Two filters. The patch-local discriminator admits a system to the core corpus when it scores at least 3 on at least one substrate. Of the 256 systems, 204 pass it, 99 of them first-pass. The other 52, 31 of them first-pass, form the contrast corpus: systems described in public as agentic or self-improving that fail the discriminator. Each second-pass contrast row gives its exclusion reason in the released CSV. The stricter composite two-loop audit then asks whether Loop 1 modifies a persisted artifact from session signal and whether Loop 2 has an explicit cross-context promotion mechanism with versioning or lineage. Representative first-pass members of both groups, contrast systems and core systems that fail the composite audit, are listed in Appendix E.1.

Systems outside the released table. Fifteen named systems have no row in the released table. Ten are treated at mechanism level in the main text and this supplement: SkillOpt, GEPA, Meta-Harness, CORAL, PACE, MOSS, RCL, FederatedSkill, SkillClaw, and ChronoMem. Each was classified from its paper after the first pass closed and is placed by that reading, in the matrix or among the research families. The other five are contrast or boundary entries, classified from their papers and marked with a dagger in Appendix E.1: ReAct, MetaGen, MetaReflection, CLIN, and ACE. The corpus counts (256, 204, 52) exclude all fifteen. SkillOpt is nonetheless one of the eleven research systems that pass the composite audit, because the audit was applied to it from its paper.

Evidence tiers. Each system carries one or more of five tiers: (T1) peer-reviewed publication; (T2) arXiv preprint; (T3) production claim with metrics; (T4) open-source artifact; (T5) industry blog or interview. T1 and T4 weigh most in cluster analysis. T3 counts where the claim is concrete, naming a benchmark or an integration target; T5 counts only when T1–T4 evidence corroborates it elsewhere. Tier flags reach cluster claims through the per-row evidence column of the released CSV, and T3 production claims are never pooled with T1 or T2 evidence in a cluster-level statistic.

Tier distribution. In the first-pass core corpus, about a quarter of the systems carry T1 evidence as their primary tier, roughly half carry T2, and the rest T3 or T4. Preprint-only systems in either pass carry T2, recorded per row in the released CSV. Some T2 work is recent enough that its ablations may yet be revised; the cluster-level claims survive dropping any single T2 system.

Scoring discipline. Each system is scored substrate by substrate on the survey's 0–5 rubric and also receives an integrated patch-local-adaptation score in $\{0, 1, 2, 2.5, 3, 3.5, 4, 4.5, 5\}$. Half-scores are used sparingly, where one criterion is clearly met and another only partly, and ambiguous cases are rounded down. Every score comes from a single rater. The full score table is therefore published, each row with a one-line justification in the released CSVs, so that any ambiguous classification can be checked independently (Appendix I).

Composite two-loop criterion (sensitivity analysis). The strict composite audit passes eleven research systems, all first-pass except SkillOpt, and three industry exemplars. Two natural relaxations change that count.

The governance relaxation counts any of versioning, lineage, or rollback as Loop-2 governance. It re-admits MAE, MetaGen, MetaReflection, and CLIN from the first pass, raising the research count to approximately fifteen, and of the second-pass promoters below it admits EvolveNet alone, whose rollback to the previous shared harness counts.

Six second-pass systems promote across users, clients, teams, or repositories with no versioning or lineage on the promoted artifact: EvolveNet, Evo2Team, XRepoSkill, ShareMem, Spotify's

recommendation loop, and Future AGI's prompt lifecycle. They fail the audit on governance alone. RASO, which retrieves skills from a public corpus, is not among them, because inbound retrieval is not a promotion (Appendix E.2.1). A weaker relaxation that asks only for an explicit promotion mechanism, with no governance property at all, would admit five of them, all landing in [Auto × Auto] except Spotify's loop, which would land in [Auto × Human]; none would open a new cell. The sixth, Future AGI's prompt lifecycle, rests on T5 evidence alone, which the tier rule above does not count without corroboration, so it is not admitted under any rule.

The optimizer-state relaxation admits as Loop 1 a session signal that updates optimizer state without persisting an artifact distinct from that state. It re-admits DSPy and TextGrad from the first pass, raising the research count to approximately thirteen. It admits no second-pass system. The second-pass contrast rows are benchmarks, evaluation and attack studies, registries, audits, and conceptual frameworks, plus one orchestrator whose adaptation lives in the weights, and none of them persists optimizer state as its only artifact.

The qualitative findings hold under every relaxation: [Auto × Auto] is well populated; [Human × Auto] and [Human × Human] are empty; [Human × None] holds only curation workflows and [None × Human] only registries, all of which fail the composite audit; the governance-first cluster sits in [None × Auto]; Letta alone passes the audit with an ungated promotion; and no system adds governed cross-tenant promotion.

D The mechanism stack and continual-learning notes

D.1 The mechanism stack

The update rule of Section 2 expands into six operations that every patch-local system performs, whether by explicit engineering or by accident.

1. **Residual capture.** Log failures, corrections, rejected outputs, user edits, evaluation regressions, and tool-call failures. This is the signal source of Loop 1; without it, z_t has no provenance.
2. **Mode discovery and credit assignment.** Cluster repeated failures into the deployment's own error modes (the catalogue of *The Architecture of Errors* [128], instantiated locally for one patch) and attribute each mode to the scaffold coordinate responsible for it (the correspondence note of Appendix A).
3. **Delta synthesis.** Turn a recurring mode into a candidate scaffold edit: an instruction rewrite, a new skill, a memory correction, a tool binding, an orchestration change, or an evaluation case. This produces z_t.
4. **Delta validation.** Estimate whether the edit lowers L_D without collateral damage on adjacent inputs. The gate G_t is Auto when the check is an algorithmic criterion and Human when it is reviewer judgment.
5. **Promotion control.** Decide whether the accepted edit stays in the local fork or climbs the scope hierarchy: PROJECT, STACK, TENANT, CORE. This is Loop 2.
6. **Drift and bloat control.** Expire stale rules, prune conflicting memories, roll back regressed skills, and detect overfitting. Without this operation, accumulated edits produce the scaffold-level analogue of plasticity loss.

Each operation is the locus of a separate engineering discipline, and surveyed systems implement

them partially and unevenly. Research full-auto systems excel at operations 3 and 4, governance-first systems at operation 5, and production hybrids at most of operations 1 to 5 but rarely at operation 6. Implementation patterns for each operation lie outside this survey.

D.2 Loop 1 and Loop 2 as descent steps of different radius

Loop 1 proposes and validates local descent steps on a single L_D. Loop 2 widens the update radius: it selects, merges, and promotes accepted local edits so that they apply to a larger set of patch losses averaged together. Acceptance at USER scope descends on a single L_D. PROJECT promotion descends on a loss averaged over the project, and STACK promotion widens the average further. TENANT promotion descends on a tenant-wide average; CORE promotion, the cross-tenant level, descends on an average over tenants. The two-loop design space is therefore more than a taxonomy of engineering patterns: it records *which loss each surveyed system is descending on*, and which estimator gates the step.

D.3 Overfitting in scaffold adaptation

In weight training, overfitting is usually a defect. In scaffold adaptation, local overfitting is partly the point. Frontier weights are trained to preserve generality across patches, so they must smooth over local conventions and rare workflow-specific failures. A deployment's patch needs the opposite: selective specialization to its own residuals. The scaffold is therefore a layer meant to overfit, but one whose overfitting is scoped, inspectable, versioned, gated, and reversible. A scaffold edit that improves one repository, hospital workflow, or legal diligence process may degrade performance elsewhere. That is not a contradiction; it is patch adaptation. The mistake is promoting that edit beyond its evidence radius.

In non-stationary patches ($D = D_t$), the same descent mechanism tracks a moving optimum rather than converging once. Recency weighting, replay, pruning, expiry, and rollback are therefore not optional hygiene; they are the scaffold-level analogues of continual-learning machinery.

E Contrast set, matrix membership, and closest approaches (May 2026 corpus)

The tables in this section describe the first-pass corpus, which closed in May 2026; the second-pass systems that occupy a cell of the Section 6 matrix are named in the cell-by-cell discussion that follows the matrix.

E.1 Contrast set

Two tables mark the two boundaries of the survey. The first lists representative members of the first-pass contrast corpus: 31 systems that fail the patch-local discriminator, scoring below 3 on every substrate. The second lists core systems, each scoring at least 3 on some substrate, that still fail the stricter composite two-loop criterion, so that the asterisked entries of the matrix in Appendix E.2.1 can be checked.

Two marks recur. An asterisk, used in the matrix, marks a system that fails at least one criterion of the composite audit and is included for contrast. A dagger has two uses: in the two contrast

tables it marks a system classified from its paper because the score table has no row for it (ReAct, MetaGen, MetaReflection, CLIN, ACE); in the matrix it marks an industry exemplar evaluated with vendor documentation as the primary evidence (Sierra Agent OS 2.0, Cognition Devin).

E.1.1 Contrast corpus: fails the patch-local discriminator

System	Year	Discriminator failure	Brief reason
Self-Refine [149]	2023	No persistent artifact	In-context iteration only
ReAct [150]†	2023	No persistent artifact	Thought-act-observe loop is intra-session
TextGrad [142]	2024	No artifact distinct from optimizer state	Textual-gradient optimization; SkillOpt [52] (Section 4; Appendix F) adds the persistent `best_skill.md`, held-out validation gate, rejected-edit memory, and textual learning-rate schedule that TextGrad lacks.
OPRO, Toolformer	2023	No persistent deployment artifact	Optimization or tool use inside a single training or inference run
Aider conventions, Lovable knowledge	2024	No feedback loop	Engineer-authored convention files
AGENTS.md (as a standard)	2025	No feedback loop	Cross-vendor artifact format as a workspace file; adoption measured, adaptation absent (Section 4). The specification's governance under the AAIF is a separate S6 entry, "AGENTS.md / AAIF"
LangChain Skills, HuggingFace Agent Skills, CrewAI Skills	2025–2026	No feedback loop	Curated skill packs without failure-triggered update
Tran & Kiela [151]	2026	Not a system	Single-agent baselines match or beat multi-agent under matched thinking-token budgets; constrains S5 claims (see Section 4)

E.1.2 Core corpus, fails the composite two-loop criterion

System	Year	Composite-criterion failure	Brief reason
Reflexion [152]	2023	No Loop 2	Per-agent episodic memory without promotion
Generative Agents	2023	No Loop 2	Per-agent reflection + importance memory
ExpeL	2023	Loop boundary blurred	Experience bank built from training tasks

System	Year	Composite-criterion failure	Brief reason
Mem0 (vanilla)	2024	No Loop 2	Per-user memory; no cross-context promotion
Zep	2024	No Loop 2	Per-agent/session knowledge graph
ChatGPT Memory	2024	No Loop 2	Per-user persistence
Claude Memory	2025	No Loop 2	Per-conversation-thread persistence
MemGPT	2023	No Loop 2	Virtual-context management per agent
DSPy	2023	No artifact distinct from optimizer state	Prompt-template optimization; SkillOpt [52] (Section 4; Appendix F) realizes the artifact-bearing form.
MetaGen†	2026	No Loop 2; no cross-session persistence	Role pool dynamic within single inference call
MAE (Multi-Agent Evolve)	2025	No Loop 2	Co-evolution within Proposer/Solver/Judge triplet
MetaReflection†	2024	No Loop 2	Per-agent semantic memory
CLIN†	2023	No cross-instance promotion	Causal abstractions per-agent-per-environment
LangGraph, AutoGen, CrewAI	2023–2024	Framework only	Orchestration frameworks; not deployed systems
W&B Weave, Agenta, Langfuse	2024	No Loop 1	Observability / eval-as-code
Mobile-Agent-E, Cradle, OS-Copilot	2024–2025	No Loop 2	Self-evolving but per-session/per-deployment
Agent S, Agent S2	2024–2025	No Loop 2	Computer-use agentic framework
ACE [143]†	2025	No Loop 2 in surveyed form	Single-context evolving playbook
Cursor Rules (vanilla)	2024	No Loop 2	Per-project artifacts
GitHub Copilot custom instructions, Windsurf, Continue.dev, Zed AI rules	2024	No Loop 1	Engineer-authored at repo base branch
Replit `replit.md`	2025	Loop 2 absent	Self-writing scaffold but per-workspace

System	Year	Composite-criterion failure	Brief reason
SkillForge [4]	2026	No Loop 2	A judge flags failed replies that drive each skill rewrite; the rewrite is committed without an acceptance test, versions stay in the scenario deployment, and engineers review replies, not skill versions
AutoAgent [74]	2026	No Loop 2	Per-agent tool and peer descriptions revised on an uncalibrated consistency judgment; never promoted or versioned
SkillRL [5]	2026	No gate on either loop	A success rate below 0.4 triggers skill proposals; the skill bank takes each one as an unconditional union

The contrast systems fail the discriminator in three main ways: they iterate in context only, they have no feedback loop, or they are not systems at all. Core systems fail the composite criterion mostly for want of a Loop 2.

E.2 Matrix membership and closest approaches

E.2.1 Two-loop matrix: cell-by-cell discussion

Both axes record who decides, as Section 2 defines: a gate is Auto when a program applies a check, Human when a person decides, and None when the step is absent or runs unchecked. An engineer-authored or registrant-supplied artifact with no failure trigger therefore has Loop 1 None in both passes, and who supplies each candidate is listed separately in the occupant table at the end of this section.

The first-pass matrix, with the marks of Appendix E.1:

	Loop-2 Auto-gated	**Loop-2 Human-gated**	**Loop-2 None**
Loop-1 Auto-gated	NanoResearch, SkillOpt, AlphaEvolve, ADAS, DGM, Voyager, SAGE, EvolveR, CoEvoSkills, AutoManual	CASCADE, Sierra Agent OS 2.0†, Cognition Devin†	AutoAgent*, MAE*, MetaGen*, MetaReflection*, CLIN*, DSPy*, TextGrad*
Loop-1 Human-gated	**[empty by engineering logic]**	**[empty in the first pass]**	*(none)*
Loop-1 None	ExpeL* (weak), Braintrust*, LangSmith Hub*, Vellum*	Anthropic skills repository*	Self-Refine*, ReAct*, Mem0*, SkillRL*, SkillForge*

The matrix places the first-pass systems audited under the composite criterion, with SkillOpt, and the contrast entries marked with an asterisk. Every matrix entry that passes the audit sits in [Auto

× Auto] or [Auto × Human]: the ten research systems of the first cell, with CASCADE, are the eleven research systems of Section 3, and Sierra Agent OS 2.0 and Cognition Devin are two of its three industry exemplars. The third exemplar, Letta's Context Repositories, is a second-pass system.

Three placements need a word. DSPy and TextGrad sit in [Auto × None] because their Loop 1 is automated but persists only optimizer state, which only the optimizer-state relaxation of Appendix C would count as a persisted artifact. DGM sits in [Auto × Auto] although its local gate is a validity check rather than a loss estimate: the archive admits every child that compiles and can still edit code, the benchmark score weights parent selection rather than deciding retention, and no person decides anything in the loop, since self-modifications are evaluated without human intervention [54]. SkillRL sits in [None × None] although its loop is failure-driven: a per-category success rate below 0.4 triggers new skill proposals, and the skill bank takes each proposal as an unconditional union, so the step that keeps an edit runs unchecked, and the bank never leaves its run. SkillForge sits beside it for the same reason: a calibrated judge selects the failed replies that drive each skill rewrite, the rewrite is committed without an acceptance test, every version stays in the scenario deployment, and its support engineers review replies, not skill versions.

COSPLAY and Cuadros et al., which are first-pass rows, and PACE, CORAL, GEPA, Meta-Harness, MOSS, and RCL, which were classified from their papers after the first pass closed, carry no composite audit, so they are neither placed in the matrix nor counted in Section 3. Second-pass systems record their gate pairs in the released score table, and the paragraphs below name those that occupy a cell.

[**Auto** × **Auto**] is the most populated cell: ten systems, each with a machine-evaluated criterion on both loops. Speed on both loops is what its members share; cross-organization governance is what every one of them lacks.

[**Auto** × **None**] holds the systems whose local loop is gated and whose promotion has no gate, whether because nothing is promoted or because everything is. One audit passer sits here: Letta's Context Repositories push shared memory across agents automatically, with git lineage and no evaluation of any kind, so the composite audit admits them on lineage and the matrix records a promotion that nothing decides (Section 6). Mem0 Platform, TOKI, Learning on the Job, and RASO sit here with no promotion at all; RASO retrieves skills from a read-only public corpus, which is inbound retrieval rather than promotion [153]. One first-pass system joins them, starred because it has no Loop 2: AutoAgent's analyzer accepts revisions to one agent's tool and peer descriptions on an uncalibrated consistency judgment, and nothing leaves that agent's store. VACE belongs here too although its gate is a held-out check: it compares a candidate harness with the incumbent against the freshly trained checkpoint of the same run, so the edit is made and kept inside one training run [154].

[**Auto** × **Human**] is the production-viable cell. Loop 1 fires at machine speed under an algorithmic criterion, and Loop 2 waits for a person before an edit propagates; the pairing reconciles the speed of automation with a reviewer's accountability. CASCADE, Sierra Agent OS 2.0, and Cognition Devin occupy it: a human-and-agent consolidation gate, human-authored expert answers, and a reviewer's merge of a pull request each put a person on promotion.

Six second-pass systems promote across users, clients, teams, or repositories with no versioning or lineage on the promoted artifact, so they sit outside the strict audit: EvolveNet, Evo2Team, XRepoSkill, ShareMem, Spotify's recommendation loop, and Future AGI's prompt lifecycle. Were the audit to ask for a promotion mechanism and no governance property at all, five of them would

enter, all landing in [Auto × Auto] except Spotify's loop in [Auto × Human], without opening a new cell; the sixth, Future AGI's prompt lifecycle, rests on T5 evidence alone and is not admitted under any rule (Appendix C).

[None × Auto] holds the governance-first prompts-as-code systems: Braintrust, LangSmith Hub, and Vellum. Their Loop 2 is real, since evaluation-gated CI blocks a promotion on regression. They have no Loop 1 in the patch-local sense, because engineers author the candidate artifacts at their desks and session feedback never triggers one, and the asterisks record exactly that: governance without inner-loop adaptation. The cell is informative because it shows what survives when only the outer loop is automated. ExpeL, starred and also marked weak, sits here for a different reason: its experience bank is built from training tasks, which blurs the boundary of its loop, so it fails the audit (Appendix E.1.2).

[Human × Auto] is empty in both passes: no surveyed system puts a person on the local gate of a failure-driven loop and then promotes the result by program. A plausible occupant would ship person-approved skills through an automatic held-out evaluation; no published system does.

[Human × Human] is empty too. The registries that a first reading puts here (Azure API Center, the agentic-community MCP Gateway and Registry, TrueFoundry, Buildrix) have no failure-driven loop: a person registers an artifact and a person, or an admission scan, approves its use, which is a promotion gate on a registrant-supplied artifact. They sit in **[None × Human]** with the Anthropic skills repository and Agentic Configuration Management, whose lifecycle states a governance role assigns, and the five public skill repositories of the field study (Section 6) show the same shape in practice. Enoch's checkpoint-and-validate protocol moves one agent's state between hosts and is neither loop. The released table records these re-codings (Appendix G.7).

[Human × None] is the curation cell, also filled only from the second pass and single-loop by construction. In VerificAgent an agent drafts, a person fact-checks, and the result is frozen; in SkillChain a judge filters the candidate edits and a person decides, at a reflection gate, which refined skills are deployed; in the Davis et al. case study one engineer both writes and accepts the governance mechanisms that failures call for (Section 6). Curation stays patch-local, rather than collapsing into an engineer editing a configuration file, because the candidate answers a failure and a person decides whether it stays.

The occupants, with the variable the matrix does not show:

Cell	Systems	Candidate source	Loop 1 gate	Promotion mechanism	Loop 2 gate
[Auto × Auto]	NanoResearch, SkillOpt, AlphaEvolve, ADAS, DGM, Voyager, SAGE, EvolveR, CoEvoSkills, AutoManual	agent proposals from trajectories or failures	evaluator, verifier, or threshold	merge into a shared bank or archive	the same evaluator, on the shared bank
[Auto × Human]	CASCADE, Sierra Agent OS 2.0, Cognition Devin	agent proposals	calibrated judge or outcome threshold	release to the team or product	reviewer, domain expert, or merge approval

Cell	Systems	Candidate source	Loop 1 gate	Promotion mechanism	Loop 2 gate
[Auto × None]	Letta Context Repositories; Mem0 Platform, TOKI, Learning on the Job, RASO, VACE; AutoAgent	agent (Letta: sleep-time subagents)	automatic write gate (RASO: validation-set acceptance; VACE: held-out comparison against the incumbent harness; AutoAgent: an uncalibrated consistency judgment)	Letta: unconditional git push to shared memory; the others: none	none
[None × Auto]	Braintrust, LangSmith Hub, Vellum, Governed Capability Evolution; EDDOps, with a sign-off after the threshold	engineers	none (no failure-driven loop)	evaluation-gated release	evaluation threshold
[None × Human]	Azure API Center, MCP Gateway and Registry, TrueFoundry, Buildrix, Anthropic skills repository, Agentic Configuration Management	registrants, contributors, engineers	none	admission to a registry, merge into a shared repository, or a lifecycle state	approval by a person (the MCP Gateway adds an automated admission scan)
[None × None]	SkillRL, SkillForge	agent proposals from a success-rate trigger or a judge's failure flags	none (the rewrite is kept unchecked)	none	none
[Human × None]	VerificAgent, SkillChain, Davis et al.	agent drafts (VerificAgent; SkillChain after a judge); the engineer (Davis et al.)	a person	none	none

E.2.2 Closest approaches to governed cross-tenant scaffold optimization

Section 8 defines the four properties: (a) an auto-gated Loop 1, (b) cross-organization promotion, (c) versioned lineage with access control, and (d) rollback. The closest approaches split them among themselves:

- **AlphaEvolve** has (a) and a partial (b) confined to one organization; it has no mechanism for promotion across organizations.
- **NanoResearch** has (a) at the research level, and its paper describes no enterprise governance layer.
- **SkillOpt** has (a) fully and a partial (c), through its exported `best_skill.md` artifact and protected slow-update field. Its reported positive transfer across models, harnesses, and benchmarks is an empirical measurement of evidence radius. It implements neither (b) nor explicit rollback (d). The walkthrough in Appendix F maps each of its design decisions onto the loop.
- **CORAL** [63] has a partial (a) through reflection-written shared memory and an automatic Loop 2 through heartbeat-based promotion; it has no (b), no formal access control for (c), and no (d). It is the only system classified before the second pass whose Loop 2 runs asynchronously across multiple agents inside one organization, through shared persistent memory. Evo2Team and ShareMem, second-pass systems, promote across teams and across users without versioning.
- **PACE** [56] has (a) at small-model production timescales and a fast/slow two-timescale pair of loops that consolidates across one deployment's own episodes, not across deployments; it has no (b) and no formal (c) or (d). It is the closest demonstration that the two-loop architecture runs under small-frozen-model constraints.
- **Meta-Harness** [11] has (a) over the harness (substrate S5), with a filesystem of candidates that records credit assignment and serves as the accumulated history H_t of Appendix A. It has no cross-organization promotion, no access control, and no rollback.
- **SkillForge** has no (a), since its calibrated judge selects failures and nothing tests the rewrite, and a partial (c) through VFS commits, but no (b): each rewritten skill version stays in the scenario deployment that produced it, and its support engineers review replies, not skill versions; (d) is undocumented in the public report.
- **Sierra Agent OS 2.0** has a partial (a). Its cross-organization promotion is claimed in vendor documentation only and not counted as evidence under the tier rule of Section 3, and its GitHub-style Workspaces give some lineage but not access control with rollback in the strict sense.
- **FederatedSkill** [68] has (a) and, with EvolveNet (Table 6 of the main article), the strongest partial (b) among the surveyed systems: heterogeneous clients, each a backbone-plus-harness profile under one federation operator, exchange semantic skill diffs, and a server-side evolution agent applies, adapts, or rejects each patch per client. An auditable decision log, recording source patch, reward, and justification for each modified path, gives a partial (c); there is no access control and no (d).
- **SkillClaw** [155] has a partial (a), through a nightly validation step, and a partial (b) across the users of one shared, synchronized skill repository; it reports no (c) or (d).
- **ChronoMem** [95] has (d) and the versioning half of (c) for the memory substrate alone: whole-memory snapshots, version histories, and natural-language rollback inside Google's Agent Development Kit. It has no Loop 1, no (b), and no access control.
- **Vellum, Braintrust, and LangSmith Hub** have (c) and (d) explicitly but neither (a) nor

(b) in the strict sense; they are governance-first by construction.

No surveyed system combines (b), (c), and (d) under any gate type. The cross-organization governance triad is unfilled on its own, whether or not Loop 1 is auto-gated.

E.2.3 Whether the missing architecture should be built

Whether the missing architecture should be built depends on the feedback-signal class, not on the application domain: safety-critical domains such as medical decision support, legal advice, and financial advice, where a reviewer's gate on promotion is a design feature rather than a defect, are judged or delayed domains with high stakes. Section 8 gives the argument class by class.

F The convergent research families, system by system

This appendix walks the convergent research families of Section 9 system by system, at the level of mechanism, so that each conclusion of that section can be checked against what the systems do. The walkthroughs were written before the second pass and updated where the audit of Appendix G revised a claim; the members the second pass adds to each family are cited in Section 9 and appear in Appendix I.

F.1 The convergent landscape

SkillOpt [52] is the clearest instance of artifact-layer descent we found, and it is one of at least eighteen independent research threads from 2025–2026 that converge on the same disciplined, gradient-like update for the scaffold layer. They converge from different starting points: GEPA from prompt optimization and Tiwari et al. from continual learning; RCL from cognitive-agent learning; Meta-Harness, AHE, MOSS, CANTANTE, and Ong et al. from harness engineering; MOCHA and Ratchet from skill-deployment constraints; ICT, Combee, CORAL, and PACE from cross-task, multi-agent, and production deployment. Each instantiates a different subset of the five conditions of Section 2, and each exposes a different gap.

By substrate (Section 3), GEPA and RCL work on the prompt and context layer (S1–S2); the harness family on orchestration (S5), with MOSS reaching into source code (S2+S5); MOCHA and Ratchet on skills under deployment constraints (S2); the transfer family on integrated memory-plus-orchestration stacks (S3+S5); and SkillOpt on an integrated stack led by S2.

No member of any family combines the four properties of Section 8: (a) auto-gated Loop 1, (b) cross-organizational Loop 2, (c) versioned lineage with role-based access control (RBAC), and (d) rollback. The loop of Section 2 is realized in fragments across the families, and every member leaves the Section 8 gap open.

F.2 SkillOpt: full correspondence

The table maps each component of SkillOpt to its quantity in Section 2, with the role that quantity plays and a pointer into the published procedure (Algorithm 1 of the SkillOpt paper).

Section 2 quantity	SkillOpt instantiation	Section 2 role explained	Algorithm pointer
θ_s (scaffold)	`best_skill.md` — a 379–1,995-token markdown document	the trainable external state of the frozen target model; persistent and exported as the deployment artifact	s_{best} (Algorithm 1, lines 1, 22, 37)
L_D (patch loss)	benchmark-native hard score / exact-match on a held-out test split	scalar quality signal that the optimizer is implicitly minimizing	EVALUATE(M, h, s, D) (lines 2, 17, 37)
Q_t (proposal)	optimizer-model reflection over failure and success minibatches; failures and successes analyzed separately, then merged with priority on failure fixes	proposal distribution conditioned on accumulated rollout evidence and current scaffold	optimizer calls O for failure analysis, success analysis, and merge (lines 9–11)
z_t (edit)	structured add/delete/replace patch on the skill (or rewrite-from-suggestions in rewrite mode)	the candidate update applied via $\oplus$ to the current scaffold	rank/clip to η_t edits, apply to obtain $\tilde{s}$ (lines 12–13)
η_t (textual learning rate)	bounded per-step edit budget; constant / linear / cosine / autonomous schedules; default cosine with floor $\eta_t = 2$	edit-magnitude controller; the step size of Section 2 in textual form	rank/clip step (line 12); schedule applied per epoch
G_t (gate)	held-out selection-split accept gate, strictly greater than current selection score	the per-edit gate G_t — accepts only when $\widehat{\Delta}_{z_t} L_D$ improves the selection score	comparison at line 20; cache lookup at lines 14–18
rejected-edit buffer	epoch-local store of failed proposals and the score drop they caused	converts rejected steps into negative feedback for later proposals within the same epoch; zero deployment cost	$\mathcal{B}$ updated at line 26; reset at line 5
slow/meta update	epoch-end longitudinal guidance written to a markup-fenced protected slow-update field, still passed through D_{sel}	second-order consolidation across adjacent epochs; step-level edits cannot overwrite the protected slow-update field	lines 28–31; gate applied to injected guidance
optimizer-side meta skill	m_{meta} — a prompt-side summary of which edits helped, which failed, and which failures persisted	training-time guidance for Q_t; never shipped with the deployed artifact	lines 33–34 (training-side only)

The mapping is one-to-one at the level of Algorithm 1: every Section 2 quantity has a named, code-level counterpart, and every SkillOpt control surface maps onto a Section 2 quantity. SkillOpt is therefore the clearest instance we found in which the loop of Section 2 is *mechanism* rather than *metaphor*: all five conditions are instantiated.

F.3 What SkillOpt does that prior textual-optimization systems do not

SkillOpt is compared with six baselines under one protocol (same target model, held-out test split, and scorer), and each baseline lacks a specific piece of the loop.

- **TextGrad** [142] backpropagates textual feedback through a compound system's variables but leaves no artifact distinct from optimizer state, which is why Appendix E.1 places it outside the core corpus; it has no step bound, held-out gate, or rejected-edit buffer. SkillOpt leads it on direct chat by +5.9 to +34.0 points across benchmarks at GPT-5.5 scale, and at all seven target-model scales.
- **GEPA** [114] evolves prompts by reflection under Pareto-style selection. Its target is a prompt rather than a persistent skill, its selection is not a strictly-greater held-out gate, and nothing bounds an accepted edit, so one edit can rewrite the prompt wholesale. SkillOpt matches or beats it in every reported cell of its main results table [52].
- **Trace2Skill** [156] distils trajectory-local lessons into reusable skills offline. No held-out loop accepts or rejects each candidate, so delta validation (operation 4 of Appendix D.1) and the rejected-edit buffer are missing. SkillOpt leads it by +5.0 to +19.0 points at GPT-5.5 direct chat.
- **EvoSkill** [157], the strongest baseline inside the Codex and Claude Code harnesses, evolves a skill folder from failure analysis. Its credit assignment is genuine, but per-step folder edits are unbounded and rejected edits are not kept. SkillOpt leads it by +14.0 inside Codex and +3.2 inside Claude Code on the GPT-5.5 five-benchmark average.
- **A human-written skill** (145–516 tokens per benchmark) persists but never updates: the no-Loop-1 ceiling.
- **A one-shot LLM skill**, generated once by GPT-5.5 from the benchmark description, never updates and proposes without rollout evidence: the untrained-scaffold baseline.

ABSTRAL [158] and EvoTest [159], cited by SkillOpt but not measured, work at the multi-agent-design and test-time-configuration layer and fall outside the comparison.

The live difference across the six is four controls: a persistent edit target distinct from optimizer state (θ_s); a held-out gate on every acceptance (G_t, with estimator $\widehat{\Delta}$ and threshold τ_t); rejected edits kept as negative feedback for Q_t; and a bounded, scheduled edit budget (η_t). SkillOpt is the first system we found that has all four, run as one offline training loop whose deployed artifact, a 379–1,995-token markdown file, needs no optimizer calls at inference.

Its follow-up, SkillOpt-Lite [15], asks which controls are necessary. Recasting skill optimization as zeroth-order search, it keeps three principles: trajectory exploration through the file system, consensus attribute mining (attributes common across heterogeneous rollouts, single-sample quirks discarded), and a held-out gate strictly disjoint from the training failures. It drops minibatch reflection pooling, the rejected-edit buffer, the slow update, and the learning-rate schedule. It still converges faster and scores higher than full SkillOpt (LiveMath +8.8 points at GPT-5.5, +25.4 at GPT-5.4-nano). The gate and a consensus form of credit assignment are load-bearing; negative-feedback retention and the bounded step are stabilizers whose value depends on the benchmark. None of this touches the governance properties of Section 8.

F.4 GEPA, Tiwari, and the reflective-evolution family

The immediate predecessor to SkillOpt's optimizer framing is **GEPA** [114]. It shows that natural-language reflection over execution trajectories can provide a richer adaptation signal than sparse scalar rewards for prompt evolution in compound AI systems. GEPA samples trajectories containing reasoning steps, tool calls, and tool outputs; reflects on them to diagnose failures and propose prompt updates; and maintains a Pareto frontier of candidate prompts rather than greedily mutating only the current best. Across six tasks, GEPA outperforms GRPO by 6% on average and by up to 20% while using up to 35× fewer rollouts, and also outperforms MIPROv2 in the reported aggregate comparison. In Section 2 vocabulary, GEPA strongly instantiates the proposal side, Q_t, through trajectory-conditioned reflection, and uses Pareto-based candidate selection as a proposal and pruning mechanism. What it does not instantiate is the full SkillOpt-style artifact lifecycle: a bounded edit budget on a reusable skill document, a strictly-greater held-out gate on each accepted skill update, a rejected-edit buffer used as negative feedback, and an exported deployment artifact with governance hooks.

Tiwari et al. [115] supply the broader framing from the optimization side, arguing that restricting learning to *either* in-context adaptation (prompt optimization) *or* in-weights adaptation (RL or fine-tuning) is unnecessary and costly. In-context adaptation is cheap and rapid but capacity-bounded; in-weights adaptation is expressive but induces catastrophic forgetting and loss of plasticity [125]. The scaffold layer of Section 9 is the substrate that absorbs what in-context adaptation can carry cheaply, leaving the weights free to generalize. Tiwari et al. reach the two-timescale motivation from optimization and Section 9 reaches it from production reliability; the two framings are independent and reinforce each other.

F.5 Reflective Context Learning: a parallel formalization

The most direct independent convergence on the Section 2 formalization comes from **Reflective Context Learning** [RCL; [116]]. Their opening observation is that "the fundamental problems of learning, including credit assignment, overfitting, forgetting, local optima, and high-variance learning signals, persist whether the learned object lies in parameter space or context space," and that current methods are "fragmented and ad hoc" precisely because they address these challenges one at a time rather than as a unified optimization problem over the context. The parallel to Section 2's argument, that scaffold maintenance needs the discipline of weight-space optimization applied to an external artifact layer, is exact and was arrived at independently.

RCL maps onto Section 2 as follows. The *learned context* is θ_s. The *reflection signal* is a form of credit assignment identifying which context coordinate caused a failure (Appendix D.1, operation 2). The *generalization penalty* that RCL introduces to prevent narrow sample-specific rules is structurally the complexity penalty $\Omega(z)$ of the gate condition. Where RCL and Section 2 differ is in Loop-2 scope: RCL is framed around single-agent cross-task generalization, while Section 2's Loop 2 is cross-deployment, cross-tenant, and cross-organizational promotion. The convergence is nonetheless significant: two research threads starting from different motivations (production reliability in Section 2, cognitive agent learning in RCL) arrive at the same formal observation that the classical learning problems do not disappear when the learned object moves outside the weights.

F.6 The harness-optimization cluster

A second family optimizes the *harness*, the code and configuration that wrap a frozen model, rather than skill documents. It names credit assignment as the central unsolved problem of scaffold

evolution, and its solutions instantiate two operations of the mechanism stack in Appendix D.1: mode discovery with credit assignment (operation 2) and delta validation (operation 4).

Meta-Harness [11] starts from the observation that system performance "depends not only on model weights, but also on their harness", the scaffold-versus-weights decomposition stated for this substrate. Existing text optimizers, it finds, "compress feedback too aggressively" for credit assignment, so it gives an outer-loop proposer filesystem access to the source code, scores, and execution traces of every prior candidate. The edit history becomes the credit-assignment signal. In Section 2 terms the filesystem is the accumulated history H_t that conditions the proposal Q_t, and the outer loop closes Loop 1 over orchestration (S5) rather than skills (S2).

Agentic Harness Engineering [AHE; [9]] comes at the problem from observability. It names three obstacles: a heterogeneous action space across editable components, voluminous trajectories that bury the actionable signal, and edits whose effect is hard to attribute. They are, in order, coordinate identification (operation 2), residual capture (operation 1), and the credit-assignment half of operation 2. AHE's answer, three observability pillars that give every editable harness component a file-level representation, builds the coordinate decomposition that Section 2 assumes. Descent needs that instrumentation as much as it needs the loop.

MOSS [117] observes that earlier self-evolving systems "confine evolution to text-mutable artifacts" and leave the harness untouched. Routing, hook ordering, state invariants, and dispatch live in code, so a whole class of structural failure is, in MOSS's phrase, "physically unreachable from the text layer". MOSS therefore extends Loop 1 to rewriting source code, evolving skills (S2) and orchestration (S5) together. The consequence for Section 8: a governed optimizer that covers only text artifacts leaves the code substrate ungoverned.

CANTANTE [10] takes up credit assignment in multi-agent scaffolds, where "scores are available only at the system level, whereas the parameters governing agent behavior are local". It contrasts rollouts of several jointly acting configurations to split a system-level reward into per-agent update signals. That gives the credit-assignment step a concrete mechanism when the scaffold is spread across agents (S5): the attribution, succeeding with non-zero probability per failure, that the correspondence note of Appendix A requires.

Harness-optimizer evaluation [14] finds that judging scaffold optimizers "solely by observing target agents' performance gains" misses intermediate erroneous edits that hinder performance, and that it is "unclear whether harness optimization is driven by optimizers' informed update actions or simply trial-and-error". That is the line between Stage 2 and Stage 3 of Section 2, measured. An optimizer that cannot tell informed updates from lucky ones operates in Stage 2, whatever its benchmark gains.

HARNESSEVO [12] turns the diagnosis into a measurement. It splits the harness into four separately evolvable slots (role, task strategy, tool and format rules, reflection and control) and attributes value by leave-one-in and leave-one-out evolution. On ALFWorld with a 7B target model nearly all of the gain sits in the reflection and control slot (leave-one-in +0.119; the other slots individually null). Spread evenly, a fixed 64-rollout budget leaves every slot below the optimizer's effective search floor, so each freezes at its seed; concentrated on the high-value slot, it recovers the gain (0.642 to 0.761). On WebShop the effect is absent. Credit assignment over coordinates has to precede budget allocation, and which coordinate carries the value depends on the task.

F.7 Multi-objective and lifecycle constraints

Descent on a single loss misses two kinds of constraint: the platform limits a deployed skill must satisfy, and the hygiene of the artifact pool from which the gate draws its signal.

MOCHA [119] observes that agent skills are "multi-field artifacts subject to hard platform constraints": descriptions are truncated for routing, instruction bodies are compacted by progressive disclosure, and co-resident skills compete for the context window. A skill must maximize task performance and satisfy those limits at once, and optimizing performance alone yields skills that pass benchmarks but degrade on real deployment infrastructure. MOCHA navigates the constraint frontier with *Chebyshev annealing*. It operationalizes the complexity penalty $\Omega(z)$ of the gate, and it shows that the scalar Ω of Section 2 simplifies what is, in production, a Pareto frontier.

Ratchet [34] places the bottleneck elsewhere. LLM-authored skills deliver +0.0 pp over no-skill baselines while human-curated ones deliver +16.2 pp, which points away from the optimizer, the mechanism that proposes and accepts edits, and toward *lifecycle management*: whether skills are written, retrieved, curated, and retired correctly. Four hygiene mechanisms (outcome-driven retirement, a bounded active cap, meta-skill authoring, and pattern canonicalization) lift the +0.0 pp floor substantially before any optimizer runs.

The finding clarifies a precondition of descent rather than refuting it. The gate needs a meaningful signal: if retrieval injects the wrong skill into every rollout, the proposal process conditions on corrupted evidence and descent degrades to random search. Ratchet's mechanisms keep residual capture and mode discovery (operations 1 and 2) working. Retirement keeps stale artifacts out of the residual signal, the active cap keeps scaffold bloat from making retrieval unreliable, and canonicalization collapses near-duplicates so that proposals condition on one clean artifact set. A system with strong hygiene and no optimizer sits at Stage 1 with a clean scaffold; one with a strong optimizer and poor hygiene sits at Stage 2 with a noisy loss. Stage 3 needs both.

SkillsVote [13] carries the lifecycle argument to an open ecosystem. It profiles a million-scale open-source skill corpus for environment requirements, quality, and verifiability, synthesizes tasks for the verifiable subset, and exposes skills through agentic library search before execution. After execution it splits each trajectory into skill-linked subtasks and attributes their outcomes to skill guidance, agent exploration, environment, or result signals; only successful, reusable discoveries pass into evidence-gated updates. Attribution thus happens at the admission gate rather than inside the optimizer, credit assignment used as a filter on residual capture. The redundancy and environment sensitivity SkillsVote measures are the bloat and drift of Section 7, arising here from many authors rather than one loop.

F.8 Cross-task transfer and evidence radius

Several papers measure what Section 2 calls the *evidence radius* of a scaffold artifact, how far an edit validated in one context keeps helping when redeployed elsewhere, without using the term.

In-context Training [ICT; [120]] tests cross-task self-improvement: whether experience distilled from one set of tasks transfers to unseen ones. Its finding that "current reflection-based methods can only self-correct within a single task instance" measures the radius at its lower bound, one task. ICT tries to widen it by distilling reflections into reusable lessons, which instruments Loop 2 empirically: does an edit that passed the gate on one task's L_D still help across a distribution of unseen tasks?

Combee [121] reaches the same measurement through scale. ACE [143] showed that a system prompt which accumulates and organizes strategies through agentic interaction improves on a static prompt for a single agent; Combee learns from the traces of many agents running in parallel. Each proposal then draws on a wider deployment distribution, and the promoted artifact is tested against a larger average of patch losses. This is not cross-organizational promotion. Combee measures how much of the evidence-radius benefit parallelism inside one organization captures.

CORAL [63] runs long-lived agents that explore, reflect, and collaborate through shared persistent memory, asynchronous execution, and heartbeat-based interventions. A finding one agent writes to the shared store becomes available to all, so the store is the cross-agent promoted artifact and the heartbeat a form of auto-gated Loop-2 promotion; there is no versioned lineage, RBAC, or rollback. CORAL is the only fully asynchronous multi-agent Loop-2 promotion mechanism among the systems classified before the second pass; in the full corpus, Evo2Team and ShareMem add cross-team and cross-user promotion without versioning (Appendix G.3).

PACE [56] is a production-oriented two-timescale design for small language-model agents. The fast timescale adapts prompts, parsers, validators, and control logic from failure signals within a deployment, as Loop 1 does; the slow timescale consolidates across the deployment's own episodes, a Loop 2 whose radius is one deployment. The two-loop structure holds on small, frozen production models, not only under frontier-model assumptions, and persists even when the optimizer and the target are the same frozen small model.

SkillClaw [155] and **FederatedSkill** [68] are the two systems of these walkthroughs whose Loop 2 crosses users rather than tasks or agents; in the full corpus, EvolveNet, ShareMem, and XRepoSkill also promote across users, clients, or repositories, each inside one operator (Section 6). SkillClaw aggregates many users' trajectories, turns recurring behavioral patterns into skill-set updates, and synchronizes the updated skills across users; its evaluation runs a day-night cycle in which simulated users generate trajectories by day and the system reflects, proposes, validates, and deploys by night. FederatedSkill exchanges semantic skill diffs, structured patches over a local library shipped with the task reward that produced them, so that clients share what they learned without sharing raw trajectories. A server-side evolution agent aggregates the diffs into a per-client capability matrix (covered, absorbing, broken, gap) and applies, adapts, or rejects each peer patch per client, logging source patch, reward, and justification for every modified path. Across 20 task families it reports up to +44.4% success and -37.5% compute against isolated self-evolution. FederatedSkill is the first surveyed system in which Loop 2 aggregates artifact updates rather than promoting a single artifact, and its decision log is the first lineage record at an aggregator. Its acceptance, though, uses the source client's reward rather than a held-out evaluation at the receiving client.

F.9 Synthesis: what the convergence establishes

Sections 4 and 9 draw four conclusions from the families, and the walkthroughs support them as follows.

1. **The scaffold-versus-weights decomposition is independently confirmed.** GEPA, Meta-Harness, MOSS, RCL, and PACE all begin from the observation that model weights and the external artifact layer are distinct adaptation surfaces with different cost, reversibility, and scope properties. No surveyed paper disputes this decomposition; the variation is in which substrate each paper targets (skill document, harness code, system prompt, shared memory).
2. **Credit assignment is the central engineering bottleneck.** AHE, CANTANTE, Meta-

Harness, the harness-optimizer evaluation of Ong et al. [14], HARNESSEVO, and SkillsVote all identify credit assignment, attributing a system-level failure to a specific scaffold coordinate, as the problem that separates disciplined Loop-1 evolution from trial-and-error. It is operation 2 of Appendix D.1, the attribution that the correspondence note of Appendix A requires, and it remains unsolved in the general multi-substrate case.

3. **Lifecycle hygiene is a precondition for descent, not an alternative to it.** Ratchet establishes that the gradient-analog mechanism of Section 2 requires a maintained scaffold. A dirty artifact pool (stale, bloated, miscurated) keeps the gate G_t from receiving a meaningful signal. Hygiene and optimization are sequential dependencies, not competing paradigms.

4. **No surveyed system combines Loop-1 automation with cross-organizational Loop-2 governance.** SkillOpt has the strongest Loop-1 closure and the most systematic evidence-radius measurement, but it does not implement cross-organizational promotion, versioned lineage with RBAC, or rollback. CORAL implements asynchronous cross-agent promotion within one system but not across organizations. PACE implements two-timescale consolidation within one deployment. Meta-Harness implements Loop-1 closure over the harness substrate but no Loop-2 promotion mechanism. FederatedSkill federates skill patches across clients under one operator, with a decision log but no RBAC or rollback, and EvolveNet adds a held-out merge gate and round rollback without lineage or RBAC. The four-property gap of Section 8 is narrowed on (b) and (d) by different systems and closed on none, and the independent convergence of several research threads on the same single-organization ceiling confirms it.

F.10 SkillOpt against Section 8's four properties

Mapped onto Section 8's four properties of governed cross-tenant scaffold optimization:

- (a) **Auto-gated Loop 1.** *Satisfied.* The held-out selection-split gate at strictly-greater margin instantiates G_t.
- (b) **Multi-tenant cross-org promotion.** *Absent.* SkillOpt is a single-organization research system; there is no governed cross-tenant aggregator.
- (c) **Versioned lineage with RBAC.** *Partial.* The exported `best_skill.md` is a versioned artifact, and the protected slow-update field separates fast and slow consolidation. There is no documented RBAC layer or audit-trail policy, and `edit_apply_report.json` is an operational log rather than an access-controlled lineage system.
- (d) **Rollback.** *Absent.* The training loop selects the best validation-gated skill and exports it; no rollback envelope on deployed artifacts is reported.

The cross-model, cross-harness, and cross-benchmark transfer experiments reported by SkillOpt are *empirical evidence-radius measurements*: a skill optimized in one (model, benchmark, harness) cell continues to improve performance when redeployed in nearby cells without further optimization. In Section 2 vocabulary, this is a measured Loop-2 *radius* for the artifact, not a *governance mechanism* for promoting it. Both readings are useful: SkillOpt is the strongest Loop-1 proof and supplies the first systematic transfer evidence at the scaffold layer, while the Section 8 cross-tenant governance triad ((b), (c), (d)) remains unfilled.

F.11 Cost and bloat

SkillOpt reports 0.6M to 46.4M training tokens per absolute test-point gain (their Table 6), with the procedural benchmarks (SpreadsheetBench, OfficeQA, LiveMath) at the cheap end and search/visual benchmarks (SearchQA, DocVQA) an order of magnitude more expensive per point. The deployed artifact remains 379–1,995 tokens regardless. This is, to our knowledge, the first published cost-per-point figure for scaffold-layer descent, and it complements the Section 7 scaffold-bloat discussion (and the Ratchet hygiene-precondition finding walked in Appendix F.7): the textual learning rate η_t with cosine decay keeps the deployed skill compact even when the optimizer explores many more edits than are accepted (only 1–4 edits survive into the final skill across the six benchmarks). The mechanism that prevents the deployed artifact from bloating is exactly the complexity penalty $\Omega(z)$ of Section 2, instantiated as a hard cap on accepted edits per step rather than a soft regularizer.

G Audit of the first-pass draft's absolute claims against the full corpus

The first-pass draft of this survey, written on the 130 systems of the first pass, made absolute claims of the form *no surveyed system*, *only X*, *first*, *absent*, and *empty cell*; each row below sets one of them against the full corpus of 256 systems and records what that corpus shows. A claim **holds** when the full corpus bears it out as worded, is **sharpened** when it is true and the corpus supports a more precise statement, is **revised** when the corpus contradicts its wording, and is **dropped** when the article no longer makes it; for sharpened and revised claims the Evidence column also states the claim as the full corpus supports it. Evidence names the systems that decide each verdict, and parenthetical numbers such as 2608.04968 are arXiv identifiers.

G.1 The four properties of Section 8

Claim	Verdict	Evidence
No public system combines (a) auto-gated Loop 1, (b) cross-organization promotion, (c) versioned lineage with role-based access control (RBAC), (d) rollback	**holds**	Across the 126 second-pass systems: no system demonstrates promotion across organizations, and EvolveNet has a partial (b) inside one operator (its clients are benchmark partitions, "organizations" only in motivation), while RASO, which retrieves from a public corpus, promotes nothing; MemClaw is single-tenant in its evaluation, and Buildrix is an open community hub whose artifacts a person approves, with no failure-driven loop; (c) appears as lineage without RBAC everywhere (EDDOps, Agentic Configuration Management (ACM), MemClaw, Letta, Mem0 history) and as RBAC without versioned promotion in the MCP control planes (agentic-community gateway, TrueFoundry); (d) is demonstrated by Governed Capability Evolution and EDDOps and claimed by EvolveNet (round rollback) and EvoUndo (recoverability as admission), none of which has (a)+(b).
"No surveyed system meets the FedAvg limit; FederatedSkill comes closest"	**sharpened**	EvolveNet (2608.04968) comes nearer: it broadcasts a shared harness, optimizes it locally at each client, merges on the server only when the merge fixes at least as many held-out items as it breaks, and rolls back to the previous shared harness after a failed round. It is still selection and merge over program edits, not averaging, so neither it nor FederatedSkill meets the FedAvg limit.
Fourth conclusion on the research families: "FederatedSkill comes closest ... narrowed on one property, closed on none"	**revised**	The two closest approaches are FederatedSkill (decision log; no gate at the receiving client) and EvolveNet (held-out merge gate and round rollback; no lineage or RBAC). The gap is narrowed on (b) and (d) by different systems and closed on none (Table 6).
"no member of the Section 9 research families provides (d), and only FederatedSkill provides even a partial (b)"	**revised**	(d) appears in EvolveNet (round rollback), SkillSmith and SkillAudit (edit-level rollback inside one agent), Growing Harness, WDA, and Co-Harness (rollback within an optimization run; Co-Harness adds a partial (c), a versioned diff registry without access control), and EvoUndo (recoverability as an admission test); EvolveNet provides a partial (b). Rollback appears inside single deployments and inside one federated round; no system rolls back a promotion that has crossed an organization boundary, because none makes such a promotion.

G.2 Matrix cells

Claim	Verdict	Evidence
"Three cells are populated; the rest are empty for structural reasons"	**revised**	Single-loop and governance-only systems occupy these cells, under the definition of Section 2 that a gate records who decides and not who supplies the candidate: [None × Human] (MCP registries: Azure API Center, agentic-community MCP Gateway & Registry, TrueFoundry; Buildrix; the Anthropic skills repository; ACM), [Human × None] (VerificAgent, SkillChain's second-stage reflection gate, Davis et al. governance conversion), [Auto × None] (Letta's ungated push; Mem0 Platform, TOKI, Learning on the Job, RASO), [None × Auto] (Governed Capability Evolution; EDDOps, with a sign-off after the threshold), [None × None] (SkillRL and SkillForge, failure-driven skill loops that run unchecked; Official MCP Registry; Enoch's migration protocol; Continual Harness's unchecked self-edit; the revocation guard). Two cells hold every audit passer whose promotion has a gate, a third holds the governance-first platforms, Letta passes with no promotion gate, and the human-gated row holds only curation (Section 6).
"[Human × Auto] is empty by engineering logic . . . no published example exists"	**holds**	No system in the full corpus puts a person on the local gate of a failure-driven loop and promotes by program. MCP Gateway & Registry adds an automated security scan to a human admission, but it has no failure-driven local loop, so it sits in [None × Human].
"[Human × Human] is empty in the surveyed corpus"	**holds, after re-coding**	The registries and the community hub that a first reading places here have no failure-driven loop and sit in [None × Human] (first row). The occupant the first-pass draft predicted, solo or small-team workflows, is documented by "Who Maintains Agent Skills?" (2609.05677): across five public skill repositories every substantive edit is human-authored and human-merged, 62% are AI co-authored, and there is 1 deprecation in 254 edits; that shape is [None × Human] too, because no loop fires from session signal (Section 6).
"[Human × None] is structurally improbable"	**revised**	VerificAgent (agent-drafted memories, human fact-check, frozen bank), SkillChain's second-stage human reflection gate, and the Davis et al. governance conversion occupy the cell. It is the human-curation cell, single-loop by construction.

Claim	Verdict	Evidence
"Nearly every auto-gated system fires at n = 1" (accepts an edit on a single evaluation)	**revised**	Batched or statistical acceptance is common, and has been since mid-2026: SAGE-gate (paired per-item test), Future AGI (paired bootstrap CI), AutoSaddler (mini-batch failure signals), MGM, ModularRSI (evidence across tasks), VaG (3-of-3 critics), EvolveNet (fixed $\geq$ broken on held-out), VACE, HackProbe, PROCTOR. Evidence is batched before acceptance, with SAGE-gate and VaG as instances (Section 4).
"no surveyed system has formalized agreement as a k-of-n acceptance rule"	**revised**	VaG requires all three heterogeneous critics to pass (3-of-3), and SkillSpec is consensus-gated, so agreement-based acceptance exists. The distinction between agreement and averaging stands: under agreement the false-accept rate decays as p^k, under averaging the variance decays as $1/k$ (Appendix B).
"Only AlphaEvolve demonstrates explicit multi-layer depth"	**sharpened**	HarnessBank, JIT-Agent, and MGM keep archives for lineage, and VaG adds tiered subset selection; none pairs a cascade evaluator with a diversity archive. AlphaEvolve is the only system that combines the two.

G.3 "Only" and "first"

Claim	Verdict	Evidence
CORAL "the only asynchronous multi-agent Loop-2 mechanism in the corpus"	**revised**	Evo2Team (2609.34135) promotes routing and communication skills across teams of 4–32 agents with target-side confirmation, and ShareMem pools memory across users. CORAL is the only such mechanism among the systems classified before the second pass; in the full corpus Evo2Team and ShareMem add cross-team and cross-user promotion without versioning (Appendix F.8).
AlphaEvolve Strassen claim	**holds** (attributed)	The claim is AlphaEvolve's reported improvement on Strassen's algorithm for 4×4 matrix multiplication. The result is taken from the AlphaEvolve report and was not independently measured; with that attribution in the same sentence, the claim holds.
"CASCADE reports the largest within-benchmark ablation gain in the corpus"	**dropped**	Not established for the 126 second-pass systems, and ablation gains measured on different benchmarks are not comparable. The article does not make the claim.

Claim	Verdict	Evidence
ChronoMem and rollback repair "the first surveyed instances of the rollback envelope"	**revised**	Memory rollback also appears in Mi-Memory (as a design) and Letta (git history), and checkpoint-rollback studies (2608.29381, 2608.15939) show rollback unsafe. VerMem keeps prior versions and soft-deleted records but has no operation that reverses an edit; the "Restore" in its tool table moves episodic fragments of the current task back into context. ChronoMem and rollback repair are among the first instances, and the checkpoint-rollback studies are the counter-evidence.
Cuadros "only surveyed memory system with full provenance + versioning + correction"	**holds**	The memory-governance systems of the second pass (TOKI, PPMF, TMA-NM, MemClaw, He & Yu) each carry provenance or lineage without correction loops or versioning, and the auto-gated memory loops (Recuris, MemGuard, Compiled Memory) carry no lineage, while VerMem keeps prior versions without an operation that reverts a write.
SkillOpt "cleanest published instance" / "first to instantiate all four simultaneously"	**sharpened**	SkillOpt-Lite (2607.03451) removes two of the four controls and matches or beats SkillOpt; VALVE instantiates a held-out gate that rejects a candidate before commit and never reverts an accepted update, and SkillSmith and SkillAudit instantiate held-out gates with edit-level rollback. SkillOpt is the first system we found that instantiates all four, and its own follow-up shows that two of them are stabilizers rather than necessities (Appendix F.3).
FederatedSkill "first ... federated aggregation ... first lineage record at the aggregator"	**holds**	EvolveNet (August 2026) postdates it and keeps no lineage record.
SkillOpt "first published cost-per-point figure"	**holds**	The full corpus contains no earlier figure.

G.4 Absences

Claim	Verdict	Evidence
"production deployments sit at Stage 1 or 2; none reaches Stage 3"	**sharpened**	Spotify (2609.30297) runs automated prompt and planning optimization gated by human launch and A/B metrics; this is Stage 3 on Loop 1 with a human Loop 2: production deployments that reach Stage 3 on the local loop promote behind a human gate or, in Letta's case alone, by an ungated push (Section 2). SkillForge (2604.08618) stops short of it: a judge that agrees with experts more than nine times in ten flags the failed replies that drive each skill rewrite, but the rewrite is committed without an acceptance test, and support engineers review replies, not skill versions.
No S1 system scores 5; "no instruction system closes both loops"	**holds**	On S1, n=15 and the maximum score is 4. Adoption studies point the same way: context files raise cost by more than 20% without raising success (Gloaguen), and context files dominate configuration while skills are rare (Galster).
"Anthropic's skills repository is one-pool with PR review but no automated eval gate"	**holds**	The study of public skill repositories (2609.05677) observes human merges and no evaluation gate.
"correctable, versioned memory is rare; production memory lacks failure-triggered correction and versioned lineage"	**revised**	Letta Context Repositories (git-backed memory, every edit committed, shared memory auto-pushed), Mem0's per-memory history and feedback endpoint, and Zep's bitemporal invalidation make lineage common in production memory. Failure-triggered correction is not common, and 2609.08258 shows revoked facts still outranking their replacements in five backends. Production memory is versioned, but correction is not enforced (Section 4).
No S5 system scores 5; "topology evolution is rare"	**sharpened**	With MAE and Memento-Skills corrected to 4 (Appendix G.7), S5 has no score of 5 in either pass. Gated topology evolution exists: TOCOMAS keeps only candidates that satisfy structural invariants; Evo2Team promotes across teams; AgentGrad and HarnessForge update S5 under an evaluation gate. OptiMAS updates S5 through an LLM optimizer that reads a running validation score and reverts an edit only when its own analysis refutes it, with no keep-if-better test (its Algorithm 2 returns the last configuration, and the only programmatic check is structural validity). Topology evolution is rare and, where it exists, gated by evaluation alone except in TOCOMAS, which adds a structural invariant, and OptiMAS, where the evaluation informs the edit rather than gating it.

Claim	Verdict	Evidence
"no system matures all six at once"	**holds**	None of the 126 second-pass systems scores 5 on any substrate.
"Failure-triggered memory is rare in production"	**holds**	One field report exists (SIx Harness, N=1, self-reported); Letta's dreaming writes from conversation history, not from failures.
"Cross-customer skill promotion is nearly absent"	**holds**	ShareMem, EvolveNet, and XRepoSkill promote across users, clients, or repositories inside one operator; none promotes across customers or organizations with evidence.
"Memory versioning is mostly unsolved outside Cuadros et al. and ChronoMem"	**revised**	The evidence is the same as for the claim that production memory lacks versioned lineage (fourth row of this table): Letta, Mem0 history, Zep bitemporal facts, TOKI audit rows, and MemClaw supersession chains all version memory. Memory versioning is common, not confined to Cuadros et al. and ChronoMem.
"Scaffold-level curriculum is unspecified in nearly every surveyed system"	**sharpened**	Task-CoEvolve (variance-weighted task sampling) and AutoSaddler (mini-batch failure signals) specify it; they remain the exception.
"Held-out evals at Loop-2 promotion are absent"	**revised**	Receiving-side evaluation exists: ShareMem (evaluation tasks belong to a different user), XRepoSkill (independent verification in $\geq$2 repositories), Evo2Team (target-side confirmation; 4 of 22 transfers still fail on new graphs). VACE's held-out gate compares two harnesses on the next checkpoint of the same training run, so it is a Loop-1 check rather than receiving-side evaluation and is not counted here. EvolveNet gates the merge on a server-side held-out slice but runs no check at the receiving client, and Federated Agent Optimization (2610.01195) proposes a receiving-client gate without implementing it. Receiving-side evaluation at promotion is rare, not absent (Section 6).
Research full-auto cluster: "enterprise governance is uniformly absent"	**holds**	No second-pass [Auto $\times$ Auto] system reports RBAC or audited promotion.
Governance-first cluster: "no patch-local inner loop"	**holds**	EDDOps, ACM, Governed Capability Evolution, and the MCP control planes all lack an inner loop.

Claim	Verdict	Evidence
"AGENTS.md in over 60,000 repositories"	**revised**	Adoption is measured directly rather than by one repository count: a 2,853-repository study (Galster), a 128,018-project adoption study (Robbes: 22–29% of coding-agent projects), and 36,710 repositories mined for loop engineering (Lulla, 2608.21884), alongside a controlled study in which context files raise cost without raising success (Gloaguen; Section 4).
"no surveyed system treats pruning or expiration as a first-class gated operation (outside Ratchet, SkillOpt, ChronoMem)"	**revised**	Retirement is a gated operation in at least eleven skill systems (Library Drift: retire when n $\geq$ 100 and contribution $\leq$ -0.10 with a cap of 50; SkillSmith; SkillOps; ASSAY; HiSME; SkeMex; MUSE-Autoskill; SkillMAS; SafeEvolve; Double Ratchet; SkillGen), and expiry is a product control in Mem0, Zep, and LangGraph Platform. Pruning and expiry are first-class operations; what remains rare is pruning of a promoted artifact's descendants, which VaG shows is structurally irreversible.
The failure surface is "largely unmitigated in production"	**holds**	Revocation is unenforced in five production backends (2609.08258), checkpoint rollback is unsafe in three frameworks (2608.29381), and 16 deployed frameworks keep no evidentiary record (Hearsay, 2609.32495).
Multi-sample agreement is "under-explored; SkillOpt-Lite nearest"	**revised**	The evidence is the same as for the two acceptance-rule claims of Appendix G.2, single-evaluation acceptance and k-of-n agreement: SAGE-gate, VaG, and SkillSpec. Agreement-based acceptance is instantiated, not under-explored (Appendix B).
"No surveyed paper disputes the scaffold-vs-weights decomposition"	**holds**	SIA (2605.27276) updates harness and weights together but keeps them as distinct levers.
Credit assignment "remains unsolved in the general multi-substrate case"	**holds**	HARNESSEVO (slot attribution), AgentGrad (intervention-based), and CANTANTE (contrastive) each solve one substrate's case.

G.5 The three predictions of Section 5

Section 5 predicts from the signal classes which cells of the design space each class can reach, and Section 6 confirms the predictions; the first two rows give the systems behind the two cells Section 5 assigns, and the third row gives those behind its claim about the reflexive gate.

Claim	Statement	Verdict	Evidence
First	An [Auto × Auto] system whose gate estimates the loss change per edit should run on a C1 signal or a simulated C2 environment; one that works on a C4 signal should prune proposals rather than estimate	**holds**	Before the second pass: DGM, Voyager, AlphaEvolve, SkillOpt, SAGE, EvolveR, CoEvoSkills (C1 or simulated C2); NanoResearch (semantic merge, C4) and CORAL (heartbeat, C4) prune. AutoManual estimates nothing: its Builder reads each episode's trajectory and binary outcome and writes or rewrites rules by its own judgment, with no test of the edit, and its Consolidator merges and deletes rules once they exceed twelve, so it decides by a model judge over a simulated outcome (C3 over C2) and prunes, and the claim does not bear on it. SkillRL and AutoAgent no longer sit in [Auto × Auto] (Appendix G.7). Second pass: EvolveNet, Evo2Team, XRepoSkill (C1 code signals or re-runnable benchmark outcomes, the simulated C2 of Section 5); ShareMem (judged consolidation, C3, measured on Mind2Web); Letta (no gate).
Second	[Auto × Human] should hold only C2 and C3 systems, because the human on Loop 2 substitutes for the exact signal that Loop 1 lacks	**holds**	CASCADE (C3), Sierra (C2/C3), Devin (C2 merge outcome), Spotify (C3 judge locally, C2 A/B behind human launch). No C1 system occupies the cell; SkillForge, also C3, promotes nothing and sits in [Auto × None].
Third	No scaffold system should yet calibrate a judged gate under the reflexive condition	**sharpened**	Three scaffold systems report calibration: SkillForge (>90% agreement), SkillChain (κ 0.72 to 0.91), and SkillCoach (trajectory-filter consistency 82→96% against human gold); none calibrates a judge that is itself being edited. On the weight side, no method that co-evolves a rubric or judge reports agreement with people (Appendix G.6).

G.6 The rubric-reward comparison

Reinforcement learning from verifiable rewards trains weights on exact signals, and extending it to writing and other open-ended tasks has produced the judged gates that a C3 scaffold loop can borrow. These are weight-side methods and the corpus does not score them; Section 5 reads one distinction from them, whether the rubric or judge is held fixed while the policy learns or moves with it. The comparison covers fourteen papers, thirteen methods and one survey. Of the methods, eight hold the rubric or judge fixed during policy optimization and five co-evolve it. Five report a human-agreement figure. Four of them hold the gate fixed, and the fifth, RLPR, co-evolves a probability reward rather than a rubric, checked against human labels. No method that co-evolves a rubric or judge reports one.

Gate during policy optimization	Papers	Members	Agreement with human judgment
Held fixed	8 methods	Rubrics as Rewards, an instance-specific checklist scored by a fixed judge [98]; Writing-Zero, a pairwise judge trained on writing principles and frozen for policy training [99]; reference-guided evaluation [100]; an essay scorer trained against a fixed rubric of binary items [101]; a study of cheap training-time verifiers [160]; Rubric-Grounded RL, the Open Rubric System, and query-specific rubrics learned from human preferences [161–163]	Reference-guided evaluation reaches 79.1% agreement with human annotations; the essay scorer reaches a quadratic weighted kappa of 0.803 against a single human rater's 0.805; the cheap-verifier study calibrates against language-model golden judges, not people
Co-evolved with the policy	5 methods	RLPR, scoring by the policy's own probability of a fixed reference [102]; Rubric-ARM, alternating updates of a rubric generator and a judge at the reward-model stage [103]; EvoRubrics, SCORE, and RubricEM, updating rubrics adversarially or from a rubric buffer within training [104–106]	RLPR only, for its probability reward against human labels; no method that co-evolves a rubric or judge reports a figure
Survey	1 survey	A survey of rubric evaluation, which names the dynamic regime as the trend [164]	—

The only co-evolving method with a human-agreement figure, RLPR, co-evolves a probability reward, not a rubric. No method that co-evolves a rubric or judge reports agreement with people. That is the reflexive-gate problem of Section 5, seen from the weight side.

G.7 Recount of the full corpus

The full corpus holds 256 unique systems, 130 from the first pass and 126 from the second; 204 form the core corpus and 52 the contrast corpus. None of the 126 second-pass systems scores 5 on any substrate. The composite two-loop audit passes 11 research systems and 3 industry exemplars: Sierra Agent OS 2.0, Cognition Devin, and Letta's Context Repositories, whose shared-memory push carries git lineage but no gate. The 6 second-pass systems that promote across contexts with no versioning or lineage on the promoted artifact are EvolveNet, Evo2Team, XRepoSkill, ShareMem, Spotify, and Future AGI. Of these, EvolveNet alone carries rollback and so enters under the governance relaxation of Appendix C; the first five enter only if the audit drops its governance requirement altogether; Future AGI's evidence is T5 only, so it enters not at all.

Two first-pass rows are corrected before the recount. The row "Governed Collaborative Memory / SSGM" conflated Cuadros et al. (2605.04264) with SSGM (Lam et al., 2603.11768); it is split, and

SSGM, a conceptual framework, scores 2. MAE and Memento-Skills were scored 5 on S5 in the released CSV while the first-pass draft's table said no orchestration system reaches 5; both close one loop only and score 4 under the uniform rubric, which moves the first-pass S5 mean from 2.68 to 2.58.

Nine second-pass rows had their gate labels re-coded on 2026-10-08 under the definition of Section 2, which records who decides and not who supplies the candidate: Letta's Loop-2 label moved from Auto to None (an ungated git push of shared memory; the row still passes the composite audit on its lineage); the Loop-1 label of Azure API Center, the agentic-community MCP Gateway & Registry, the Official MCP Registry, TrueFoundry, Buildrix, the revocation guard, and "Who Maintains Agent Skills?" moved from Human to None (artifacts registered or authored by people, with no failure-driven loop), and Enoch's checkpoint-and-validate protocol moved to None on both loops (a migration of one agent's state, not a promotion). The MCP Gateway's Loop-2 label now records its automated admission scan beside the human approval, and the field study's Loop-2 label records that maintainers merge. No score changed; the cells of Section 6 and Appendix E.2.1 follow the re-coded labels.

Five more rows were corrected on 2026-10-08 after a reading of the full texts against the same definition. RASO's Loop-2 label moved from Auto to None: it accepts a rewritten skill only if the rewrite beats the incumbent on a validation split, but the optimized skill never leaves the task, and its use of a public corpus is inbound retrieval, not a promotion. Three first-pass systems leave the composite audit. SkillForge moves from [Auto × Human] to [Auto × None]: its judge flags the failed replies that drive each skill rewrite, the rewrite is committed without an acceptance test, its support engineers review replies and not skill versions, and no skill version leaves the scenario deployment. On the final consistency pass SkillForge moved on to [None × None], since nothing tests its rewrite, the same reading that places SkillRL there. SkillRL's Loop-1 label is None: a per-category success rate below 0.4 triggers new skill proposals, and the skill bank takes each one as an unconditional union, a step that runs unchecked. AutoAgent moves from [Auto × Auto] to [Auto × None]: its analyzer accepts text revisions on an uncalibrated consistency judgment, and each revision stays, unversioned, in one agent's own store. Co-Harness's governance entry now records a partial (c), its versioned diff registry, and (d) within one run, and its signal class records C1 for the held-out gate beside C2 for attribution. The audit therefore passes 11 research systems rather than 14, [Auto × Auto] holds ten research systems rather than twelve, the [Auto × Human] cell of Appendix E.2.1 holds CASCADE, Sierra Agent OS 2.0, and Cognition Devin, and six systems rather than seven promote without lineage. No score changed.

The recount runs over all substrate-level rows and gives the counts of Table 2 of the main article; its last line is the first pass alone (systems / core / contrast) after the two corrections.

```
rows: 269  unique systems: 256  core: 204  contrast: 52
substrate            n   mean  score-5 systems
INTEGRATED           52   3.81  AutoAgent, NanoResearch, SkillRL
S1_Instructions      15   3.07  (none)
S2_Skills            50   3.34  COSPLAY, SAGE
S3_Memory            50   3.06  Governed Collaborative Memory (Cuadros et al.)
S4_Tools             20   3.65  Harvey, Hippocratic AI, MCP
S5_Orchestration     37   3.08  (none)
S6_Governance        45   3.02  AGENTS.md / AAIF, Braintrust, LangSmith Hub, Vellum
First pass alone (after corrections): 130 / 99 / 31; S5 mean 2.58.
```

H Extended related work

The main article makes four design commitments that another survey could make differently: it classifies systems by whether a feedback loop closes over a persisted artifact, rather than by which mechanism is edited; it labels the gate on each of two loops as algorithmic, human, or absent; it classifies the signal that feeds the gate by verifiability; and it audits governance against four named properties. Each commitment meets a separate line of prior work. This section traces what each line established, the setting it assumed, and the condition that changes here, so that a reader can see the differences without the article having to argue for them. Every description below is taken from the cited work's own abstract, method section, or documentation; the evidence record is released with the article.

H.1 Loop closure as the unit of classification

The surveys of self-evolving agents organize the field by what is evolved and when. Gao et al. [129] ask what, when, how, and where to evolve, with a timing axis that separates adaptation inside a test-time episode from adaptation between episodes, and they note that a human in the loop can validate a newly created tool. Fang et al. [130] give a unified frame of system inputs, agent system, environment, and optimizers, and cover single-agent, multi-agent, and domain-specific settings. Zhou et al. [131] taxonomize 108 self-evolving coding methods by evolution target, timing, and code-specific signal, and they come closest to a gate: their admission rule accepts an update when it improves a metric by a margin and passes a check, and they observe that a passing test "checks only selected behavior". Ren et al. [165] describe an agent as a foundation model times an operational scaffold and treat self-improvement as a self-induced update operator organized by target and driving signal. Surveys of agent memory [166–168], context engineering [137, 169], and agent skills [135, 136, 170] do the same for their substrate: each catalogues mechanisms and, in the skills surveys, lifecycle stages.

What this line established is the inventory. The substrates this article scores are the ones these surveys name, and the mechanisms it finds in the corpus are the ones they catalogue. The setting they assume is a mechanism-centered one: a system is described by what it edits. The condition that changes here is the question asked of each system. Whether a loop closes at all, whether it closes under a gate, and whether a second loop promotes the result are properties that the mechanism inventory does not determine, and they are what the two-loop audit of Section 3 measures. Zhou et al.'s admission rule is the nearest precedent for the gate; it is stated for one domain, code, where the check is a test, and this article asks what the rule becomes when the check is a judge or a delayed outcome.

H.2 The gate as algorithmic, human, or absent

Three lines bear on gate type. The first is prompt and program optimization, which a reader will expect to see named. APE treats the instruction as a program and selects among LLM-proposed candidates [171]; ProTeGi forms natural-language "gradients" from minibatches that criticize the current prompt and searches edits with beam search [172]; OPRO has the language model propose new instructions from a prompt that lists previous candidates with their scores and reports gains of up to 8% on GSM8K and up to 50% on Big-Bench Hard over human-designed prompts [138]; EvoPrompt applies evolutionary operators to a prompt population and improves it against a development set, with gains of up to 25% on BBH across 31 datasets [139]; Promptbreeder evolves task prompts and the mutation prompts that modify them, evaluated on arithmetic and

commonsense reasoning and on hate-speech classification [140]. DSPy compiles pipelines of language-model calls into text-transformation graphs and optimizes each pipeline to a metric by creating and collecting demonstrations [141], MIPRO proposes instructions and demonstrations jointly for multi-stage programs under a stochastic mini-batch surrogate [173], TextGrad backpropagates textual feedback through the variables of a compound system [142], and Trace generalizes the pattern to any computation whose execution trace can be shown to an optimizer [174]. This line established that a development-set score is a usable gate for an algorithmic proposer and that the proposer can be the model itself. The setting it assumes is a single optimization run against a fixed task distribution. The condition that changes here is persistence and scope: these optimizers produce a prompt or program as the optimizer's output, not a versioned deployment artifact edited from session signal and promoted under a second gate. That is why OPRO and TextGrad score below the core-corpus threshold and DSPy, which persists only optimizer state, closes no second loop; Appendix E.1 lists all three. SkillOpt is the point where the lineage crosses into the corpus, by adding a persistent skill document, a held-out acceptance gate, and a rejected-edit buffer to the same proposer (Sections 4 and 9).

The second line optimizes the scaffold as a whole rather than one prompt, and it is the nearest precedent for the loop of Section 2. Agent symbolic learning treats prompts, tools, and the pipeline that connects them as symbolic parameters and updates them with textual losses and gradients [144]; STOP has a scaffolding program recursively rewrite the code that improves it [145]; a self-improving coding agent edits its own harness and raises its benchmark score as a result [146]; and workflow search optimizes the orchestration graph itself, by node and edge optimization over agent graphs [20] or by tree search over code-represented workflows [21]. These established that the whole scaffold is an optimizable object and that textual feedback can drive its updates. Each runs within one deployment under a benchmark score, with no promotion to other deployments and no governance of the artifact it produces; the two-loop audit of Section 3 is the measurement of exactly those two absences.

The third line is the classic single-loop agents and their memories: Reflexion's verbal reinforcement into episodic memory [152], the reflection-and-importance memory of generative agents [175], Voyager's growing skill library [70], the experience bank ExpeL extracts from training tasks [176], MemGPT's virtual context management [177], Agent Workflow Memory's induced and reused workflows [178], the Zettelkasten-style organization of A-MEM [179], the persistent test-time memory of Dynamic Cheatsheet [180], the evolving context playbook of agentic context engineering [143], and the in-context iteration of Self-Refine and ReAct [149, 150]. They established that a frozen model can improve across episodes through an artifact it writes, and the cognitive-architecture framing of Sumers et al. [3] gave that artifact the episodic, semantic, and procedural division this article's memory substrate uses. Reflexion, generative agents, Voyager, ExpeL, MemGPT, Agent Workflow Memory, A-MEM, Dynamic Cheatsheet, and the context playbook persist the artifact and close Loop 1; Self-Refine and ReAct do not persist it. None promotes the artifact beyond the agent that wrote it, which is the absence the second loop of this article is defined to measure.

H.3 Verifiability of the feedback signal

No surveyed work classifies systems by the verifiability of the signal that feeds the gate, but three lines come near it. Guo et al. [134] note that software engineering provides mechanical feedback, in the form of builds, unit tests, linters, and execution logs, that most other domains lack; Zhou et al. [131] note that even a passing test is incomplete feedback. The judged end of the spectrum has its own measurements: language-model judges agree with human preference at rates that make them

usable evaluators of open-ended answers [2], a proxy reward optimized past a point lowers the true score [59], and people who grade model outputs revise their own criteria as they grade, so that the evaluator drifts along with the evaluated [60]. The second line is the weight-side effort to extend reinforcement learning from verifiable rewards into domains without verifiers. Rubrics as Rewards generates an instance-specific checklist before training and scores responses with a fixed judge [98]; Writing-Zero trains a pairwise generative judge on writing principles and holds it fixed during policy training [99]; reference-guided evaluation reaches 79.1% agreement with human annotations across the datasets its authors test [100]; an essay scorer trained against a fixed rubric reaches a quadratic weighted kappa of 0.803 against a single human rater's 0.805 [101]. A second group moves the gate with the policy: RLPR scores by the policy's own probability of a fixed reference [102], Rubric-ARM alternates updates of a rubric generator and a judge [103], and EvoRubrics, SCORE, and RubricEM update rubrics adversarially or from a rubric buffer within training [104–106]. The groups part on human calibration: among the methods that move the gate with the policy, only RLPR, whose reward is a probability rather than a rubric, reports agreement with human labels, and no method that co-evolves a rubric or judge does (Appendix G.6 gives the tally). One study calibrates its verifier against language-model golden judges rather than people [160], and a survey of rubric evaluation names the dynamic regime as the trend [164].

This line established the gate constructions for judged signals: checklists, references, pairwise judges, and rubric buffers. The setting it assumes is weight-space optimization with the policy as the learner. The condition that changes here is the learner: when the scaffold is what learns, the rubric is itself a scaffold coordinate, so the fixed-versus-co-evolved distinction becomes the question of whether a gate under edit can be calibrated at all. The article borrows the constructions and the distinction; it does not score the methods.

H.4 Governance: lineage, access control, rollback, and promotion across organizations

Three lines bear on the four-property audit. The first is governance surveys. Ding et al. [132] code 435 works on six dimensions per state item, including provenance and recoverability, and find the literature "concentrates more heavily on accumulating and retrieving state than on governing, recovering, or relinquishing it". Lin et al. [133] organize memory security by lifecycle phase and name five primitives, among them principal-scoped retrieval and rollbackability. Neither carries a promotion dimension, and neither scores whether an inner loop fires from session signal.

The second line is the systems that come closest to the four properties. FederatedSkill exchanges semantic skill diffs between clients that are agent instances under one operator, and a server-side evolution agent applies, adapts, or rejects each patch per client while writing an auditable decision log of source patch, reward, and justification; it reports up to 44.4% higher success than isolated self-evolution across 20 task families [68]. EvolveNet broadcasts a shared harness, optimizes it locally at each client, and merges only when a merge fixes at least as many held-out items as it breaks, rolling back to the previous shared harness after a failed round; its authors describe clients as spanning users and organizations, evaluate on benchmark partitions, and note that a program edit can itself encode client information [67]. The registries supply the remaining properties without an inner loop: EDDOps promotes through automatic evaluation thresholds and human sign-off with supersession lineage and emergency rollback, and does not implement role-based access control [80]; Agentic Configuration Management keeps immutable revisions and baselines that enable rollback [81]; Governed Capability Evolution gates activation and rolls back on monitoring, recovering in 72.4% of drift cases in its proof of concept [82]; and the MCP control planes add role-based access

and quarantine without an evaluation gate [77–79]. Letta's Context Repositories commit every memory edit to git and push shared memory across agents without a gate [39].

The third line is federated learning. FedAvg averages client updates in parameter space [110], DP-FedAvg adds user-level differential privacy [111], the field's open problems are catalogued in Kairouz et al. [112], and PromptFL moves the federated object from weights to a shared prompt that participants learn cooperatively [113]. The setting is numeric updates that are comparable by construction. The condition that changes here is the artifact: scaffold promotion is selection and merge over symbolic objects, and the nearest systems aggregate program edits with a language model rather than by averaging, which is why the article uses the federated analogy only as a limit.

H.5 What is and is not new

The ingredients have precedent, and the article says so. The six substrates are the inventory of the self-evolving-agent, memory, context, skill, and harness surveys, and the memory taxonomy is CoALA's. The gate as a development-set score is the prompt-optimization lineage; the scaffold as an optimizable object is agent symbolic learning, STOP, and workflow search; the persisted artifact that closes Loop 1 is Reflexion, Voyager, and the agent-memory line; the admission rule with a margin and a check is Zhou et al.'s; the gate constructions for judged signals are the rubric-reward line; provenance, recoverability, and rollback as governance dimensions are Ding et al. and Lin et al.; federated aggregation of artifact patches is FederatedSkill and EvolveNet. What is new is the combination and the measurement: a corpus of 256 systems scored on one rubric and released; a classification of each by which of two loops it closes and under which gate type; a classification of the signal that feeds the gate into four verifiability classes, with the demonstration that the class fixes the reachable cells of the design space and the dominant failure; and an audit of every system against four governance properties, with the finding that none meets all four and that the two nearest, FederatedSkill and EvolveNet, each narrow different properties without meeting the set. These are the claims of the abstract and the introduction. Three limits bound them: scoring is single-rater over public evidence of uneven strength (Section 3), the corpus closes at the October 2026 cutoff, and the evidence is the published behavior of the 256 systems rather than new experiments (Section 1).

I Full score table

Every scored (substrate, system) row of the corpus: 130 systems scored by May 2026 and 126 added by the October 2026 sweep. Scores follow the 0–5 rubric of the main text (Table 1); Loop 1 / Loop 2 gate types and signal classes are recorded for the October rows (the gate pairs of the first-pass systems audited under the composite criterion are in the matrix of Appendix E.2.1). Substrate codes: INT integrated, S1 instructions, S2 skills, S3 memory, S4 tools, S5 orchestration, S6 governance. A year of "—" marks a classical system that predates the survey window. The CSV files with per-row justifications are released as ancillary files with this preprint.

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
INT	ADAS / Meta Agent Search	4	2024	arXiv:2408.08435				May 26
INT	Agent S / S2	4.5	2024	arXiv:2410.08164				May 26
INT	Agent Workflow Memory (AWM)	3.5	2024	arXiv:2409.07429				May 26
INT	AlphaEvolve	4	2025	arXiv:2506.13131				May 26
INT	AutoAgent	5	2026	arXiv:2603.09716				May 26
INT	AutoManual	3.5	2024	NeurIPS 2024				May 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
INT	AutoSaddler	4	2026	arXiv:2608.23041	Auto	None	C2	Oct 26
INT	CaptainAgent	2.5	2024	arXiv:2405.19425				May 26
INT	CASCADE	4.5	2025	arXiv:2512.23880				May 26
INT	ChatDev	3	2023	arXiv:2307.07924				May 26
INT	Co-Harness	4	2026	arXiv:2607.22688	Auto	None	C1,C2	Oct 26
INT	CoEvoSkills	3.5	2026	arXiv:2604.01687				May 26
INT	Cognition Devin	3	2024	cognition.ai				May 26
INT	Continual Harness	4	2026	arXiv:2605.09998	None	None	C4	Oct 26
INT	Cradle	4	2024	arXiv:2403.03186				May 26
INT	Darwin Gödel Machine	4	2025	arXiv:2505.22954				May 26
INT	Double Ratchet	4	2026	arXiv:2607.12790	Auto	None	C3 anchored by C1	Oct 26
INT	EvolveNet	4	2026	arXiv:2608.04968	Auto	Auto	C1,C2	Oct 26
INT	EvolveR	4	2025	arXiv:2510.16079				May 26
INT	ExpeL	3.5	2023	arXiv:2308.10144				May 26
INT	Generative Agents	4	2023	arXiv:2304.03442				May 26
INT	GenericAgent	4	2026	arXiv:2604.17091	Auto	None	C2	Oct 26
INT	GITM	3	2023	arXiv:2305.17144				May 26
INT	HarnessBank	4	2026	arXiv:2607.13683	Auto	None	C2	Oct 26
INT	HarnessX	4	2026	arXiv:2606.14249	Auto	None	C2	Oct 26
INT	Hippocratic AI Polaris	3.5	2024	hippocraticai.com				May 26
INT	JARVIS-1	3.5	2023	arXiv:2311.05997				May 26
INT	JIT-Agent	4	2026	arXiv:2608.25593	Auto	None	C2	Oct 26
INT	Life-Harness	4	2026	arXiv:2605.22166	Auto	None	C2	Oct 26
INT	Mendel Gödel Machine (MGM)	4	2026	arXiv:2608.07645	Auto	None	C1	Oct 26
INT	MetaGPT	3.5	2023	arXiv:2308.00352				May 26
INT	Mobile-Agent-E	4.5	2025	arXiv:2501.11733				May 26
INT	ModularRSI	4	2026	arXiv:2609.14857	Auto	None	C2	Oct 26
INT	NanoResearch	5	2026	arXiv:2605.10813				May 26
INT	OpenCore	4	2025	github.com/sibmike/opencore				May 26
INT	OpenHands / CodeAct	3	2024	arXiv:2407.16741				May 26
INT	OS-Copilot / FRIDAY	4	2024	arXiv:2402.07456				May 26
INT	Plan4MC	2.5	2023	GitHub/arXiv				May 26
INT	SAGE	4	2025	arXiv:2512.17102				May 26
INT	Sakana AI Scientist	3	2024	arXiv:2408.06292				May 26
INT	ScholarEvolve	4	2026	arXiv:2609.40169	Auto	None	C2	Oct 26
INT	SIA	4	2026	arXiv:2605.27276	Auto	None	C2	Oct 26
INT	Sierra Agent OS	4	2024	sierra.ai				May 26
INT	Sierra Agent OS 2.0	3	2025	sierra.ai/blog/agent-os-2-0				May 26
INT	SkillForge	4	2026	arXiv:2604.08618				May 26
INT	SkillMAS	4	2026	arXiv:2605.09341	Auto	None	C2	Oct 26
INT	SkillOS	3	2026	SemanticScholar:2653e4c7				May 26
INT	SkillRL	5	2026	arXiv:2602.08234				May 26
INT	TTHE	4	2026	arXiv:2607.08124	Auto	None	C3	Oct 26
INT	VACE	4	2026	arXiv:2609.37105	Auto	None	C2	Oct 26
INT	Voyager	4	2023	arXiv:2305.16291				May 26
INT	Workspace Optimization (DreamTeam)	3	2026	arXiv:2605.09650	Auto	None	C1	Oct 26
S1	AGENTS.md	2	2026	arXiv:2601.20404				May 26
S1	Aider Conventions	2	2024	aider.chat				May 26
S1	Claude Code CLAUDE.md	4	2025	docs.anthropic.com				May 26
S1	Cody / Sourcegraph	2	2024	sourcegraph.com				May 26
S1	Continue.dev Rules	3	2024	docs.continue.dev				May 26
S1	Cursor Rules	3	2024	docs.cursor.com/rules				May 26
S1	GitHub Copilot Custom Instructions	3	2024	docs.github.com				May 26
S1	GRACE	4	2026	arXiv:2607.09175	Auto	None	C2	Oct 26
S1	Lovable Knowledge	2	2024	docs.lovable.dev				May 26
S1	Meta Context Engineering (MCE)	4	2026	arXiv:2601.21557	Auto	None	C2	Oct 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S1	Replit replit.md	4	2025	docs.replit.com				May 26
S1	Spotify conversational-recommendation loop	4	2026	arXiv:2609.30297	Auto	Human	C3	Oct 26
S1	VS Code Copilot Instructions	3	2024	code.visualstudio.com				May 26
S1	Windsurf Cascade Rules	3	2024	docs.windsurf.com				May 26
S1	Zed AI Rules	3	2024	zed.dev/docs				May 26
S2	AgentVerse	1	2023	arXiv:2308.10848				May 26
S2	Anthropic Claude Agent Skills	3	2025	anthropic.com/ engineering				May 26
S2	Anthropic/skills GitHub repo	2	2025	github.com/ anthropics/skills				May 26
S2	ASSAY	4	2026	arXiv:2606.15390	Auto	None	C2	Oct 26
S2	AutoSkill	3	2026	arXiv:2603.01145	Auto	None	C3	Oct 26
S2	Buildrix	3	2026	arXiv:2606.25139	None	Human	C3	Oct 26
S2	ChatGPT Plugins (legacy)	1	2023	openai.com				May 26
S2	Claude Code Skills	3	2025	docs.anthropic.com				May 26
S2	COSPLAY	5	2026	arXiv:2604.20987				May 26
S2	CrewAI Skills	2	2025	docs.crewai.com				May 26
S2	HiSME	4	2026	arXiv:2605.28390	Auto	None	C1,C3	Oct 26
S2	HuggingFace Agent Skills	2	2025	huggingface.co				May 26
S2	LangChain Skills	2	2025	langchain.com				May 26
S2	Library Drift curator	4	2026	arXiv:2605.19576	Auto	None	C1	Oct 26
S2	Meta Context Engineering (MCE)	4	2026	arXiv:2601.21557	Auto	None	C2	Oct 26
S2	MetaGPT (SOP-as-skill)	3	2023	arXiv:2308.00352				May 26
S2	MetaSkill-Evolve	4	2026	arXiv:2607.05297	Auto	None	C2	Oct 26
S2	Microsoft Semantic Kernel Plugins	2	2024	learn.microsoft.com				May 26
S2	MUSE-Autoskill	4	2026	arXiv:2605.27366	Auto	None	C1,C2	Oct 26
S2	NanoResearch Skill Bank	4	2026	arXiv:2605.10813				May 26
S2	OpenAI Custom GPTs / Actions	2	2023	platform.openai.com				May 26
S2	Personalized developer skills (Huang-Du-Lan)	3	2026	arXiv:2608.10319	Auto	None	C3	Oct 26
S2	RASO	4	2026	arXiv:2609.38024	Auto	None	C2	Oct 26
S2	SAGE	5	2025	arXiv:2512.17102				May 26
S2	SCAFFOLD	4	2026	arXiv:2609.05511	Auto	None	C2	Oct 26
S2	SE-GoS	4	2026	arXiv:2609.08228	Auto	None	C2	Oct 26
S2	SkeMex	4	2026	arXiv:2606.09365	Auto	None	C2	Oct 26
S2	Skill Issues (Bauplan)	4	2026	arXiv:2606.01185	Auto	None	C1,C3	Oct 26
S2	SKILL.md smells (Hong-Imani-Ahmed)	1	2026	arXiv:2607.01456	None	None	-	Oct 26
S2	SkillAudit	4	2026	arXiv:2606.14239	Auto	None	C3,C1-partial	Oct 26
S2	SkillBoost	4	2026	arXiv:2607.26643	Auto	None	C1	Oct 26
S2	SkillChain	4	2026	arXiv:2606.12984	Auto->Huma	None	C3,C2	Oct 26
S2	SkillComposer	4	2026	arXiv:2606.06079	Auto	None	C1,C2	Oct 26
S2	SkillEvoReg	4	2026	arXiv:2609.30861	Auto	None	C1,C2	Oct 26
S2	SkillFoundry	3	2026	arXiv:2604.03964	Auto	None	C1	Oct 26
S2	SkillGen	4	2026	arXiv:2605.10999	Auto	None	C2	Oct 26
S2	SkillGLoW	4	2026	arXiv:2609.02217	Auto	None	C2	Oct 26
S2	SkillLift	4	2026	arXiv:2609.15396	Auto	None	C3 inner, C2 outer	Oct 26
S2	SkillOps	4	2026	arXiv:2605.13716	Auto	None	C2	Oct 26
S2	SkillSmith	4	2026	arXiv:2606.01314	Auto	None	C1,C2	Oct 26
S2	SkillSpec (consensus-gated)	4	2026	arXiv:2610.00704	Auto	None	C1,C2	Oct 26
S2	SuperAGI Toolkits	1	2024	github.com/ TransformerOptimus				May 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S2	TRUSS	3	2026	arXiv:2608.17588	Auto	None	C1,C2	Oct 26
S2	VALVE	4	2026	arXiv:2609.32990	Auto	None	C1	Oct 26
S2	VCE-Skill	4	2026	arXiv:2608.16544	Auto	None	C1,C2	Oct 26
S2	Voyager	4	2023	arXiv:2305.16291				May 26
S2	Web2BigTable	3	2026	arXiv:2604.27221	Auto	None	C1,C2	Oct 26
S2	Who Maintains Agent Skills? (empirical)	2	2026	arXiv:2609.05677	None	Human	C3	Oct 26
S2	WikiSkill	4	2026	arXiv:2608.27454	Auto	None	C2	Oct 26
S2	XRepoSkill	4	2026	arXiv:2609.36807	Auto	Auto-parti	C1	Oct 26
S3	A-MEM	3	2025	NeurIPS 2025				May 26
S3	Agent-native memory system evaluation (Zhou et a	2	2026	arXiv:2606.24775	n/a	n/a	n/a	Oct 26
S3	AIM privacy-aware interoperable memory (Johnson	3	2026	arXiv:2609.12320	Auto	None	C3	Oct 26
S3	Beyond the Harness (enterprise Text-to-SQL cards	3	2026	arXiv:2608.22830	Auto	None	C2	Oct 26
S3	ChatGPT Memory	3	2024	openai.com				May 26
S3	Claude Memory	3	2025	anthropic.com				May 26
S3	Cognee	3	2024	github.com/ topoteretes/cognee				May 26
S3	Compiled Memory (Rhodes & Kang)	4	2026	arXiv:2603.15666	Auto-gated	None	C1	Oct 26
S3	Control-plane placement / ForgetEval (Yang)	2	2026	arXiv:2606.15903	n/a	n/a	n/a	Oct 26
S3	FRESH failure-aware heterogeneous graph memory (	4	2026	arXiv:2609.28003	Auto	None	C2	Oct 26
S3	GateMem multi-principal memory-governance benchm	2	2026	arXiv:2606.18829	n/a	n/a	n/a	Oct 26
S3	Gemini Code Assist Memory	4	2025	cloud.google.com				May 26
S3	Generative Agents	3	2023	arXiv:2304.03442				May 26
S3	Governed Collaborative Memory (Cuadros et al.)	5	2026	arXiv:2605.04264				May 26
S3	HippoRAG	2	2024	arXiv				May 26
S3	Honest Lying mitigation / programmatic failure s	4	2026	arXiv:2605.29463	Auto	None	C1/C2	Oct 26
S3	KV-cache rollback consistency (Zhang & Yang)	2	2026	arXiv:2608.15939	n/a	n/a	n/a	Oct 26
S3	LangChain Primitives (legacy)	1	2023	langchain.com				May 26
S3	LangGraph Long-Term Memory	3	2024	langchain.com				May 26
S3	LangGraph Store / LangMem (2026 refresh)	3	2026	docs.langchain.com/ oss/python	None	None	n/a	Oct 26
S3	Learning on the Job	4	2026	arXiv:2607.22157	Auto+ Human	None	C2,C3	Oct 26
S3	Letta MemFS + Context Repositories + Dreaming (2	4	2026	docs.letta.com/ letta-code	Auto	None	C3	Oct 26
S3	Mem0	3	2025	arXiv				May 26
S3	Mem0 Platform (2026 refresh: history, feedback,	3	2026	docs.mem0.ai/ api-reference	Auto	None	C3	Oct 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S3	MemAudit post-hoc poisoned-memory audit (Tan et	3	2026	arXiv:2605.23723	Auto	None	C2	Oct 26
S3	MemClaw / Governed Shared Memory (Margalit et al	3	2026	arXiv:2606.24535	Auto	None	C3	Oct 26
S3	MemGPT / Letta	3	2023	arXiv:2310.08560				May 26
S3	MemGuard persisted verifier signals (Wang et al.	4	2026	arXiv:2608.21867	Auto-gated	None	C3	Oct 26
S3	MemoryBank	3	2023	Zhong et al.				May 26
S3	Mi-Memory lifecycle framework (Liu et al.; Xiaom	3	2026	arXiv:2607.18975	Auto	None	C3	Oct 26
S3	MIRIX	3	2025	arXiv				May 26
S3	MPBench memory-poisoning study (Dash et al.)	2	2026	arXiv:2606.04329	n/a	n/a	n/a	Oct 26
S3	NanoResearch Memory Module	3	2026	arXiv:2605.10813				May 26
S3	PPMF provenance-preserving memory firewall (Xu e	3	2026	arXiv:2607.29167	Auto	None	C1	Oct 26
S3	REALM retrieval-driven memory reconsolidation (S	4	2026	arXiv:2609.16053	Auto	None	C2/C3	Oct 26
S3	Recuris recursive experiential-working memory (Y	4	2026	arXiv:2608.24876	Auto-gated	None	C2	Oct 26
S3	Reflexion	4	2023	arXiv:2303.11366				May 26
S3	Revocation guard / Revoked but Still Authoritati	3	2026	arXiv:2609.08258	None	None	C1	Oct 26
S3	Safe to Resume? checkpoint-rollback continuity (	2	2026	arXiv:2608.29381	n/a	n/a	n/a	Oct 26
S3	ShareMem	4	2026	arXiv:2609.32511	Auto	Auto-parti	C3	Oct 26
S3	SIx Harness / Memory as Infrastructure (Helwig)	3	2026	arXiv:2609.05510	Auto	None	C2	Oct 26
S3	SOAR Cognitive Architecture	4	—	Newell/Laird				May 26
S3	SSGM stability-and-safety governed memory (Lam e	2	2026	arXiv:2603.11768	None	None	n/a	Oct 26
S3	Stale constraints in inherited memory (Nakayashi	2	2026	arXiv:2608.25553	n/a	n/a	n/a	Oct 26
S3	TMA-NM origin-bound memory authority (Louck)	3	2026	arXiv:2606.24322	Auto	None	C1	Oct 26
S3	TOKI bitemporal operator algebra (Wang)	3	2026	arXiv:2606.06240	Auto or Hu	None	C3	Oct 26
S3	Trust-scored memory sanitization (Devarangadi Su	2	2026	arXiv:2601.05504	Auto gate	None	n/a	Oct 26
S3	Typed provenance + assertion guardrails / Stored	3	2026	arXiv:2609.02127	Auto	None	C1	Oct 26
S3	Utility Under Attack (Karunanidhi)	2	2026	arXiv:2608.21230	n/a	n/a	n/a	Oct 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S3	VerificAgent human-verified frozen memory (Nguye	3	2025	arXiv:2506.02539	Human	None	C3	Oct 26
S3	VerMem verifiable memory management (Sun et al.)	4	2026	arXiv:2608.03137	Auto	None	C2/C3	Oct 26
S3	Voyager Skill Library	4	2023	arXiv:2305.16291				May 26
S3	Zep / Graphiti	3	2025	arXiv				May 26
S3	Zep / Graphiti bitemporal facts (2026 refresh)	3	2026	help.getzep.com	Auto	None	C3	Oct 26
S4	Anthropic Tool Use API	4	2024	docs.anthropic.com				May 26
S4	Browser Tools (Playwright/Computer Use)	3	2024	anthropic.com				May 26
S4	Claude Code Tool API	4	2024	docs.anthropic.com				May 26
S4	Code Interpreter	3	2023	openai.com				May 26
S4	Cursor MCP Integration	4	2024	docs.cursor.com				May 26
S4	Devin / SWE-agent Tool Harness	4	2024	cognition.ai				May 26
S4	Gemini Function Calling	3	2024	ai.google.dev				May 26
S4	GPT Actions	3	2023	platform.openai.com				May 26
S4	Harvey (Legal Vertical)	5	2024	harvey.ai				May 26
S4	Hippocratic AI	5	2024	hippocraticai.com				May 26
S4	LangChain Tools Abstraction	3	2023	langchain.com				May 26
S4	LangGraph Tool Nodes	4	2024	langchain.com				May 26
S4	LlamaIndex Tool Ecosystem	4	2023	llamaindex.ai				May 26
S4	Model Context Protocol (MCP)	5	2024	modelcontextprotocol.io				May 26
S4	OpenAI Assistants API	3	2023	platform.openai.com				May 26
S4	OpenAI Function Calling	3	2023	platform.openai.com				May 26
S4	Sierra	4	2024	sierra.ai				May 26
S4	SkillSmith	4	2026	arXiv:2606.01314	Auto	None	C1	Oct 26
S4	Toolformer	2	2023	arXiv:2302.04761				May 26
S4	ToolLLM / ToolBench	3	2023	arXiv:2307.16789				May 26
S5	AgentConductor	0	2026	arXiv:2602.17100	None	None	C2	Oct 26
S5	AgentGrad	4	2026	arXiv:2609.08572	Auto	None	C2	Oct 26
S5	AgentVerse	2	2023	arXiv:2308.10848				May 26
S5	AutoGen / MagenticOne	3	2024	arXiv:2308.08155				May 26
S5	AWS Bedrock Multi-Agent	2	2024	aws.amazon.com				May 26
S5	Bilevel Autoresearch	4	2026	arXiv:2603.23420	Auto	None	C1	Oct 26
S5	CAMEL	1	2023	arXiv				May 26
S5	ChatDev	2	2023	arXiv:2307.07924				May 26
S5	ContextEvo	4	2026	arXiv:2609.34649	Auto	None	C2	Oct 26
S5	Control System for Frozen Agents (Paul)	4	2026	arXiv:2607.25415	Auto	None	C2	Oct 26
S5	CrewAI	3	2025	docs.crewai.com				May 26
S5	Evo2Team	4	2026	arXiv:2609.34135	Auto	Auto	C2	Oct 26
S5	Growing Harness	4	2026	arXiv:2609.26760	Auto	None	C2	Oct 26
S5	HARBOR	4	2026	arXiv:2604.20938	Auto	None	C2	Oct 26
S5	HarnessForge	4	2026	arXiv:2606.01779	Auto	None	C2	Oct 26
S5	Inspect AI	2	2024	ukgovernmentbeis.github.io				May 26
S5	Institutional MAS architectures (Fei et al.)	1	2026	arXiv:2604.27691	None	None	C2	Oct 26
S5	LangGraph	3	2024	langchain.com				May 26
S5	Letta	3	2025	letta.com				May 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S5	LlamaIndex Workflows	2	2024	llamaindex.ai				May 26
S5	Memento-Skills	4	2026	arXiv				May 26
S5	MetaGPT	2	2023	arXiv:2308.00352				May 26
S5	Microsoft Semantic Kernel	2	2024	learn.microsoft.com				May 26
S5	Multi-Agent Evolve (MAE)	4	2025	arXiv				May 26
S5	OpenAI AgentKit	4	2025	openai.com				May 26
S5	OpenAI Swarm	1	2024	openai.com				May 26
S5	OpenHands / CodeAct	3	2024	arXiv:2407.16741				May 26
S5	OptiMAS	4	2026	arXiv:2608.21918	None	None	C2	Oct 26
S5	PluginRSI	4	2026	arXiv:2609.32423	Auto	None	C2	Oct 26
S5	Pydantic AI	2	2024	pydantic.dev				May 26
S5	Recursive Harness Self-Improvement (RHI)	4	2026	arXiv:2607.15524	Auto	None	C3	Oct 26
S5	Skill-MAS (meta-skill)	4	2026	arXiv:2606.18837	Auto	None	C2	Oct 26
S5	TACO	4	2026	arXiv:2604.19572	Auto	None	C2	Oct 26
S5	Task-CoEvolve	4	2026	arXiv:2608.20169	Auto	None	C2	Oct 26
S5	TOCOMAS	4	2026	arXiv:2609.37953	Auto	None	C2	Oct 26
S5	Vertex AI Agent Builder	4	2024	cloud.google.com				May 26
S5	Workflow-Designing Agents (WDA)	4	2026	arXiv:2609.32458	Auto	None	C2	Oct 26
S6	Agenta	3	2024	agenta.ai				May 26
S6	Agentic Configuration Management ACM (Quessada-V	3	2026	arXiv:2608.11166	None	Human	n/a	Oct 26
S6	AGENTS.md / AAIF	5	2025	agents.md				May 26
S6	Anthropic skills repo	4	2025	github.com/ anthropics/skills				May 26
S6	Azure API Center MCP registry	3	2026	learn.microsoft.com	None	Human	-	Oct 26
S6	Braintrust	5	2024	braintrust.dev				May 26
S6	Cleric.ai	2	2025	cleric.ai				May 26
S6	Competence Not Accuracy	2	2026	arXiv:2608.18719	None	None	-	Oct 26
S6	Constitutional AI	2	2022	Anthropic				May 26
S6	Cursor Rules	3	2024	cursor.com				May 26
S6	Decagon	2	2026	decagon.ai				May 26
S6	DSPy compile	3	2024	dspy.ai				May 26
S6	EvolveNet	4	2026	arXiv:2608.04968	Auto	Auto	C2	Oct 26
S6	EvoUndo	4	2026	arXiv:2608.28363	Auto	None	C1	Oct 26
S6	From Prompts to Contracts	2	2026	arXiv:2607.08028	None	None	C1	Oct 26
S6	Future AGI prompt lifecycle	4	2026	futureagi.com/blog	Auto	Auto	C3	Oct 26
S6	GitHub Copilot enterprise	4	2026	github.com				May 26
S6	Governance conversion (Davis et al. case study)	3	2026	arXiv:2607.01087	Human	None	C2	Oct 26
S6	Governed Capability Evolution (Qin et al.)	3	2026	arXiv:2604.08059	None	Auto	C2	Oct 26
S6	HackProbe	3	2026	arXiv:2609.04665	Auto	None	C1	Oct 26
S6	Inspect AI	3	2024	ukgovernmentbeis. github.io				May 26
S6	LangFuse	3	2024	langfuse.com				May 26
S6	LangSmith Hub	5	2024	smith.langchain.com				May 26
S6	MCP Gateway & Registry (agentic-community)	3	2026	github.com/ agentic-community	None	Human	C1	Oct 26
S6	NanoResearch	2	2026	arXiv:2605.10813				May 26
S6	Official MCP Registry	2	2026	modelcontextprotocol. io	None	None	-	Oct 26
S6	OpenAI AgentKit Evals	3	2025	openai.com				May 26

Sub.	System	Score	Year	Source	Loop 1	Loop 2	Class	Added
S6	OPRO	2	2023	arXiv				May 26
S6	Organizational Control Layer (Shi et al.)	2	2026	arXiv:2606.04306	None	None	n/a	Oct 26
S6	Registry-Governed Agent Lifecycle / EDDOps on AW	3	2026	arXiv:2607.00345	None	Auto+ Human	C3	Oct 26
S6	Reward-Free Pairwise Validator	4	2026	arXiv:2607.14408	Auto	None	C3	Oct 26
S6	Runtime Skill Audit	2	2026	arXiv:2606.11671	None	None	-	Oct 26
S6	Runtime-Independent Persistent Agents (Enoch)	3	2026	arXiv:2609.00546	None	None	C1	Oct 26
S6	SafeEvolve	4	2026	arXiv:2608.12851	Auto	None	C3,C2	Oct 26
S6	SAGE statistical acceptance gate (≠ SAGE 2512.17	4	2026	arXiv:2609.36043	Auto	None	C2	Oct 26
S6	Self-Refine	1	2023	Madaan et al.				May 26
S6	Sierra Ghostwriter	2	2026	sierra.ai				May 26
S6	SkillCoach	3	2026	arXiv:2607.01874	Auto	None	C3,C1	Oct 26
S6	SkillProbe	2	2026	arXiv:2603.21019	None	None	-	Oct 26
S6	SkillSpec (intent-masked correctness)	2	2026	arXiv:2609.06052	None	None	-	Oct 26
S6	TextGrad	2	2024	arXiv				May 26
S6	TrueFoundry MCP Gateway	3	2026	truefoundry.com/docs	None	Human	-	Oct 26
S6	Vellum	5	2024	vellum.ai				May 26
S6	Verifier-as-Gatekeeper (VaG)	4	2026	arXiv:2608.05810	Auto	None	C2	Oct 26
S6	W&B Weave (Prompts)	3	2024	weave-docs.wandb.ai				May 26