
The Architecture of Errors: From Universal Impossibility to Patch-Local LLM Reliability

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Universal LLM reliability is not a finite-library problem: across all possible tasks,
2 tools, schemas, knowledge sources, and evaluator expectations, new intervention-
3 distinguishable failure modes can appear without bound, so no finite interven-
4 tion dictionary can guarantee bounded residual error universally. But deployed
5 systems do not operate over the whole universe; they operate inside bounded
6 patches (legal review, medical RAG, code repair, customer-support agents, con-
7 tract extraction) with recurring tasks, schemas, tools, and evaluator expectations.
8 Within such patches, empirical evidence suggests failures are sparse, repetitive,
9 and concentrated in a small recurring catalogue, so reliability becomes a local
10 catalogue-discovery and intervention-coverage problem rather than an exponential
11 token-length problem. We formalize this transition with a negative result showing
12 that no universal finite intervention dictionary exists, an inverse-discovery corollary
13 showing that logarithmic mode discovery implies exponential data cost for linearly
14 expanding the discovered tail, and a positive patch-local result showing that under
15 logarithmic or capped mode discovery the intervention budget grows polylogarith-
16 mically in sequence length and becomes domain-constant once the patch catalogue
17 saturates. The framework relocates rather than dissolves long-context difficulty:
18 where the number of hard decisions itself grows with task length, reliability remains
19 hard; the contribution is to identify the on-axis intervention rather than to make
20 those regimes easy.

21 1 Introduction

22 The standard worry about long-context autoregressive generation is exponential. If every token
23 has independent error probability e , then after n tokens the chance of a fully correct output is
24 $(1 - e)^n$, which collapses to zero for any nontrivial e and large enough n [LeCun, 2023, Dziri et al.,
25 2023]. Earlier work in this series [Anonymous, 2025] argued that this worry is misplaced because
26 errors are not distributed uniformly across tokens: only 5–10% of tokens are “key” (genuinely
27 dependent on long-range context [Fang et al., 2025]), while the remaining majority become nearly
28 deterministic once enough context accumulates. Replacing $(1 - e)^n$ with the two-rate model
29 $P(\text{correct}) = (1 - e_{\text{key}})^k (1 - e_{\text{non}})^{n-k}$ with $k \ll n$ scaling sublinearly in n recovers the observed
30 long-context coherence and converts the reliability question from “how does n grow?” to “how does
31 k grow?”.

32 **From sparse tokens to recurring patterns.** This paper takes the next step. Sparsity tells us *where*
33 errors live; the natural follow-up is *what* they are. Recent failure-mode atlases supply a striking
34 empirical answer: errors are not only sparse, they are also *repetitive*. ErrorAtlas [Ashury-Tahan et al.,

2026], covering 83 models across 35 datasets, organises observed failures into 17 named categories sorted by prevalence, with a long-tailed distribution heavily concentrated in the head. In code, two error types (`AssertionError` and `NameError`) cover 86.35% of failures on HumanEval across 14 LLMs, replicated across 23 models on HumanEval Pro and MBPP Pro [Wen et al., 2024]. In math, MWPEs-300K [Sun et al., 2025] categorises 304,865 errors from 15 LLMs across four math word-problem datasets and reports that dataset characteristics shape error patterns systematically. The same picture appears in multi-hop QA [Zhang et al., 2026], agentic tool use [Cemri et al., 2025], and retrieval-augmented generation [Wood and Forbes, 2024]. These observations suggest a third layer in the reliability architecture, complementing key-token sparsity (Layer 1) and within-key stratified-manifold structure (Layer 2 in Anonymous 2025). Within the small set of key tokens, only a fraction β produce *hard* failures, and inside bounded deployment patches those hard failures appear to cluster into a finite or effectively capped catalogue of recurring modes whose size grows much more slowly than the number of observed events.

Contributions. The central contribution is a shift in the reliability object: universal LLM reliability is not a finite-library problem, but patch-local reliability can be treated as catalogue discovery and intervention coverage. We formalise this transition with two propositions and one corollary. Proposition 1 gives the negative result: if an unbounded domain keeps producing intervention-distinguishable failures, no finite intervention dictionary can guarantee bounded residual error universally. Corollary 1 gives the inverse economic implication of logarithmic discovery: linearly expanding the discovered-mode catalogue requires exponentially more observed hard-failure events, so open-domain tail discovery has rapidly diminishing returns. Proposition 2 gives the corresponding positive result inside a fixed deployment patch: under logarithmic or capped mode discovery, the intervention budget required to bound per-hard-token residual error satisfies $m \geq \lceil |C_{\text{eff}}|^{1-\varepsilon/e_{\text{hard}}} \rceil$, with a doubly-logarithmic sequence-length rate as the optimistic pre-cap special case and a domain-constant budget once the patch catalogue saturates. The sequence-level analogue is strictly tighter and approaches full-catalogue coverage as k grows; we surface that asymmetry rather than bury it.

Around this transition the paper does four supporting things. A three-layer framework, comprising sparsity (α), hard-token stratification (β), and patch-local mode catalogue ($|C_D|$), separates *where* errors occur, *what* recurring forms they take, and *which* capability interventions address them (§3). Logarithmic mode discovery is stated as an empirical postulate, not a theorem, calibrated against ErrorAtlas, HumanEval, and MWPEs with $\sigma \in [0.87, 1.85]$ across anchors and $\sigma \approx 1.85$ carried as a conservative planning value. Section 4 synthesises evidence for failure clustering, cluster-selective interventions, and sublinear length scaling, drawing on ~ 60 prior published results including the six-axis capability-elimination harvest of 28 quantitatively-anchored citations stratified into Patterns A/B/C (Appendix B). And the most-cited steep-decay counter-evidence (Appendix C) is re-audited and shown to decay primarily over task-structure variables, including compositional graph size, fact count, log-time horizon, capacity threshold, and evidence scope, rather than raw token length.

2 Related Work

Failure-mode taxonomies. ErrorAtlas [Ashury-Tahan et al., 2026], MWPEs-300K [Sun et al., 2025], the HumanEval categorisation of Wen et al. [2024], and the RFMDataset [Guo et al., 2025] together suggest that, at any given corpus size, the number of distinct named failure modes is small (typically 8–20) and the cumulative coverage of the top modes is high. Domain-specific taxonomies for multi-hop QA [Zhang et al., 2026], multi-agent systems [Cemri et al., 2025], and tool-augmented agents [Yao et al., 2024] report the same pattern.

Targeted interventions. Each named cluster has been the subject of focused intervention research. Python execution closes most of the arithmetic-error cluster [Gao et al., 2023a, Chen et al., 2023, Gou et al., 2024b]; constrained decoding eliminates format violations [Suresh et al., 2025, Wang et al., 2025b, Dong et al., 2025]; execution feedback closes most code-logic errors [Shinn et al., 2023, Huang et al., 2023]; process reward models reduce step-level reasoning errors [Wang et al., 2024b, Lightman et al., 2024]; RAG strongly reduces hallucinations [Wood and Forbes, 2024]; preference optimisation closes over-refusal [Karaman et al., 2024]; structured uncertainty fixes most tool-call failures [Suri et al., 2025]. Two structural observations recur: each intervention is *cluster-selective* (residuals belong to a different named cluster, not to the targeted one), and *additivity is approximate but not perfect* [Patel et al., 2026, Le, 2026].

89 **Length-decay benchmarks.** A parallel literature measures the *shape* of the reliability decay curve.
 90 Mild-decay results (Loong [Wang et al., 2024a], GSM- ∞ [Zhou et al., 2025], RULER [Hsieh et al.,
 91 2024], anchor-based LLMs [Pang et al., 2024]) report log-linear, sigmoidal, or threshold-like decay
 92 inconsistent with smooth $(1 - \varepsilon)^n$. Steep-decay results [Dziri et al., 2023, Kuratov et al., 2024,
 93 Kwa et al., 2025, Wan et al., 2026] are sometimes cited as exponential-compounding evidence. The
 94 apparent tension is reduced on careful reading: every steep-decay paper decays over a variable distinct
 95 from raw token length, which Appendix C documents.

96 **Gap.** What is missing is a quantitative bridge between the small, repeating taxonomy catalogue and
 97 a polylog intervention budget. We provide that bridge in §3.

98 3 Theoretical Framework

99 **Four levels of $|C|$.** Before any proposition, a definitional setup. LLM error analysis routinely
 100 conflates four objects that the rest of this paper needs to keep apart:

- 101 • **L1: failure events.** Concrete LLM errors observed in a benchmark; the raw data behind any
 102 taxonomy.
- 103 • **L2: empirical taxonomy categories.** Researcher-chosen labels grouping L1 events. ErrorAt-
 104 las’s 17 categories, MWPE’s top-4 classes, and HumanEval’s AssertionError/NameError
 105 partition all live here. L2 is observable but depends on the taxonomer’s resolution choices.
- 106 • **L3: latent failure modes.** The unobserved underlying clusters that L2 approximates. L3 is the
 107 quantity Postulate 1 is morally about, but it cannot be measured directly; we treat L2 as a noisy
 108 proxy for L3.
- 109 • **L4: capability axes / interventions.** The engineering unit: a Python interpreter, a constrained
 110 decoder, a retrieval-augmented generator. L4 is coarser than L2 in practice; one capability axis
 111 typically targets several L2 categories at once (§4, Claim B).

112 Throughout this section, $|C|$ refers to the **L2** count: that is what published taxonomies report.
 113 Postulate 1 is therefore *engineering*-relevant because L4 is coarser than L2 (so a small intervention
 114 library can sweep across many categories at once), and *epistemically* conditional because L2 is a noisy
 115 proxy for the actual L3 mode catalogue whose faithfulness has not been independently measured.
 116 Tagging the formal claims at the right level keeps the framework honest where casual error-talk slips.

117 **Roadmap.** The framework separates three questions that are often conflated. First, *where* do errors
 118 occur? Only a sparse subset of key decisions. Second, *what* kinds of errors recur? A local catalogue
 119 of failure modes. Third, *what fixes them?* A smaller library of capability interventions. The two
 120 propositions of §3.4 formalise the transition from universal impossibility to patch-local tractability:
 121 there is no finite intervention dictionary that covers all possible deployments, but inside a fixed
 122 deployment patch the required library is small and slowly growing.

123 3.1 β -stratification of key tokens

124 The two-rate model of Anonymous [2025] distinguishes k key tokens (error rate e_{key}) from $n - k$
 125 non-key tokens ($e_{\text{non}} \ll e_{\text{key}}$). Empirical atlases suggest that even within the key-token class errors
 126 are not uniformly distributed: most key tokens are “decisions” for which the model has stable
 127 representations; only a fraction concentrate the actual failures.

128 **Definition 1** (Hard fraction). Partition the k key tokens of a sequence into easy and hard subsets,

$$k_{\text{hard}} = \beta k, \quad k_{\text{easy}} = (1 - \beta)k, \quad \beta \in (0, 1),$$

129 where hard key tokens have an elevated error rate e_{hard} corresponding to manifold-transition decisions
 130 in the sense of Anonymous [2025, §3.2], and easy key tokens have $e_{\text{easy}} \approx e_{\text{non}}$.

131 Under Definition 1, the composed sequence-level reliability becomes

$$P(\text{correct}) = (1 - e_{\text{hard}})^{\beta k} (1 - e_{\text{easy}})^{(1-\beta)k} (1 - e_{\text{non}})^{n-k}. \quad (1)$$

132 We do not assume iid token errors: the three rates summarise *conditional per-decision failure*
 133 *hazards* after conditioning on prior decisions being correct, and grouping by stratum yields the

134 multiplicative survival expression. Because $e_{\text{easy}} \approx e_{\text{non}}$ once context accumulates, the failure rate
 135 is dominated by the βk hard tokens, and we treat e_{hard} as the load-bearing quantity. The parameter
 136 β is *latent*: failure-mode atlases report $\Pr(\text{category} \mid \text{error occurred})$, which does not identify
 137 $\Pr(\text{hard} \mid \text{key-token decision})$. We carry β symbolically.

138 3.2 Empirical postulate: logarithmic mode discovery

139 Stratification names a target, the βk hard-token decisions, but does not yet tell us how many distinct
 140 ways those hard decisions can fail. The empirical answer from Section 1 is that failures repeat: a
 141 small number of recurring patterns covers most of the mass. The next question is how the catalogue’s
 142 size grows with the number of observed failures, because that is the quantity an intervention library
 143 has to keep up with.

144 Two candidates exist in the type–token literature: Heaps’ law (power-law, $|C| \approx K \cdot k_{\text{hard}}^b$ with
 145 $b \in [0.4, 0.6]$ [Manning et al., 2008]), and logarithmic discovery. **Zipfian rank-frequency does not**
 146 **imply logarithmic growth.** Heaps’ law is the correct type–token consequence of Zipf-distributed
 147 events, and it is power-law, not logarithmic. We therefore state logarithmic mode discovery as an
 148 empirical postulate, defensible by direct measurement of error taxonomies, not as a theorem.

149 Two sample-size variables matter, and they are easy to conflate. Let $h(n) = \beta k(n)$ denote the
 150 number of hard-token decisions in a single sequence of length n ; let T denote the number of observed
 151 hard-failure events in a corpus used to discover the domain catalogue. These are not the same object:
 152 $h(n)$ controls per-sequence exposure, T controls empirical discovery. We write C_D for the full
 153 reachable catalogue inside domain patch D , $C_{\text{seen},D}(T)$ for the subset discovered after T sampled
 154 failures, and $C_{\text{active},D}(n)$ for the subset a single sequence of length n can activate.

155 **Postulate 1** (Patch-indexed catalogue discovery). Within a fixed application domain D , the number
 156 of named recurring failure modes discovered after T sampled hard-failure events is bounded by

$$|C_{\text{seen},D}(T)| \leq \min(A_D + \sigma_D \ln T, |C_D|), \quad \sigma_D > 0, A_D \geq 0.$$

157 The cap $|C_D|$ is the domain-imposed ceiling. The postulate concerns catalogue *discovery*, not the
 158 number of hard decisions inside a single sequence.

159 **Assumption 2** (Per-sequence active-mode exposure). For a single sequence of length n in domain D ,

$$|C_{\text{active},D}(n)| \leq \min(A'_D + \sigma'_D \ln h(n), |C_D|).$$

160 Primed constants are distinct from those of Postulate 1: corpus discovery and per-sequence activation
 161 need not share the same rate.

162 The two bounds answer different questions, and the rest of the paper picks the right one at each point.
 163 Proposition 2’s sequence-length scaling rate is governed by $C_{\text{active},D}(n)$ via Assumption 2: how much
 164 of the catalogue can a single sequence touch? Full domain-library budgeting, the practical engineering
 165 target for a deployed system, uses C_D : how large does the library need to be to cover the domain’s
 166 reachable failure modes? The cap $|C_D|$ enters both: once enough has been sampled (or activated) to
 167 saturate the domain ceiling, further growth stops and the budget becomes a domain-constant function
 168 of $|C_D|$ alone.

169 **Empirical calibration.** Existing LLM error taxonomies provide endpoint anchors for small named
 170 catalogues at observed corpus scales. ErrorAtlas, HumanEval-style code taxonomies [Wen et al.,
 171 2024], and MWPEs [Sun et al., 2025] place $|C|$ roughly in the 8–20 range across corpora ranging
 172 from approximately 10^4 to 3×10^5 observed failures, yielding $\sigma \in [0.87, 1.85]$ under the simple
 173 $A = 0$ logarithmic calibration. We use $\sigma \approx 1.85$ as a conservative planning value, not as a fitted law.
 174 These endpoint counts are not discovery curves: the missing test is a subsample-vs-discovered-modes
 175 measurement within a fixed deployment patch. Full calibration details are in Appendix E.

176 **Capabilities are coarser than error categories.** The category count $|C|$ is conservative for engi-
 177 neering because deployed interventions operate at the capability level, not the label level. A Python
 178 interpreter can remove execution-error components across arithmetic, unit conversion, counting, list
 179 manipulation, and date arithmetic; a constrained decoder can eliminate several structural output
 180 failures at once. The formal coverage problem is therefore closer to weighted set cover over failure
 181 mass than one-intervention-per-category counting. The ranked-category model used in Proposition 2
 182 is a tractable proxy; the full capability harvest appears in Appendix B.

183 **Domain patches cap the engineering problem.** A fixed deployment patch restricts the task family,
 184 schemas, tools, workflows, and evaluator expectations. Prior work on localized representations [Park
 185 et al., 2024, Li and Sarwate, 2025], cross-domain performance variation [Wang et al., 2024c], and
 186 long-tail knowledge [Mallen et al., 2023, Kandpal et al., 2023] supports the weaker claim that model
 187 behaviour is strongly domain-dependent; it does not directly measure $|C_D|$. We therefore treat the
 188 patch’s reachable catalogue as finite or effectively capped as a modelling assumption motivated by
 189 domain heterogeneity, not as a theorem derived from boundedness. The full evidence breakdown is
 190 in Appendix D.

191 **Corollary 1** (Inverse Discovery Cost). *The logarithmic discovery postulate inverts to a sample-budget*
 192 *cost on distinct-mode discovery. If $q(T) \approx A_D + \sigma_D \ln T$ with $\sigma_D > 0$, then discovering q distinct*
 193 *modes requires*

$$T(q) \approx \exp\left(\frac{q - A_D}{\sigma_D}\right) \quad (2)$$

194 *observed hard-failure events, and each additional Δq modes raises the sample budget by a factor of*
 195 *$\exp(\Delta q / \sigma_D)$.*

196 **Engineering meaning.** Linearly expanding the discovered-mode catalogue therefore requires
 197 multiplicatively more data. At the conservative calibration $\sigma_D \approx 1.85$, five extra discovered modes
 198 need roughly $15\times$ more observed hard failures; ten extra modes need roughly $220\times$ more. The
 199 claim is about *newly distinguishable* tail modes, not ordinary failures, which may remain common.
 200 Combined with the head-heavy coverage model of §3.3, this gives the economic interpretation:
 201 generic open-domain training spends increasing data on low-mass tail discovery, while patch-local
 202 reliability directly targets the high-mass local head. Full derivation, the Heaps polynomial variant,
 203 and a mode-mediated capability-gain sub-corollary are in Appendix A.6.

204 3.3 Coverage by a targeted intervention library

205 With the catalogue defined, the next question is how much residual error a library of m targeted
 206 interventions actually closes off. Given a catalogue of $|C|$ recurring modes, let $p_1 \geq p_2 \geq \dots \geq p_{|C|}$
 207 denote the ranked empirical hard-error mass of those modes with $\sum_i p_i = 1$, so a library covering
 208 the top m modes removes cumulative error mass $\sum_{i=1}^m p_i$. We adopt the log-coverage form

$$F(m; |C|) = \min\left(1, \frac{\ln m}{\ln |C|}\right) \quad (3)$$

209 as a *postulated* parametric approximation to that cumulative top- m mass, not as a raw count of
 210 categories covered. Reading the formula: the first few interventions cover disproportionate mass
 211 (the head of the distribution is dense), then incremental coverage falls off as the library saturates the
 212 catalogue. For our ErrorAtlas anchor ($|C| = 17, m = 5$), $F(5; 17) \approx 56.8\%$; doubling the library to
 213 $m = 10$ pushes it to $F(10; 17) \approx 81.3\%$.

214 This is an empirical best-fit choice, not derived from Zipf. ErrorAtlas sorts its 17 categories by
 215 prevalence but does not publish a single cross-domain cumulative top- m curve at this writing;
 216 alternative cumulatives (Zipf–Mandelbrot, truncated power law, saturating exponential) match the
 217 same qualitative shape with comparable freedom. Appendix A.5 shows the polylog conclusion of
 218 §3.4 survives across the family.

219 Two caveats matter for what comes next. (a) F tracks cumulative error mass across ranked L2
 220 categories, not L4 capability axes deployed; in real deployments, where one capability often sweeps
 221 several L2 categories, residual error at a given m is generally lower than (3) predicts. (b) The form is
 222 intended for $m \geq 2, |C| \geq 2$; we adopt $F(0; |C|) := 0$ and $F(1; |C|) := p_1$ (the mass of the single
 223 most prevalent mode) since the log expression is not anchored at $m = 1$. After deploying the top- m
 224 library, the residual per-hard-token error rate is

$$e_{\text{res}}(m) = (1 - F(m; |C|)) e_{\text{hard}}. \quad (4)$$

225 3.4 From universal impossibility to patch-local reliability

226 The finite-catalogue claim is patch-local. It does not say that all possible LLM failures can be covered
 227 by a universal list of fixes. Across all possible tasks, tools, schemas, knowledge sources, workflows,

and evaluator expectations, new intervention-distinguishable failures can keep appearing. In that setting, a finite intervention dictionary is not a well-posed reliability target.

The positive result begins only after a deployment patch D has been fixed. Inside a patch, the task family, schemas, tools, evaluator expectations, and admissible workflows recur. The engineering question then changes from “Can one library cover all possible LLM failures?” to “How large must the local library be to cover enough of the reachable failure catalogue C_D ?”

Proposition 1 (No Universal Finite Intervention Dictionary). *Fix a residual-error tolerance ε . If a domain D contains an infinite sequence of failures where each new failure is not covered by any finite intervention dictionary covering all earlier failures, then the ε -resolution failure catalogue C_D^ε is infinite. Consequently, no finite intervention dictionary can guarantee residual error below ε over all of D .*

Proof sketch. Each new failure is intervention-distinguishable from the failures before it. The sequence therefore generates infinitely many distinct intervention modes, and any finite dictionary must miss some later mode. The full proof, including the precise definition of mode-level coverage, is in Appendix A.1.

Engineering meaning. Universal reliability is not a finite-library problem. A fixed list of interventions cannot cover open-ended LLM use in general. The rest of the paper is therefore about *patch-local* reliability, not universal reliability. Proposition 1 is the inoculation against a misread: we are not claiming a fixed list of ≈ 50 named patterns covers LLM use in general.

The positive result composes β -stratification (Definition 1), the active-mode bound (Assumption 2), and the log-coverage form (Eq. 3) into a local budget statement. The result is conditional engineering math, not an unconditional theorem about LLMs: it depends on the log-coverage approximation, the active-mode exposure assumption, and the patch hypothesis that C_D is finite or effectively capped. Appendix A.2 states these conditions explicitly before the derivation.

Proposition 2 (Patch-Local Intervention Budget). *Fix a deployment patch D . Let $\varepsilon \in (0, e_{\text{hard}})$ be the target per-hard-token residual error rate. Under the log-coverage approximation of §3.3, a library covering the dominant local modes satisfies the target once*

$$m \geq \left\lceil |C_{\text{eff}}|^{1-\varepsilon/e_{\text{hard}}} \right\rceil, \quad (5)$$

where C_{eff} is either the active catalogue $C_{\text{active},D}(n)$ touched by a single sequence, or the full reachable catalogue C_D of the deployment patch. If $C_{\text{active},D}(n)$ grows logarithmically with the number of hard decisions and $k(n) = \Theta(\log n)$, then the pre-cap intervention budget grows doubly-logarithmically in sequence length. Once the patch catalogue saturates, the required budget becomes domain-constant in $|C_D|$.

Proof sketch. The library removes cumulative hard-error mass $F(m; |C|)$. The residual hard-token error rate is therefore $(1-F)e_{\text{hard}}$. Requiring this residual to be at most ε rearranges, after substituting the log-coverage form, to the stated bound. The full step-by-step derivation, including the pre-cap and cap regimes, is in Appendix A.2.

Engineering meaning. Once the patch is fixed, the problem changes. The relevant engineering question is no longer whether arbitrary future failures exist, but how quickly the reachable local catalogue is discovered and how much hard-error mass is removed by the top interventions. The intervention budget for per-hard-token reliability is small and slowly growing, domain-constant in the cap regime, with the exact size determined by the local discovery and rank-coverage curves rather than by a universal prior.

Sequence-level caveat. Proposition 2 bounds residual error per hard decision. A one-shot sequence-level target is strictly stricter because many hard decisions occur in one output: as k grows the allowable residual per hard decision shrinks, and the required library approaches full-catalogue coverage. In some regimes the non-hard-token error mass alone already exceeds the sequence-level budget, so hard-token interventions cannot meet the SLA by themselves. The full three-regime analysis (and the per-hard-token tolerance τ_{seq} that converts a sequence target into the Proposition 2 bound) is in Appendix A.3.

A common engineering rule of thumb, “fix ≈ 50 patterns per domain $\rightarrow \approx 90\%$ error reduction,” is best understood as a per-hard-token planning prior. The analogous sequence-level claim requires correspondingly more interventions and tighter coverage of the tail.

We call these statements propositions rather than theorems because their force is conditional: they formalise the consequences of the paper’s modelling assumptions rather than deriving a universal law of LLM reliability. The doubly-logarithmic rate is the optimistic special case under logarithmic active-mode exposure; Appendix A.5 gives the Heaps and saturation variants. Appendix A.4 expands on the conditional reading.

4 Empirical Evidence

We collect evidence for three load-bearing claims: errors cluster into a small recurring set (A); each cluster is addressable by one targeted intervention (B); reliability decays sublinearly with output length (C).

Claim A: Failure-mode clustering. The strongest single result is ErrorAtlas [Ashury-Tahan et al., 2026]: 83 models $\times$ 35 datasets, $\approx 7,000$ failures, 17 named categories with a long-tailed, head-concentrated prevalence ordering. Domain-specific Paretos reproduce the shape inside each domain: math (MWPEs top-4 categories dominate per model with strong cross-model overlap [Sun et al., 2025]), code (AssertionError+NameError = 86.35% on HumanEval, replicated across 23 models [Wen et al., 2024]), multi-hop QA (a single dominant “missing-evidence” mode [Zhang et al., 2026]), agentic tool use (under 20 recurring modes across two studies [Cemri et al., 2025, Yao et al., 2024]), and RAG (a handful of distinguishable hallucination types [Wood and Forbes, 2024]). The catalogue is also *stable* across models: ErrorAtlas reports a fixed category hierarchy across 83 models; RFMDataset [Guo et al., 2025] finds “strikingly similar failure mode distributions” across ten advanced reasoning models; EDIT [Dai et al., 2025] measures a $\approx 4.7\%$ key-step fraction stable across models. Schaeffer et al. [2025] provide the closest theoretical anchor, proving that the observed aggregate power-law scaling of LLM evaluations requires the per-problem success-rate distribution to be heavy-tailed near $p = 0$, structurally related to Postulate 1.

Claim B: Cluster-selective capability interventions. A dedicated harvest yields 28 capability-elimination citations across six independent axes (arithmetic, code execution, format/structure, perception/grounding, knowledge/RAG, verification), each axis independently confirmed by between three and nine citations. The full table and the three structural patterns (Pattern A by-construction, B strong-empirical-with-class-shift, C moderate-with-shift) appear in Appendix B.

Arithmetic is the cleanest case: PAL [Gao et al., 2023a] lifts GSM-Hard from 20.1% to 61.5%, and the residuals are problem-comprehension errors rather than arithmetic ones. Format violations admit something stronger still: constrained decoding zeroes invalid-token probability by construction [Suresh et al., 2025, Wang et al., 2025b, Dong et al., 2025]. Code-logic errors yield to execution feedback (Reflexion + AgentCoder push HumanEval pass@1 from 80% to 96.3%, with residuals in spec-misinterpretation [Shinn et al., 2023, Huang et al., 2023]); reasoning steps to process supervision (Math-Shepherd: Mistral-7B MATH 28.6% $\rightarrow$ 43.5% [Wang et al., 2024b]); hallucinations to RAG, where Acurai [Wood and Forbes, 2024] reports 100% elimination on RAGTruth, a benchmark-conditional empirical result rather than a general by-construction guarantee. Two further clusters complete the picture: POROver lifts the not-overrefused rate from 57.6% to 82.1% [Karaman et al., 2024], SAGE-Agent raises When2Call accuracy from 36.5% to 65.2% [Suri et al., 2025].

The class-shift signature, in which post-intervention residuals belong to structurally different classes, is the empirical content of the cluster-orthogonality claim that underwrites Proposition 2’s composition step. Capabilities are coarser than error classes: a single Python interpreter removes execution-error components from five named clusters; constrained decoding eliminates format and the structural component of “missing required element” jointly. A negative control sharpens the selectivity claim. DebugBench [Tian et al., 2024] reports that execution feedback works for syntax and reference errors and is explicitly “unhelpful for logic errors,” exactly where the framework predicts capability provisioning fails.

Claim C: Sublinear length scaling. Direct evidence for sublinear (rather than exponential) decay shows up across very different measurement designs. Loong [Wang et al., 2024a] reports GPT-4o on

Chain-of-Reasoning declining $81.6\% \rightarrow 32.9\%$ across the $10K \rightarrow >200K$ context bins; log-linear over $\log L$ and not well explained by smooth independent per-token exponential decay. GSM- ∞ [Zhou et al., 2025] attacks the independent-error model from a different angle: exponentially more inference compute yields only *linear* AUC gains, and DeepSeek-R1 maintains 10%+ accuracy at 130 reasoning operations where $(1 - \varepsilon)^{130}$ for any reasonable ε predicts near-zero. Press et al. [2023] find a $\approx 40\%$ compositionality gap that is approximately *constant* across the GPT-3 family, direct evidence against per-step independent compounding (which would predict the gap to grow with chain length).

Structural correlates of the same phenomenon appear elsewhere. Pang et al. [2024]’s Anchor LLMs achieve $\approx 99\%$ K/V reduction with only minor accuracy compromise: the effective number of distinct attention-key vectors a model needs is far smaller than n . Think-Prune-Train [Costello et al., 2025] raises Pass@1 while Pass@20 plateaus, the manifold-transition signature. RULER [Hsieh et al., 2024] finds threshold behaviour inconsistent with smooth $(1 - \varepsilon)^n$. METR [Kwa et al., 2025] fits logistic-in-log-length, which is mathematically sublinear in length itself.

Self-consistency [Wang et al., 2023] is consistent with but does not *prove* clustered structure: Condorcet aggregation of iid Bernoulli chains can yield similar gain magnitudes, so we treat it as suggestive rather than diagnostic. Prominent steep-decay counter-evidence [Dziri et al., 2023, Kuratov et al., 2024, Wan et al., 2026, Karpinska et al., 2024] decays over variables distinct from raw token length; per-paper re-audits showing the relocation pattern are in Appendix C.

5 Practical Implications

Reliability engineering is local patch coverage. Within a fixed domain patch, Proposition 2 makes reliability a small-catalogue engineering problem rather than an asymptotic scaling problem. A team should initially budget for a library on the order of tens of interventions, then refine from the local mode-discovery curve $C_{\text{seen},D}(T)$ and the empirical rank-frequency distribution. As a rule of thumb consistent with the head-mass figures in ErrorAtlas, HumanEval, and MWPEs, ≈ 50 named interventions cover the bulk of the per-hard-token failure distribution in many measured domains. This is a planning prior calibrated to current taxonomies, not a universal constant; the local mode-discovery curve sets the actual library size for any given deployment. The same base model in cardiology RAG, legal contract drafting, and code review yields three different libraries because C_D , A_D , σ_D all change with the deployment domain.

Per-hard-token vs. sequence-level targets matter for SLA design. Per-hard-token residual error, the right metric for continuous-correction systems, has a polylog budget. Sequence-level failure probability, the right metric for one-shot systems, is strictly stricter and approaches full-catalogue coverage as k grows. Production deployments should design SLAs to match the actual cost structure, not read the per-token result as the production SLA.

Library design as engineering, not asymptotic theory. An intervention library can be assembled module-by-module, in any order, as long as modules target distinct failure classes. Layer-separated interventions (constrained decoding + retrieval + process supervision + tool call) compose approximately additively in our data; the Le [2026] caveat applies only to interventions sharing a prompt channel. Library construction is therefore highly parallelisable.

Capability libraries are coarser than error-class libraries. Five named math classes (arithmetic, units, counting, list manipulation, date arithmetic) collapse under one Python interpreter; three code classes (SyntaxError, NameError, most TypeError) collapse under one execution-feedback loop; format violations and the structural component of “missing required element” (together more than 20% of ErrorAtlas) collapse under one constrained-decoder deployment. A team budgeting for ≈ 50 interventions per domain can expect to need fewer than 50 capability-axis interventions while covering ≈ 50 named clusters. The six axes of Appendix B are closer to the right unit of engineering accounting than the twelve clusters.

6 Discussion

By-construction elimination as a boundary case. A subset of the capability-elimination harvest deserves separate accounting. Seven of the 28 citations (Appendix B; constrained decoders [Suresh et al., 2025, Dong et al., 2025, Li et al., 2026, Zhang et al., 2023, OpenAI, 2024], proof kernels [Ren et al., 2025], static syntax checks) achieve residual error rate equal to zero *by mathematical construction*, not statistically. Constrained decoders set $P(\text{invalid token}) = 0$ at every generation step; the class of grammar-violating outputs is mathematically empty. In this regime the polylog bound of Proposition 2 is loose: covered clusters contribute exactly zero to residual error, not a small positive quantity. The result strengthens, rather than refutes, the framework. And Pattern A applies precisely to structural/verifiable classes (format, syntax, schema, invalid proofs), which sit at the head of the heavy-tailed cluster frequency distribution: the strongest mechanism lands exactly where the catalogue is densest.

Counter-evidence relocates, not dissolves. Five prominent steep-decay papers [Dziri et al., 2023, Kuratov et al., 2024, Kwa et al., 2025, Wan et al., 2026, Karpinska et al., 2024] are routinely cited as exponential-compounding evidence. Every one decays over a variable distinct from raw token length: compositional graph size [Dziri et al., 2023], number of supporting facts [Kuratov et al., 2024], log-time horizon [Kwa et al., 2025], capacity threshold [Wan et al., 2026], evidence scope [Karpinska et al., 2024]. The apparent rapid decay is, in each case, in a quantity our framework already concentrates the action in (k_{hard} and $|C|$), not in raw sequence length: a relocation, not a dissolution. These regimes remain hard; the framework’s value is directing intervention toward capability provisioning along the actual decay axis rather than toward context-window expansion that does not help. Appendix C walks through each paper.

Limitations. Postulate 1 is empirical, not derived; a domain with genuinely Heaps-power-law mode discovery would invalidate the doubly-logarithmic special case while preserving the qualitative polylog conclusion. The coverage form of §3.3 is an empirical best-fit; readers should not interpret the numerical match against anchors as theoretical confirmation. Cluster-orthogonality is approximate; Le [2026]’s caveat on prompt-channel interference tightens the bound in shared-channel settings. Domain narrowness: Postulate 1’s empirical anchor rests on three taxonomies (ErrorAtlas, HumanEval, MWPEs), all published 2025–2026 and covering general/code/math; the framework is untested on agentic workflows, long-running scientific reasoning, and multi-turn tool use over millions of tokens, precisely the regimes where $|C|$ may grow faster than logarithmically. The $\sigma \approx 1.85$ estimate is a single-point calibration, not a discovery-curve fit; a subsample-vs-distinct-modes plot has not been published for any LLM error taxonomy at this writing. β is latent and we report no empirical range. Taxonomy granularity (the L2–L3 gap) is unmeasured. Patch-shift between deployments changes $A_D, \sigma_D, \beta_D, |C_D|$ simultaneously; libraries calibrated against one patch under-cover the next. Finally, the framework *relocates* the difficulty of long-context reliability rather than resolving it: where k_{hard} grows with task length (adversarial compositional structure, multi-hop chains, long agent horizons), reliability remains hard. The contribution is to name the on-axis intervention, not to make those regimes easy. A falsifiability test the framework passes: DebugBench [Tian et al., 2024] explicitly catalogues which classes execution-feedback can and cannot repair, and the empirical split (works for syntax/reference; “unhelpful for logic errors”) is exactly what the framework’s selectivity claim predicts.

7 Conclusion

LLM reliability is often framed as an asymptotic scaling problem. Universal reliability is not a finite-library problem (Proposition 1): an unbounded domain that keeps producing intervention-distinguishable failures cannot be covered by any finite dictionary. The contribution of this paper is to show that, once a bounded deployment patch is fixed, reliability becomes a local engineering problem. Prior work in this line [Anonymous, 2025] argued that the standard asymptotic framing rests on a false uniformity assumption: errors concentrate at $\approx 5\text{--}10\%$ of tokens. This paper takes the next step: within that sparse set, errors are not only concentrated but also *repetitive*, clustering into a finite catalogue whose size grows logarithmically with observed failures under the empirical postulate of §3.2, or as a small power under the conservative Heaps alternative (Appendix A.5). The conditional consequence (Proposition 2) is that the intervention budget required to bound per-hard-token residual

error scales polylogarithmically in sequence length within a fixed domain patch, and becomes a domain-constant once the patch ceiling $|C_D|$ is reached. Sequence-level reliability targets are strictly tighter. Available evidence is consistent with libraries on the order of tens of interventions covering the head of the per-hard-token failure distribution in many fixed domains. The exact number is patch-indexed and should be estimated from local discovery and rank-coverage curves, not assumed from cross-domain priors. This reframes the question from “can we bound the growth of n -token error?” to “have we catalogued enough failure modes *inside the deployment patch*?” The latter is finite and addressable once a bounded deployment patch has been fixed.

The same logarithmic discovery postulate carries an inverse interpretation (Corollary 1, Appendix A.6): discovering linearly more distinct tail modes requires exponentially more observed hard-failure events. Combined with head-heavy failure mass, this explains why generic frontier post-training can face diminishing reliability returns in open-ended deployment settings — the high-mass head is discovered early, while the tail becomes increasingly expensive to find and contributes less marginal residual-error reduction. The framework therefore does not say frontier scaling is useless. It says frontier-only reliability is economically misaligned with fixed deployment patches: once D is known, local adaptation, tools, validators, retrieval, constrained decoding, and process supervision can target recurring failure mass directly, where generic post-training would have to rediscover the same local repairs indirectly across an open-ended task universe.

The open empirical question this paper invites is direct measurement of the mode-rate constant σ on new domains (agentic, scientific, code-in-production), and a corresponding sequence-level validation that measures S_{base} , β , and active-mode exposure directly in deployed systems. The conceptual follow-on goes in a different direction. This paper names the engineering object: a patch-local failure catalogue and the intervention budget that covers its head. It does not say how production systems actually accumulate and govern that library over time. That question, how the deployment-time scaffold (instructions, tools, retrieval, memory, orchestration, governance) learns to cover the local failure topology, is the subject of follow-on work in this line. Reliability engineering, in that frame, is governance of the scaffold rather than scaling of the weights.

References

- Anonymous. Beyond exponential decay: Rethinking error accumulation in large language models. *arXiv preprint*, 2025. Author field anonymized for double-blind submission; arXiv ID public.
- Akari Asai, Zeqiu Wu, Yizhong Wang, Avirup Sil, and Hannaneh Hajishirzi. Self-RAG: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations (ICLR)*, 2024.
- Shir Ashury-Tahan, Yifan Mai, Elron Bandel, Michal Shmueli-Scheuer, and Leshem Choshen. ErrorMap and ErrorAtlas: Charting the failure landscape of large language models. *arXiv preprint*, 2026.
- Mert Cemri, Melissa Z. Pan, Shuyi Yang, Lakshya A. Agrawal, Bhavya Chopra, Rishabh Tiwari, Kurt Keutzer, Aditya Parameswaran, Dan Klein, Kannan Ramchandran, Matei Zaharia, Joseph E. Gonzalez, and Ion Stoica. Why do multi-agent LLM systems fail? *arXiv preprint*, 2025. Introduces the MAST taxonomy (Multi-Agent System Failure Taxonomy).
- Wenhu Chen, Xueguang Ma, Xinyi Wang, and William W. Cohen. Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks. *Transactions on Machine Learning Research (TMLR)*, 2023.
- Kanzhi Cheng, Qiushi Sun, Yougang Chu, Fangzhi Xu, Yantao Li, Jianbing Zhang, and Zhiyong Wu. SeeClick: Harnessing GUI grounding for advanced visual GUI agents. In *Proceedings of ACL*, 2024.
- Caia Costello, Simon Guo, Anna Goldie, and Azalia Mirhoseini. Think, prune, train, improve: Scaling reasoning without scaling models. *arXiv preprint*, 2025.
- Chengwei Dai et al. Capture the key in reasoning to enhance CoT distillation generalization. In *Proceedings of ACL*, 2025. Earlier arXiv version titled “Beyond Imitation: Learning Key Reasoning Steps from Dual Chain-of-Thoughts in Reasoning Distillation”.

481 Alex Dantart. Reliability by design: Quantifying and eliminating fabrication risk in LLMs. from
482 generative to consultative AI: A comparative analysis in the legal domain and lessons for high-
483 stakes knowledge bases. *arXiv preprint*, 2026.

484 Yixin Dong, Charlie F. Ruan, Yaxing Cai, Ruihang Lai, Ziyi Xu, Yilong Zhao, and Tianqi Chen.
485 XGrammar: Flexible and efficient structured generation engine for large language models. In
486 *Proceedings of MLSys*, 2025.

487 Nouha Dziri, Ximing Lu, Melanie Sclar, Xiang Lorraine Li, Liwei Jiang, Bill Yuchen Lin, Peter West,
488 Chandra Bhagavatula, Ronan Le Bras, Jena D. Hwang, et al. Faith and fate: Limits of transformers
489 on compositionality. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

490 Lizhe Fang, Yifei Wang, Zhaoyang Liu, Chenyang Zhang, Stefanie Jegelka, Jianfeng Gao, Bolin
491 Ding, and Yisen Wang. What is wrong with perplexity for long-context language modeling? In
492 *International Conference on Learning Representations (ICLR)*, 2025.

493 Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and
494 Graham Neubig. PAL: Program-aided language models. In *Proceedings of ICML*, 2023a.

495 Tianyu Gao, Howard Yen, Jiatong Yu, and Danqi Chen. Enabling large language models to generate
496 text with citations. In *Proceedings of EMNLP*, 2023b. Introduces the ALCE benchmark.

497 Alex J. Goodell, Simon N. Chu, Dara Rouholiman, and Larry F. Chu. Large language model
498 agents can use tools to perform clinical calculations. *npj Digital Medicine*, 8(1):163, 2025. doi:
499 10.1038/s41746-025-01475-8.

500 Boyu Gou, Ruohan Wang, Boyuan Zheng, Yanan Xie, Cheng Chang, Yiheng Shu, Huan Sun, and
501 Yu Su. Navigating the digital world as humans do: Universal visual grounding for GUI agents. In
502 *International Conference on Learning Representations (ICLR)*, 2025. ICLR 2025 Oral. Introduces
503 the UGround model.

504 Zhibin Gou, Zhihong Shao, Yeyun Gong, Yelong Shen, Yujiu Yang, Nan Duan, and Weizhu Chen.
505 CRITIC: Large language models can self-correct with tool-interactive critiquing. In *International*
506 *Conference on Learning Representations (ICLR)*, 2024a.

507 Zhibin Gou, Zhihong Shao, Yeyun Gong, Yelong Shen, Yujiu Yang, Minlie Huang, Nan Duan, and
508 Weizhu Chen. ToRA: A tool-integrated reasoning agent for mathematical problem solving. In
509 *International Conference on Learning Representations (ICLR)*, 2024b.

510 Dadi Guo, Jiayu Liu, Zhiyuan Fan, Zhitao He, Haoran Li, Yuxin Li, Yumeng Wang, and Yi R. Fung.
511 Mathematical proof as a litmus test: Revealing failure modes of advanced large reasoning models.
512 *arXiv preprint*, 2025. Introduces the RFMDataset (Reveal Failure Modes).

513 Cheng-Ping Hsieh, Simeng Sun, Samuel Kriman, Shantanu Acharya, Dima Rekesh, Fei Jia, Yang
514 Zhang, and Boris Ginsburg. RULER: What’s the real context size of your long-context language
515 models? In *Conference on Language Modeling (COLM)*, 2024.

516 Dong Huang, Qingwen Bu, Jie M. Zhang, Michael Luck, and Heming Cui. AgentCoder: Multi-agent-
517 based code generation with iterative testing and optimisation. *arXiv preprint*, 2023.

518 Nikhil Kandpal, Haikang Deng, Adam Roberts, Eric Wallace, and Colin Raffel. Large language
519 models struggle to learn long-tail knowledge. In *Proceedings of ICML*, 2023.

520 Batuhan K. Karaman, Ishmam Zabir, Alon Benhaim, Vishrav Chaudhary, Mert R. Sabuncu, and
521 Xia Song. POROver: Improving safety and reducing overrefusal in large language models with
522 overgeneration and preference optimization. *arXiv preprint*, 2024.

523 Marzena Karpinska, Katherine Thai, Kyle Lo, Tanya Goyal, and Mohit Iyyer. One thousand and one
524 pairs: A “novel” challenge for long-context language models. In *Proceedings of EMNLP*, 2024.

525 Yuri Kuratov, Aydar Bulatov, Pavel Anokhin, Ivan Rodkin, Dmitry Sorokin, Artyom Sorokin, and
526 Mikhail Burtsev. BABILong: Testing the limits of LLMs with long context reasoning-in-a-haystack.
527 In *NeurIPS Datasets and Benchmarks Track*, 2024.

528 Thomas Kwa et al. Measuring AI ability to complete long software tasks. In *Advances in Neural*
529 *Information Processing Systems (NeurIPS)*, 2025. METR technical report.

530 Yifan Le. Schema key wording as an instruction channel in structured generation under constrained
531 decoding. *arXiv preprint*, 2026.

532 Yann LeCun. Auto-regressive LLMs are doomed. Tweet, March 26, 2023. [https://x.com/](https://x.com/ylecun/status/1640122342570336267)
533 [ylecun/status/1640122342570336267](https://x.com/ylecun/status/1640122342570336267), 2023.

534 Linzhang Li, Yixin Dong, Guan jie Wang, Ziyi Xu, Alexander Jiang, and Tianqi Chen. XGrammar-2:
535 Efficient dynamic structured generation engine for agentic LLMs. *arXiv preprint*, 2026.

536 Xuyang Li and Anand D. Sarwate. Unraveling the localized latents: Learning stratified manifold
537 structures in LLM embedding space with sparse mixture-of-experts. *arXiv preprint*, 2025.

538 Yujia Li et al. Competition-level code generation with AlphaCode. *Science*, 378(6624):1092–1097,
539 2022. doi: 10.1126/science.abq1158.

540 Hunter Lightman, Vineet Kosaraju, Yura Burda, Harri Edwards, Bowen Baker, Teddy Lee, Jan
541 Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *International*
542 *Conference on Learning Representations (ICLR)*, 2024.

543 Fangyu Liu, Julian Martin Eisenschlos, Francesco Piccinno, Syrine Krichene, Chenxi Pang, Kenton
544 Lee, Mandar Joshi, Wenhua Chen, Nigel Collier, and Yasemin Altun. DePlot: One-shot visual
545 language reasoning by plot-to-table translation. In *Findings of ACL*, 2023.

546 Yadong Lu, Jianwei Yang, Yelong Shen, and Ahmed Awadallah. OmniParser for pure vision based
547 GUI agent. *arXiv preprint*, 2024.

548 Alex Mallen, Akari Asai, Victor Zhong, Rajarshi Das, Daniel Khashabi, and Hannaneh Hajishirzi.
549 When not to trust language models: Investigating effectiveness of parametric and non-parametric
550 memories. In *Proceedings of ACL*, 2023. Introduces the PopQA benchmark.

551 Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. *Introduction to Information*
552 *Retrieval*. Cambridge University Press, 2008. Section on Heaps’ law: [https://nlp.stanford.](https://nlp.stanford.edu/IR-book/html/htmledition/heaps-law-estimating-the-number-of-terms-1.html)
553 [edu/IR-book/html/htmledition/heaps-law-estimating-the-number-of-terms-1.](https://nlp.stanford.edu/IR-book/html/htmledition/heaps-law-estimating-the-number-of-terms-1.html)
554 [html](https://nlp.stanford.edu/IR-book/html/htmledition/heaps-law-estimating-the-number-of-terms-1.html).

555 Ayana Niwa and Hayate Iso. AmbigNLG: Addressing task ambiguity in instruction for NLG. In
556 *Proceedings of EMNLP*, 2024.

557 OpenAI. Introducing structured outputs in the API. OpenAI Engineering Blog. [https://openai.](https://openai.com/index/introducing-structured-outputs-in-the-api/)
558 [com/index/introducing-structured-outputs-in-the-api/](https://openai.com/index/introducing-structured-outputs-in-the-api/), 2024. Published August 6,
559 2024.

560 OpenAI. GPT-5 system card. <https://cdn.openai.com/gpt-5-system-card.pdf>, 2025. Pub-
561 lished August 7, 2025.

562 Jianhui Pang, Fanghua Ye, Derek F. Wong, Xun He, Wenxiang Chen, and Longyue Wang. Anchor-
563 based large language models. In *Findings of ACL*, 2024.

564 Kiho Park, Yo Joong Choe, and Victor Veitch. The linear representation hypothesis and the geometry
565 of large language models. In *Proceedings of ICML*, 2024.

566 Khush Patel, Siva Surendira, Jithin George, and Shreyas Kapale. The six sigma agent: Achieving
567 enterprise-grade reliability in LLM systems through consensus-driven decomposed execution.
568 *arXiv preprint*, 2026.

569 Ofir Press, Muru Zhang, Sewon Min, Ludwig Schmidt, Noah A. Smith, and Mike Lewis. Measuring
570 and narrowing the compositionality gap in language models. In *Findings of EMNLP*, 2023.

571 Z. Z. Ren et al. DeepSeek-Prover-V2: Advancing formal mathematical reasoning via reinforcement
572 learning for subgoal decomposition. *arXiv preprint*, 2025.

573 Rylan Schaeffer et al. How do large language monkeys get their power (laws)? In *International*
574 *Conference on Machine Learning (ICML)*, 2025.

575 Yu Shang, Yu Li, Keyu Zhao, Likai Ma, Jiahe Liu, Fengli Xu, and Yong Li. AgentSquare: Automatic
576 LLM agent search in modular design space. *arXiv preprint*, 2024.

577 Yuling Shi, Songsong Wang, Chengcheng Wan, Min Wang, and Xiaodong Gu. From code to
578 correctness: Closing the last mile of code generation with hierarchical debugging. *arXiv preprint*,
579 2024. Introduces the MGDebugger method.

580 Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu
581 Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural*
582 *Information Processing Systems*, volume 36, 2023.

583 Yuhong Sun, Zhangyue Yin, Xuanjing Huang, Xipeng Qiu, and Hui Zhao. Error classification of
584 large language models on math word problems: A dynamically adaptive framework. In *arXiv*
585 *preprint*, 2025. Introduces the MWPEs-300K dataset (304,865 error samples from 15 LLMs
586 across four MWP datasets).

587 Tarun Suresh et al. DINGO: Constrained inference for diffusion LLMs. *arXiv preprint*, 2025.

588 Manan Suri, Puneet Mathur, Nedim Lipka, Franck Dernoncourt, Ryan A. Rossi, and Dinesh Manocha.
589 Structured uncertainty guided clarification for LLM agents. *arXiv preprint*, 2025. Introduces the
590 SAGE-Agent method.

591 Runchu Tian, Yining Ye, Yujia Qin, Xin Cong, Yankai Lin, Zhiyuan Liu, and Maosong Sun. De-
592 bugBench: Evaluating debugging capability of large language models. In *Findings of ACL*,
593 2024.

594 Akihiko Wada et al. Retrieval-augmented generation elevates local LLM quality in radiology contrast
595 media consultation. *npj Digital Medicine*, 8(1):395, 2025. doi: 10.1038/s41746-025-01802-z.

596 Kai Wan et al. A Fano-style accuracy upper bound for LLM single-pass reasoning in multi-hop QA.
597 In *International Conference on Learning Representations (ICLR)*, 2026.

598 Benlu Wang et al. From scores to steps: Diagnosing and improving LLM performance in evidence-
599 based medical calculations. In *Proceedings of EMNLP*, 2025a. Introduces the MedRaC benchmark.

600 Darren Yow-Bang Wang, Zhengyuan Shen, Soumya Smruti Mishra, Zhichao Xu, Yifei Teng, and
601 Haibo Ding. SLOT: Structuring the output of large language models. *arXiv preprint*, 2025b. SLOT
602 = Structured LLM Output Transformer.

603 Minzheng Wang et al. Leave no document behind: Benchmarking long-context LLMs with extended
604 multi-doc QA. In *Proceedings of EMNLP*, 2024a.

605 Peiyi Wang, Lei Li, Zhihong Shao, R. X. Xu, Damai Dai, Yifei Li, Deli Chen, Y. Wu, and Zhifang
606 Sui. Math-Shepherd: Verify and reinforce LLMs step-by-step without human annotations. In
607 *Proceedings of ACL*, 2024b.

608 Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdh-
609 ury, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models.
610 In *International Conference on Learning Representations (ICLR)*, 2023.

611 Yubo Wang et al. MMLU-Pro: A more robust and challenging multi-task language understanding
612 benchmark. In *NeurIPS Datasets and Benchmarks Track (Spotlight)*, 2024c.

613 Hao Wen et al. Fixing function-level code generation errors for foundation large language models.
614 *arXiv preprint*, 2024. Introduces the LlmFix method. Reports AssertionError 63.64% + NameError
615 22.71% = 86.35% of HumanEval failures across 14 LLMs.

616 Michael C. Wood and Adam A. Forbes. 100% elimination of hallucinations on RAGTruth for GPT-4
617 and GPT-3.5 Turbo. *arXiv preprint*, 2024. Introduces the Acurai method.

- 618 Tianbao Xie et al. Scaling computer-use grounding via user interface decomposition and synthesis.
619 In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. NeurIPS 2025 Spotlight.
620 Introduces the Jedi grounding dataset and OSWorld-G benchmark.
- 621 Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A benchmark for
622 tool-agent-user interaction in real-world domains. *arXiv preprint*, 2024.
- 623 Cyril Zakka, Rohan Shad, Akash Chaurasia, Alex R. Dalal, Jennifer L. Kim, Michael Moor, Robyn
624 Fong, Curran Phillips, Kevin Alexander, Euan Ashley, Jack Boyd, Kathleen Boyd, Karen Hirsch,
625 Curtis Langlotz, Rita Lee, Joanna Melia, Joanna Nelson, Karim Sallam, Stacey Tullis, Melissa Ann
626 Vogelsong, John Patrick Cunningham, and William Hiesinger. Almanac — retrieval-augmented
627 language models for clinical medicine. *NEJM AI*, 1(2):AIoa2300068, 2024. doi: 10.1056/
628 AIOa2300068.
- 629 Kexun Zhang et al. Don’t fine-tune, decode: Syntax error-free tool use via constrained decoding.
630 *arXiv preprint*, 2023. Introduces the ToolDec method.
- 631 Meiru Zhang, Zaiqiao Meng, and Nigel Collier. Failure modes in multi-hop QA: The weakest link
632 effect and the recognition bottleneck. *arXiv preprint*, 2026.
- 633 Yang Zhou, Hongyi Liu, Zhuoyan Chen, Yuandong Tian, and Beidi Chen. GSM-Infinite: How do
634 your LLMs behave over infinitely increasing context length and reasoning complexity? *arXiv*
635 *preprint*, 2025.

636 A Formal Proofs, Derivations, and Sensitivity Analysis

637 A.1 Proof of Proposition 1: No Universal Finite Intervention Dictionary

638 Fix a residual-error tolerance ε . Say that an intervention *covers* a failure mode if it reduces the residual
639 error of that mode below ε . Coverage is mode-level: an intervention that covers a mode covers every
640 failure event in that mode. The two conditions “ D is intervention-unbounded” and “ $|C_D^\varepsilon| = \infty$ ” are
641 equivalent under mode-level coverage; an unbounded witnessing sequence is constructed by taking
642 one representative per mode, and an infinite catalogue forces the existence of such a sequence.

643 Let D be a domain. Call D *intervention-unbounded* if it contains an infinite sequence of failures
644 $f_1, f_2, f_3, \dots$ such that each new f_j is not covered by any finite intervention dictionary that covers
645 all earlier failures $\{f_1, \dots, f_{j-1}\}$.

646 **Step 1: assume the opposite.** Suppose, for contradiction, that C_D^ε is finite. Then there are only
647 finitely many intervention-distinguishable failure modes in D . Write them as $C_D^\varepsilon = \{c_1, \dots, c_M\}$
648 for some finite M .

649 **Step 2: what finiteness means.** If there are only M intervention-distinguishable modes, then after
650 all M modes have appeared in the sequence, every later failure must belong to one of the already-seen
651 modes.

652 **Step 3: same mode means same intervention class.** Modes are defined at intervention resolution
653 ε . If a later failure belongs to the same mode as an earlier failure, then the intervention dictionary that
654 covers the earlier representative of that mode also covers the later failure below residual tolerance ε .

655 **Step 4: contradiction.** Intervention-unboundedness says exactly the opposite: each new f_j is not
656 covered by any finite dictionary that covers $\{f_1, \dots, f_{j-1}\}$. Therefore f_j cannot belong to any earlier
657 intervention mode, so each f_j introduces a new intervention-distinguishable mode. The sequence
658 $f_1, f_2, f_3, \dots$ induces infinitely many such modes, contradicting the assumption that C_D^ε is finite.
659 Hence $|C_D^\varepsilon| = \infty$.

660 **Step 5: no finite dictionary covers the domain.** Suppose, again for contradiction, that some finite
661 intervention dictionary $\mathcal{I}$ covers all of D . Then $\mathcal{I}$ covers every finite prefix $\{f_1, \dots, f_{j-1}\}$ for every
662 j . By intervention-unboundedness, any dictionary covering that prefix fails to cover f_j . Therefore $\mathcal{I}$
663 does not cover f_j , contradicting the assumption that $\mathcal{I}$ covers all of D . $\square$

664 A.2 Derivation of Proposition 2: Patch-Local Intervention Budget

665 **Conditions used.** Proposition 2 depends on four assumptions:

666 1. **Coverage model.** The cumulative hard-error mass covered by the top m modes is approximated
667 by

$$F(m; |C|) = \min\left(1, \frac{\ln m}{\ln |C|}\right)$$

668 on the declared domain $m \geq 2, |C| \geq 2$.

669 2. **Non-trivial, attainable target.** The residual target satisfies $\varepsilon \in (0, e_{\text{hard}})$. If $\varepsilon \geq e_{\text{hard}}$ no
670 intervention is needed; if $\varepsilon \leq 0$ the target is unattainable unless all residual hard-token error is
671 eliminated.

672 3. **Patch-local catalogue.** The result applies only after a deployment patch D has been fixed and
673 its reachable catalogue is modelled as finite or effectively capped.

674 4. **Sequence-length claim.** The doubly-logarithmic rate further requires Assumption 2 and $k(n) =$
675 $\Theta(\log n)$. Without these, Proposition 2 still gives a catalogue-size budget but not the same
676 sequence-length scaling.

677 **Step 1: define the target.** Let e_{hard} be the baseline hard-token error rate and $\varepsilon \in (0, e_{\text{hard}})$ the target
678 after intervention.

679 **Step 2: define the effective catalogue.** For per-sequence scaling, set $C_{\text{eff}} = C_{\text{active}, D}(n)$: how
680 many failure modes can a single sequence of length n activate? For full deployment-library budgeting,
681 set $C_{\text{eff}} = C_D$: how large must the library be to cover the recurring failure modes reachable inside
682 D ?

683 **Step 3: cumulative coverage.** Let the ranked local failure modes have hard-error masses $p_1 \geq$
684 $p_2 \geq \dots \geq p_{|C_{\text{eff}}|}$ with $\sum_i p_i = 1$. A library covering the top m modes removes cumulative
685 hard-error mass $F(m; |C_{\text{eff}}|) = \sum_{i=1}^m p_i$, approximated by the log-coverage form of §3.3.

686 **Step 4: residual after intervention.** The uncovered mass fraction is $1 - F(m; |C_{\text{eff}}|)$, so the
687 residual per-hard-token error rate is $e_{\text{res}}(m) = (1 - F(m; |C_{\text{eff}}|)) e_{\text{hard}}$.

688 **Step 5: impose the target.** Requiring $e_{\text{res}}(m) \leq \varepsilon$ and dividing by $e_{\text{hard}} > 0$ gives

$$1 - F(m; |C_{\text{eff}}|) \leq \frac{\varepsilon}{e_{\text{hard}}}, \quad \text{i.e.,} \quad F(m; |C_{\text{eff}}|) \geq 1 - \frac{\varepsilon}{e_{\text{hard}}}.$$

689 **Step 6: substitute the log-coverage form.** Since $\varepsilon < e_{\text{hard}}$, the required coverage $1 - \varepsilon/e_{\text{hard}} \in$
690 $(0, 1)$, so the cap in F is inactive before saturation. Using $F = \ln m / \ln |C_{\text{eff}}|$:

$$\frac{\ln m}{\ln |C_{\text{eff}}|} \geq 1 - \frac{\varepsilon}{e_{\text{hard}}}.$$

691 **Step 7: solve for m .** Multiplying by $\ln |C_{\text{eff}}| > 0$ and exponentiating,

$$m \geq \exp\left[\left(1 - \frac{\varepsilon}{e_{\text{hard}}}\right) \ln |C_{\text{eff}}|\right] = |C_{\text{eff}}|^{1 - \varepsilon/e_{\text{hard}}}.$$

692 Taking the ceiling (since m is integer-valued) yields $m \geq \lceil |C_{\text{eff}}|^{1 - \varepsilon/e_{\text{hard}}} \rceil$, the bound stated in Eq. (5).

693 **Step 8: sequence-length rate.** For per-sequence scaling, set $C_{\text{eff}} = C_{\text{active}, D}(n)$. Assumption 2
694 bounds $|C_{\text{active}, D}(n)| \leq \min(A'_D + \sigma'_D \ln h(n), |C_D|)$. In the pre-cap regime, i.e., while $A'_D +$
695 $\sigma'_D \ln h(n) < |C_D|$, the ceiling has not been reached and $|C_{\text{active}, D}(n)| = O(\ln h(n))$. With $h(n) =$
696 $\beta k(n)$ and $k(n) = \Theta(\log n)$, we get $h(n) = \Theta(\log n)$ and $\ln h(n) = \Theta(\log \log n)$, hence

$$|C_{\text{active}, D}(n)| = O(\log \log n),$$

697 and substituting into the budget,

$$m = O((\log \log n)^{1 - \varepsilon/e_{\text{hard}}}).$$

698 This is the doubly-logarithmic pre-cap special case.

699 **Step 9: cap regime.** Once active-mode discovery saturates the patch catalogue, $|C_{\text{active},D}(n)| =$
700 $|C_D|$, so $m \geq \lceil |C_D|^{1-\varepsilon/e_{\text{hard}}} \rceil$, which no longer depends on n . The intervention budget is then
701 domain-constant. $\square$

702 A.3 Sequence-Level Reliability Derivation

703 Proposition 2 gives a per-hard-token residual target. Production systems often care about a stricter
704 target: the probability that the entire sequence is correct.

705 Starting from the composed reliability of Eq. (1), and writing the post-intervention hard-token rate as
706 e_{res} ,

$$P(\text{correct}) = (1 - e_{\text{res}})^{\beta k} (1 - e_{\text{easy}})^{(1-\beta)k} (1 - e_{\text{non}})^{n-k}.$$

707 Define the non-hard-token survival factor

$$S_{\text{base}} = (1 - e_{\text{easy}})^{(1-\beta)k} (1 - e_{\text{non}})^{n-k}, \quad (6)$$

708 so that $P(\text{correct}) = (1 - e_{\text{res}})^{\beta k} S_{\text{base}}$. The sequence-level target $P(\text{correct}) \geq 1 - \varepsilon_{\text{seq}}$ becomes

$$(1 - e_{\text{res}})^{\beta k} S_{\text{base}} \geq 1 - \varepsilon_{\text{seq}}, \quad \text{i.e.,} \quad (1 - e_{\text{res}})^{\beta k} \geq \frac{1 - \varepsilon_{\text{seq}}}{S_{\text{base}}}.$$

709 Three regimes follow.

710 **Regime (i): non-hard-token errors already violate the target.** If $S_{\text{base}} < 1 - \varepsilon_{\text{seq}}$, then even
711 setting $e_{\text{res}} = 0$ cannot meet the target, because the maximum possible survival after eliminating all
712 hard-token failures is only S_{base} . Hard-token interventions alone cannot meet the sequence-level SLA.
713 This is the route by which the exponential-in- n concern re-enters and should be diagnosed before any
714 catalogue-budgeting exercise.

715 **Regime (ii): hard-token residual error determines feasibility.** If $S_{\text{base}} \geq 1 - \varepsilon_{\text{seq}}$, the target may
716 be feasible. Taking the $(1/(\beta k))$ -th power of the rearranged inequality and isolating e_{res} ,

$$e_{\text{res}} \leq 1 - \left(\frac{1 - \varepsilon_{\text{seq}}}{S_{\text{base}}} \right)^{1/(\beta k)} =: \tau_{\text{seq}}.$$

717 Applying Proposition 2 with ε replaced by τ_{seq} ,

$$m \geq \left\lceil |C_{\text{eff}}|^{1-\tau_{\text{seq}}/e_{\text{hard}}} \right\rceil.$$

718 As the sequence-level target becomes stricter (ε_{seq} shrinks), $\tau_{\text{seq}} \rightarrow 0$ and the exponent $1 -$
719 $\tau_{\text{seq}}/e_{\text{hard}} \rightarrow 1$, so $m \rightarrow |C_{\text{eff}}|$. Strict one-shot sequence-level reliability pushes the system to-
720 ward full-catalogue coverage.

721 **Regime (iii): baseline hard-token error is already acceptable.** If $\tau_{\text{seq}} \geq e_{\text{hard}}$, the baseline hard-
722 token rate already satisfies the sequence target (no intervention is required because $e_{\text{res}}(0) = e_{\text{hard}}$
723 from Eq. (4)).

724 **Conclusion.** Per-hard-token reliability is easier than sequence-level reliability. A library that gives
725 a large reduction in residual hard-token error may still be insufficient for one-shot sequence-level
726 guarantees when many hard decisions occur in a single output. This is why the main paper treats the
727 “ ≈ 50 patterns” rule of thumb as a per-hard-token planning prior, not a one-shot sequence-level SLA.

728 A.4 Why these are propositions rather than unconditional theorems

729 Proposition 1 is a definitional impossibility result: once intervention-unboundedness is assumed, an
730 infinite intervention-resolution catalogue follows. Its role is not to prove that every unbounded domain
731 necessarily has infinite failure modes, but to show that open-ended domains cannot be assumed to
732 admit finite dictionaries.

733 Proposition 2 is conditional engineering math. It does not prove that LLM failures universally obey
734 logarithmic mode discovery. It proves that *if* a bounded patch has a finite or effectively capped

735 reachable catalogue, *and if* cumulative intervention coverage follows the stated head-heavy form, *then*
 736 the required per-hard-token intervention budget grows slowly and becomes constant after catalogue
 737 saturation.

738 The empirical burden therefore lies not in the algebra but in measuring, for each deployment patch, the
 739 local discovery curve $C_{\text{seen},D}(T)$, the per-sequence activation $C_{\text{active},D}(n)$, the cumulative coverage
 740 $F(m; |C_D|)$, the hard-token fraction β_D , and the baseline hard-token rate e_{hard} . This is why the paper
 741 frames the results as a reliability-engineering scaffold rather than as universal theorems about LLM
 742 behaviour.

743 A.5 Heaps Power-Law Variant and Cluster-Count Sensitivity

744 A reader who prefers to derive cluster-count growth from standard Heaps’ law rather than from
 745 Postulate 1 obtains a qualitatively similar result. Let $|C|(k_{\text{hard}}) \approx K \cdot k_{\text{hard}}^b$ with $b \in (0, 1)$. Canonical
 746 fits give $b \approx 0.5$ for natural-language vocabularies; failure-mode taxonomies plateau much more
 747 sharply (ErrorAtlas stabilises at $|C| = 17$ across 10^4 failures), implying a smaller effective $b \approx 0.2$
 748 in our setting. Composing with $k = \Theta(\log n)$, we get $|C| = O((\log n)^b)$, and Proposition 2 becomes

$$m = O\left((\log n)^{b \cdot (1 - \varepsilon / e_{\text{hard}})}\right).$$

749 This is still polylogarithmic in n for any $b \in (0, 1)$ and any $\varepsilon < e_{\text{hard}}$. The paper’s qualitative claim
 750 survives either choice of cluster-count law.

751 **Symbolic-form sensitivity across candidate laws.** The polylog conclusion depends on which
 752 cluster-count law one accepts; available evidence is consistent with multiple candidates because no
 753 subsample-discovery curve has been published for any LLM failure-mode taxonomy at this writing.
 754 We therefore report symbolic rates rather than fitted constants. With $h(n) = \beta k(n)$:

- 755 • **Logarithmic:** $|C_{\text{active},D}(n)| = O(\log h(n))$. If $k(n) = \Theta(\log n)$, then $m =$
 756 $O((\log \log n)^{1 - \varepsilon / e_{\text{hard}}})$.
- 757 • **Heaps:** $|C_{\text{active},D}(n)| = O(h(n)^b)$ with $b \in (0, 1)$. If $k(n) = \Theta(\log n)$, then $m =$
 758 $O((\log n)^{b(1 - \varepsilon / e_{\text{hard}})})$.
- 759 • **Saturating:** $|C_{\text{active},D}(n)| \leq |C_D|$. Then $m = O(|C_D|^{1 - \varepsilon / e_{\text{hard}}})$, constant in n once the
 760 patch ceiling is reached.

761 The qualitative conclusion is robust **under the $k(n) = \Theta(\log n)$ regime**: m grows more slowly than
 762 any positive power of n under every candidate, and the directional claim (“a small library covers the
 763 head of the failure distribution in the per-hard-token regime”) survives. Only the exponent shifts:
 764 doubly-logarithmic under logarithmic discovery, $(\log n)^b$ with small b under Heaps, constant in the
 765 cap regime. The doubly-logarithmic rate is the optimistic special case. If $k(n)$ grows as a positive
 766 power of n , the Heaps variant inherits that power and the polylog-in- n language fails along that axis;
 767 the framework’s intervention prescription still applies, but its asymptotic-rate framing does not.

768 Numerical constants require a measured discovery curve $C_{\text{seen},D}(T)$ or $C_{\text{active},D}(n)$. Existing tax-
 769 onomies provide endpoint category counts at a single corpus scale (ErrorAtlas at $|C| = 17$ for
 770 $\approx 10^4$ failures), not discovery curves. The subsample-discovery measurement remains the explicit
 771 empirical test that would either tighten the postulate or fall back to the Heaps variant. Until that
 772 measurement exists, the framework’s headline rate should be read as *polylogarithmic in the pre-cap*
 773 *regime, domain-constant in the cap regime*, with the specific exponent flagged as a falsifiability test
 774 rather than a fitted prediction.

775 A.6 Inverse Discovery Cost

776 The body Corollary 1 inverts the logarithmic discovery postulate to read out a sample-budget cost
 777 for distinct-mode discovery. This appendix gives the algebra, the numerical anchors, sensitivity to
 778 the Heaps cluster-count law of §A.5, and a separate mode-mediated gain corollary that connects
 779 discovery cost to broad capability proxies.

780 **Setup.** Let $q(T) = |C_{\text{seen},D}(T)|$ be the number of distinct failure modes discovered in patch D
 781 after T observed hard-failure events. Postulate 1 approximates discovery as

$$q(T) \approx A_D + \sigma_D \ln T, \quad \sigma_D > 0.$$

782 **Inversion.** Subtract A_D , divide by σ_D , and exponentiate. The sample budget required to discover
 783 at least q distinct modes is

$$T(q) \approx \exp\left(\frac{q - A_D}{\sigma_D}\right).$$

784 **Multiplicative cost per additional mode.** Comparing the budget at $q + \Delta q$ to the budget at q ,

$$\frac{T(q + \Delta q)}{T(q)} \approx \exp\left(\frac{\Delta q}{\sigma_D}\right),$$

785 so each block of Δq additional modes costs a *multiplicative* sample increase. At the conservative
 786 calibration $\sigma_D \approx 1.85$ (§3.2), $\exp(5/1.85) \approx 14.9$ and $\exp(10/1.85) \approx 222$: five extra modes cost
 787 roughly $15\times$ more observed hard failures, ten extra modes cost roughly $220\times$.

788 **What this is and is not.** The corollary describes *new distinct-mode discovery*. Ordinary failures
 789 inside already-discovered modes may remain common and cheap to observe; the exponential cost
 790 lives on the category-novelty axis. Conflating ordinary failure rate with novel-mode arrival rate would
 791 over-claim the result.

792 **Heaps alternative.** Under the Heaps cluster-count law of §A.5, $q(T) = KT^b$ with $b \in (0, 1)$, the
 793 inverse cost is polynomial:

$$T(q) = (q/K)^{1/b}.$$

794 The exponential inverse-cost claim is specific to the logarithmic postulate; the broader qualitative
 795 claim that tail discovery has diminishing returns survives under any concave discovery curve. Sensi-
 796 tivity is therefore: exponential under logarithmic, polynomial under Heaps, undefined past the patch
 797 ceiling.

798 **Saturation regime.** Once $q(T) \leq |C_D|$ has been saturated, discovery stops; the inversion applies
 799 only in the pre-cap regime. Inside saturated patches the corollary’s multiplicative-cost claim is
 800 vacuous because no novel modes remain to discover, which is itself a property of the patch and not a
 801 failure of the corollary.

802 **Mode-mediated capability gain (sub-corollary).** Suppose broad capability or reliability gain G
 803 inside the patch is approximately linear in the number of useful discovered modes, $G(q) = G_0 + \gamma q$
 804 for some $\gamma > 0$. Composing with logarithmic mode discovery,

$$G(T) = G_0 + \gamma A_D + \gamma \sigma_D \ln T,$$

805 so G grows logarithmically in observed hard-failure exposure. Inverting,

$$T(G) = \exp\left(\frac{G - G_0 - \gamma A_D}{\gamma \sigma_D}\right).$$

806 A linear gain in mode-mediated broad reliability therefore requires exponential growth in observed
 807 hard-failure exposure under the postulate. We deliberately phrase G as a *mode-mediated capa-*
 808 *bility/reliability* proxy, not as “intelligence”: the corollary does not say frontier scaling is useless
 809 or that intelligence requires exponential data in any general sense. It says, more narrowly, that
 810 for fixed deployment reliability where improvement is mediated by discovering new useful modes,
 811 generic open-domain training pays a heavy data tax relative to direct patch-local measurement and
 812 intervention.

813 **Engineering reading.** The corollary explains, without invoking new mechanisms, two empirical
 814 signals: (a) why generic post-training shows diminishing reliability returns once a domain’s head
 815 modes are covered, and (b) why patch-local measurement combined with targeted tools, retrieval,
 816 validators, constrained decoding, and process supervision often outperforms more frontier-scale
 817 data on the deployment SLA. Frontier scaling and patch-local engineering solve different problems:
 818 scaling improves the substrate; patch-local engineering removes recurring deployment failure mass.
 819 □

820 B Full Failure-Mode Taxonomy and the Capability-Elimination Harvest

821 The intervention literature provides at least one targeted countermeasure for each named cluster of §4.
822 Two organising granularities exist: at the capability level (six axes) the cluster-orthogonality property
823 is most clearly visible; at the error-class level (twelve named clusters) the evidence aligns with the
824 taxonomies of §4 but is finer-grained than the underlying capability mechanisms.

825 **Six capability axes and three structural patterns.** A dedicated harvest yields **28 quantitatively-**
826 **anchored citations** across **six independent capability axes**: Arithmetic (Python/symbolic execution),
827 Code Execution (REPL/sandbox feedback), Format/Structure (constrained decoding, FSMs, grammar
828 engines), Perception/Grounding (visual grounding for GUI, charts, tables), Knowledge/RAG (dense
829 retrieval and citation grounding), Verification (proof checkers, learned verifiers, classifier rerouting,
830 process supervision). Each axis is independently confirmed by between three and nine citations. We
831 stratify the 28 by kind of evidence into three patterns:

832 **Pattern A: hard guarantees (by-construction).** Seven citations achieve residual error rate equal
833 to zero *by construction*, restricted strictly to structural/verifiable classes: constrained decoders set
834 $P(\text{invalid token}) = 0$ at every step [Suresh et al., 2025, Zhang et al., 2023, Dong et al., 2025,
835 Li et al., 2026, OpenAI, 2024]; static syntax checks reject programs with a `SyntaxError` before
836 execution [Wen et al., 2024]; proof kernels reject any output failing type-checking [Ren et al., 2025].
837 The class of grammar-violating outputs is mathematically empty under these mechanisms.

838 **Pattern B: strong empirical reductions with class-shift signature.** Roughly fourteen citations
839 report empirical reductions of 80–100% in a named error class, with the post-intervention failure
840 log dominated by structurally different residual classes. Program-of-Thoughts on GSM8K [Chen
841 et al., 2023]: calculation errors drop from 30% of failures to 0%, residuals are 62% reasoning + 36%
842 misunderstanding. OpenMedCalc [Goodell et al., 2025]: “only interpretation errors were identified.”
843 Acurai [Wood and Forbes, 2024]: 100% hallucination elimination on RAGTruth (95% CI 91–100%),
844 strong empirical rather than by-construction.

845 **Pattern C: moderate reductions (60–80%) with residuals shifting outside the target class.**
846 Legal RAG [Dantart, 2026]: fabricated citations $> 30\% \rightarrow < 0.2\%$. GPT-5 SimpleQA with web
847 access [OpenAI, 2025]: $47\% \rightarrow 9.6\%$ (inter-condition, not within-condition). CRITIC [Gou et al.,
848 2024a]: toxic generation -79.2% .

849 **The 28 citations, organised by axis and pattern.**

Axis	Citation & targeted class	Pre → Post	Pattern
Verification	Ren et al. [2025]: Invalid Lean proof	any → 0% by constr.	A
Format	OpenAI [2024]: JSON schema violation	>60% → 0% by constr.	A
Format	Zhang et al. [2023]: Tool-call syntax	21–100% → 0%	A
Format	Suresh et al. [2025]: JSON parse failure	13–82% → 0%	A
Format	Dong et al. [2025]: Multi-format errors	20–38% → 0%	A
Format	Li et al. [2026]: Malformed tool calls	33–78% → 0%	A
Arithmetic	Chen et al. [2023]: Calculation on GSM8K	30% → 0% of failures	B
Arithmetic	Goodell et al. [2025]: Clinical arithmetic	“only interpretation errors”	B
Knowledge/RAG	Wada et al. [2025]: RAG hallucinations	8% → 0% ($p = 0.012$)	B
Knowledge/RAG	Wood and Forbes [2024]: Context-conflict hallucinations	100% → 0% on subset	B
Knowledge/RAG	Dantart [2026]: Fabricated legal citations	>30% → <0.2%	C
Code Exec	Wen et al. [2024]: <code>SyntaxError</code> (HumanEval)	5.76% → 0.01%	A
Code Exec	Li et al. [2022]: False-positive submissions	62% → 4%	B
Code Exec	Shi et al. [2024]: 5 code-bug classes	100% repair (5/6 classes)	B
Knowledge/RAG	Gao et al. [2023b]: Citation grounding (ELI5)	≈50% w/o complete support	C
Knowledge/RAG	OpenAI [2025]: Factual errors	47% → 9.6% w/web	C
Knowledge/RAG	Zakka et al. [2024]: Clinical citation errors	44% → 9%	C
Arithmetic	Wang et al. [2025a]: Arithmetic in medical reasoning	426 → 74 errors	B
Code Exec	Wen et al. [2024]: <code>NameError</code> repair	22.7% → 2.3%	C
Perception	Gou et al. [2025]: GUI grounding errors	16.2% → 73.3% acc.	B
Perception	Xie et al. [2025]: GUI task failures	5% → 27% SR ($5.4\times$)	B
Perception	Cheng et al. [2024]: Element mislocation	5.2% → 53.4% ($10.3\times$)	B
Perception	Liu et al. [2023]: Chart-reading errors	38.2% → 67.6% (+29.4 pp)	B
Verification	Gou et al. [2024a]: Toxic generation	−79.2%	C
Verification	Wang et al. [2024b]: Reasoning step errors	28.6% → 43.5% w/rerank	B
Knowledge/RAG	Asai et al. [2024]: Unsupported facts	55.5% → 22% error	B
Perception	Lu et al. [2024]: Icon/element errors	70.5% → 93.8% (79% red.)	B
Knowledge/RAG	Mallen et al. [2023]: Long-tail entity errors	≈80% → ≈50% failure	B

Capabilities are coarser than error classes. Several entries in the harvest reveal “two-for-one” reductions where one capability addresses multiple named clusters: a single Python interpreter removes the execution-error component of arithmetic, unit conversion, simple counting, list manipulation, and date arithmetic; constrained decoding eliminates by construction both format violations and the structural component of “missing required element”; code execution feedback strongly reduces `SyntaxError`, `NameError`, and most `TypeError` together; RAG strongly reduces factual hallucinations, fabricated citations, and outdated information jointly. The practical capability library required is therefore smaller than the count of named error categories.

The twelve named clusters.

Cluster	Axis	Intervention	Before → After	Patt.
A. Arithmetic	Arithmetic	PAL Python [Gao et al., 2023a]	20.1% → 61.5%	B
B. Unit conversion	Arithmetic	Folded into A (Python)	n/a	n/a
C. Counting	Arithmetic	Explicit counter (gap)	n/a	gap
D. Format/schema	Format	DINGO/SLOT [Suresh et al., 2025, Wang et al., 2025b]	up to +68 pp; 99.5% acc.	A
E. Code logic/type	Code Exec	Reflexion/AgentCoder [Shinn et al., 2023, Huang et al., 2023]	80% → 96.3% pass@1	A/B
F. Multi-hop drift	Knowledge	Entity-grounded rewriter	≈5–10 pp	C
G. Reasoning step	Verification	Math-Shepherd [Wang et al., 2024b]	28.6% → 43.5%	B
H. Spec misinterpret.	Verification	Clarification [Niwa and Iso, 2024]	ROUGE-L +15	C
I.a Struct. missing	Format	Subsumed by D (constr. decode)	n/a	A
I.b Semantic missing	Verification	Subsumed by G/H	n/a	C
J. Hallucination	Knowledge	RAG [Wood and Forbes, 2024]	non-zero → 0%	B/C
K. Refusal	Verification	POROver [Karaman et al., 2024]	57.6% → 82.1%	B
L. Tool/API	Format+Verif.	SAGE-Agent [Suri et al., 2025]	36.5% → 65.2%	B

Eight of twelve categories have a strong citation with double-digit percentage-point improvement; three (B, C, I) lack a clean single-cluster ablation; one (F) has consistent medium-strength evidence. Category I (missing required elements) splits mechanistically into I.a (structural, absorbed by D’s constrained decoders) and I.b (semantic, absorbed by G/H). Categories B (units) and C (counting) remain technical gaps: B is naturally folded into A; C is the smallest residual gap.

Additivity and its limits. Patel et al. [2026] demonstrate that stacking interventions targeting orthogonal failure modes can produce large compound reliability gains: their parallel-consensus framework yields a $14,700\times$ improvement over single-pass baseline, evidence for compound gains from decomposition plus consensus aggregation under their specific parallel-voting regime rather than a direct demonstration that heterogeneous cluster-targeted interventions compose additively without voting. Shang et al. [2024] similarly show supra-additive gains when modules target distinct failure modes. Le [2026] reports that schema-level and prompt-level instructions interact *non-additively* when sharing a prompt channel: interventions on orthogonal processing layers (decoding constraint vs. retrieval vs. training-signal vs. inference-time tool call) compose near-additively, while those sharing a channel may interfere.

The irreducible-semantic residual. Of the 17 ErrorAtlas categories, 13 are addressed under Patterns A, B, or C by one of the six capability axes; four are not: residuals of inappropriate refusal beyond preference optimisation, specification misinterpretation not closed by clarification, problem-decomposition reasoning bottlenecks, and a “user wanted something different” semantic remainder. These are classes where the failure is in choosing what to do, not executing it. Proposition 2’s prediction of $O(\log)$ rather than $O(0)$ residual reliability reflects exactly this irreducible core. The framework does not claim 100% coverage of all failures; it claims polylog-bounded *capability-eliminable* failures, with the named semantic residual as the floor.

885 C Counter-Evidence Re-Audits

Five prominent papers are routinely cited as evidence that LLM reliability decays steeply with length [Dziri et al., 2023, Kuratov et al., 2024, Kwa et al., 2025, Wan et al., 2026, Karpinska et al., 2024]. A careful re-reading shows that *every one* of these papers decays over a variable distinct from raw token length n . Identifying the decay axis is not the same as dissolving the practical concern: where compositional graph size, fact count, or evidence scope grow with problem length, the framework predicts steep failure curves. The contribution is to identify which interventions help (capability provisioning along the actual decay axis) and which do not.

Dziri et al. [2023], Faith and Fate. GPT-4 multi-digit multiplication accuracy drops from 59% (3-digit) to 4% (4-digit) zero-shot, with the authors theorising “probability of incorrect predictions

converges exponentially to ≈ 1 for abstract compositional tasks.” The decay variable is *compositional graph size* N , not raw token length: the 3×3 graph has on the order of d^2 partial products plus carries, and multi-digit multiplication is engineered so that $k_{\text{hard}} \approx N$, with every node in the computation graph a hard decision. For natural-language tasks where $k_{\text{hard}} \ll n$, Dziri’s regime is the *boundary case* in which our framework reduces to their result. A clean $(1 - \varepsilon)^N$ exponential cannot simultaneously reproduce 59% at 3-digit and 4% at 4-digit for any single per-node ε : the observed drop is locally steeper than per-node iid exponential, consistent with a finite catalogue of failure modes exhausting as N grows. *Where the framework concedes*: in adversarial compositional tasks engineered so $k_{\text{hard}} \approx N$, the framework reduces to Dziri’s regime and does not relieve it; the relocation is informational, not magical.

Kuratov et al. [2024]. The abstract reports “performance declines sharply with increased reasoning complexity” and models “effectively utilise only 10–20% of the context.” Read in detail, BABILong varies two axes: context length n (0K to 10M tokens) and number of supporting facts k (QA1 = 1 fact to QA3 = 3 facts). The sharp decline is in k , not n . For QA1 (single-fact), most models “perform well up to 4,000 tokens”: a plateau, not exponential decay. The famous RAG-flat-across-length result (60% on single-fact QA *independent of context length*) is the cleanest possible demonstration that when relevant evidence is in window, length does not matter. Recurrent Memory Transformers maintaining performance to 50M tokens further confirms effective k is determined by architecture, not raw n . The framework’s concession is narrow but real: for multi-hop tasks whose required fact count grows with task complexity, BABILong’s sharp k -axis decay is exactly what the framework predicts happens, not a counter-example.

Kwa et al. [2025] (METR). Per-model success fits a logistic in $\log(\text{human-task-duration})$: $S(t) = \sigma(\beta(\log t - \log h))$. Logistic-in-log-length is *mathematically sublinear* in length itself. The 80%-horizon being 4–6 $\times$ shorter than the 50%-horizon is a steep within-model cliff but compatible by construction with the polylog result: a cliff at a specific capacity threshold is a manifold-transition signature. The famous exponential (capability doubling every seven months) is an *inter-model* claim about how the horizon h moves across generations, orthogonal to within-model decay shape. An honest qualifier: the within-model logistic-in-log cliff is steep on a practical scale. Calling it “sublinear in length” is technically correct but engineering-useful only for systems designed to operate well below the 50% horizon.

Wan et al. [2026], Fano-style upper bound. This paper theorises super-linear information-demand growth and identifies an “accuracy cliff” at capacity overflow: “when the task’s information demand surpasses the model’s output capacity, performance does not degrade gracefully but instead collapses sharply.” The behavioural prediction is a *threshold*, not smooth $(1 - \varepsilon)^n$. A cliff at a specific capacity threshold is exactly the manifold-transition behaviour our framework predicts at the boundary between covered and uncovered modes; Wan et al.’s mechanism (capacity overflow) is one specific cause, complementary to ours. *Where the framework concedes*: Wan documents a mechanism by which $|C|$ effectively exceeds the model’s coverage in a single forward pass; Pattern A interventions (constrained decoding, formal verification) are the kind of response the framework expects but does not automatically supply.

Karpinska et al. [2024]. GPT-4o achieves 55.8% pair accuracy across 1,001 minimally-different true/false claim pairs about long fictional books (mean length 127K tokens). The paper reports performance by *evidence scope*, not by context length: 59.8% on sentence-level retrieval, 47.6% on passage-level, 41.6% on global reasoning. No per-context-length curve is reported. The decay axis is the number of evidence pieces that must be integrated, a k -axis quantity rather than an n -axis one. NoCha is therefore not counter-evidence to a sublinear-in- n claim. Evidence-scope decay is exactly the k -axis observation the framework predicts cannot be addressed by scaling raw context length; it requires retrieval or process supervision along the actual decay axis. The relocation is informational, not magical.

Unifying observation. Every counter-paper decays over a variable that is not n : compositional graph size, fact count, log-time horizon, capacity threshold, or evidence scope. The apparent rapid decay is, in each case, in a quantity our framework already concentrates the action in (k_{hard} and $|C|$), not in raw sequence length. This is a relocation, not a dissolution. Where k_{hard} grows with task length (adversarial compositional structure, multi-hop fact chains, long horizons that force more

949 decisions), reliability remains hard; the framework’s value is directing intervention toward capability
 950 provisioning along the actual decay axis rather than toward context-window or compute-budget
 951 expansion that does not help.

952 D Patch Evidence Detail

953 The body’s claim that domain patches cap the engineering problem rests on a triangulation of indirect
 954 evidence. None of the items below *measures* the patch ceiling $|C_D|$ directly; they support the weaker
 955 claim that model behaviour is strongly domain-dependent and that operational neighbourhoods
 956 occupy bounded regions of the model’s behavioural space.

Evidence type	What it supports	What it does not prove
Low intrinsic-dimensional structure in representations [Park et al., 2024, Li and Sarwate, 2025]	Domains occupy localised structure in model representation space	Does not measure failure-catalogue size $ C_D $
958 Cross-domain accuracy spreads on a single base model [Wang et al., 2024c]	Same model behaves differently across domains; performance is patch-dependent	Does not imply finite failure modes
Long-tail / popularity thresholds in knowledge tasks [Mallen et al., 2023, Kandpal et al., 2023]	Coverage depends on domain frequency and training exposure	Does not prove logarithmic mode discovery
Patch-specific intervention performance (across §4, Appendix B)	Local tools, schemas, and capability libraries change residual error structure	Does not imply universal cross-patch transfer of the same intervention library

959 The purpose of this evidence is motivational. It supports patch-indexing of $\sigma_D, A_D, \beta_D, |C_D|$, but
 960 the finite (or effectively capped) reachable catalogue remains an empirical modelling assumption to
 961 be measured per deployment, not derived from any of the rows above.

962 E Empirical Calibration Detail

963 Three published LLM error taxonomies anchor the mode-rate parameter σ in the body’s $\sigma \in$
 964 $[0.87, 1.85]$ range. The simple calibration uses $|C| \approx A + \sigma \ln T$ with $A = 0$; this is a deliberately
 965 conservative readout that ignores any positive intercept and uses endpoint counts rather than discovery
 966 curves.

Source	Approx. observed failures T	Named categories $ C $	Implied σ at $A = 0$	Role
ErrorAtlas [Ashury-Tahan et al., 2026]	$\gtrsim 10^4$	17	≈ 1.85	Conservative cross-domain anchor (used as planning value)
968 HumanEval categorisation [Wen et al., 2024]	comparable scale	8–12	≈ 0.87 – 1.30	Code-domain anchor
MWPES-300K [Sun et al., 2025]	$\approx 3 \times 10^5$	15–20	≈ 1.2 – 1.6	Math-domain anchor (largest corpus)

969 These are endpoint counts, not discovery curves: they tell us *how many* named categories a taxonomer
 970 assigned at a single corpus scale, not how the count grew with T . They do not prove logarithmic
 971 mode discovery. Their role in this paper is twofold: they motivate the empirical postulate of §3.2, and
 972 they identify the explicit empirical test that would either tighten the postulate or move the analysis to
 973 the Heaps variant of Appendix A.5: repeatedly subsample failures from a fixed deployment patch D
 974 and plot discovered modes $|C_{\text{seen},D}(T)|$ against T .