
Telegraph English: Semantic Prompt Compression via Structured Symbolic Rewriting

Anonymous Author(s)

Affiliation

Address

email

Abstract

We introduce Telegraph English (TE), a prompt-compression protocol that rewrites natural language into a symbol-rich, formally-structured dialect. Where token-deletion methods such as LLMingua-2 train a classifier to delete low-importance tokens at a fixed ratio, TE performs a full semantic rewrite: it decomposes the input into atomic fact lines, substitutes verbose phrases with ~ 40 logical and relational symbols, and lets the compression ratio adapt to each document’s information density. A consequence of the line-structure rule is that compression and semantic chunking become the same operation—each output line is an independently addressable fact, so the compressed representation is simultaneously a semantic index. We evaluate TE on 4,081 question-answer pairs from LongBench-v2 across five OpenAI models and two difficulty levels. At roughly 50% token reduction, TE preserves 99.1% accuracy on key facts with GPT-4.1 and outperforms LLMingua-2 at matched compression ratios on every model and task tested. The gap widens on smaller models—up to 11 percentage points on fine-detail tasks—suggesting that explicit relational structure compensates for limited model capacity. We release the grammar specification, compression prompt, benchmark data, and reference implementation.

1 Introduction

Large language models are increasingly embedded in retrieval-augmented generation (RAG), multi-agent orchestration, and long-context reasoning pipelines. Input cost scales linearly with token count, so prompt compression—feeding fewer tokens to the model while preserving the information it needs—has become a practical lever for controlling latency and cost.

Two families of approach exist. *Extractive* methods select a subset of tokens or sentences from the input [Jiang et al., 2023, Pan et al., 2024]; *abstractive* methods paraphrase or summarise it [Chevalier et al., 2023]. LLMingua-2 [Pan et al., 2024], currently the strongest published baseline, trains a GPT-4-distilled XLM-RoBERTa-large classifier to delete tokens below an importance threshold at a user-specified ratio.

Token deletion works, but it has structural limits that become visible once one looks past the compression ratio. The ratio is fixed regardless of input density. Deleting tokens can sever co-reference chains and destroy logical connectives, leaving the downstream model to hallucinate the relationships between surviving fragments. Token-deletion methods are input-only preprocessors—they compress the initial prompt, but generated output passes uncompressed to the next pipeline stage, so multi-step agent systems cannot compound the savings. Most consequentially, token deletion produces no structure: the output is a degraded copy of the input, unable to be indexed, selectively pruned, or dynamically updated.

36 We propose Telegraph English (TE), a different kind of compression. Rather than selecting which
37 tokens to keep, TE rewrites the passage into a compact, formally-structured dialect. The original
38 sentence

39 *“According to research by Author A and colleagues (2023), the application of machine*
40 *learning techniques to medical diagnostics resulted in a 27.5% increase in early detection*
41 *rates while simultaneously reducing false positives by approximately 12% compared to*
42 *traditional methods.”*

43 becomes, under TE:

44 `ML→MEDICAL-DIAGNOSTICS: EARLY-DETECTION+27.5% ^ FALSE-POSITIVE-12%`
45 `[AUTHOR-A:2023]`

46 Sixty-eight tokens become fourteen. The causal relationship, both quantitative claims, and the citation
47 are each on record as separate, addressable units—and the phrase “application of...resulted in” has
48 collapsed into a single symbol.

49 What makes TE architecturally distinctive is a property that emerges from the grammar’s line-structure
50 rule: compression and semantic chunking are the same operation. Every TE output line contains
51 exactly one atomic fact—one claim, one relationship, one datum. This is not a post-processing step
52 but a consequence of how the grammar defines a legal output. The result is a representation that is
53 simultaneously compressed, retrieval-ready, and amenable to dynamic management: atomic lines are
54 individually embeddable; tagged sections support hierarchical context budgeting; and facts can be
55 updated, merged, or pruned without re-running the compressor.

56 **Contributions.** (1) A formal grammar specification for structured prompt compression (§3). (2) A
57 unified compression-and-chunking framework where semantic compression, retrieval-ready indexing,
58 and dynamic context management emerge from a single rewriting pass (§3, Appendix A). (3) A
59 large-scale empirical comparison against LLMingua-2 on 4,081 key-fact and 801 fine-detail QA
60 pairs across five models (§5). (4) Evidence that the advantage of semantic rewriting over token
61 deletion grows on smaller models and on detail-intensive tasks (§6). (5) A reference implementation
62 with CLI tools for compression, benchmarking, and error analysis.

63 2 Related Work

64 **Prompt compression.** LLMingua [Jiang et al., 2023] introduced budget-constrained prompt
65 compression using perplexity-based token selection. LLMingua-2 [Pan et al., 2024] improved on
66 this with a data-distillation approach: GPT-4 labels token importance on the MeetingBank corpus,
67 and an XLM-RoBERTa-large classifier learns to predict which tokens to delete. The architectural
68 constraint is that the output remains a degraded subset of the input tokens—no new structure is
69 introduced.

70 **Abstractive compression.** AutoCompressors [Chevalier et al., 2023] train summary tokens that
71 substitute for long contexts; RECOMP [Xu et al., 2024] compresses retrieved documents into
72 extractive or abstractive summaries tailored to the downstream task. Both are effective but lossy by
73 design—they discard information that cannot be recovered, and neither produces a structured output
74 that supports selective manipulation.

75 **Structured representations and agent context.** Chain-of-thought prompting [Wei et al., 2022]
76 demonstrates that imposing structure on LLM inputs improves reasoning. TE extends this insight
77 to compression: explicit logical and relational operators help downstream models reconstruct the
78 intended meaning more reliably than degraded natural language. For long-running agents, MemGPT
79 [Packer et al., 2023] addresses context-window growth via virtual memory hierarchies, and Reflexion
80 [Shinn et al., 2023] maintains explicit self-reflection buffers; both operate on natural-language
81 representations. TE offers a complementary strategy of structured, fact-level representations that can
82 be selectively updated and pruned without further LLM calls.

83 **Semantic chunking for RAG.** Standard RAG pipelines split documents using fixed token windows
84 or sentence-boundary heuristics [Lewis et al., 2020, Gao et al., 2023]. TE sidesteps the chunking

question: compression produces atomic fact lines as a structural by-product, so no separate chunking stage is needed. Conceptually adjacent are controlled natural languages such as Attempto Controlled English [Fuchs et al., 2008], but those are designed for knowledge representation and formal reasoning rather than compression, and consumed by formal reasoners rather than LLMs.

3 The Telegraph English Grammar

The grammar (version 5) lives in a 430-line specification document that doubles as the system prompt for the LLM-based compressor. We summarise its key design principles here; the full specification is supplementary material.

3.1 Foundations

Four principles govern the grammar, in strict priority order: (i) **fidelity over brevity**—no information may be dropped unless inferable from what remains; (ii) **atomic line structure**—each line contains exactly one claim, step, event, or question; (iii) **upper-case default**, except where case carries information (proper names, code, SI symbols); (iv) **target compression $\sim 5\times$** when feasible, but correctness, auditability, and reversibility take strict priority over token reduction.

3.2 Symbol vocabulary

TE defines a fixed vocabulary of relational and logical operators. The full set numbers roughly 40; Table 1 shows the core symbols that appear in most compressions. Each symbol has a single, non-interchangeable meaning. The grammar caps symbol density at three consecutive symbols per line—a readability constraint learned from early iterations where dense symbol chains became opaque even to GPT-4.

Table 1: Core relational and logical operators in the TE symbol vocabulary. The full vocabulary contains roughly 40 symbols organised by function (causal, logical, comparative, modal).

Symbol	Meaning	Example
=	Definition / equality	VELOCITY=DISTANCE/TIME
→	Causation / flow	HEAT→EXPANSION
⇒	Logical implication	RAIN⇒WETNESS
∴	Therefore / conclusion	X>Y ∧ Y>Z ∴ X>Z
∵	Because / reason	MOTOR-FAILURE ∵ OVERLOAD
↑/↓	Increase / decrease	TEMPERATURE↑
∧/∨/¬	And / or / not	A∧B, ¬EVIDENCE
≈/≠	Approximate / not equal	COST≈USD10M
VS	Contrast (never causal)	MODEL-A VS MODEL-B

3.3 Tags and domain conventions

Beyond the symbol vocabulary, the grammar provides a tagging system and a set of domain-specific formatting rules. Tags handle the framing that natural language carries through verbose syntactic constructions: temporal state (PAST:, NOW:, FUTURE:), modality (LIKELY:, POSSIBLE:, CONF=0.87), roles (AGENT:, PATIENT:, INSTRUMENT:), scope (CTX: for shared context), and structured content types (DEF:, Q:/A:). Each tag does double duty: it collapses verbose framing into a single token and provides the structural handle that downstream systems use for selective retrieval and context management.

Domain conventions standardise the surface forms that vary most across writers: quantities (VAR=VALUEUNIT), citations ([AUTH:YEAR], DOI:, ARXIV:), financial data (USD10.5M, Y/Y+5%, +2.5PT), and URLs. Locking these down at the grammar level removes a class of factual-error failure modes the compressor would otherwise need to handle case by case.

117 3.4 Compressor self-verification

118 Two mechanisms keep the compressor honest within a single LLM call: a quality gate and a prescribed
119 distillation sequence. The quality gate is a 12-point checklist covering formatting consistency, symbol
120 precision, abbreviation policy, number formatting, information preservation, and citation integrity.
121 It is embedded directly in the compression prompt, so the compressor self-verifies output before
122 returning it. The distillation sequence prescribes a six-pass reasoning order: (1) concept identification,
123 (2) claim extraction, (3) relation mapping, (4) redundancy elimination, (5) numerical verification,
124 (6) citation cross-checking. This is a chain-of-thought scaffold inside a single inference, not a
125 multi-call pipeline. Ordering matters: numerical verification before citation cross-checking, because
126 citations sometimes attach to numerical claims that must be confirmed first.

127 3.5 Compression as semantic chunking

128 Compression and semantic chunking are not separate stages—they are the same operation. Every
129 TE output line is an atomic fact, every section is tagged, and every CTX: block defines a scope. The
130 structure falls out of the grammar’s line-structure rule, not from any additional processing, and it
131 enables three things that token-deleted text cannot support: **selective retrieval** (a query about adverse
132 events retrieves exactly the relevant line and its scope, no chunking heuristic required); **graduated**
133 **compression-on-read** (a context-assembly system can keep the most relevant lines at full fidelity,
134 retain only heading tags for moderately relevant sections, and drop irrelevant sections entirely—no
135 LLM call needed); and **continuous state refinement** (facts can be updated in-place, merged, or
136 pruned during a session). We call this the *compress-once, manage-continuously* principle. A worked
137 example and a more detailed treatment appear in Appendix A.

138 4 Experimental Setup

139 4.1 Dataset

140 LongBench-v2 [Bai et al., 2025] contributes 503 long-context multiple-choice items; we use their
141 source documents as our corpus, generating our own QA pairs per chunk (§4) rather than reusing
142 the bundled questions. We filter to three categories suitable for factual QA—Single-Document QA,
143 Multi-Document QA, and Long-Dialogue History Understanding—which leaves 339 documents.
144 NLTK sentence tokenisation chunks each one into segments of at most 1,000 words, producing
145 4,081 chunk-level evaluation units. The categories span technical reports, multi-source narrative
146 synthesis, and conversational history—three regimes where compression methods fail differently.
147 The 1,000-word chunk cap matches the practical input size for which prompt compression actually
148 saves money.

149 4.2 Compression

150 Each chunk is compressed into TE using the v5 grammar prompt with OpenAI’s o4-mini model.
151 Token counts are measured with tiktoken (cl100k_base). The mean compression ratio is 0.585—a
152 41.5% token reduction—with a range from 0.13 to 1.57. The upper end deserves explanation: rare,
153 very short inputs that are already informationally dense occasionally *expand* under TE, because the
154 grammar’s fidelity-first principle prohibits dropping information even when doing so would reduce
155 token count. This is a feature, not a failure. The full distribution is shown in Figures 1 and 2.

156 For the LLMLingua-2 baseline, the same chunks are compressed using the publicly available
157 llmlingua package at two retention rates: 0.50 (50% kept) and 0.33 (33% kept).

158 4.3 QA evaluation protocol

159 We design a multiple-choice protocol that isolates comprehension: can a model answer a factual
160 question correctly when reading compressed text instead of the original? GPT-4.1 generates a verbatim
161 QA pair from the original chunk, plus a semantically equivalent “modified answer” that prevents
162 simple string matching from inflating scores. GPT-4.1 (temperature 0.7) generates three plausible
163 distractors matched in style, length, and specificity. The modified answer and three distractors are
164 shuffled into a four-option question. The evaluation model sees the original, then the compressed text,




Figure 1: Distribution of compression ratios across 4,081 LongBench-v2 chunks compressed with TE (o4-mini, tiktoken cl100k_base). Mean 0.585, range 0.13–1.57. The right tail above 1.0 corresponds to short, dense inputs that expand under TE.




Figure 2: Distribution of per-chunk compression rate ($1 - \text{ratio}$) over the same corpus. The median chunk loses roughly 43% of its tokens; the bottom decile loses very little, reflecting TE’s adaptive behaviour on already-dense inputs.

165 and selects an answer in each setting. Accuracy is the fraction of correct selections; an *error* is a case
 166 where the model answered correctly on the original but incorrectly on the compressed version.

167 4.4 Test suites and models

168 Two suites probe different levels of information preservation. **key_facts** (4,081 QA pairs) targets core
 169 concepts—headline findings, main claims, central arguments—with generically plausible distractors.
 170 **fine_facts** (801 QA pairs) is adversarially designed to target information that lossy compression is
 171 most likely to destroy: precise numerical qualifiers, conditional statements, boundary conditions,
 172 secondary details. Distractors are near-miss variants—e.g. changing 4.8% to 4.3%—that can only be
 173 distinguished with access to the exact original detail.

174 We evaluate five OpenAI models: GPT-4.1, GPT-4o, GPT-4o-mini, GPT-4.1-nano, and a fine-tuned
 175 GPT-4o variant. GPT-4.1 also generates the QA pairs and distractors. Different suites use different
 176 model subsets: **key_facts** is run on GPT-4.1, GPT-4o-mini, and GPT-4.1-nano; **fine_facts** on GPT-4o
 177 and GPT-4o-mini. The fine-tuned variant is not used as a separate accuracy benchmark—it serves as
 178 a sanity check that fine-tuning on the original distribution does not change comparative behaviour at
 179 compression-decoded inputs.

180 5 Results

181 5.1 Key facts accuracy

Table 2: Accuracy on the **key_facts** suite (4,081 QA pairs). TE is Telegraph English at ~50% compression; LLML2-50 is LLMingua-2 at 50% retention. Drop is in percentage points (pp) relative to original. **Bold** = best compressed.

Model	Original	TE	LLML2-50	TE Drop	LLML2-50 Drop
GPT-4.1	1.000	0.991	0.990	−0.9	−1.0
GPT-4o-mini	0.991	0.957	0.946	−3.4	−4.5
GPT-4.1-nano	0.980	0.950	0.949	−3.0	−3.1

182 On headline facts (Table 2), TE matches or edges out LLMingua-2 across the board. The accuracy
 183 loss is negligible for GPT-4.1—less than a percentage point while halving the token count. The gap
 184 widens on smaller models: 1.1 pp on GPT-4o-mini, with the same direction at GPT-4.1-nano. Not
 185 dramatic. But consistent—the direction never reverses across configurations.

Table 3: Accuracy on the adversarial `fine_facts` suite (801 QA pairs). Fine-detail tasks expose larger compression effects; TE preserves more than LLMLingua-2 at matched retention.

Model	Original	TE	LLML2-50	TE Drop	LLML2-50 Drop
GPT-4o	0.996	0.965	0.933	−3.1	−6.3
GPT-4o-mini	0.938	0.843	0.820	−9.5	−11.8

186 5.2 Fine facts accuracy

187 Fine details are harder (Table 3). Compression loss runs 3–4 $\times$ higher than on key facts, regardless of
 188 method. TE holds an advantage of 3.2 pp over LLMLingua-2 on GPT-4o and 2.3 pp on GPT-4o-mini
 189 at matched 50% retention. The fine-facts regime is where token deletion starts to break down: the
 190 very tokens the questions probe—numerical qualifiers, conditional caveats, secondary details—are
 191 precisely the ones an importance classifier flags as low-value.

192 5.3 Accuracy hierarchy and compression statistics

193 Across all models and tasks the ranking holds without exception: original > TE > LLML2-50 >
 194 LLML2-33. TE’s mean compression ratio of 0.585 (std = 0.254) hides a wide spread: half the corpus
 195 sits between 0.41 and 0.74, with median 0.57. Documents dense with technical content or data tables
 196 resist compression; verbose narrative text yields ratios of 5:1 or better. Fidelity-first design means the
 197 ratio is an outcome, not a parameter.

198 5.4 Error analysis

199 Of the 4,081 `key_facts` items, 187 (4.6%) were correct on the original and incorrect on TE for
 200 GPT-4.1-nano. These error cases have a mean compression ratio of 0.531, slightly more compressed
 201 than the population mean—aggressive compression and error risk are correlated. Failures cluster
 202 around fine details: dates, units, conditional qualifications, and numerical relationships where TE
 203 either abbreviates a critical modifier or collapses a distinction the question specifically probes. One
 204 characteristic failure: a legal-document chunk where TE compressed “no later than 30 calendar days
 205 after receipt of written notice” into DEADLINE=30D-AFTER-NOTICE, and the question asked whether
 206 the deadline was in calendar or business days. The 30D abbreviation does not distinguish. This is a
 207 limitation of the symbol vocabulary, not a compressor error.

208 6 Analysis

209 6.1 Why semantic rewriting outperforms token deletion

210 Four mechanisms explain the pattern in the results. They are not ranked; different mechanisms
 211 dominate in different regimes. **Semantic-unit preservation**: token deletion operates at the token
 212 level and can split multi-word expressions, sever noun-modifier pairs, strand a number from its unit;
 213 TE works one level up, with related concepts grouped into hyphenated compounds and complete
 214 claims occupying single lines. **Explicit logical structure**: when LLMLingua-2 deletes a connective
 215 like “therefore” or “in contrast to,” the downstream model has to guess the relationship; TE refuses to
 216 offer the guess, with `∴`, `VS`, `→` each unambiguous and preserved regardless of what else is removed.
 217 **Co-reference stability**: TE’s one-claim-per-line discipline and upper-case entity naming eliminate
 218 pronoun resolution ambiguity; token deletion can strand a pronoun whose antecedent has been
 219 removed. **Adaptive compression**: a fixed-ratio method compresses dense and verbose passages
 220 identically; TE does not—dense passages emerge at ratios near 1.0, verbose ones below 0.2. The four
 221 mechanisms interlock, which is why LLMLingua-2 cannot match TE by adopting any single one of
 222 them.

223 6.2 The small-model effect

224 The TE advantage grows as model capacity shrinks. GPT-4.1 barely notices the difference between
 225 TE and LLMLingua-2 on key facts; GPT-4.1-nano and GPT-4o-mini show a wider gap, and on fine

facts the divergence becomes substantial. The likely explanation is capacity-dependent. Smaller models have less ability to reconstruct implicit relationships from token-deleted fragments; TE compensates by offloading that reconstruction work to the compression stage—the evaluation model receives a representation where the relationships are already marked, rather than having to hallucinate them from sparse clues. This has practical weight: smaller models are precisely the ones deployed in cost-sensitive production pipelines, which is where prompt compression earns its keep.

6.3 The fine-facts gap

Key facts survive both compression methods reasonably well. Central claims are often redundantly signalled, and even aggressive token deletion tends to preserve them. Fine details are stubborn in a different way: precise numerical qualifiers, conditional caveats, and secondary attributions are exactly the tokens an entropy-based classifier flags as low-importance in isolation. A number like “4.8%” may look dispensable next to surrounding prose. But if the question asks whether the figure was 4.8% or 4.3%, that token is the entire answer. TE’s claim-level decomposition and explicit numerical formatting (+27.5%, CONF=0.87, Y/Y+12.3%) are designed to preserve these details: numbers are never abbreviated, always attached to their units, and always placed in a structured format the downstream model can parse unambiguously.

6.4 Pipeline-level cost

There is a structural difference between the two methods that the accuracy comparison alone obscures: LLMingua-2 operates as an input-only preprocessor. It compresses the initial prompt; generated output passes uncompressed to subsequent stages. TE can persist as a native format throughout a pipeline. Consider a five-step agent pipeline with 2,000 tokens of initial context and five generation steps averaging 400 tokens each, at \$10 per million tokens (Table 4). The savings compound because each stage operates on TE-formatted text. LLMingua-2 compresses only the first stage’s input; the remaining four stages process uncompressed output at full token cost. A more architectural treatment of dynamic context management appears in Appendix B.

Table 4: Pipeline-level cost for a five-step agent pipeline (2,000-token initial context, five 400-token generations) at \$10 per million tokens. TE persists across stages; LLMingua-2 compresses only the first stage.

Method	Total tokens	Cost / 1K calls	Savings
Original	4,000	\$40	—
LLMingua-2	~3,300	\$33	\$7
Telegraph English	~1,600	\$16	\$24

7 Implementation

The reference implementation is a Python package with five pipeline stages: synchronous and asynchronous (Batch API) compression of LongBench-v2 documents using the TE grammar prompt; automated quality review via Claude (structured JSON scores 0–10 with strengths, weaknesses, and example pairs); end-to-end QA benchmarking (generation, distractor creation, MC evaluation); LLMingua-2 baseline evaluation against the same QA pairs; and error analysis with case-level output. Each stage is accessible as both a CLI command and an importable library function.

8 Limitations

LLM-dependent compression. TE requires an LLM call per chunk, adding latency and cost at compression time. This is amortised when compressed text is reused, but TE is poorly suited for compressing ephemeral inputs that will be read once and discarded. **Proprietary evaluation models.** Our benchmark relies on OpenAI models that are not open-weight, limiting reproducibility; future work should extend evaluation to open models. **English only.** The grammar and benchmarks are English; adapting the symbol vocabulary to other languages—particularly agglutinative or logographic ones—is non-trivial. **Compressor model sensitivity.** TE quality depends on the model performing

the rewrite; we have not yet mapped this sensitivity curve. **QA generation bias.** Both QA pairs and evaluations are produced by OpenAI models; an ideal evaluation would include human-written questions or a diverse set of QA generators. **Dynamic context management is not yet benchmarked.** The semantic chunking and dynamic state-management capabilities described in §3 and Appendix B are architectural arguments, not empirical results from a multi-turn evaluation. We have demonstrated format-level feasibility; we have not measured downstream effects over extended sessions. This is the most important gap in the current evaluation. **Comparison scope.** We benchmark against LLMingua-2 only—currently the strongest published baseline at our compression ratios. A broader comparison against AutoCompressors, RECOMP, and more recent methods would strengthen the claims.

9 Conclusion

Telegraph English demonstrates that structured semantic rewriting is a viable alternative to token deletion for prompt compression—and, on the evidence presented here, a better one. The advantage is largest where compression matters most practically: on smaller, cheaper models and on fine-grained details. The quantitative comparison may not be the most interesting part of this work. Token-deletion methods produce a smaller copy with no internal organisation; TE produces a representation where every line is an identified fact, every section is tagged, every relationship is marked with an explicit symbol. That structure makes the output not just smaller but more *useful*—more retrievable, more auditable, more maintainable over time. The compress-once, manage-continuously principle is, at this stage, an architectural argument rather than an empirical result; validating it in production agent systems is the obvious next step.

References

- Yushi Bai et al. Longbench v2: Towards deeper understanding and reasoning on realistic long-context multitasks. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*, pages 3639–3664, 2025.
- Alexis Chevalier, Alexander Wettig, Anirudh Ajith, and Danqi Chen. Adapting language models to compress contexts. In *Proceedings of EMNLP*, 2023.
- Norbert E. Fuchs, Kaarel Kaljurand, and Tobias Kuhn. Attempto controlled english for knowledge representation. In *Reasoning Web*, volume 5224 of *Lecture Notes in Computer Science*, pages 104–124. 2008.
- Yunfan Gao et al. Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*, 2023.
- Huiqiang Jiang, Qianhui Wu, Chin-Yew Lin, Yuqing Yang, and Lili Qiu. Llmilingua: Compressing prompts for accelerated inference of large language models. In *Proceedings of EMNLP*, 2023.
- Patrick Lewis et al. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, 2020.
- Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. Memgpt: Towards llms as operating systems. *arXiv preprint arXiv:2310.08560*, 2023.
- Zhuoshi Pan, Qianhui Wu, Huiqiang Jiang, Menglin Xia, Xufang Luo, Jue Zhang, Qingwei Lin, Victor Rühle, Yuqing Yang, Chin-Yew Lin, H. Vicky Zhao, Lili Qiu, and Dongmei Zhang. LLMingua-2: Data distillation for efficient and faithful task-agnostic prompt compression. In *Findings of ACL*, 2024.
- Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- Jason Wei et al. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems*, volume 35, 2022.

313 Fangyuan Xu, Weijia Shi, and Eunsol Choi. RECOMP: Improving retrieval-augmented LMs with
 314 context compression and selective augmentation. In *International Conference on Learning Repre-*
 315 *sentations (ICLR)*, 2024.

316 A Compression as Semantic Chunking: Worked Example

317 To make §3 concrete, consider a multi-paragraph clinical-trial summary compressed into TE:

```

318 H1: CLINICAL-TRIAL OUTCOMES
319 CTX: PHASE-III RANDOMISED CONTROLLED-TRIAL (RCT); N=2400
320 PRIMARY-ENDPOINT: MORTALITY↓23% VS PLACEBO; p<0.001
321 [AUTHOR-B:2024]
322 SECONDARY-ENDPOINT: HOSPITALIZATION↓18%; p=0.003
323 ADVERSE-EVENTS: NAUSEA=12% ^ HEADACHE=8% ^ SERIOUS=2.1%
324 H1: SUBGROUP-ANALYSIS
325 AGE>65: MORTALITY↓31% (STRONGER-EFFECT)
326 AGE<65: MORTALITY↓14% (WEAKER-EFFECT)
327 CONF=0.92 FOR INTERACTION-EFFECT
328 H1: LIMITATIONS
329 FOLLOW-UP=18 MONTHS; LONG-TERM-EFFECTS UNKNOWN
330 EXCLUSION: PATIENTS WITH RENAL-IMPAIRMENT

```

331 Each line is a fact; each heading is a section boundary; each CTX: block defines a scope. The structure
 332 falls out of the grammar’s line-structure rule rather than from any additional processing, and it enables
 333 three things that token-deleted text cannot support.

334 **Selective retrieval.** A query about adverse events retrieves exactly the ADVERSE-EVENTS line and
 335 its CTX: scope. No sliding-window heuristic, no overlap parameter, no risk of splitting a relevant fact
 336 across chunk boundaries; the semantic boundaries are intrinsic to the format.

337 **Graduated compression-on-read.** When assembling a prompt under a tight token budget, an agent
 338 can apply different policies to different sections: keep the lines most relevant to the current query
 339 at full fidelity; retain only the heading tags (H1: LIMITATIONS) for moderately relevant sections,
 340 preserving topic structure at near-zero cost; drop irrelevant sections entirely. This second-stage
 341 compression is semantically principled—it operates on identified sections, not on token positions.

342 **Continuous state refinement.** During a conversation, facts from earlier turns can be revised without
 343 re-compressing the source: *update* (replace a corrected figure in place), *merge* (combine related facts
 344 when the distinction no longer matters), *prune* (remove claims that have moved past relevance), and
 345 *promote/demote* (expand a heading-collapsed section, or collapse a fully expanded one).

346 B Beyond Static Compression: Dynamic Context Architecture

347 The results in §5 measure TE as a static compression method—compress once, read once, evaluate.
 348 This is the fair comparison against LLMLingua-2 and where the benchmark numbers live. But the
 349 more consequential property of TE may not be the compression ratio; it is the structure of the output.

350 **Unifying compression and chunking.** Conventional RAG systems run documents through two
 351 stages: chunking (splitting into fixed-size segments for embedding) and optional compression
 352 (reducing each chunk’s token count). These stages have different objectives and can interfere—a
 353 chunk boundary splits a sentence, then compression deletes the tokens needed to reconstruct it. TE
 354 collapses both stages into one. Each output line is a complete semantic unit; the chunking boundaries
 355 *are* the compression output. A TE-compressed document is immediately embeddable at the line
 356 level. The practical consequence for retrieval precision: fixed-window chunking inevitably includes
 357 irrelevant context within each chunk and risks splitting relevant information; TE surfaces exactly the
 358 facts a query matches, at the granularity of individual claims.

Hierarchical context budgeting. Because TE output is tagged with headings, context scopes, and role markers, a context-assembly system can make graduated decisions about inclusion. For a given token budget: full-fidelity inclusion of all atomic lines for the most relevant sections; heading-only retention for moderately relevant sections, preserving topic structure at near-zero cost; omission of irrelevant sections entirely. This graduated policy can achieve very high total compression (10–50×) when only a fraction of the document is relevant, while maintaining full detail where it matters. The policy operates on the TE output’s structure—no LLM call needed.

Dynamic state in agentic sessions. Long-running agent sessions accumulate context over many exchanges. The standard solutions are blunt: hard truncation drops the oldest tokens regardless of relevance; periodic summarisation requires an LLM call and is irreversible. TE enables something finer. Because context is already decomposed into tagged atomic facts, an agent can maintain a living state: *fact updates* replace the old line in place rather than appending alongside it; *redundancy pruning* removes facts whose information has been absorbed by later ones; *scope closure* collapses an entire CTX: block to a heading once a topic is resolved; *priority re-ranking* reorders facts by current relevance, placing the most important context where transformer attention is strongest. Context growth is controlled by continuously refining the active fact set, not by discarding the oldest tokens. This is cheap (string manipulation, no LLM calls) and semantically principled.

The cost profile is asymmetric by design: one expensive LLM rewrite per document, then indefinite cheap manipulation of the structured output.

C Full Results Tables

Table 5: Complete key_facts results with compression statistics.

Model	n	Original	TE	LLML2-50	TE Drop	LLML2-50 Drop	Mean ratio
GPT-4.1	4,081	1.000	0.991	0.990	−0.9	−1.0	0.585
GPT-4o-mini	4,081	0.991	0.957	0.946	−3.4	−4.5	0.585
GPT-4.1-nano	4,081	0.980	0.950	0.949	−3.0	−3.1	0.585

Table 6: Complete fine_facts results.

Model	n	Original	TE	LLML2-50	TE Drop	LLML2-50 Drop
GPT-4o	801	0.996	0.965	0.933	−3.1	−6.3
GPT-4o-mini	801	0.938	0.843	0.820	−9.5	−11.8

Table 7: Compression ratio statistics (tiktoken cl100k_base, $n = 4,081$ chunks).

Statistic	Value
Mean	0.585
Std	0.254
Min	0.000
25th percentile	0.407
Median	0.570
75th percentile	0.739
Max	1.567

Table 8: Error analysis: key_facts items correct on original, incorrect on TE (GPT-4.1-nano).

Statistic	Value
Total items	4,081
Error items	187
Error rate	4.58%
Mean compression ratio (errors)	0.531
Mean compression ratio (all)	0.585

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and §1 state four main claims: (i) TE matches or outperforms LLMingua-2 at matched compression ratios on every model and task tested; (ii) the gap widens on smaller models, up to ~11 pp on fine-detail tasks; (iii) the line-structure rule makes compression and semantic chunking the same operation; (iv) the reference implementation, grammar, and benchmark are released. Claims (i) and (ii) are supported by Tables 2 and 3 (§5) and Appendix C. Claim (iii) is the architectural argument supported in §3 and Appendix A; we mark it explicitly as architectural rather than empirical and reiterate this scoping in §Limitations. Claim (iv) is the artifact release listed in §Implementation.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: §Limitations enumerates seven specific limitations: LLM-dependent compression cost, reliance on proprietary OpenAI models for evaluation, English-only grammar, compressor-model sensitivity not yet characterised, QA-generation bias, no benchmark of dynamic context management (the most important gap), and comparison scope (LLMingua-2 only).

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.

- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper is empirical and architectural; it contains no formal theorems or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: §4 specifies the source corpus (LongBench-v2), filtering criteria (three QA categories, 339 documents), chunking method (NLTK sentence tokenisation, 1,000-word cap), tokenizer (tiktoken cl100k_base), compressor model (OpenAI o4-mini), grammar prompt version (v5), baseline (LLMLingua-2 at 50% and 33% retention), evaluation models (GPT-4.1, GPT-4o, GPT-4o-mini, GPT-4.1-nano, fine-tuned GPT-4o), and the five-step QA evaluation protocol. The reference implementation reproduces all stages end-to-end (§Implementation).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.

- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We release the grammar specification (CC-BY-SA-4.0), the v5 compression prompt, the LongBench-v2-derived benchmark data (CC0), and a Python reference implementation (MIT License) covering all five pipeline stages. An anonymized snapshot of the code and benchmark data is provided with the supplementary material; the public repository link will be disclosed in the camera-ready version.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: §4 specifies the dataset split (3 LongBench-v2 categories, 339 documents, 4,081 chunks of $\leq 1,000$ words each); the QA generation protocol (GPT-4.1, temperature 0.7 for distractors); the LLMingua-2 retention rates (0.50, 0.33); and the multiple-choice evaluation protocol. No model training is performed, so there are no training-time hyperparameters or optimizer choices.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [No]

Justification: We report point estimates with sample size n (4,081 for key_facts, 801 for fine_facts) but do not report explicit error bars or confidence intervals. At these sample sizes the binomial standard error of a proportion is small (e.g. $\sigma \leq 0.005$ at $n = 4,081$, $p = 0.95$), and the gaps we report between TE and LLMingua-2 are several standard errors wide; evaluation calls use temperature 0 and are deterministic conditional on distractor placement. We did not run multiple seeds of distractor randomization to characterise that variance, which is a limitation we will address in revisions.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: All experiments use OpenAI API endpoints (no local GPUs); the reference implementation runs on consumer CPUs with the OpenAI Batch API doing the heavy lifting. We discuss per-pipeline token cost in §6.4 but do not break out wall-clock time or total token volume per experiment in the main paper. These details will be added in the camera-ready version.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We have reviewed the NeurIPS Code of Ethics. The work involves no human subjects, no privacy-sensitive data, no deployed system, and no scraping beyond the publicly released LongBench-v2 corpus (used under its license). Anonymity is preserved in this submission.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: TE is a foundational efficiency method for LLM input compression and does not involve a deployed system, generative content, or new data collection. Positive impact is reduced compute / cost / latency for retrieval-augmented and agent pipelines; negative impacts are no greater than those generic to any LLM efficiency improvement (e.g. enabling lower-cost downstream applications, both beneficial and harmful). We do not believe the method introduces a unique societal risk warranting a dedicated impact section.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.

- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: We release no pre-trained models, no image-generation assets, and no scraped data. The released artifacts are a grammar specification, a compression prompt, and a benchmark derived from the publicly available LongBench-v2 corpus.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: LongBench-v2 [Bai et al., 2025] is cited and used under its release terms. LLMingua-2 [Pan et al., 2024] is used via the publicly available llmlingua Python package, with citation. OpenAI models (GPT-4.1, GPT-4o, GPT-4o-mini, GPT-4.1-nano, fine-tuned GPT-4o) are accessed through the public OpenAI API in compliance with its terms of service.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The grammar specification (CC-BY-SA-4.0), v5 compression prompt, derived benchmark data (CC0), and reference implementation (MIT License) are documented in §Implementation and provided in the supplementary materials with usage instructions, license files, and per-stage CLI commands.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The work involves no crowdsourcing and no human subjects. All evaluation is automated through API calls to OpenAI models on the LongBench-v2 corpus.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: No human subjects are involved; IRB review is not applicable.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [\[Yes\]](#)

Justification: LLMs are central to the methodology and are explicitly described in §3 and §4: OpenAI’s o4-mini performs the TE compression rewrite using the v5 grammar prompt; GPT-4.1 generates the QA pairs and distractors; GPT-4.1, GPT-4o, GPT-4o-mini, GPT-4.1-nano, and a fine-tuned GPT-4o variant evaluate on compressed and uncompressed inputs. All model identifiers, prompt versions, and roles are specified in the paper.

Guidelines:

- The answer [\[N/A\]](#) means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.